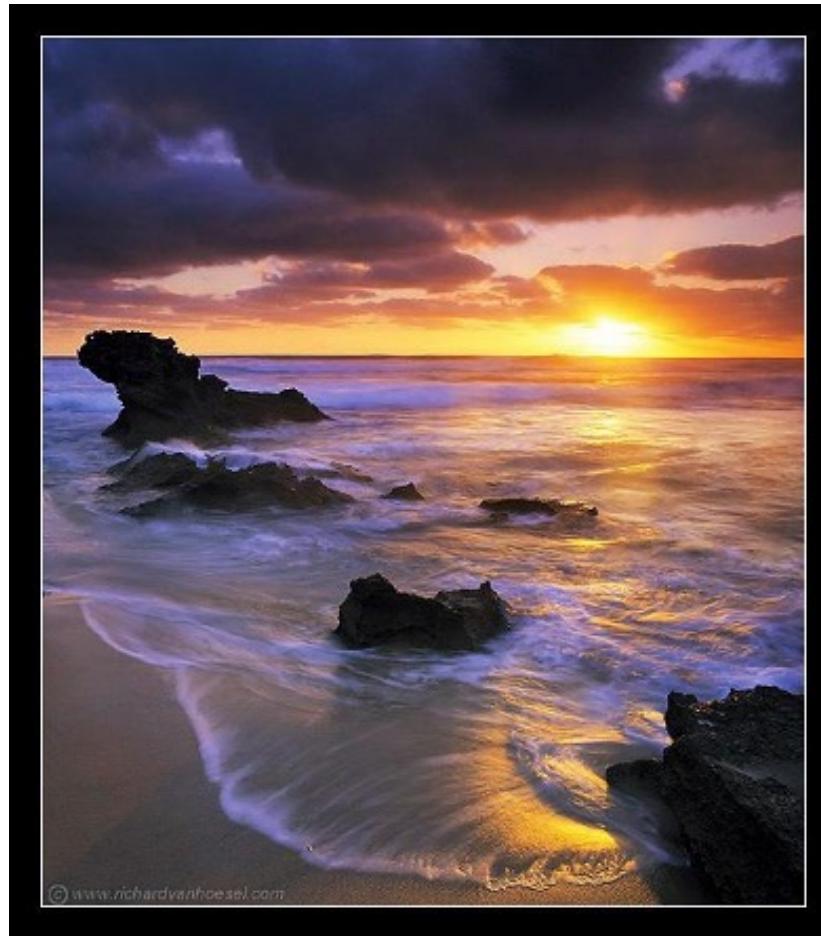


Capturing Light... in human and machine



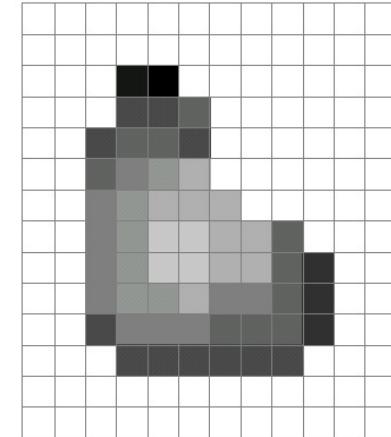
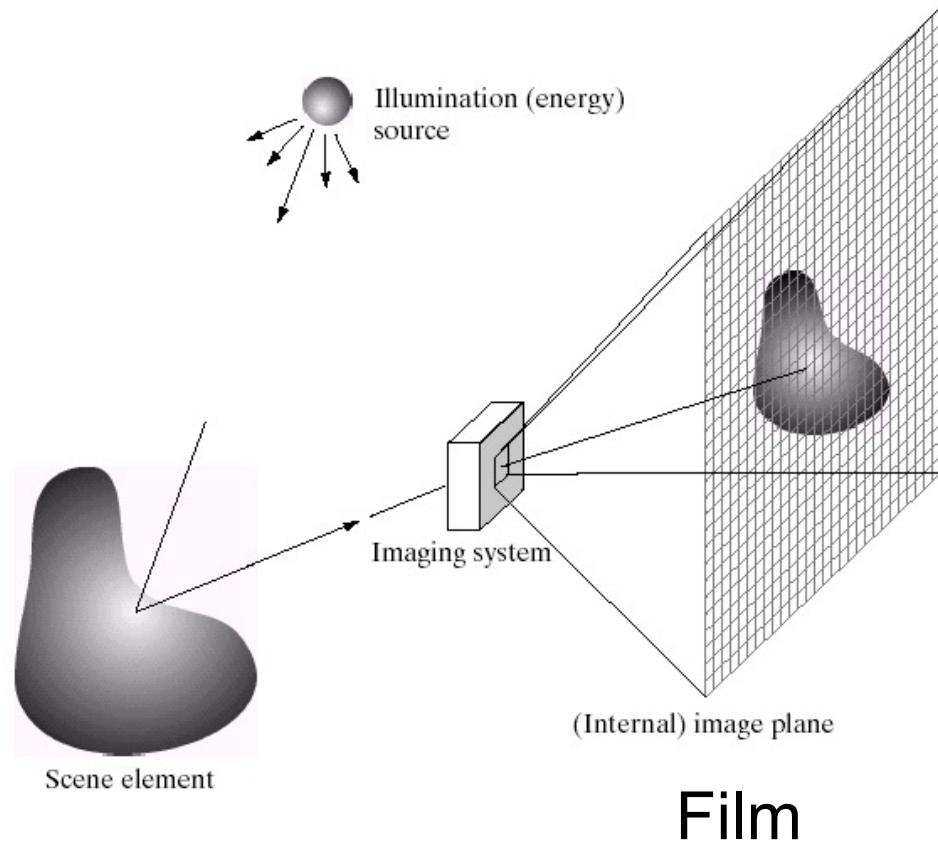
CS194-26: Computer Vision and Comp. Photography
Alexei Efros & Angjoo Kanazawa, UC Berkeley, Fall 2022

Etymology

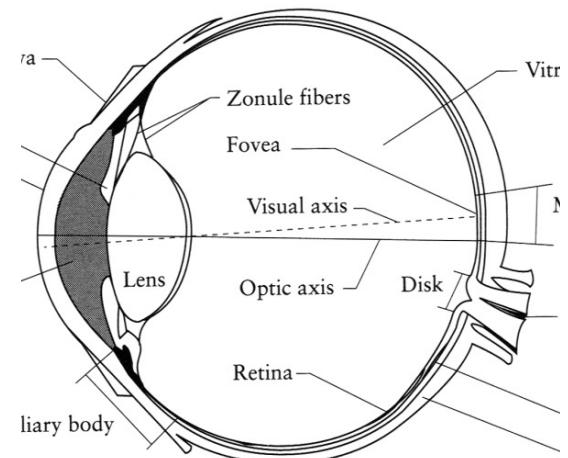
PHOTOGRAPHY

The word "PHOTOGRAPHY" is displayed in large, bold, black capital letters. It is visually divided into two main components by a red curly brace underneath the prefix "PHOTO-" and a blue curly brace underneath the suffix "-GRAPHY". Below the red brace, the word "light" is written in red capital letters. Below the blue brace, the words "drawing" and "/ writing" are written in blue capital letters, stacked vertically.

Image Formation

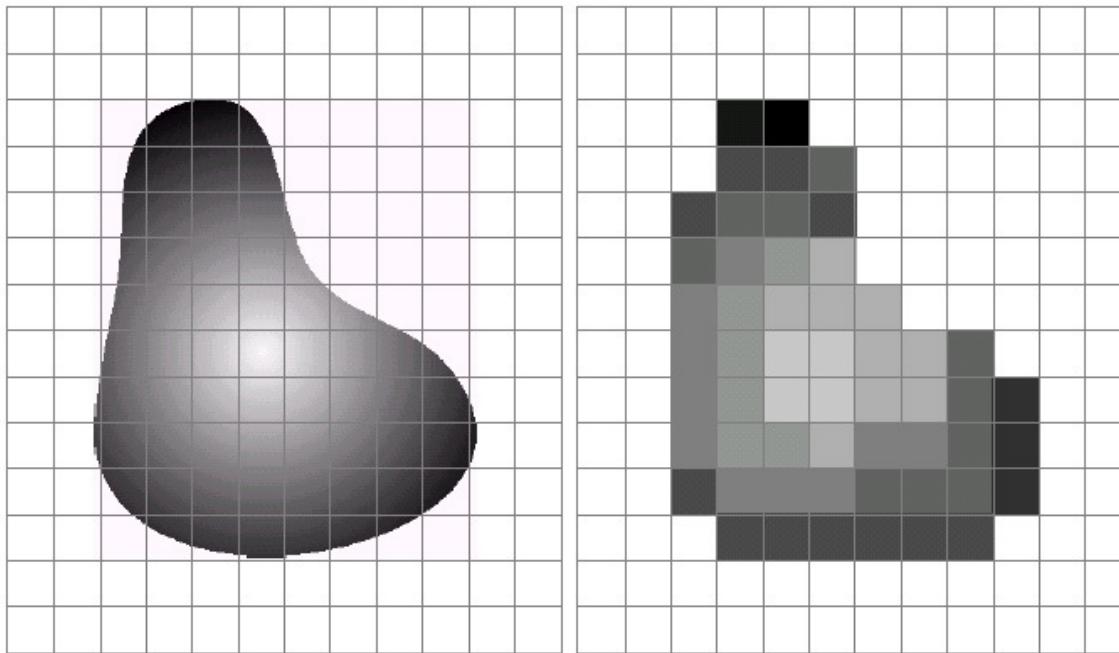


Digital Camera



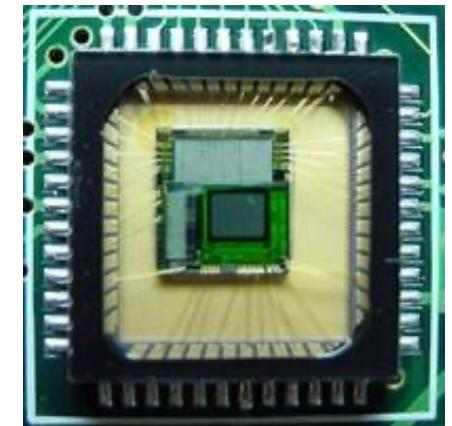
The Eye

Sensor Array



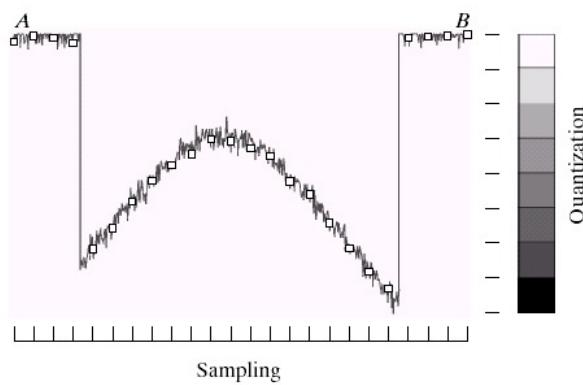
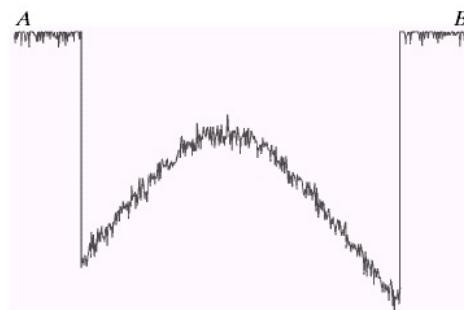
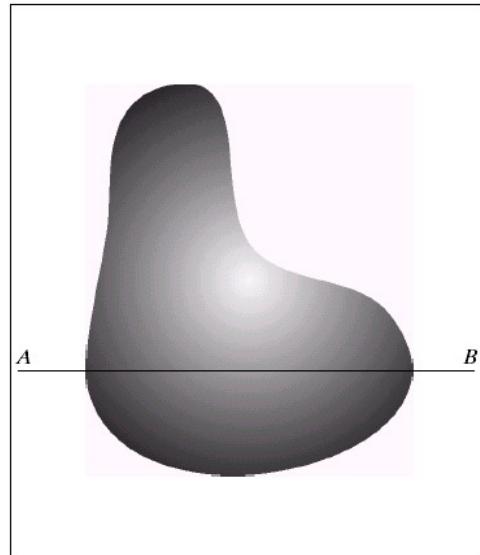
a b

FIGURE 2.17 (a) Continuos image projected onto a sensor array. (b) Result of image sampling and quantization.



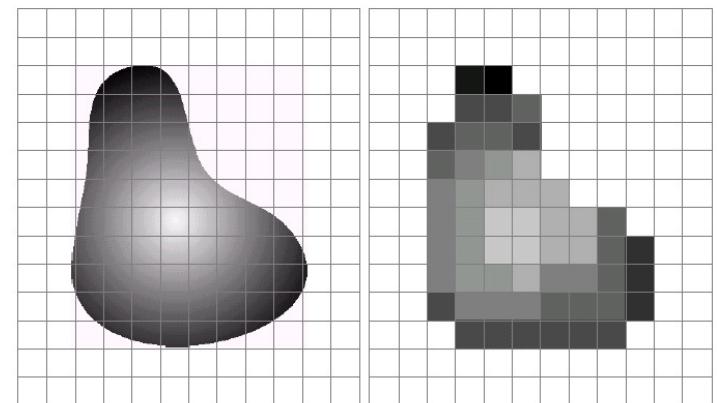
CMOS sensor

Sampling and Quantization



a
b
c
d

FIGURE 2.16 Generating a digital image. (a) Continuous image. (b) A scan line from A to B in the continuous image, used to illustrate the concepts of sampling and quantization. (c) Sampling and quantization. (d) Digital scan line.

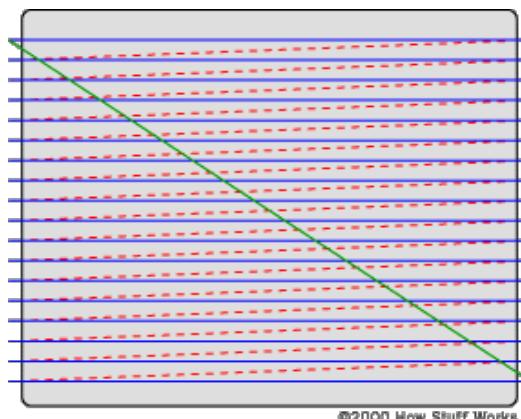
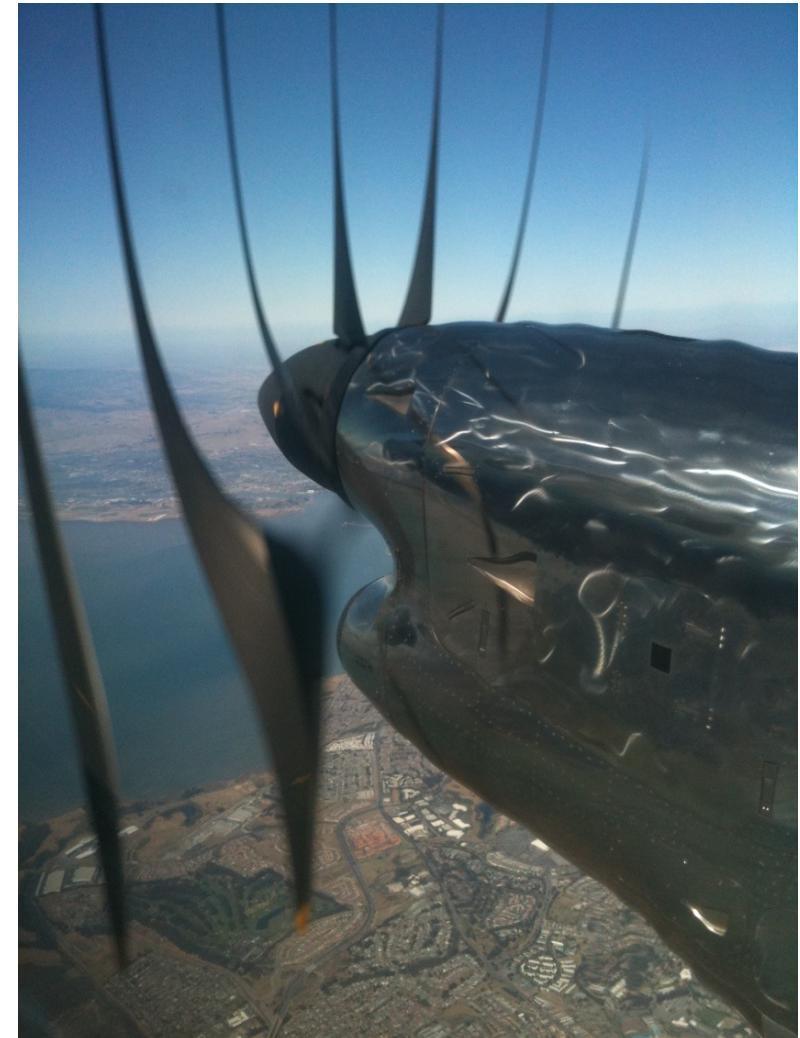


a
b

FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

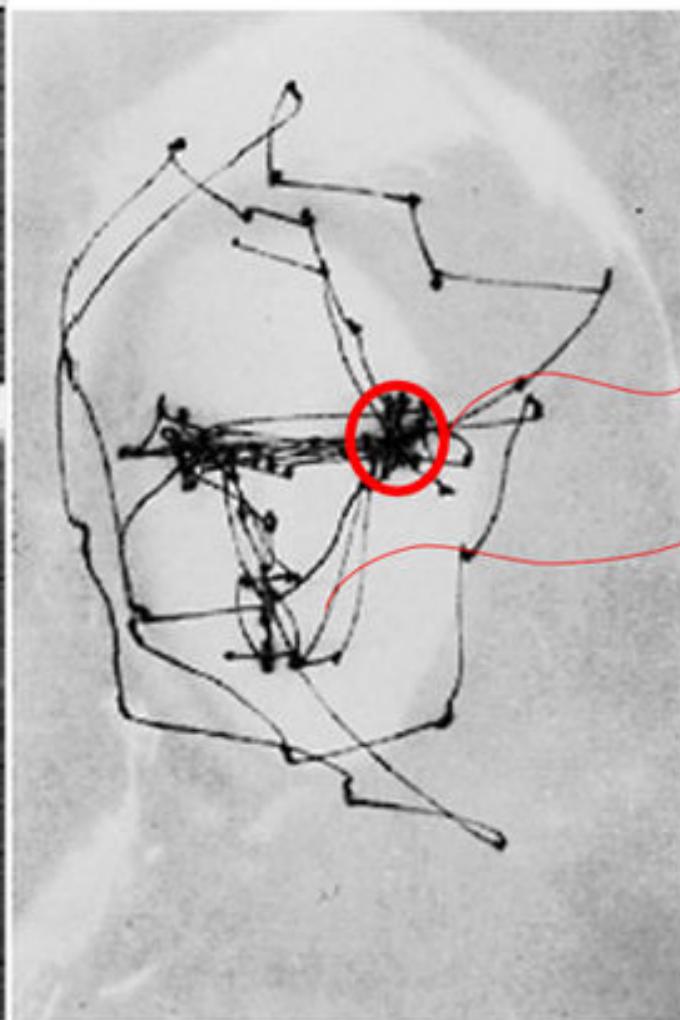
c
d

Rolling Shutter



http://en.wikipedia.org/wiki/Rolling_shutter

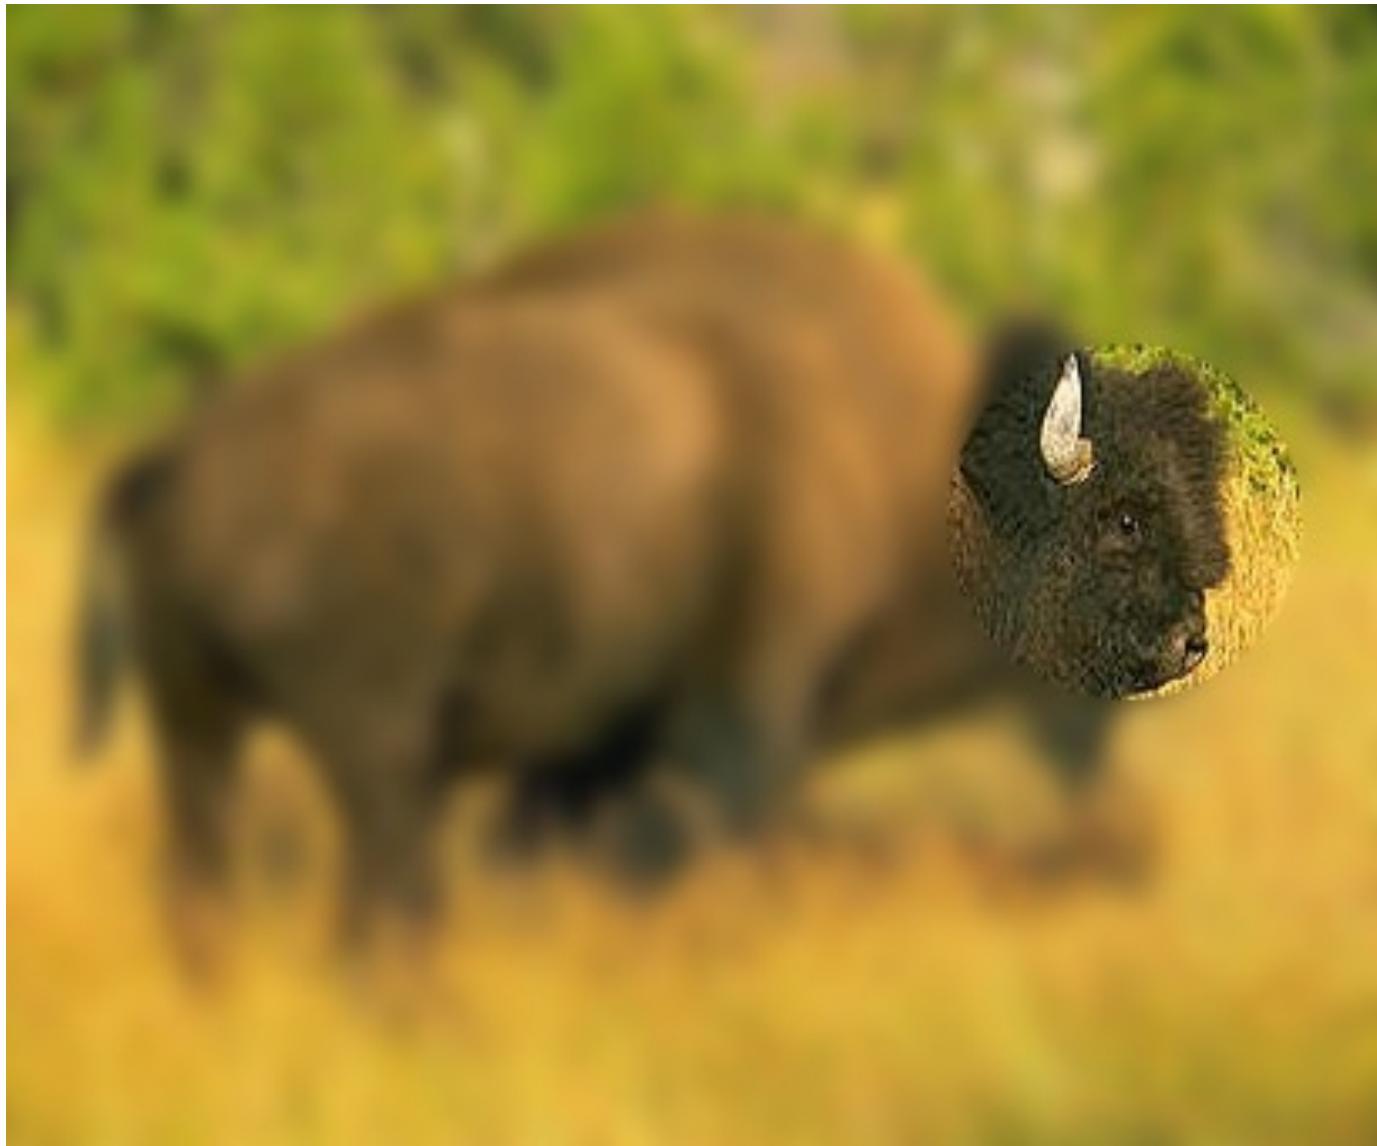
Saccadic eye movement



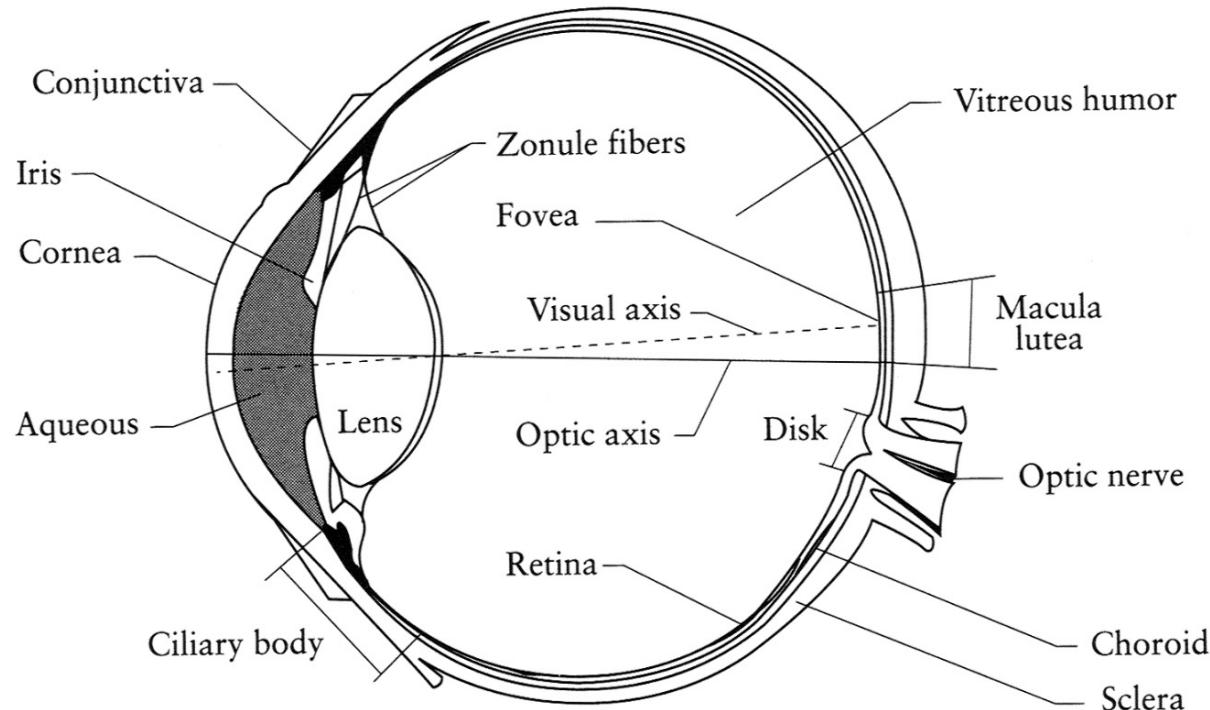
Micro-saccadic movements

Large-saccadic movements

Saccadic eye movement



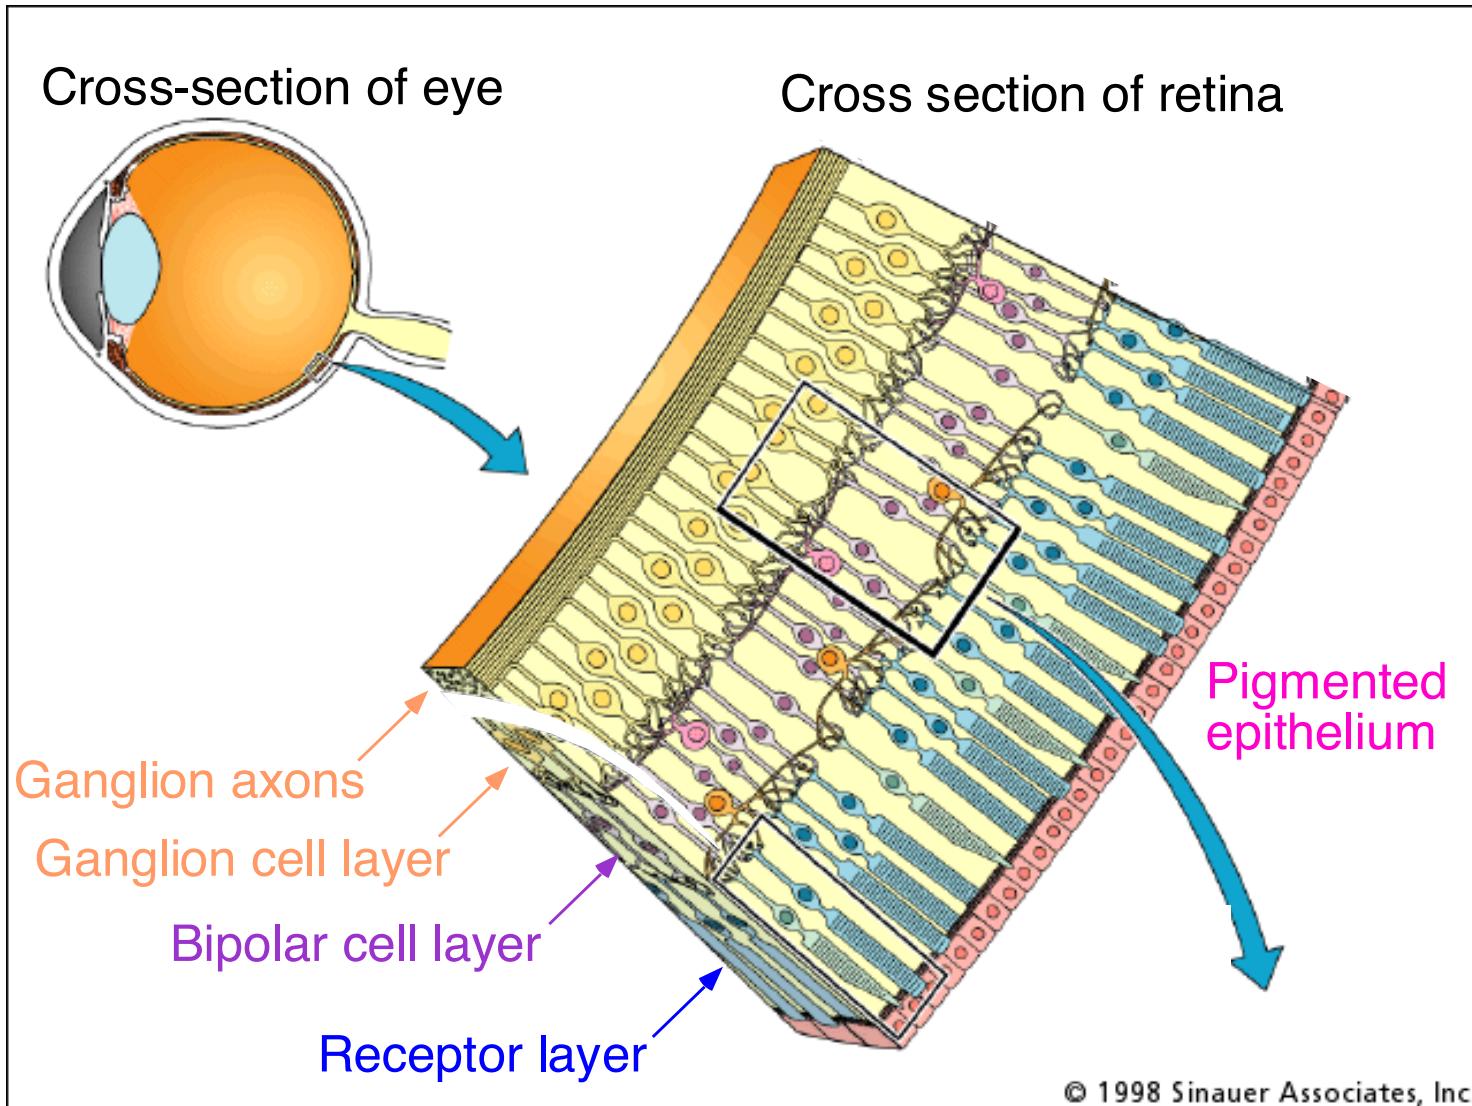
The Eye



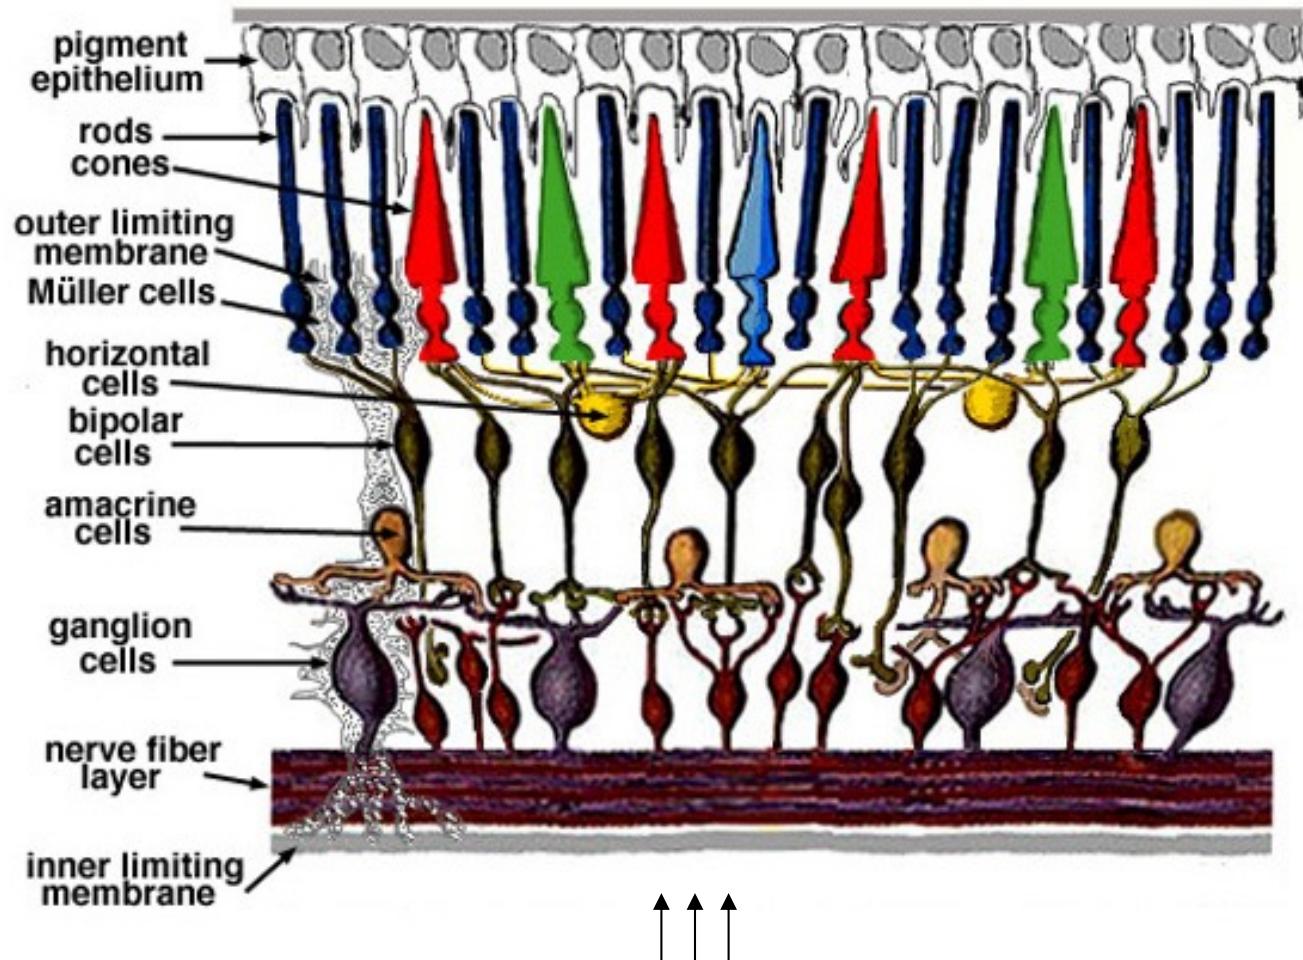
The human eye is a camera!

- **Iris** - colored annulus with radial muscles
- **Pupil** - the hole (aperture) whose size is controlled by the iris
- What's the “film”?
 - photoreceptor cells (rods and cones) in the **retina**

The Retina



Retina up-close



Light

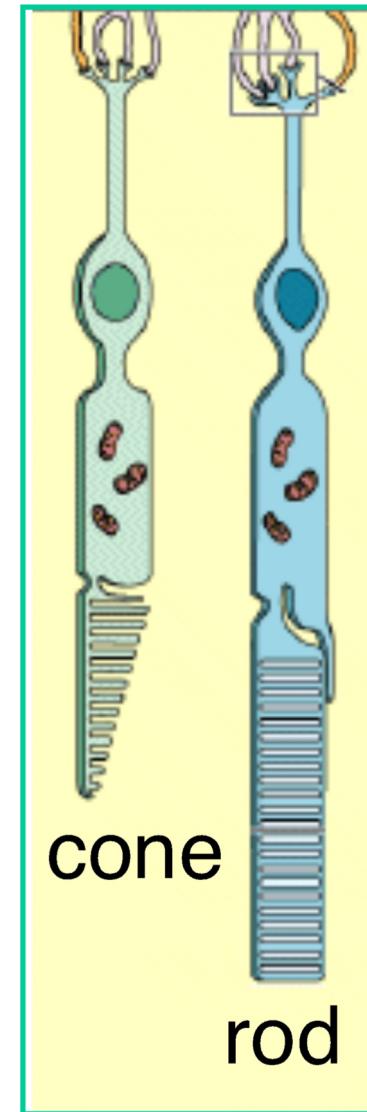
Two types of light-sensitive receptors

Cones

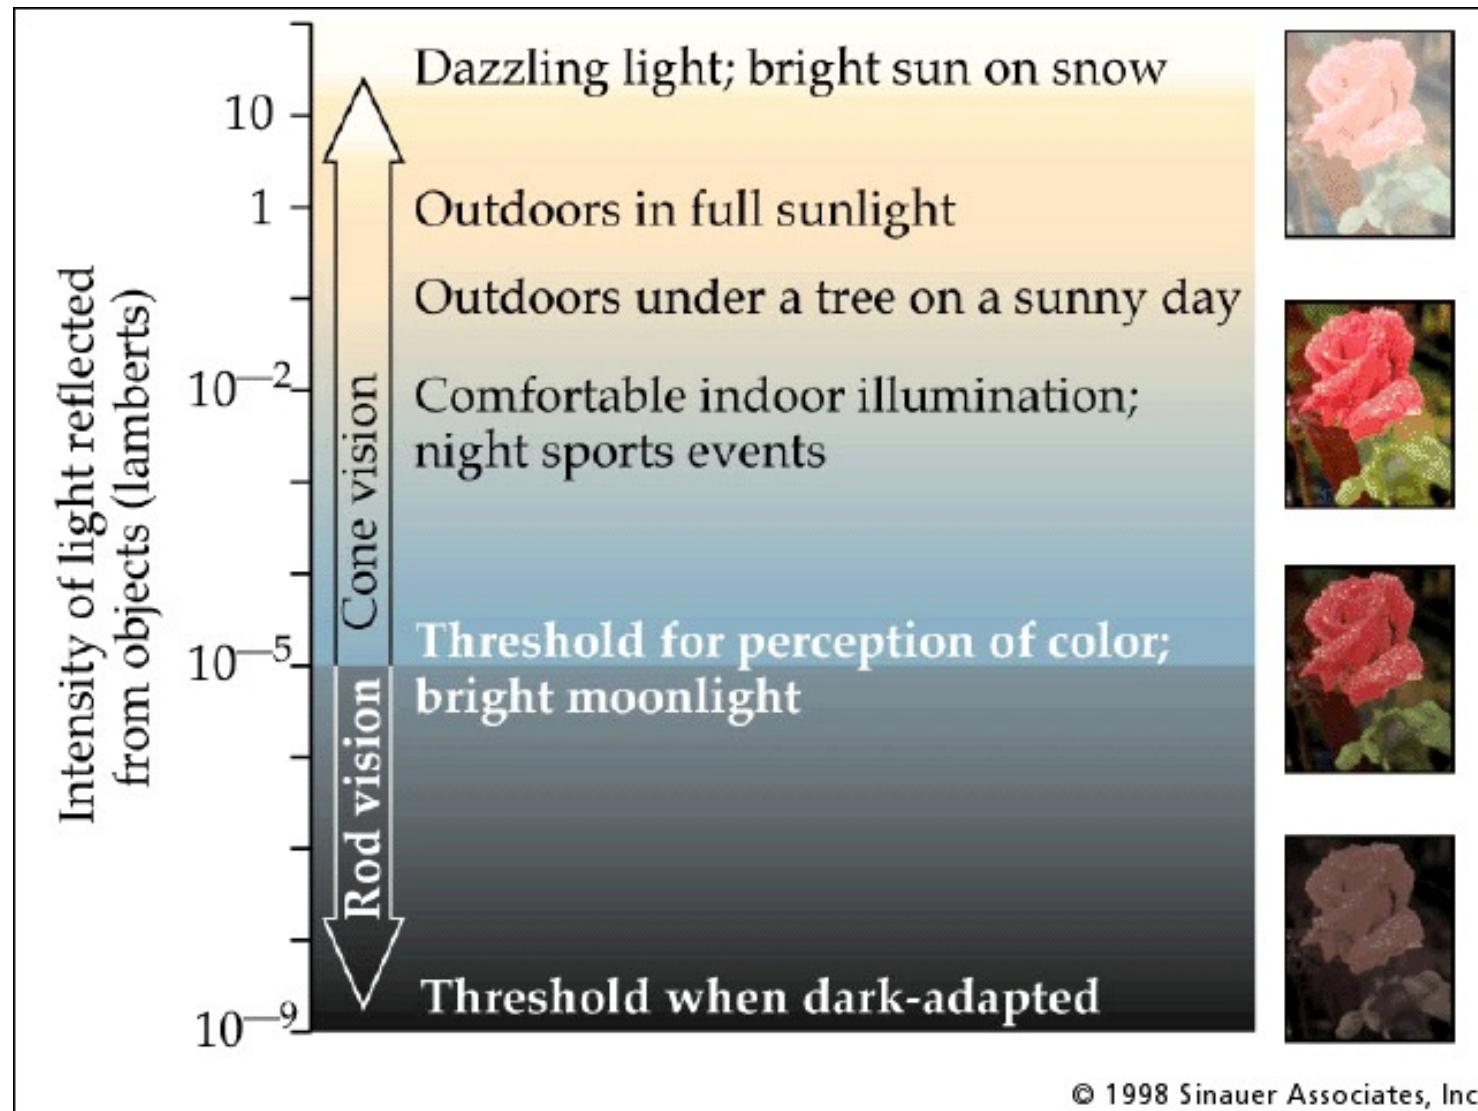
cone-shaped
less sensitive
operate in high light
color vision

Rods

rod-shaped
highly sensitive
operate at night
gray-scale vision

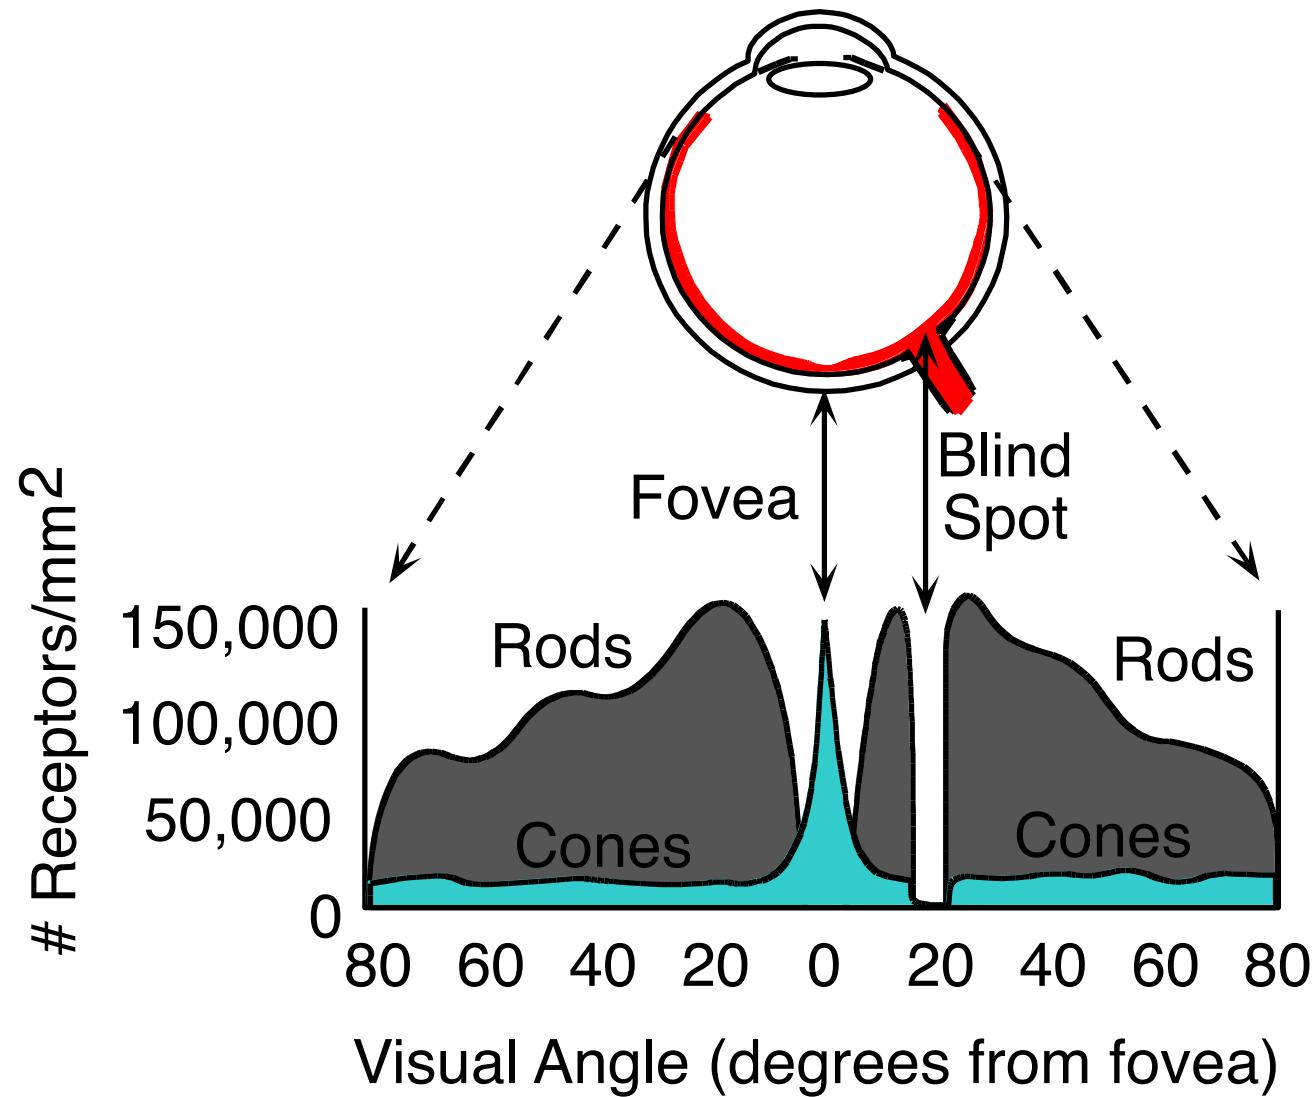


Rod / Cone sensitivity

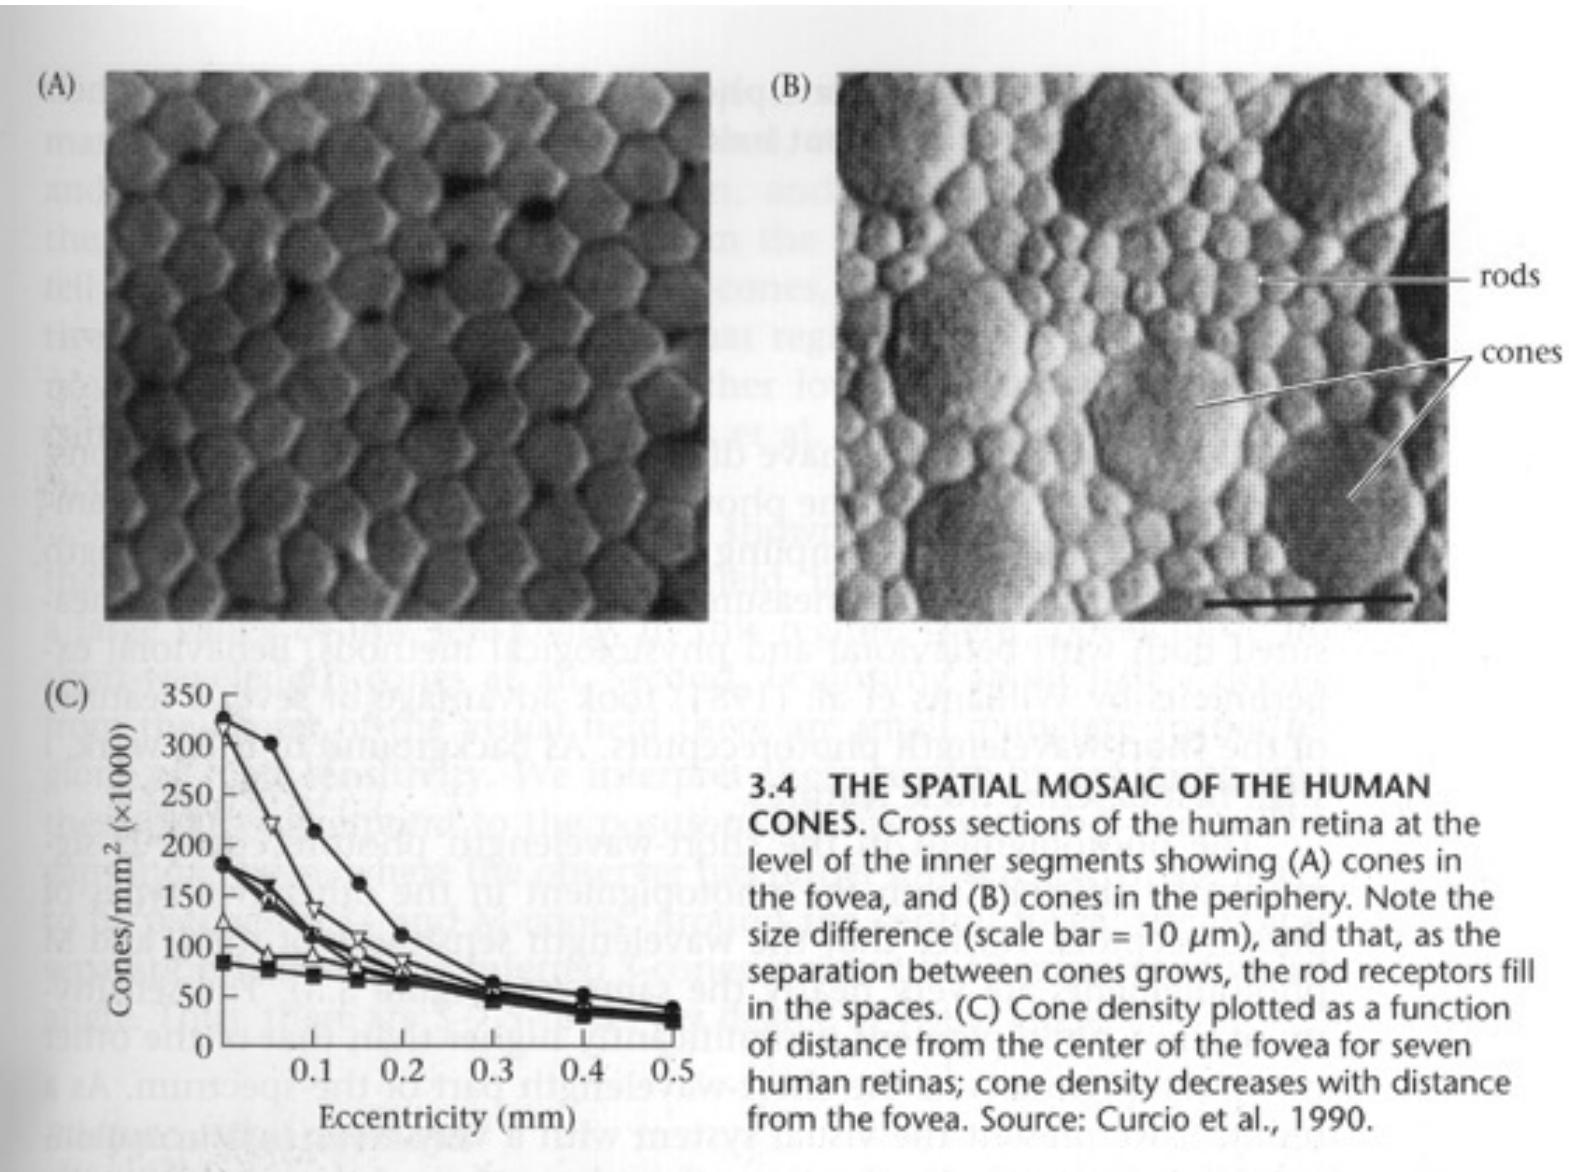


The famous sock-matching problem...

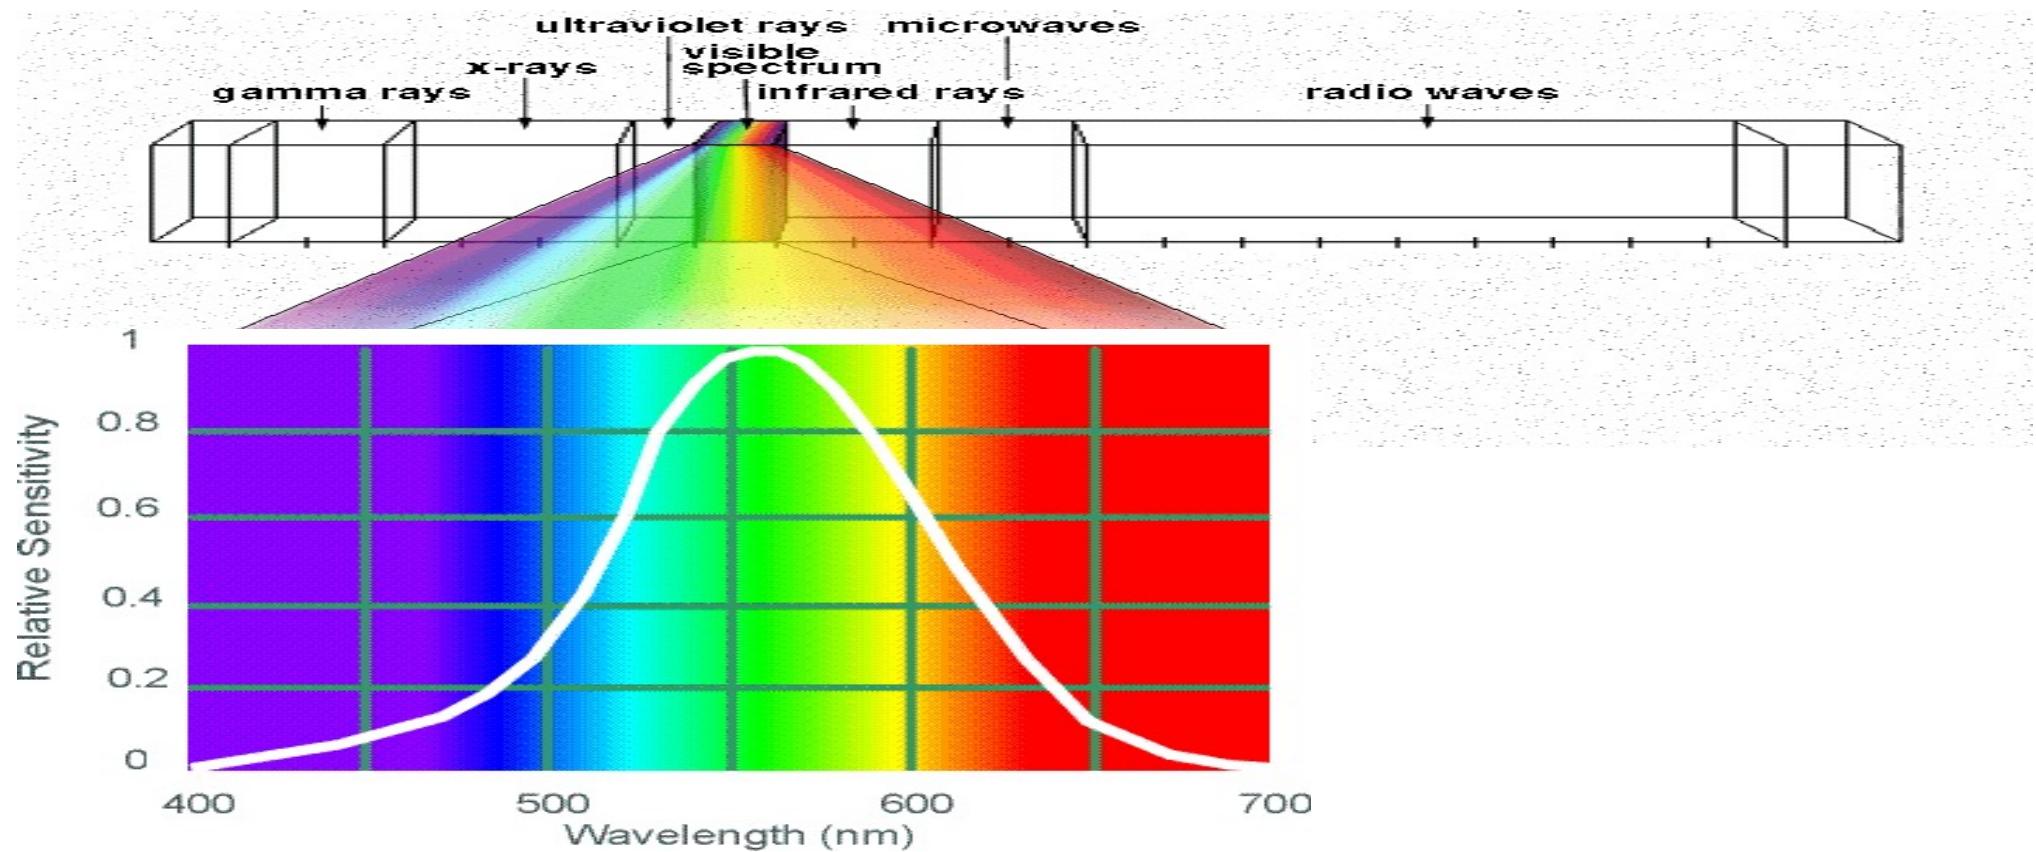
Distribution of Rods and Cones



Night Sky: why are there more stars off-center?



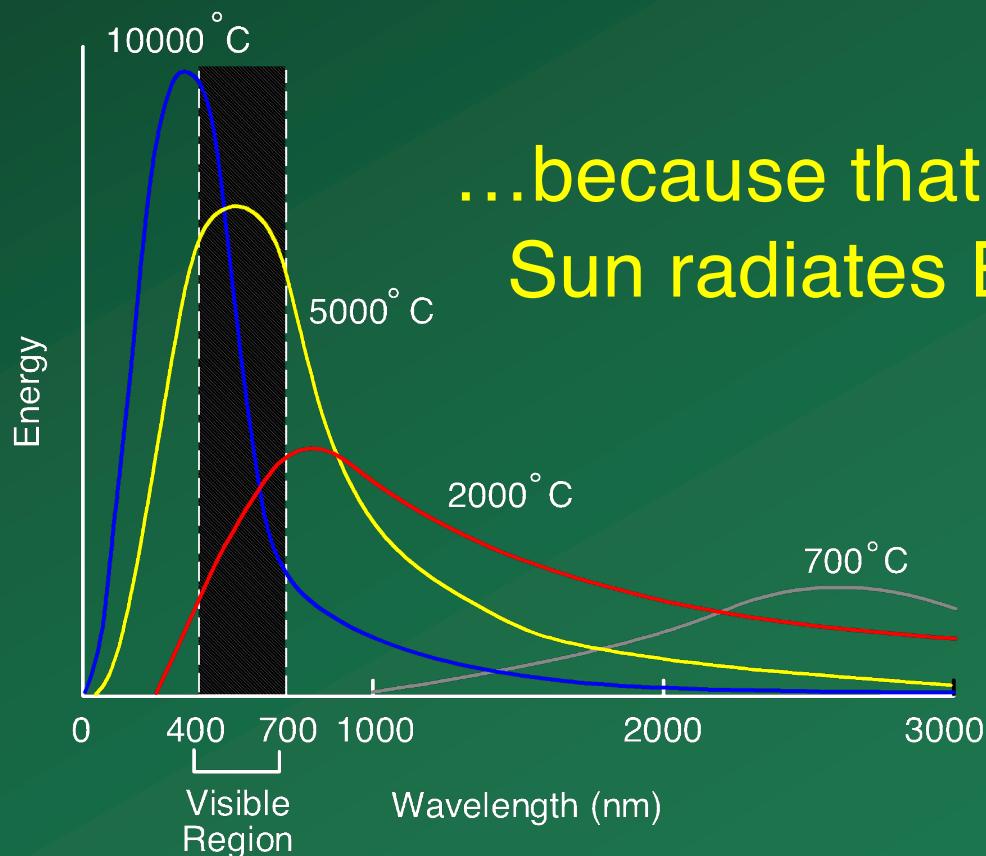
Electromagnetic Spectrum



Human Luminance Sensitivity Function

Visible Light

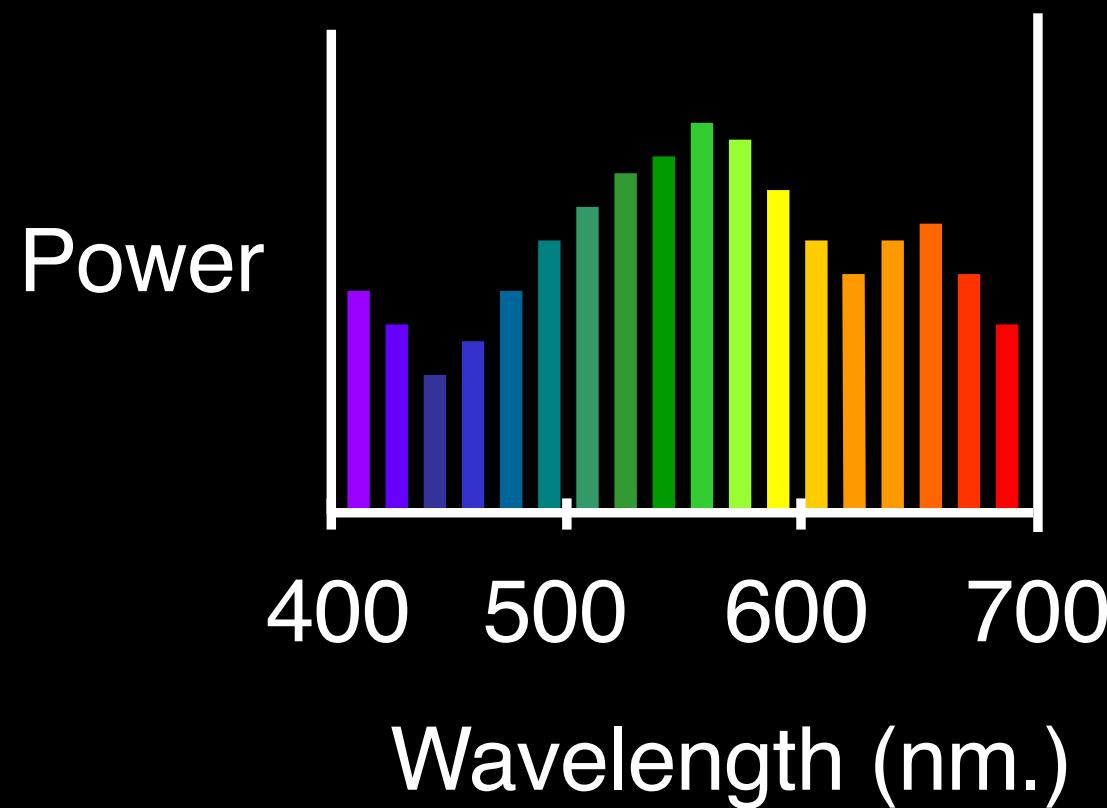
Why do we see light of these wavelengths?



...because that's where the Sun radiates EM energy

The Physics of Light

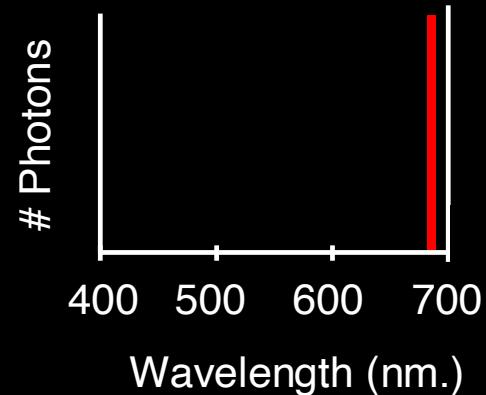
Any patch of light can be completely described physically by its spectrum: the number of photons (per time unit) at each wavelength 400 - 700 nm.



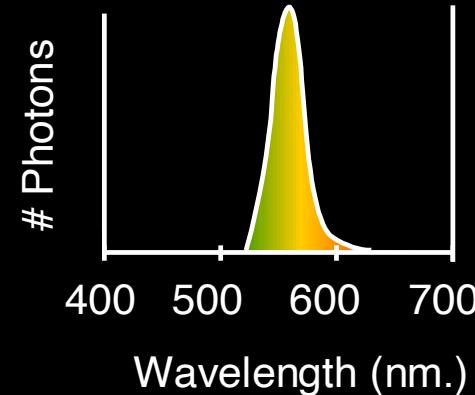
The Physics of Light

Some examples of the spectra of light sources

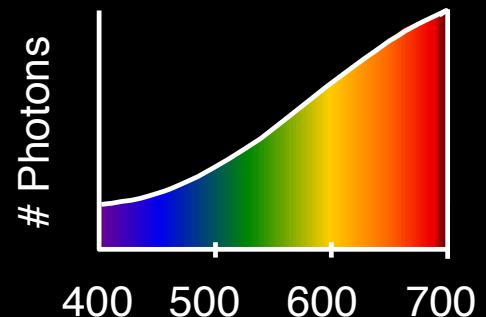
A. Ruby Laser



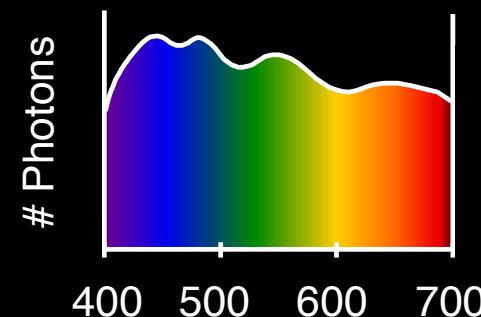
B. Gallium Phosphide Crystal



C. Tungsten Lightbulb



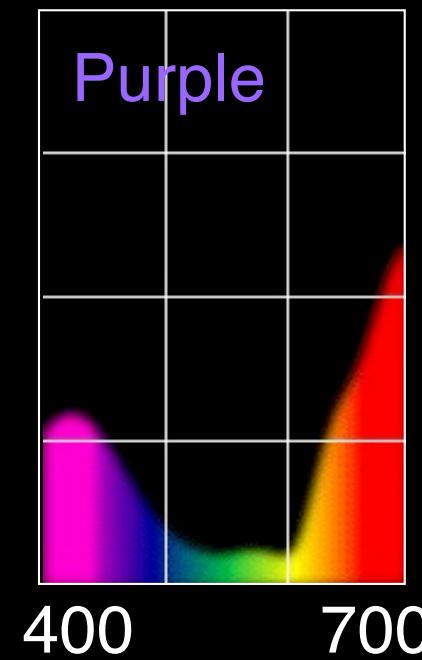
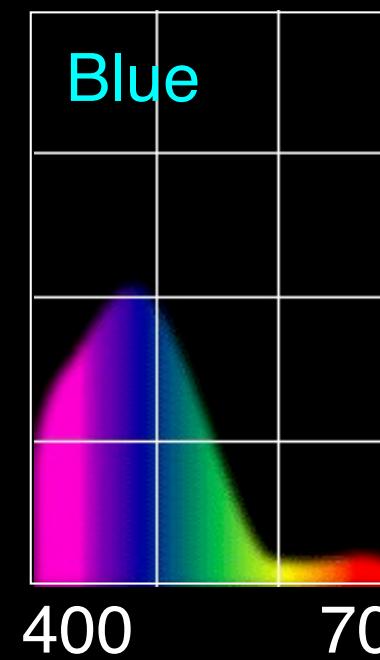
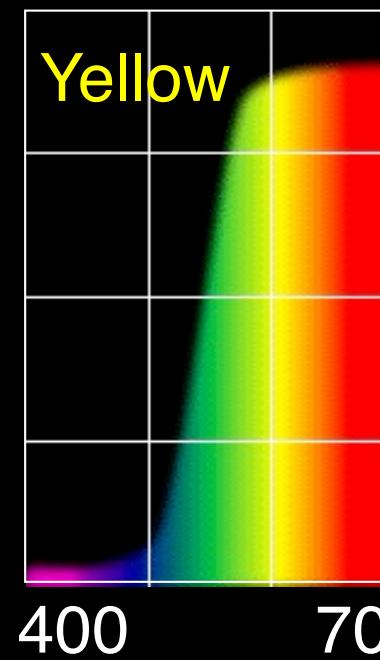
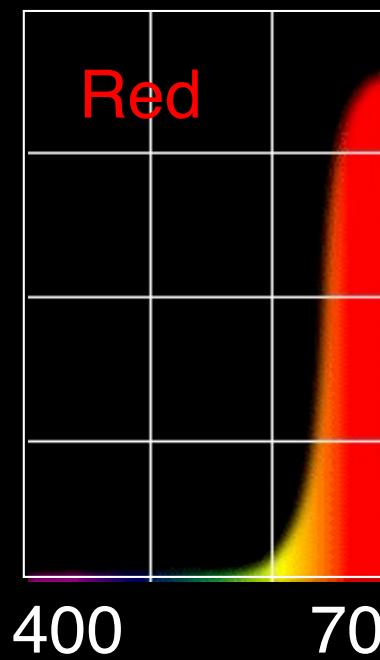
D. Normal Daylight



The Physics of Light

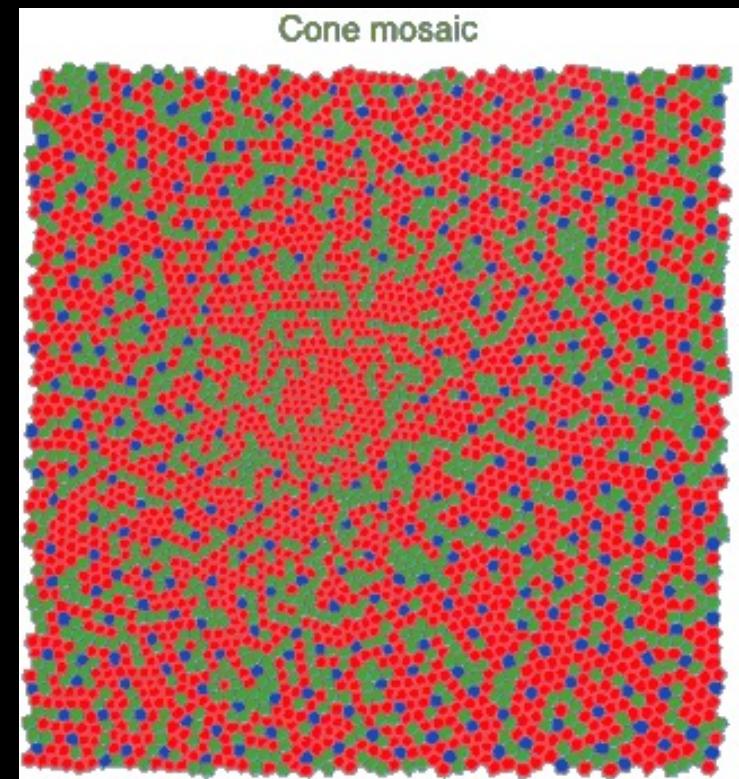
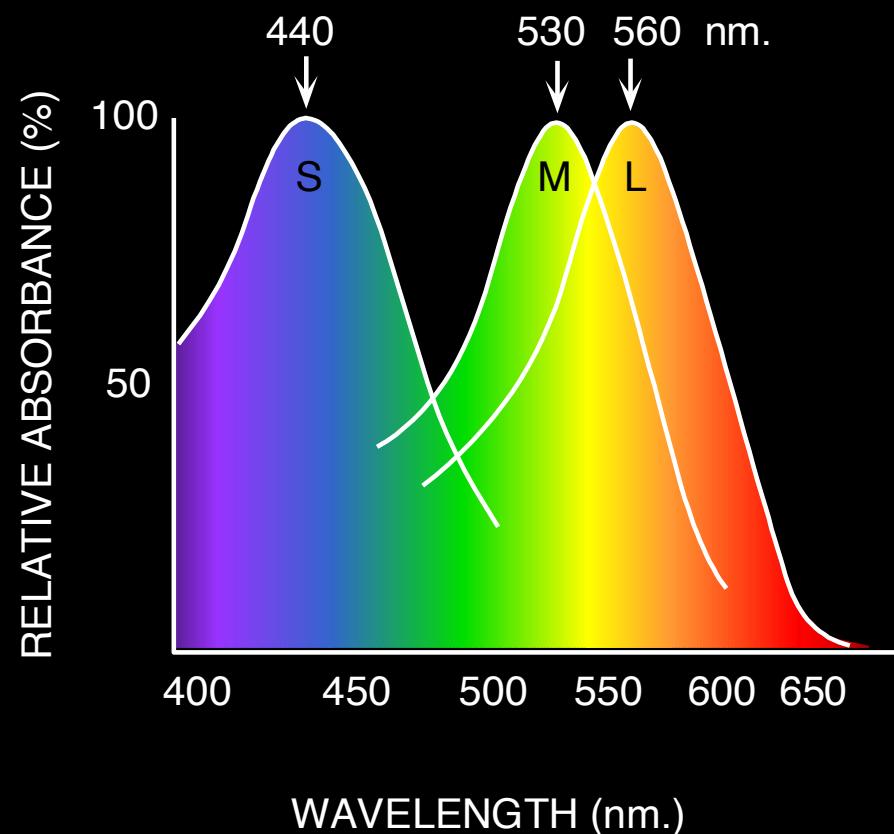
Some examples of the reflectance spectra of surfaces

% Photons Reflected



Physiology of Color Vision

Three kinds of cones:



- Why are M and L cones so close?
- Why are there 3?

Not everyone is trichromat

- Types of color blindness:
 - Deutanopia: missing M cones
 - Protanopia: missing L cones
 - Tritanopia: missing S cones
- “M” and “L” on the X-chromosome
 - Why men are more likely to be color blind
 - “L” has high variation, so some women are tetrachromatic
- Some animals have
 - 1 (night animals)
 - 2 (e.g., dogs)
 - 4 (fish, birds)
 - 5 (pigeons, some reptiles/amphibians)
 - 12 (mantis shrimp)



NORMAL VISION



DEUTERANOMALIA

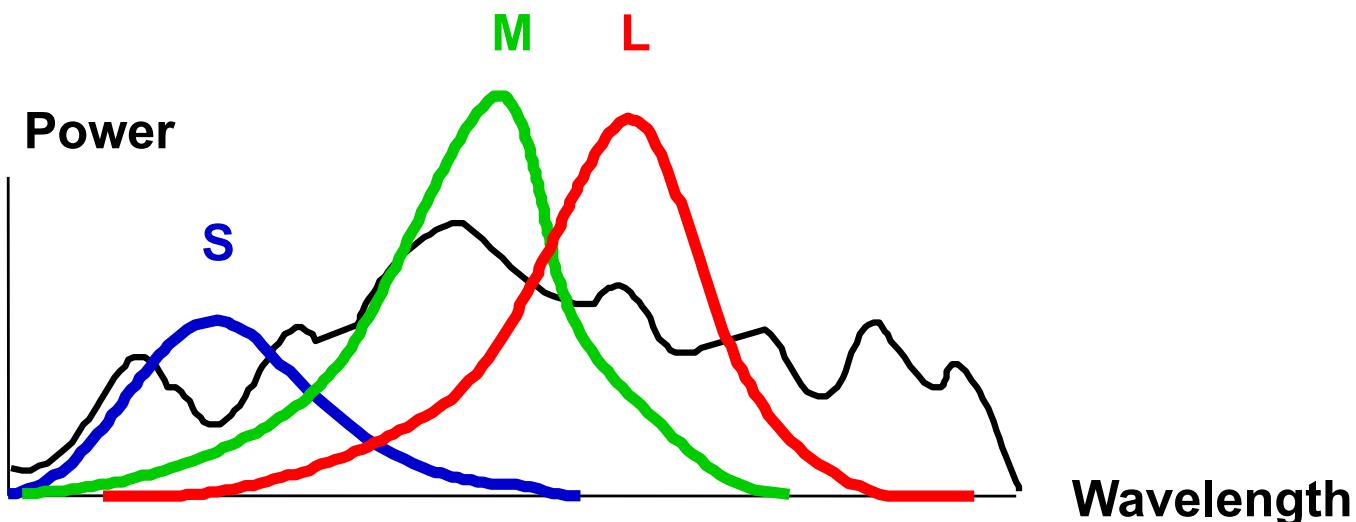


PROTANOPIA



TRITANOPIA

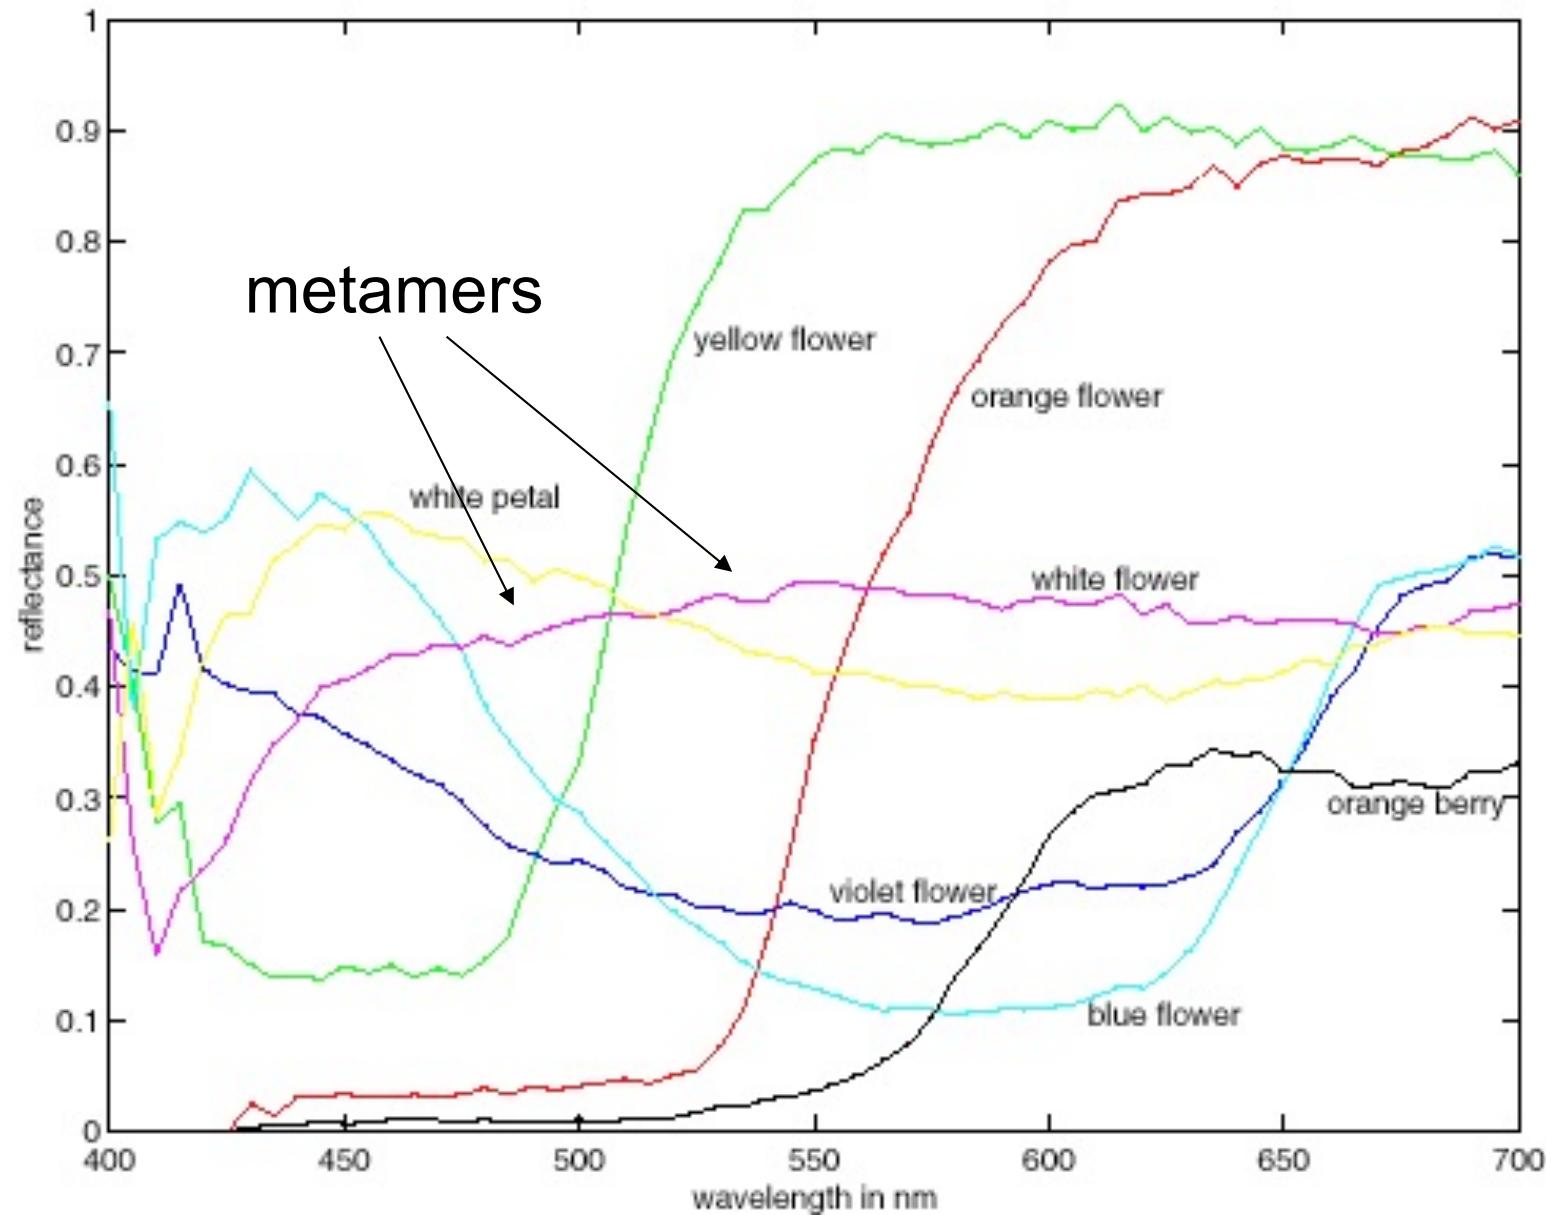
Trichromacy



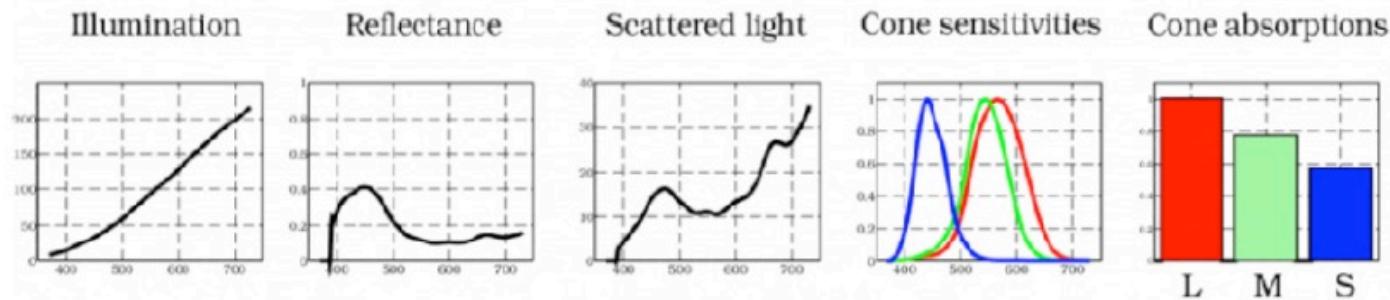
Rods and cones act as *filters* on the spectrum

- To get the output of a filter, multiply its response curve by the spectrum, integrate over all wavelengths
 - Each cone yields one number
- How can we represent an entire spectrum with 3 numbers?
- We can't! Most of the information is lost
 - As a result, two different spectra may appear indistinguishable
 - » such spectra are known as **metamers**

We are all color-blind!



Spectral Radiometry: Color Image Formation



incoming illumination \times reflectance = outgoing scene radiance

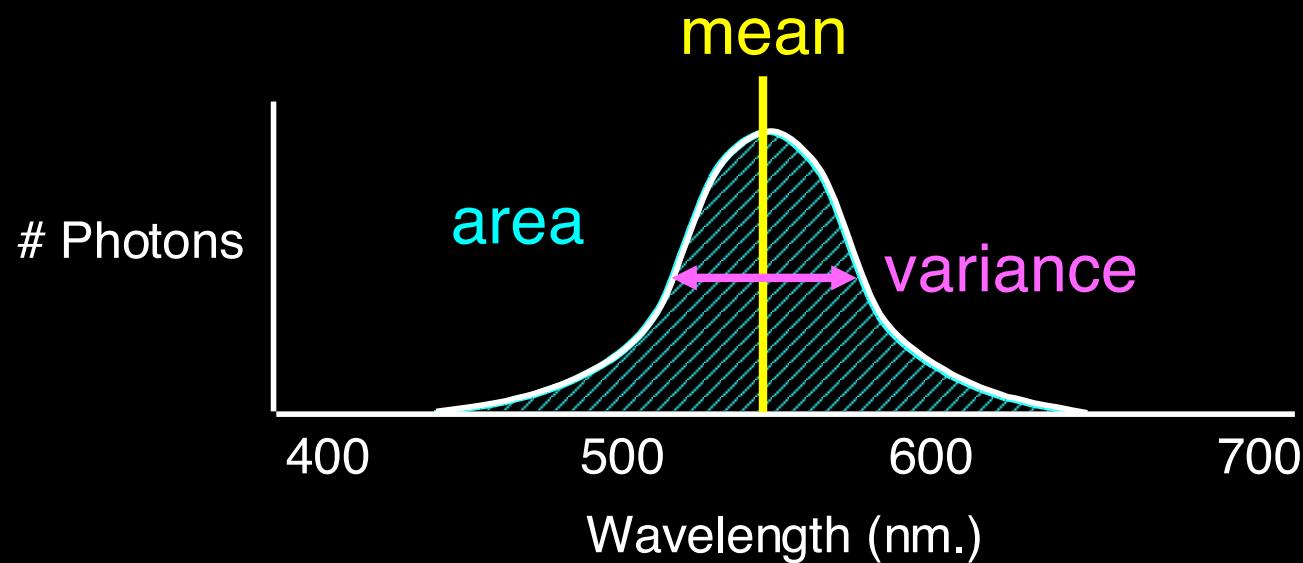
$$\int_{\lambda} \text{scene radiance} \times \text{cone sensitivities } d\lambda = \text{cone absorptions}$$

The Psychophysics of “Color”

There is no simple functional description for the perceived color of all lights under all viewing conditions, but

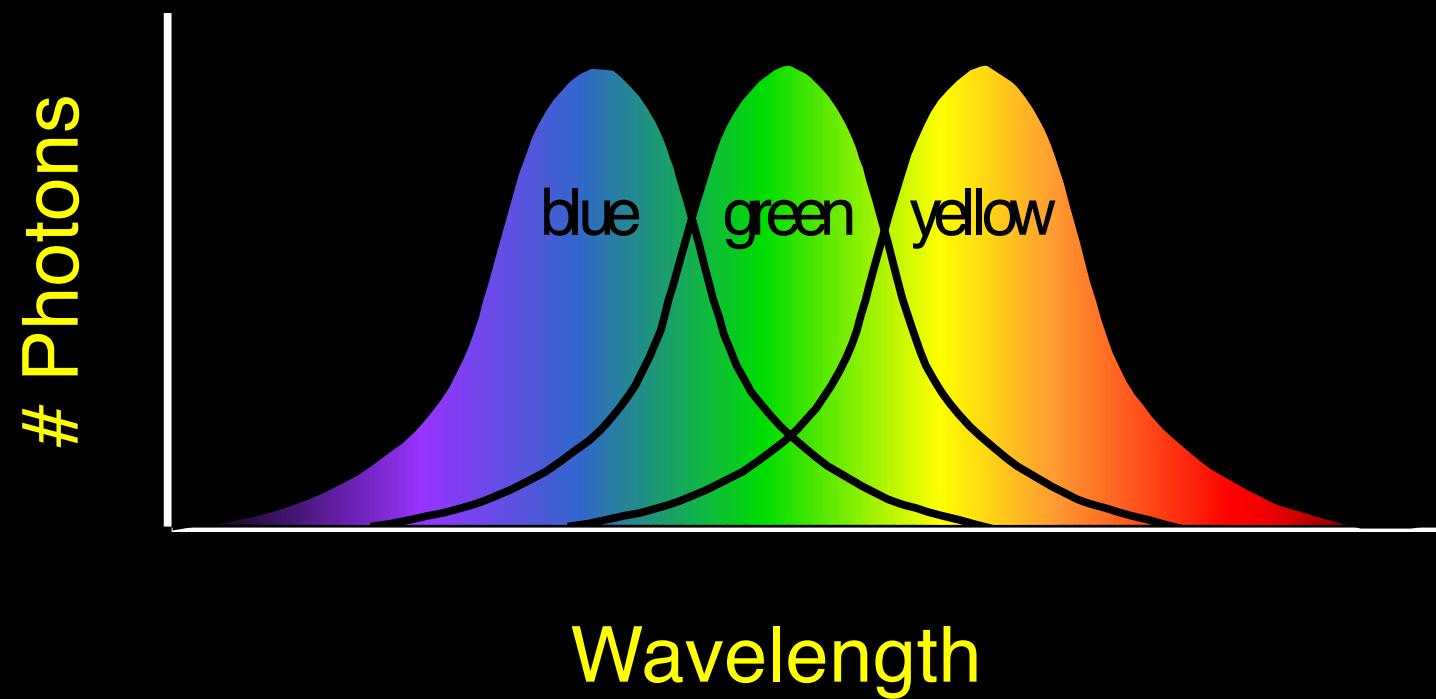
A helpful constraint:

Consider only physical spectra with normal distributions



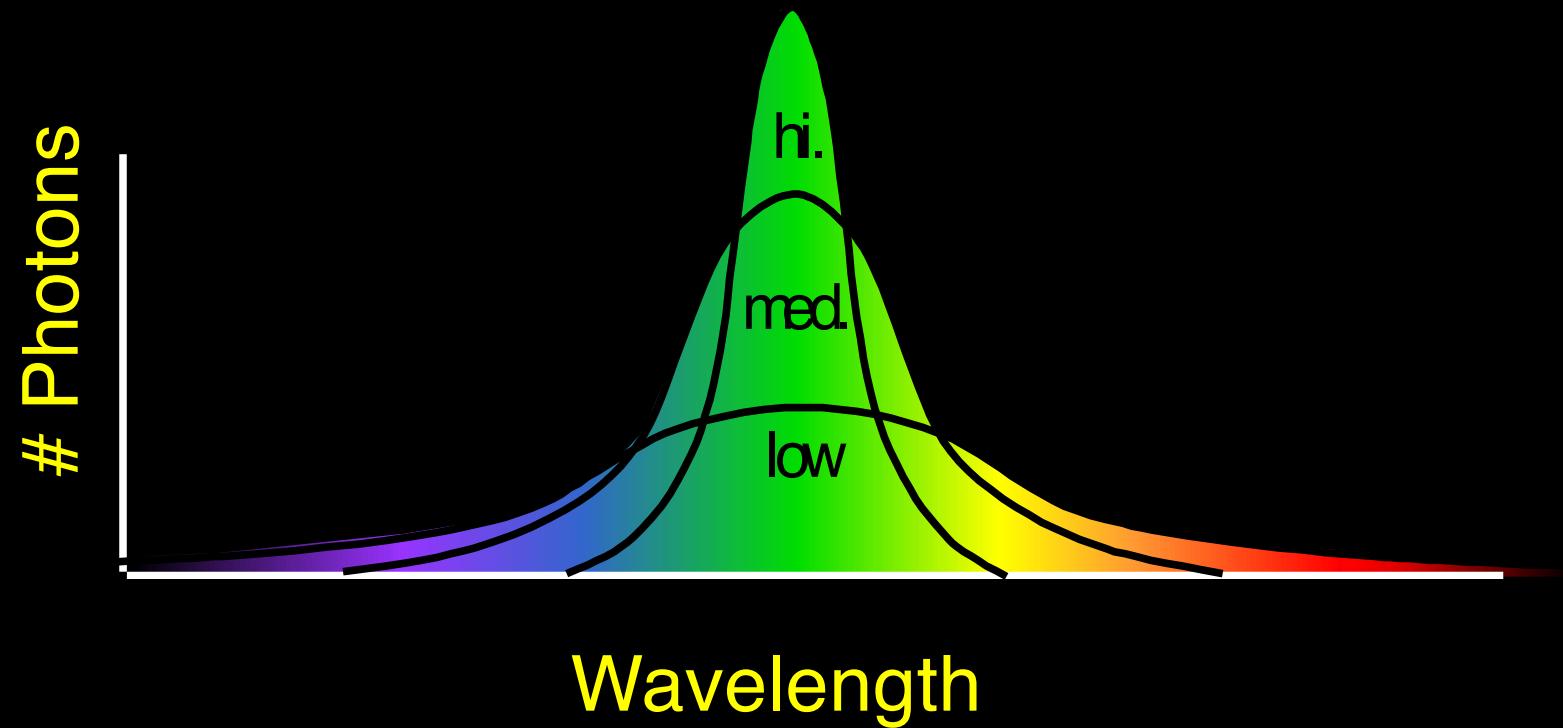
The Psychophysical Correspondence

Mean \longleftrightarrow Hue



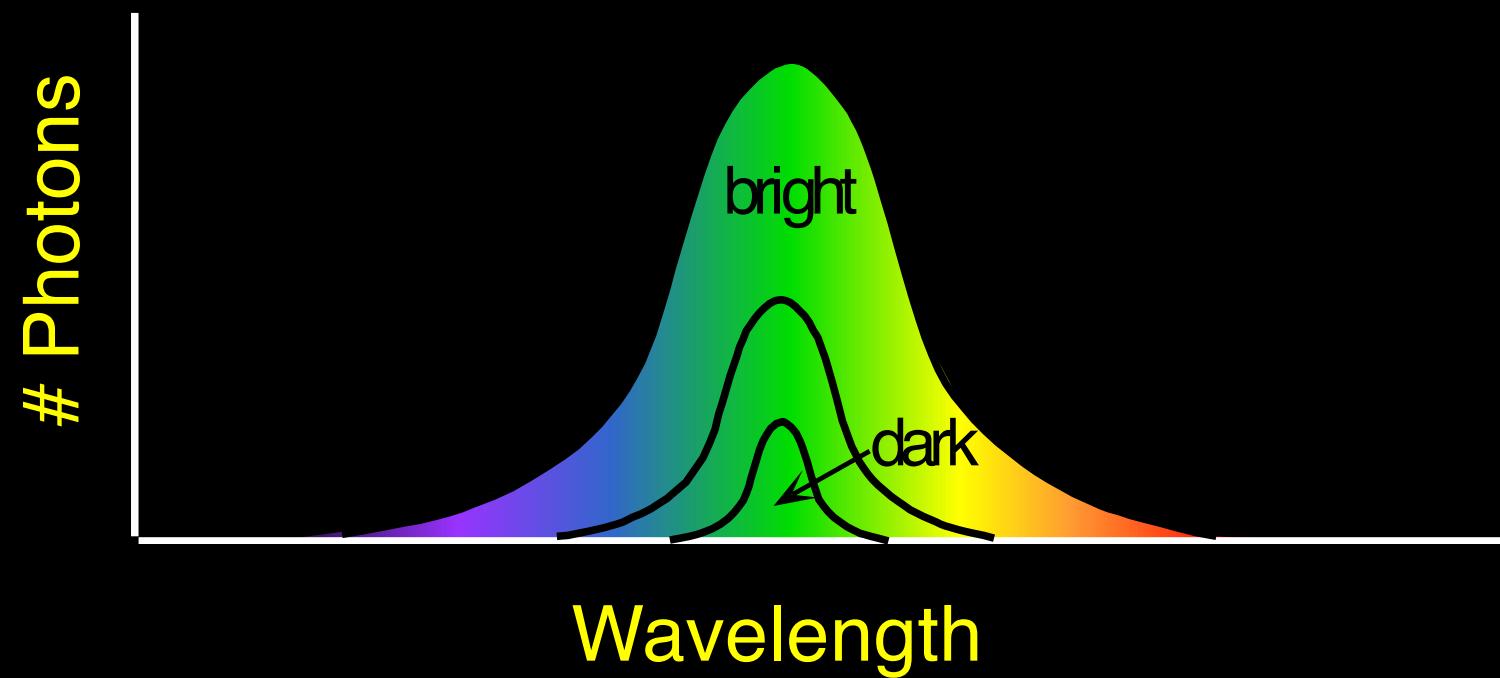
The Psychophysical Correspondence

Variance \longleftrightarrow Saturation



The Psychophysical Correspondence

Area \longleftrightarrow Brightness



Color Constancy

The “photometer metaphor” of color perception:
Color perception is determined by the spectrum of light
on each retinal receptor (as measured by a photometer).



Color Constancy

The “photometer metaphor” of color perception:
Color perception is determined by the spectrum of light
on each retinal receptor (as measured by a photometer).



Color Constancy

~~The “photometer metaphor” of color perception:
Color perception is determined by the spectrum of light
on each retinal receptor (as measured by a photometer).~~

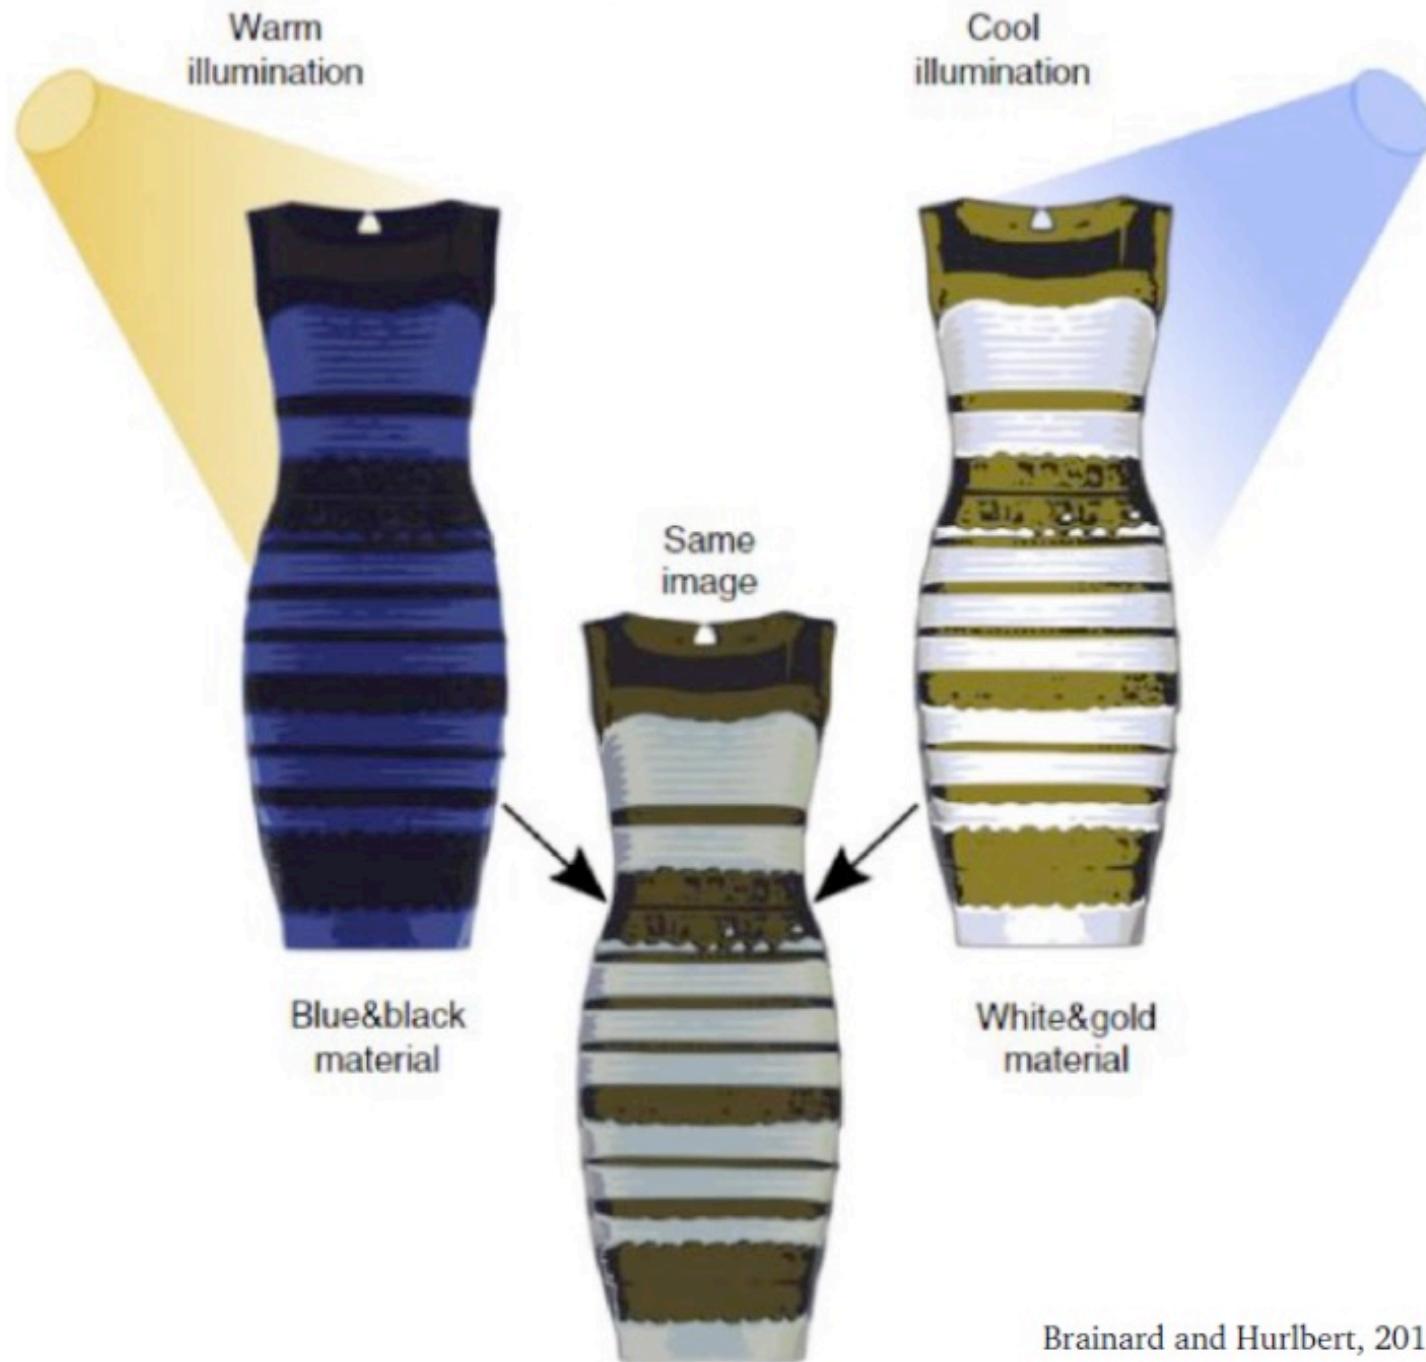


What color is the “The Dress”?



https://en.wikipedia.org/wiki/The_dress

Two Scene Interpretations of #thedress



Color Constancy

~~Do we have constancy over
all global color transformations?~~



60% blue filter



Complete inversion

Color Constancy

Color Constancy: the ability to perceive the invariant color of a surface despite ecological Variations in the conditions of observation.

Another of these hard inverse problems:
Physics of light emission and surface reflection
underdetermine perception of surface color

Camera White Balancing

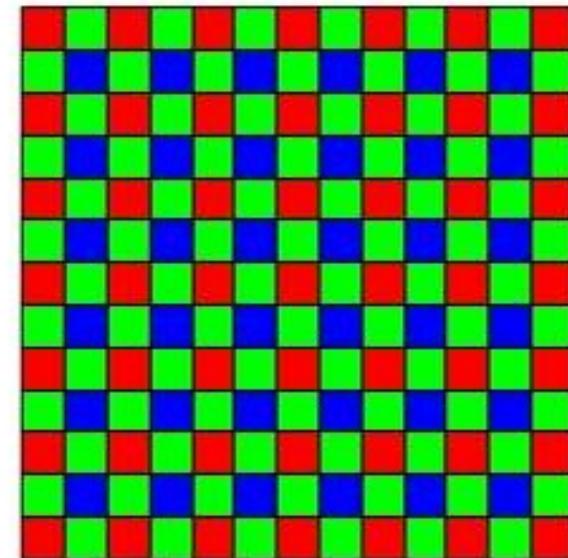
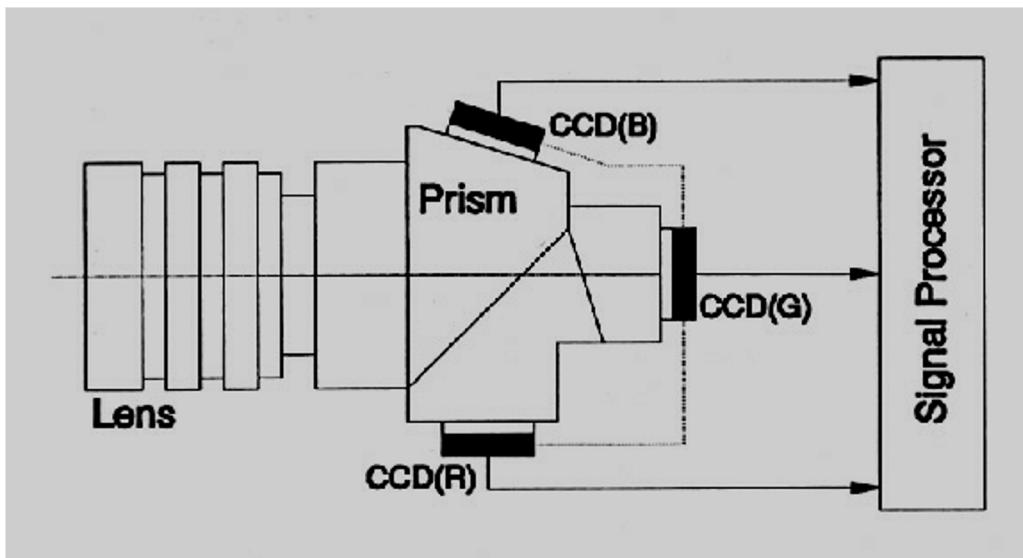
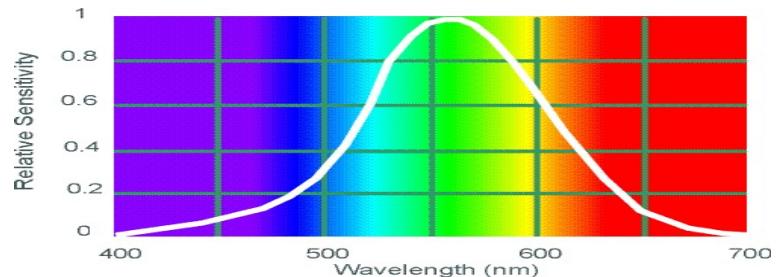


- Manual
 - Choose color-neutral object in the photos and normalize
- Automatic (AWB)
 - Grey World: force average color of scene to grey
 - White World: force brightest object to white

Color Sensing in Camera (RGB)

3-chip vs. 1-chip: quality vs. cost

Why more green?



Why 3 colors?

Ruff Works

<http://www.cooldictionary.com/words/Bayer-filter.wikipedia>

Slide by Steve Seitz

Green is in!



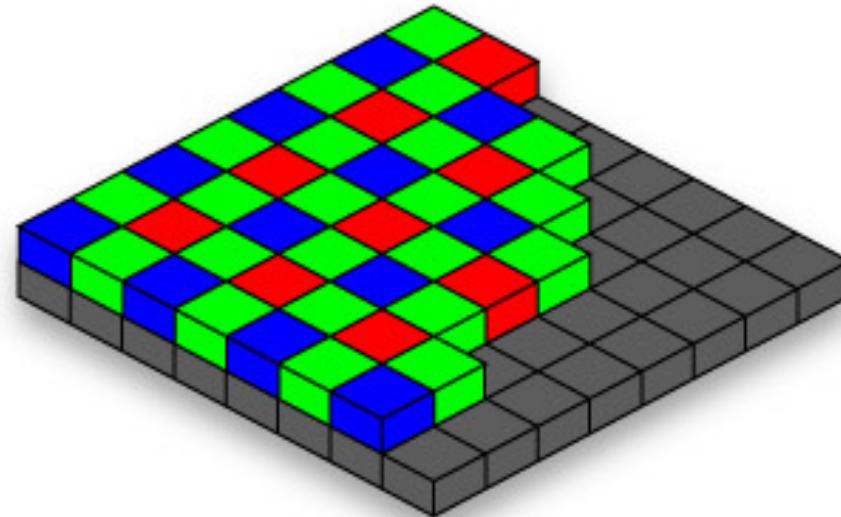
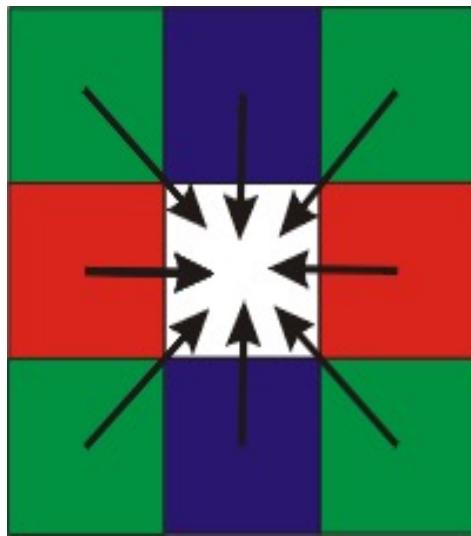
R

G

B

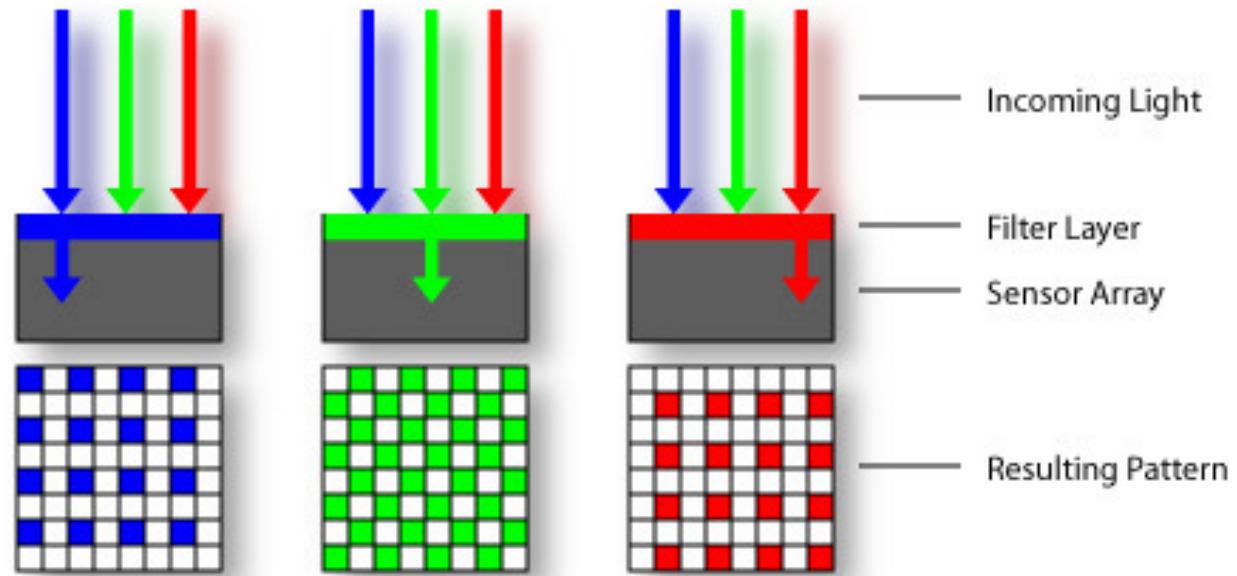


Practical Color Sensing: Bayer Grid

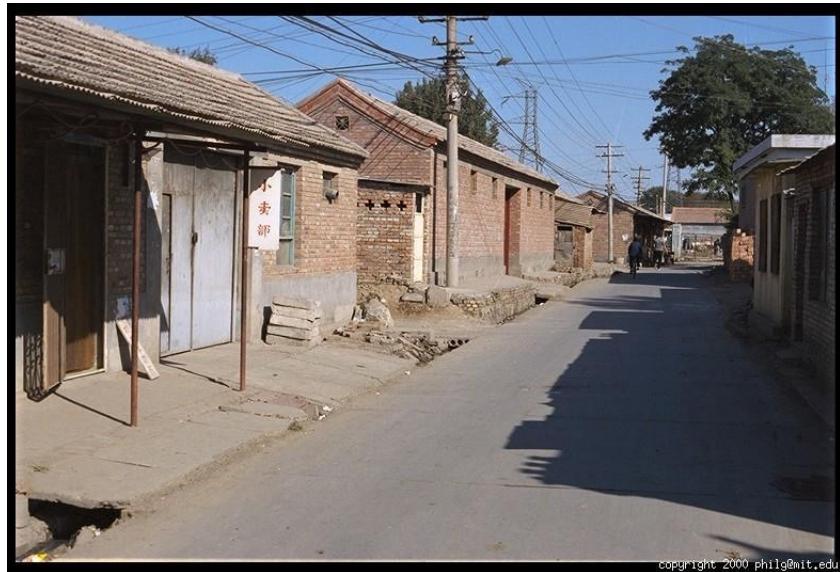


Estimate RGB
at 'G' cells from
neighboring
values

[http://www.cooldictionary.com/
words/Bayer-filter.wikipedia](http://www.cooldictionary.com/words/Bayer-filter.wikipedia)



Color Image



R



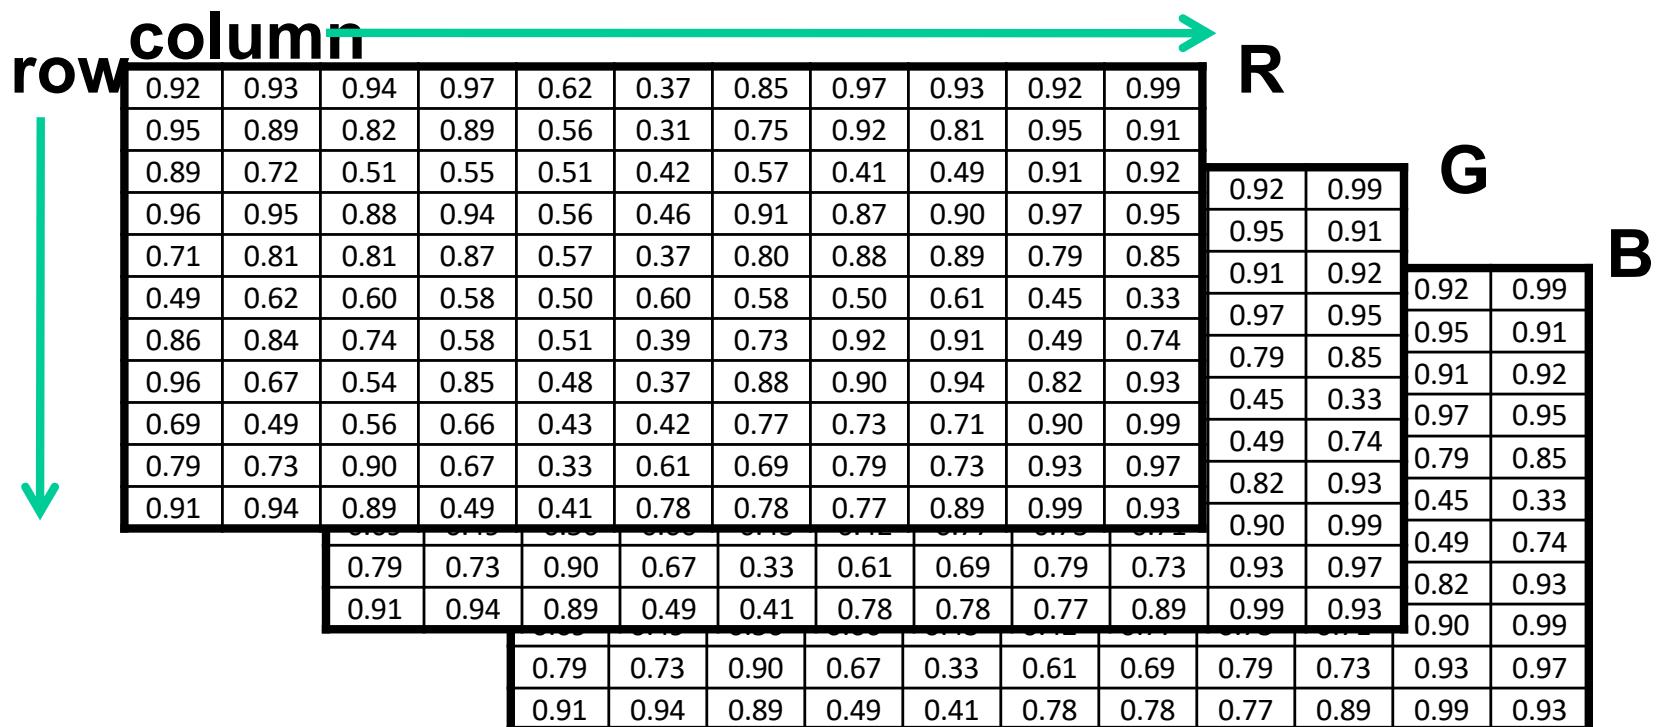
G



B

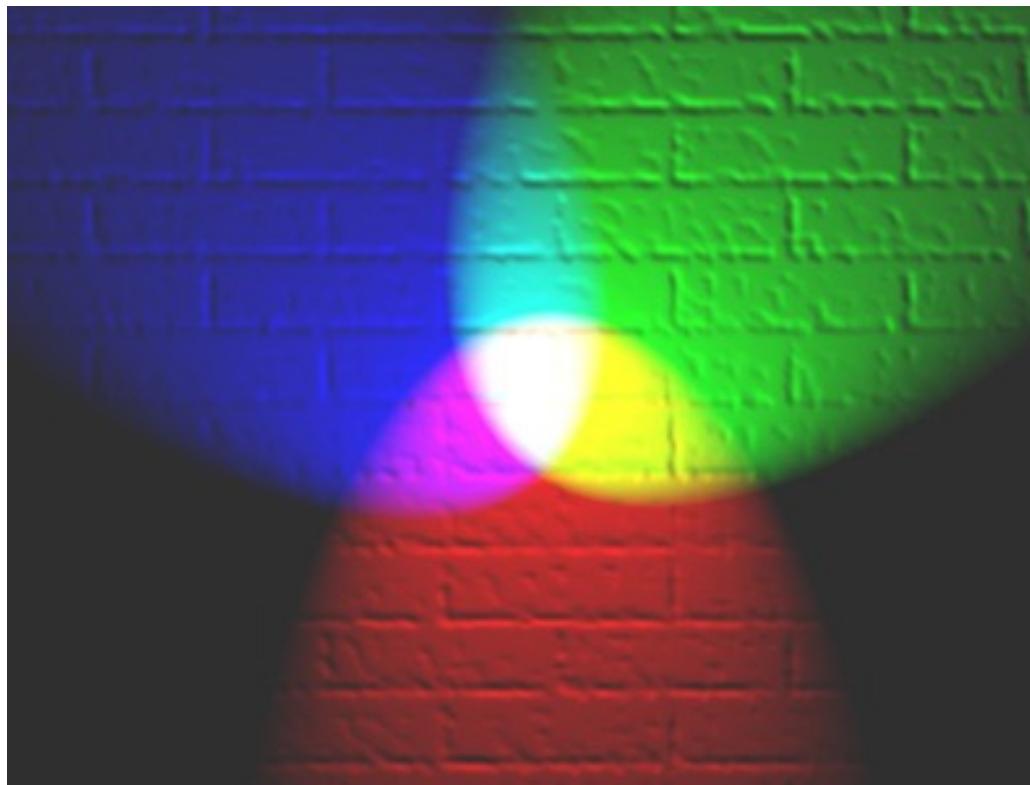
Image representation

- Images represented as a matrix
 - Suppose we have a NxM RGB image called “im”
 - $\text{im}(1,1,1)$ = top-left pixel value in R-channel
 - $\text{im}(y, x, b)$ = y pixels down, x pixels to right in the b^{th} channel
 - $\text{im}(N, M, 3)$ = bottom-right pixel in B-channel
 - **imread(filename)** returns a uint8 image (values 0 to 255)
 - Convert to double format (values 0 to 1) with **im2double**



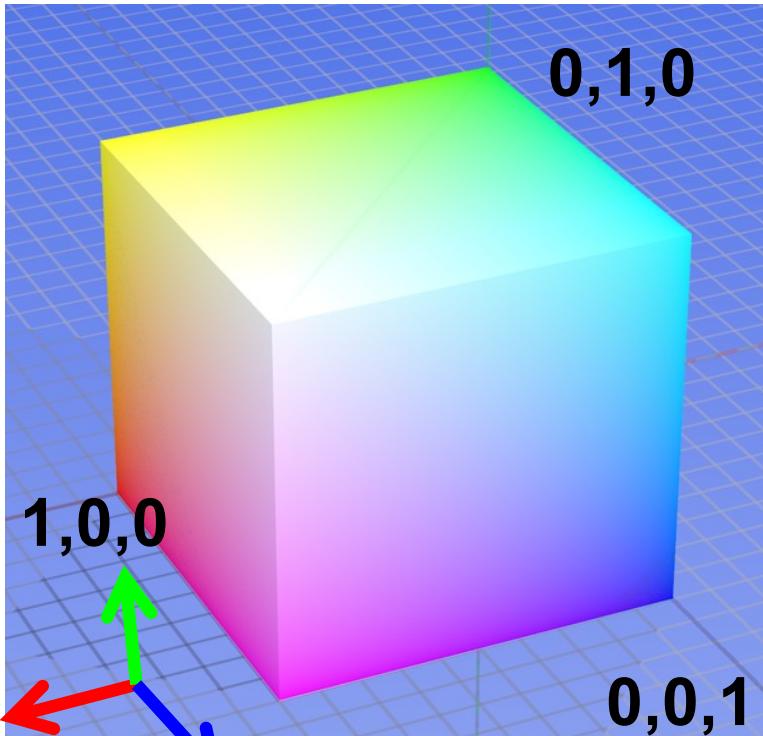
Color spaces

How can we represent color?



Color spaces: RGB

Default color space



RGB cube

- Easy for devices
- But not perceptual
- Where do the grays live?
- Where is hue and saturation?

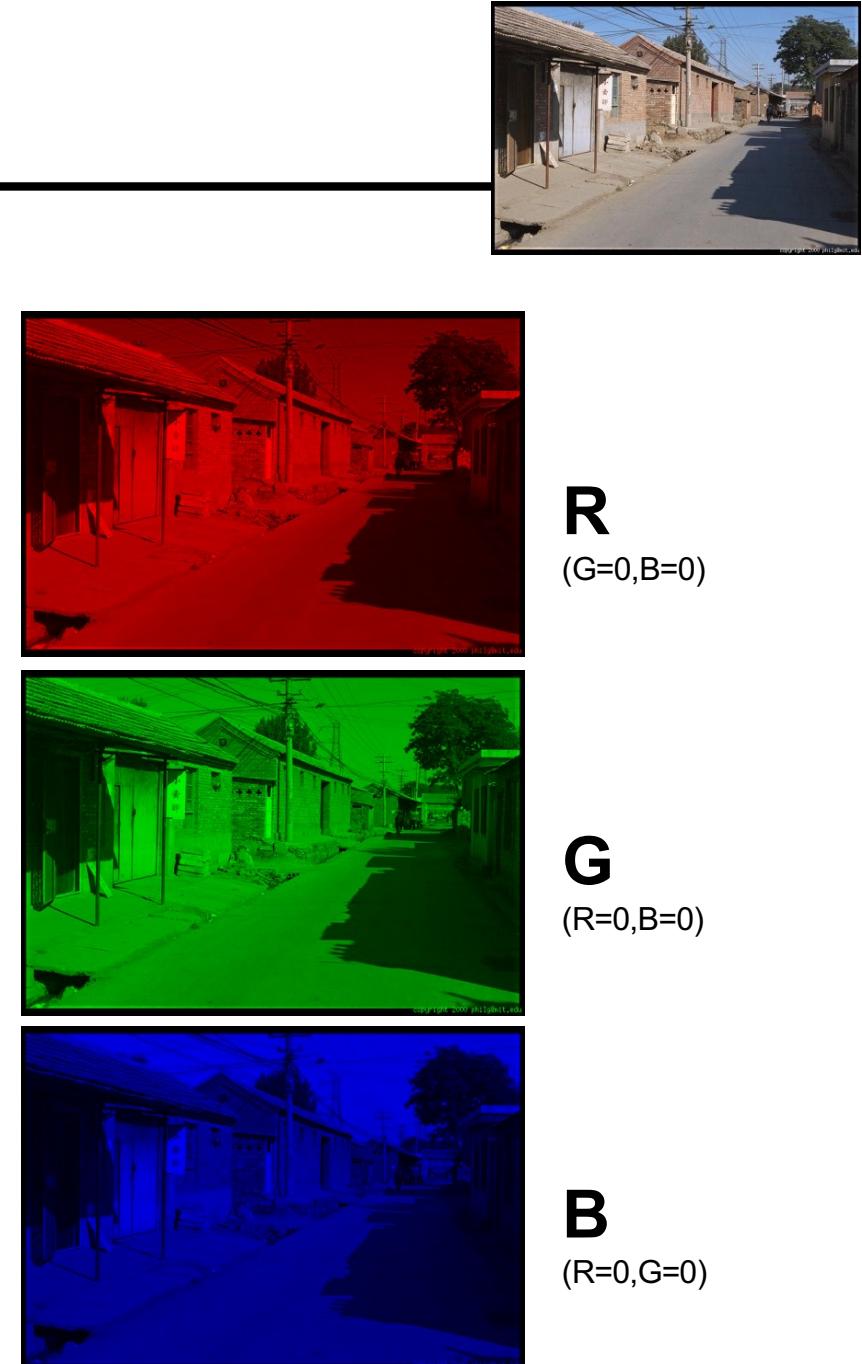


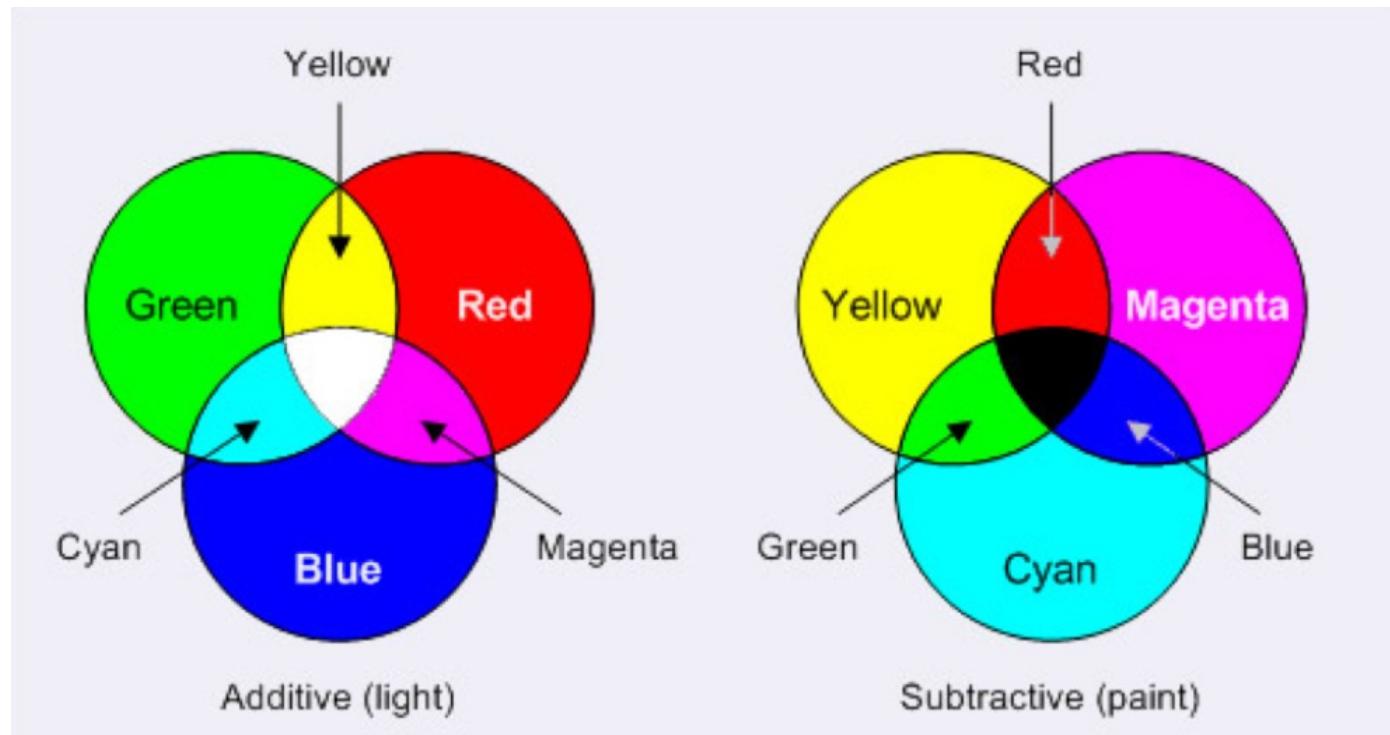
Image from: http://en.wikipedia.org/wiki/File:RGB_color_solid_cube.png

Color spaces: CMYK

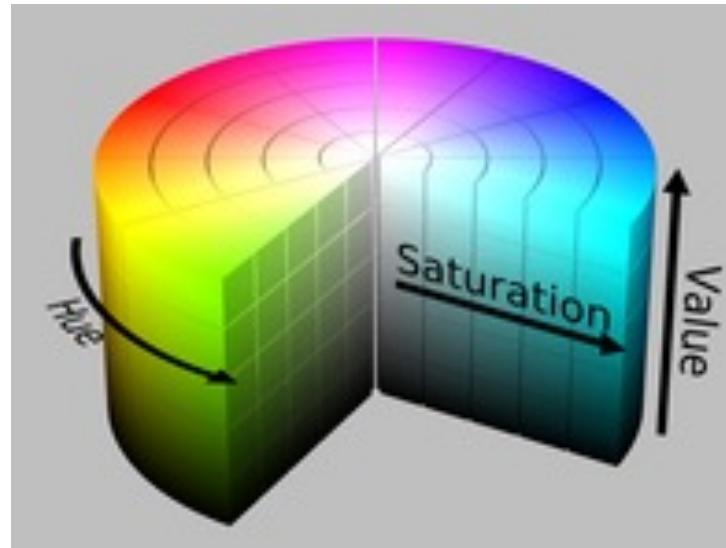
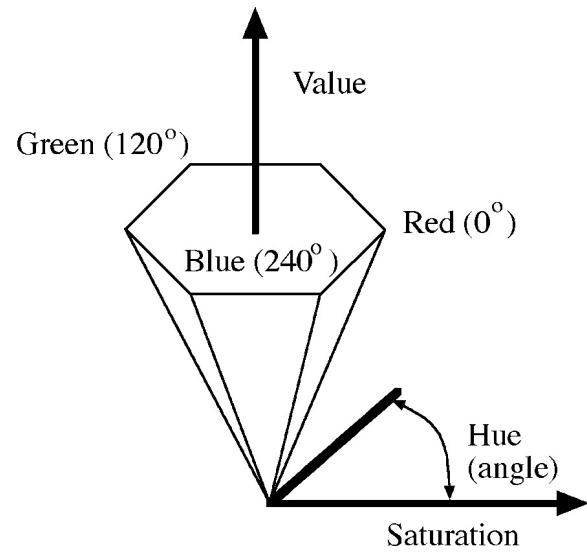
How paint works

White is the background color (all colors!)

Opposite of RGB (stage light)



HSV



Perceptual dimensions of color:

Hue: the “kind” of color, regardless of its attributes

Saturation: Purity, “colorfulness”

Value (or lightness): total amount of light

Use `rgb2HSV()` and `HSV2RGB()` in Matlab, in Python w/`skimage`

Color spaces: L*a*b*

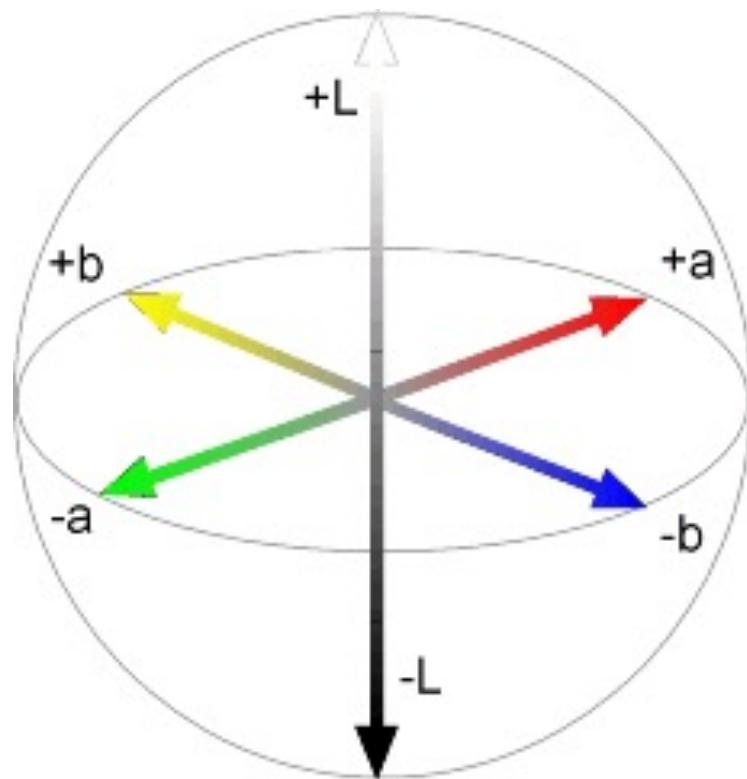


“Perceptually uniform”* color space

L = Luma (Lightness)

a = Red to Green

b = Blue to yellow



L
($a=0, b=0$)



a
($L=65, b=0$)



b
($L=65, a=0$)

Example: Original Image



Slide credit: Ren Ng

Example: Y only



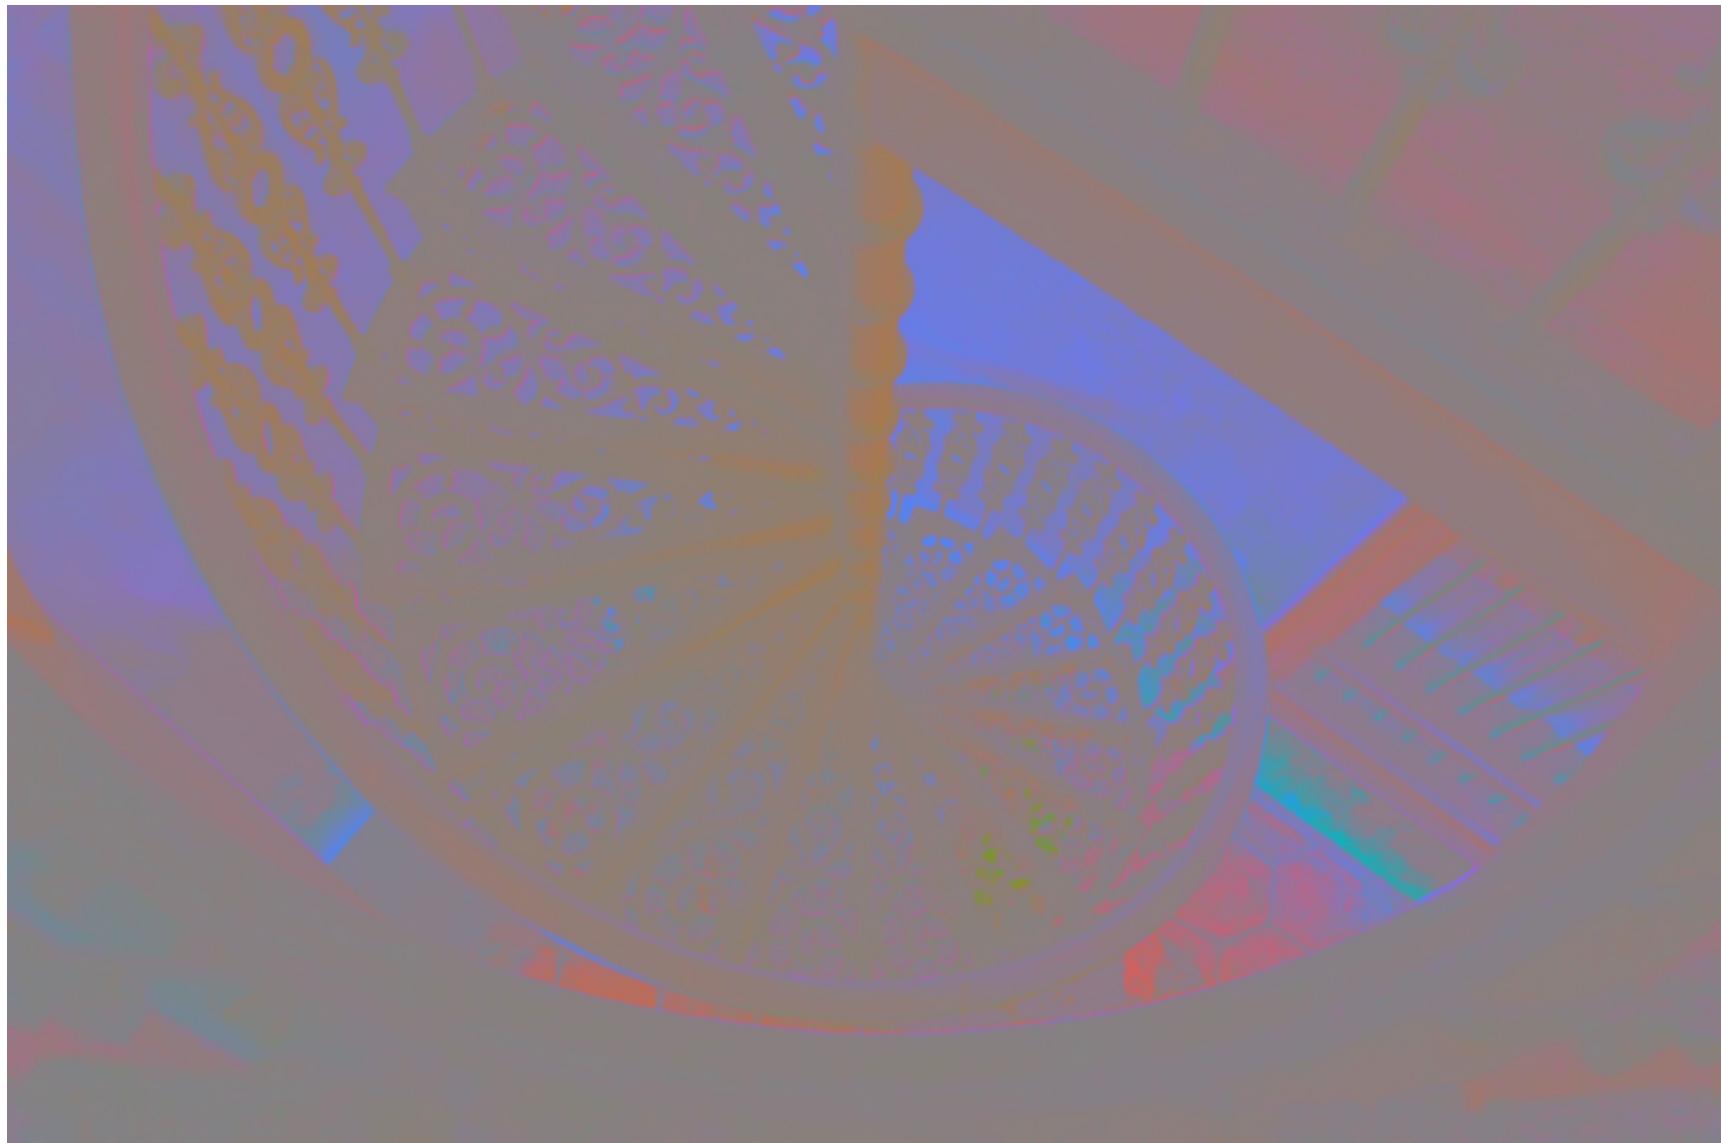
Slide credit: Ren Ng

Example: Y only – Downsampled by 4x4



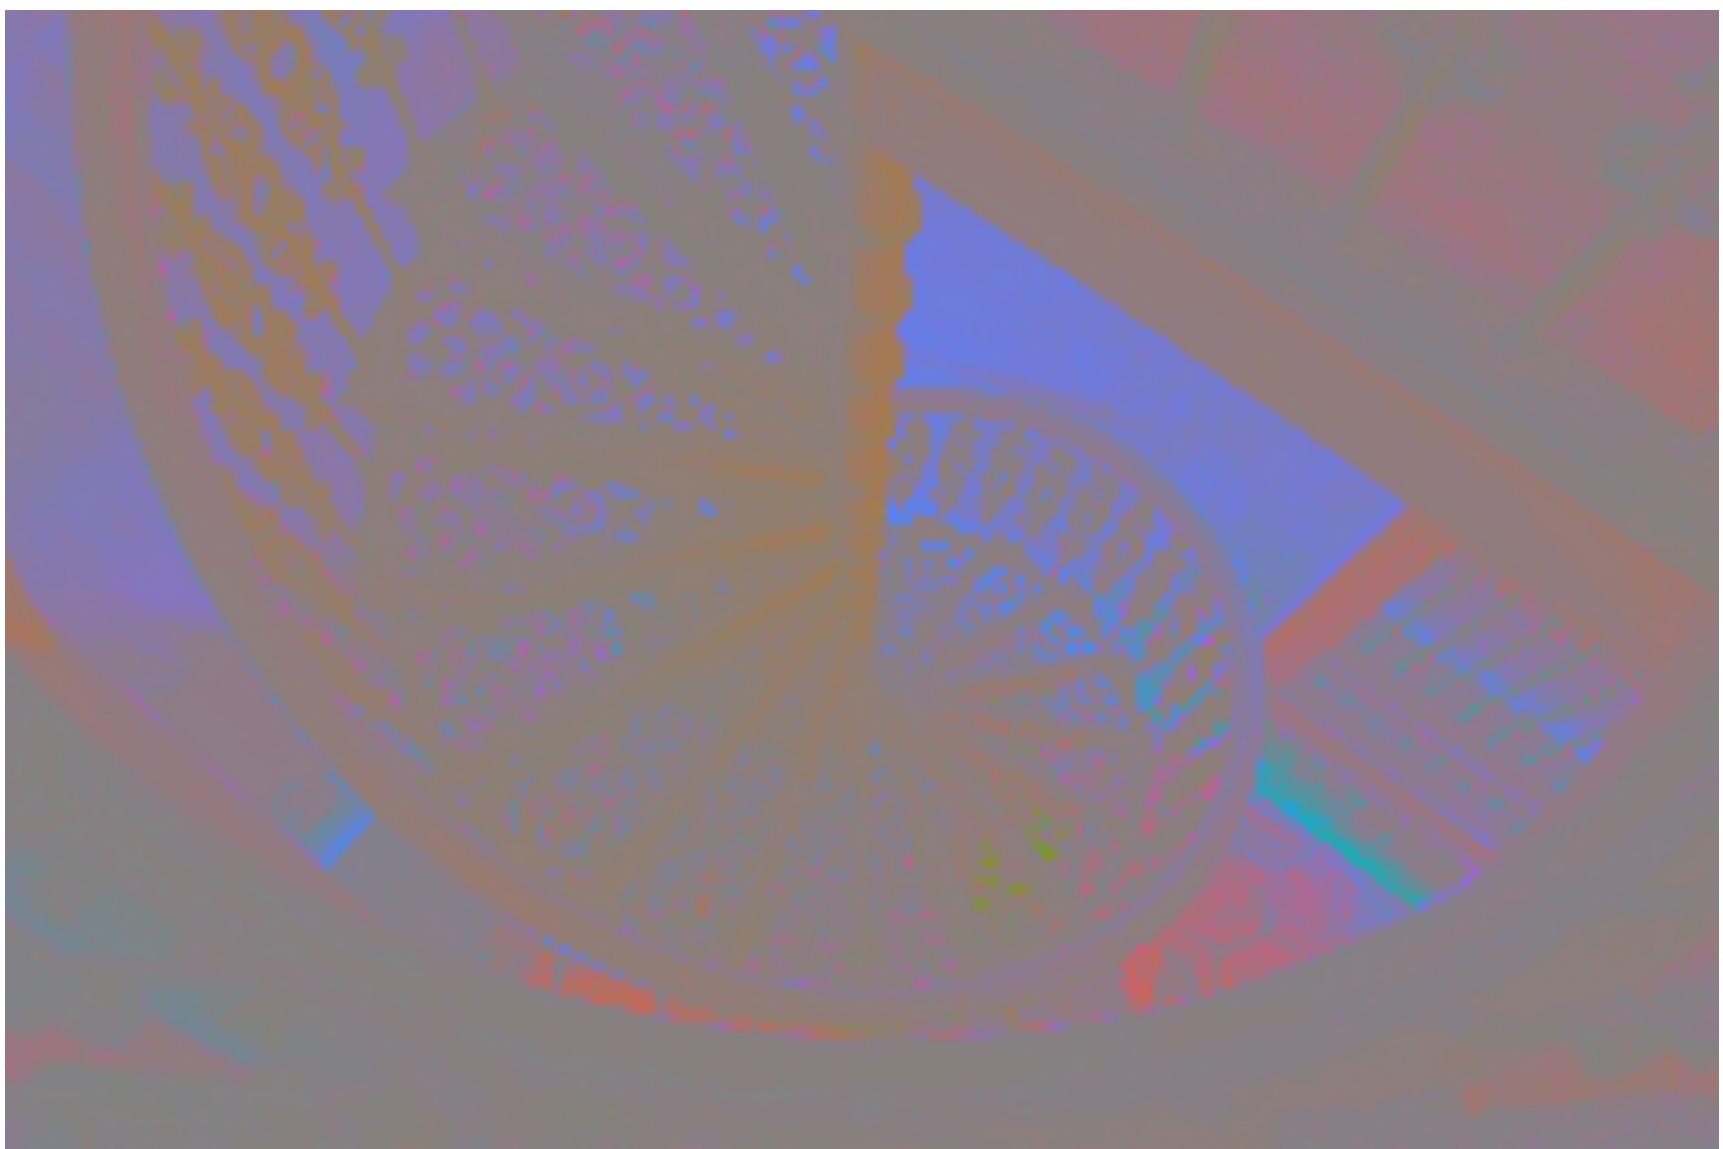
Slide credit: Ren Ng

Example: Chroma only (ab)



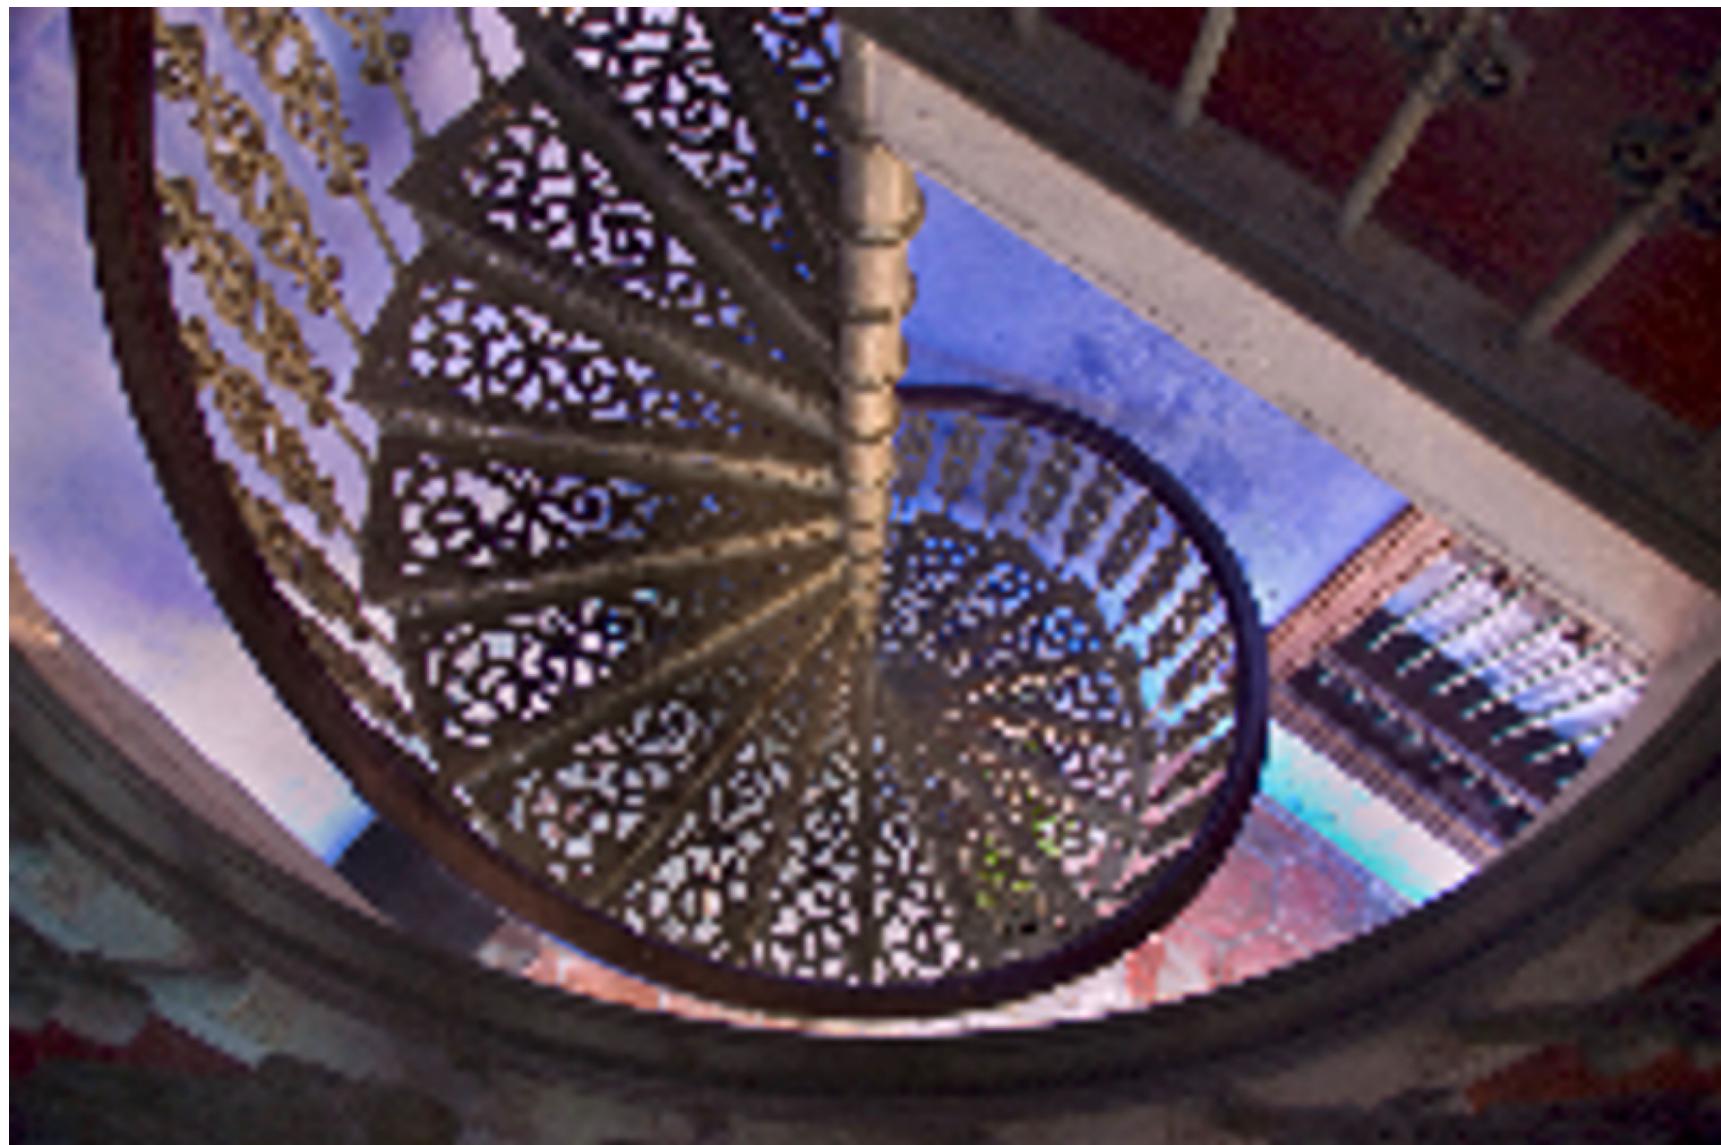
Slide credit: Ren Ng

Example: Chroma only – Downsampled 4x4



Slide credit: Ren Ng

Example: Compression in Luma



Down-sampled L, full resolution ab

Slide credit: Ren Ng

Example: Compression in Chroma channels (ab)



Full resolution L, downsampled ab

Slide credit: Ren Ng

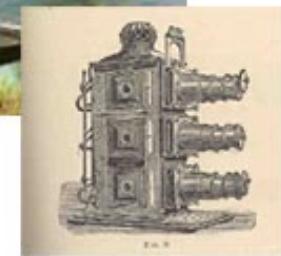
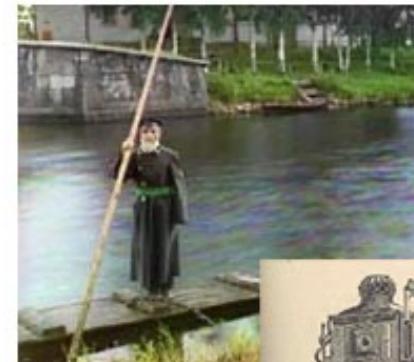
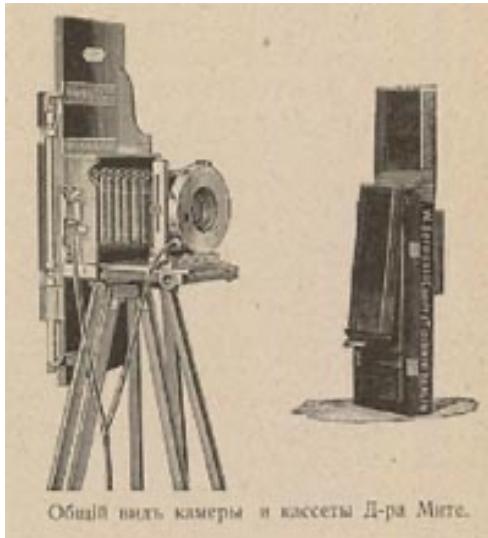
Example: Original Image



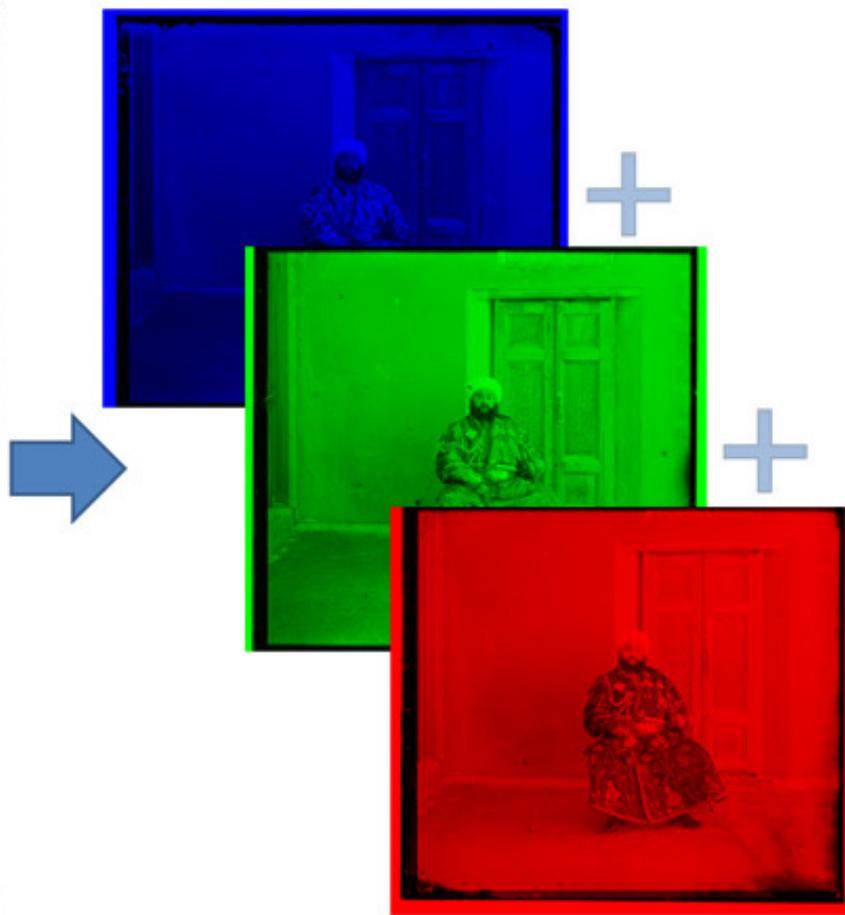
Slide credit: Ren Ng

Programming Project #1

Prokudin-Gorskii's Color Photography (1907)



Programming Project #1



Programming Project #1

- How to compare R,G,B channels?
- No right answer
 - Sum of Squared Differences (SSD):

$$ssd(u, v) = \sum_{(x,y) \in N} [I(u+x, v+y) - P(x, y)]^2$$

- Normalized Correlation (NCC):

$$ncc(u, v) = \frac{\sum_{(x,y) \in N} [I(u+x, v+y) - \bar{I}] [P(x, y) - \bar{P}]}{\sqrt{\sum_{(x,y) \in N} [I(u+x, v+y) - \bar{I}]^2} \sum_{(x,y) \in N} [P(x, y) - \bar{P}]^2}$$



Watch these 5 min videos!

Color in 5 min: <https://youtu.be/6tTNgvAl1y4>

Displays in 5 min: <https://youtu.be/1albYPL9Cfg>