

Point Processing and Filtering



CS194: Intro to Comp. Vision, and Comp. Photo
Angjoo Kanazawa and Alexei Efros, UC Berkeley, Fall 2022

About me

Angjoo Kanazawa

- First name pronounced like “Andrew”
- Grew up in Kobe, Japan



Saito et al. PIFu: Pixel-Aligned Implicit Function for High-Resolution Clothed Human ICCV 2019

Recitation & Project Party

**Python Intro: Thursday 11am-12pm
BWW 1st floor 1216**

Project 1 Party: Friday 3-4:30pm

All in the class calendar.

Piazza

What is an image?

We can think of an **image** as a function, f , from \mathbf{R}^2 to \mathbf{R} :

- $f(x, y)$ gives the **intensity** at position (x, y)
- Realistically, we expect the image only to be defined over a rectangle, with a finite range:
 - $f: [a,b] \times [c,d] \rightarrow [0,1]$

A color image is just three functions pasted together.
We can write this as a “vector-valued” function:

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

Images as functions

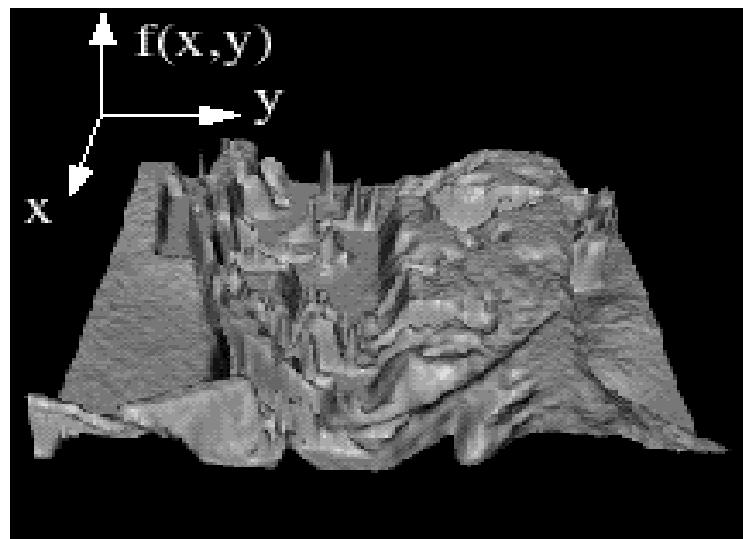
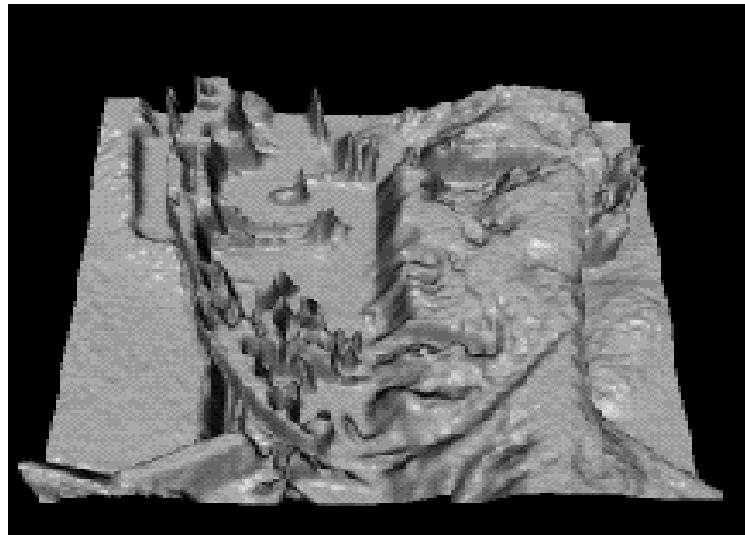


Image Formation

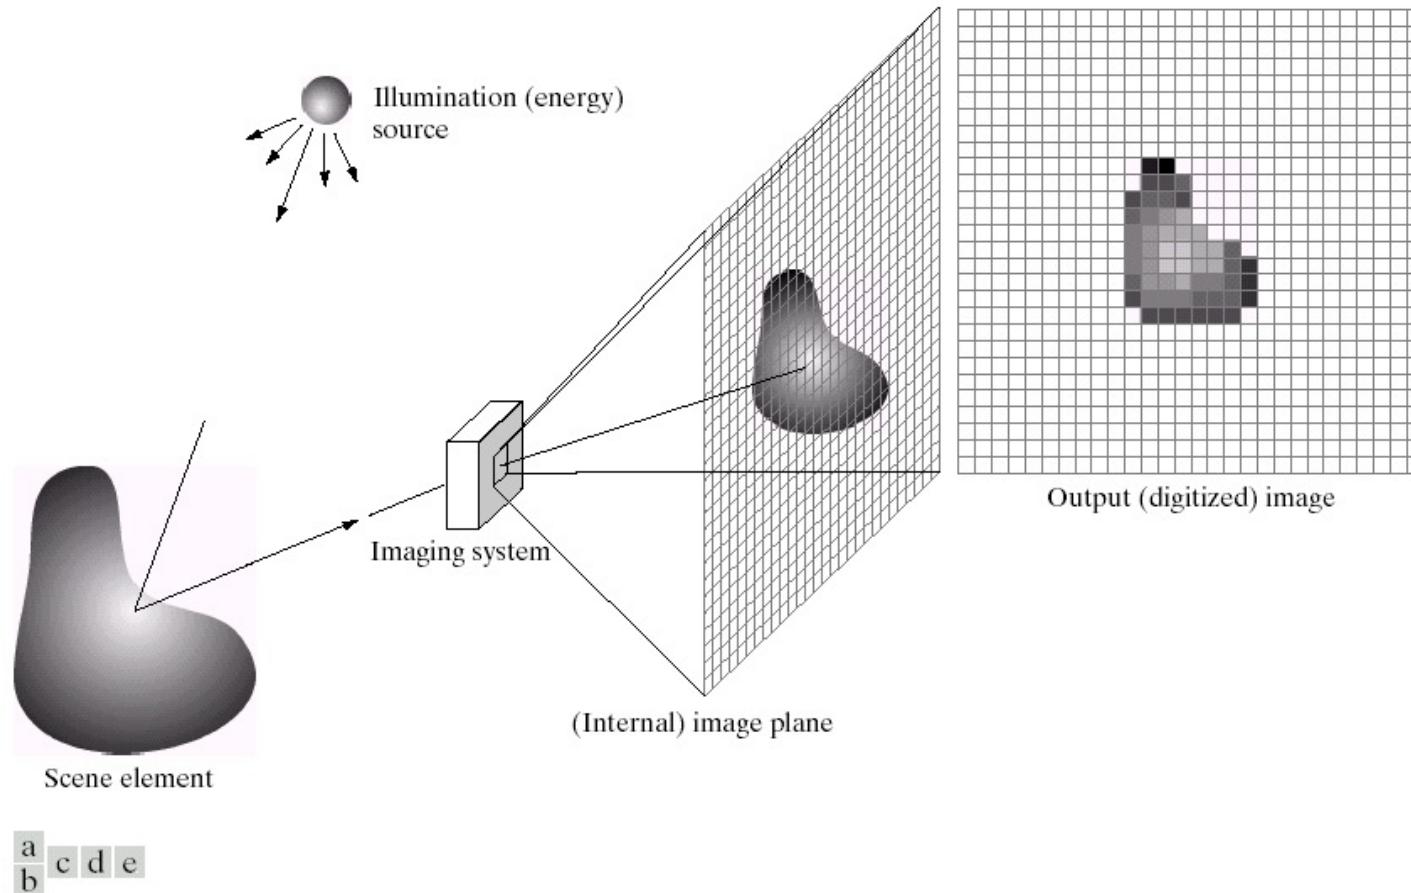


FIGURE 2.15 An example of the digital image acquisition process. (a) Energy (“illumination”) source. (b) An element of a scene. (c) Imaging system. (d) Projection of the scene onto the image plane. (e) Digitized image.

$$f(x,y) = \text{reflectance}(x,y) * \text{illumination}(x,y)$$

Reflectance in $[0, 1]$, illumination in $[0, \infty]$

Problem: Dynamic Range



1



1500



25,000



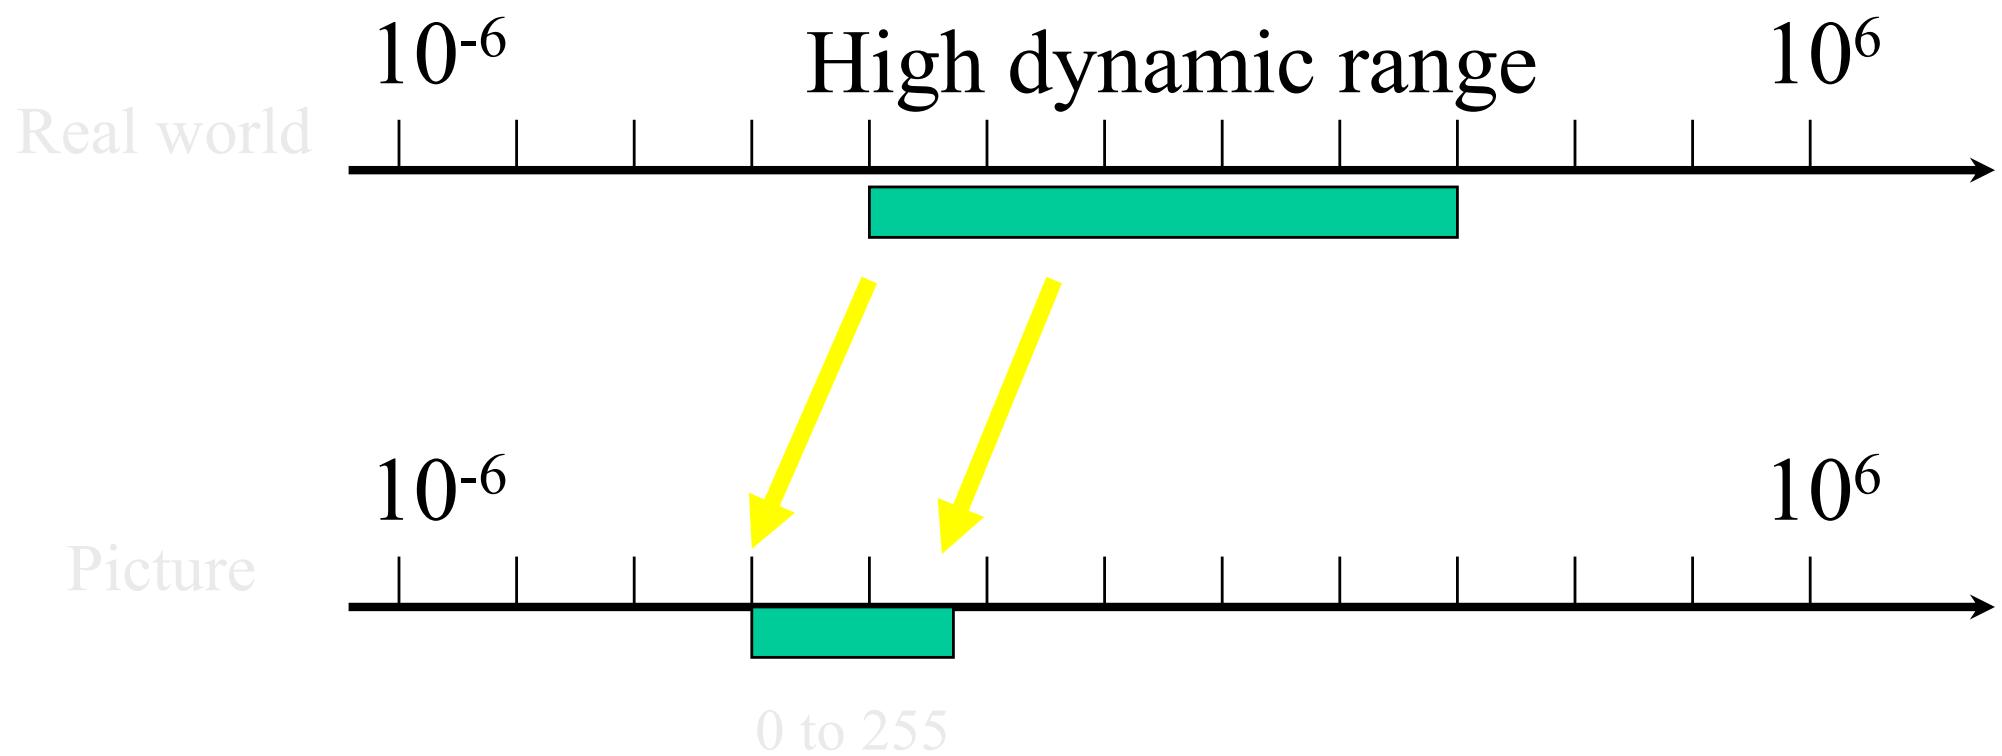
The real world is
High dynamic range

400,000



2,000,000,000

Long Exposure



Short Exposure

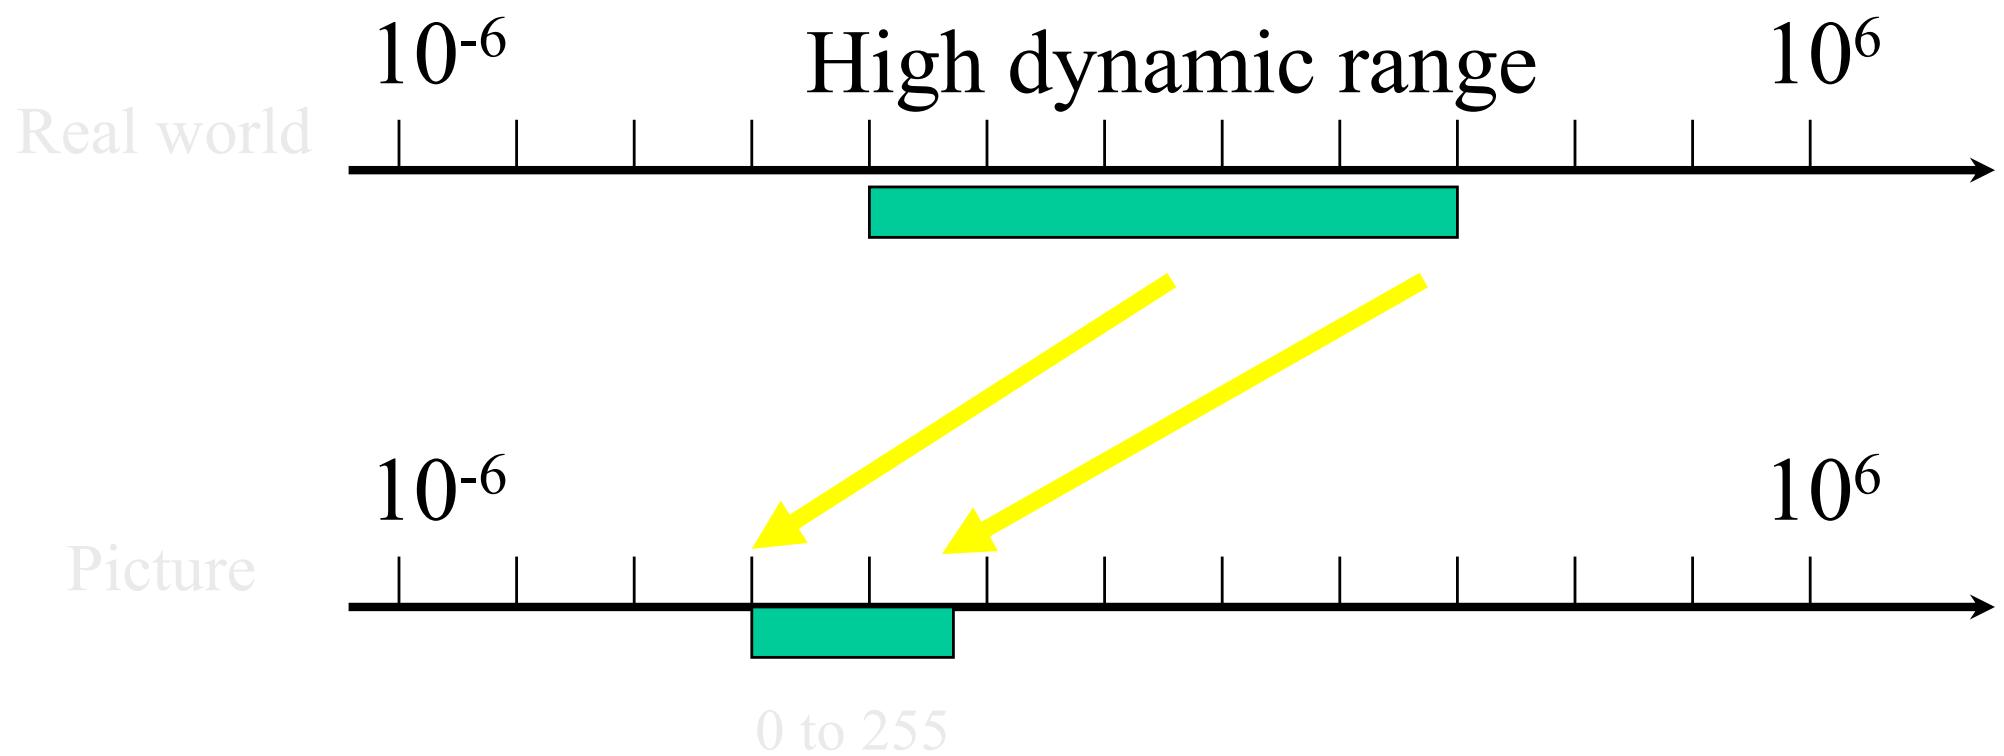
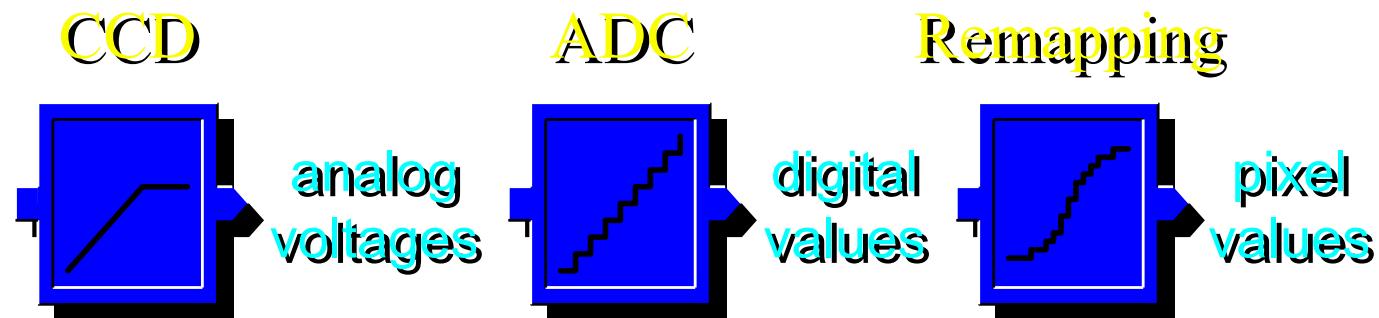
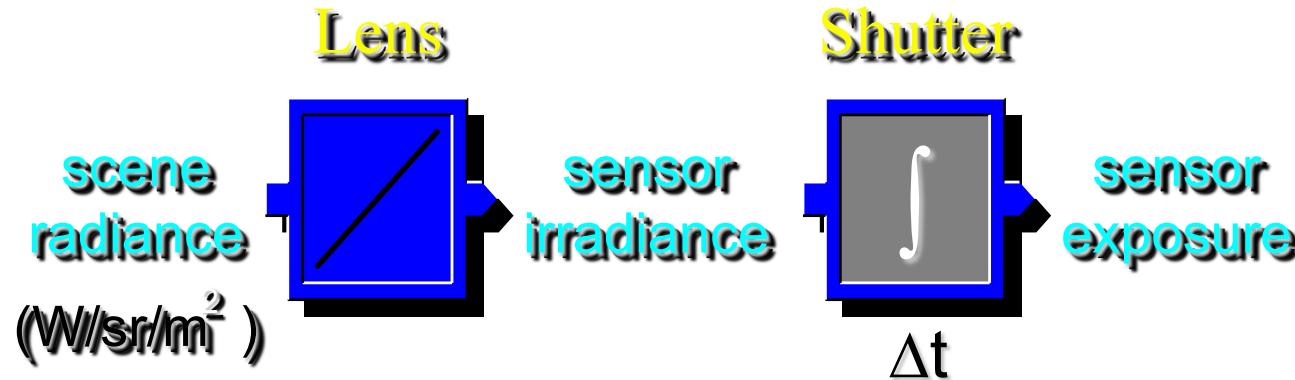


Image Acquisition Pipeline

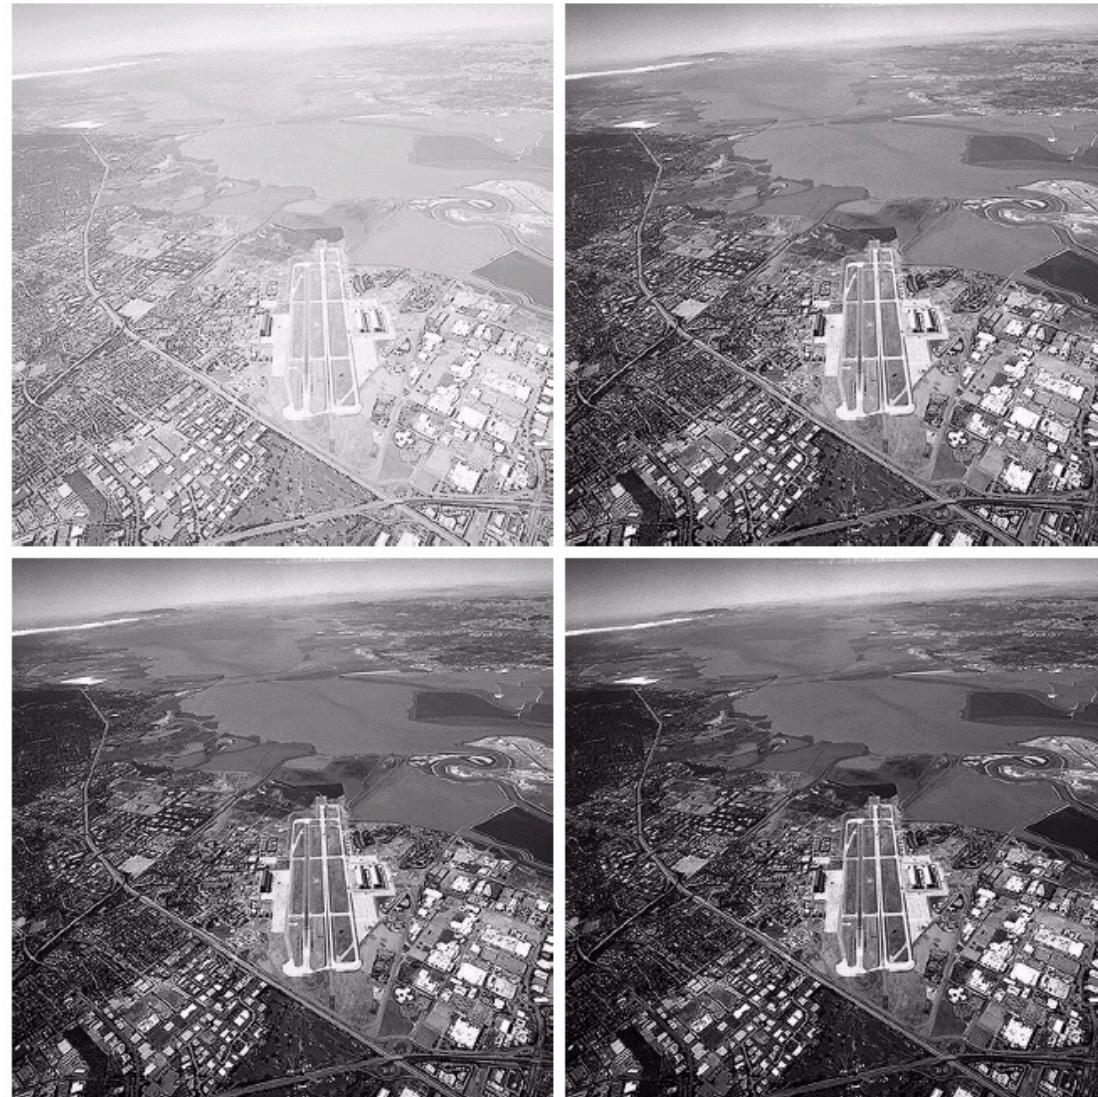


Simple Point Processing: Enhancement

a	b
c	d

FIGURE 3.9

(a) Aerial image.
(b)–(d) Results of applying the transformation in Eq. (3.2-3) with $c = 1$ and $\gamma = 3.0, 4.0$, and 5.0 , respectively. (Original image for this example courtesy of NASA.)



Power-law transformations

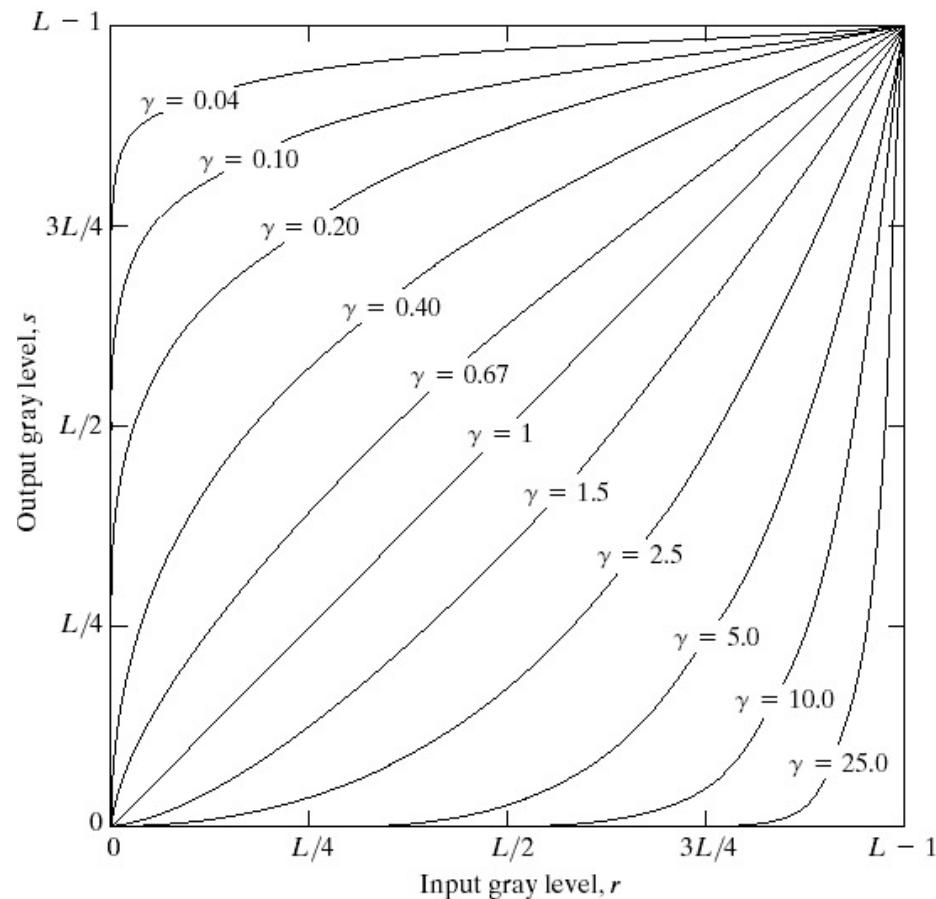
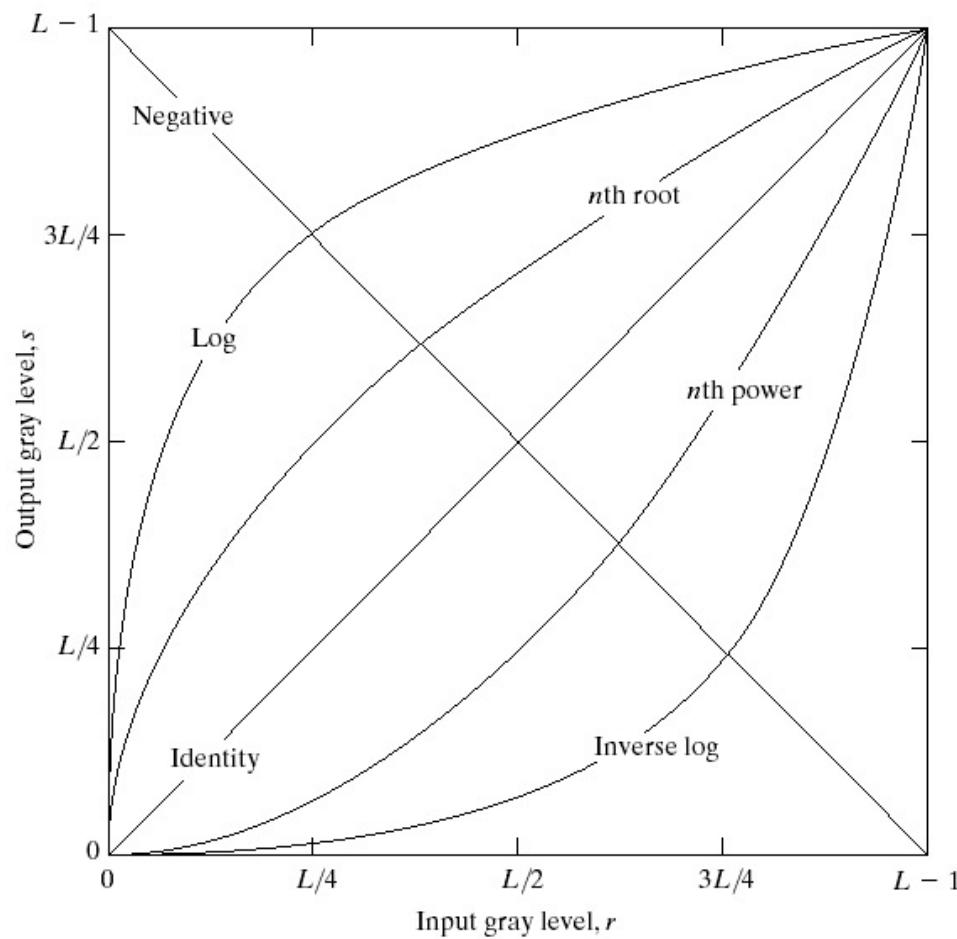


FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

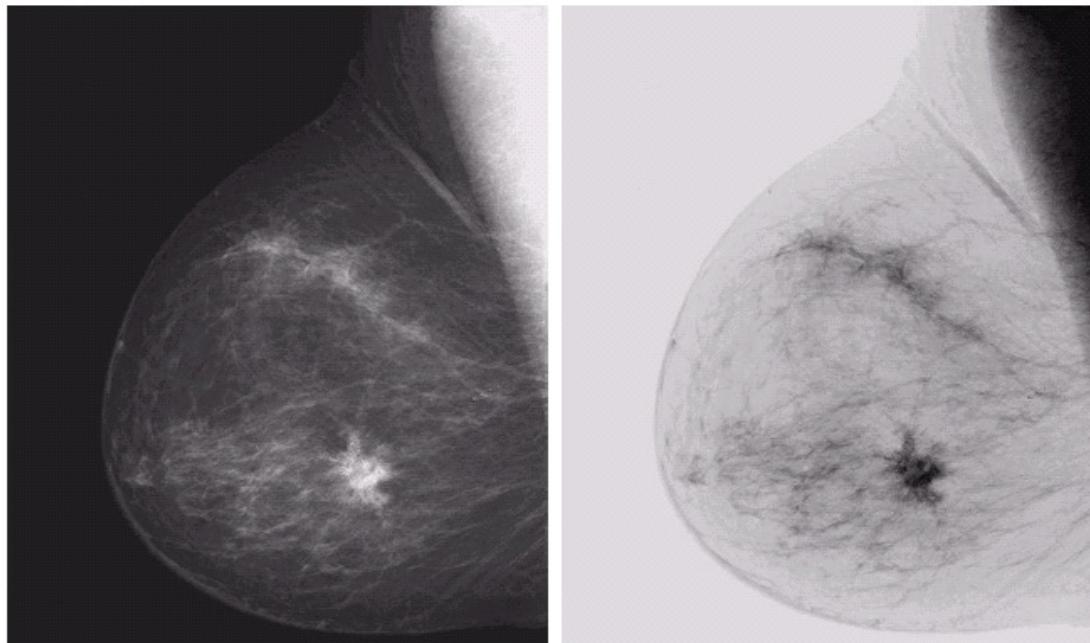
$$s = cr^\gamma$$

Basic Point Processing

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.



Negative



a b

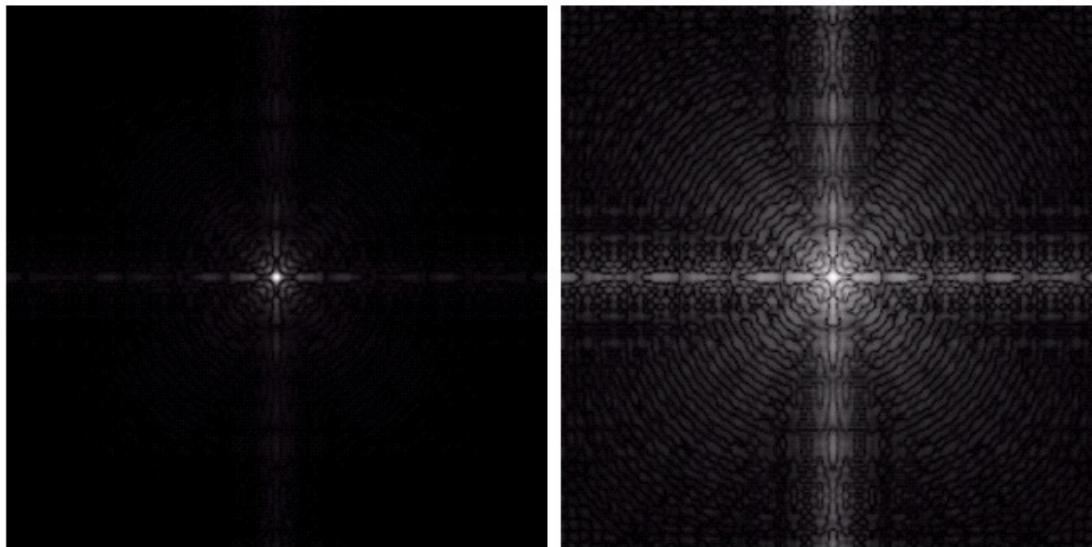
FIGURE 3.4
(a) Original
digital
mammogram.
(b) Negative
image obtained
using the negative
transformation in
Eq. (3.2-1).
(Courtesy of G.E.
Medical Systems.)

Log

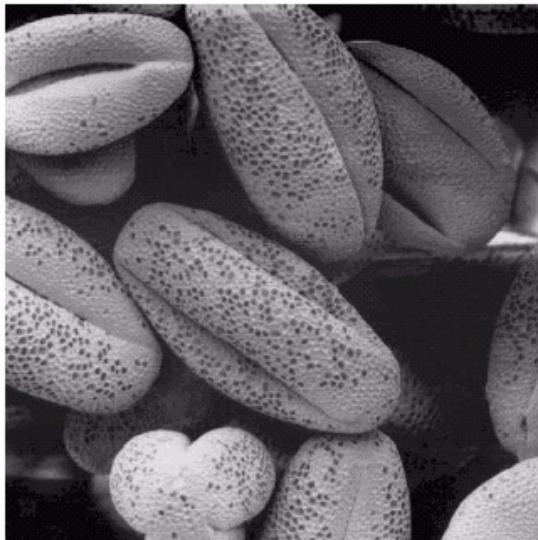
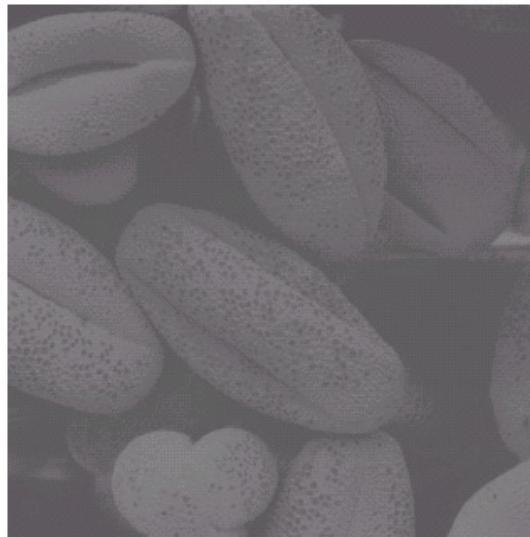
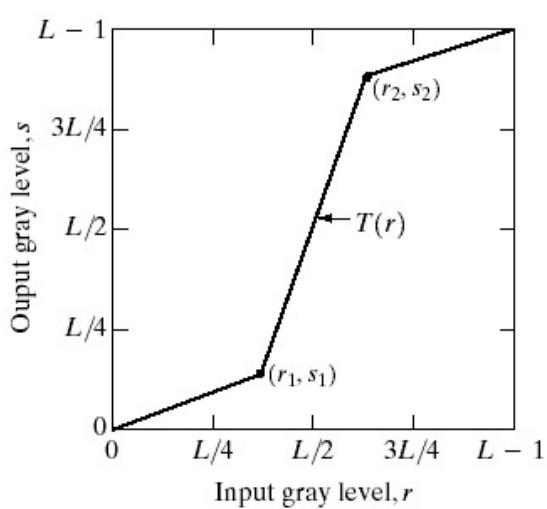
a b

FIGURE 3.5

(a) Fourier spectrum.
(b) Result of applying the log transformation given in Eq. (3.2-2) with $c = 1$.



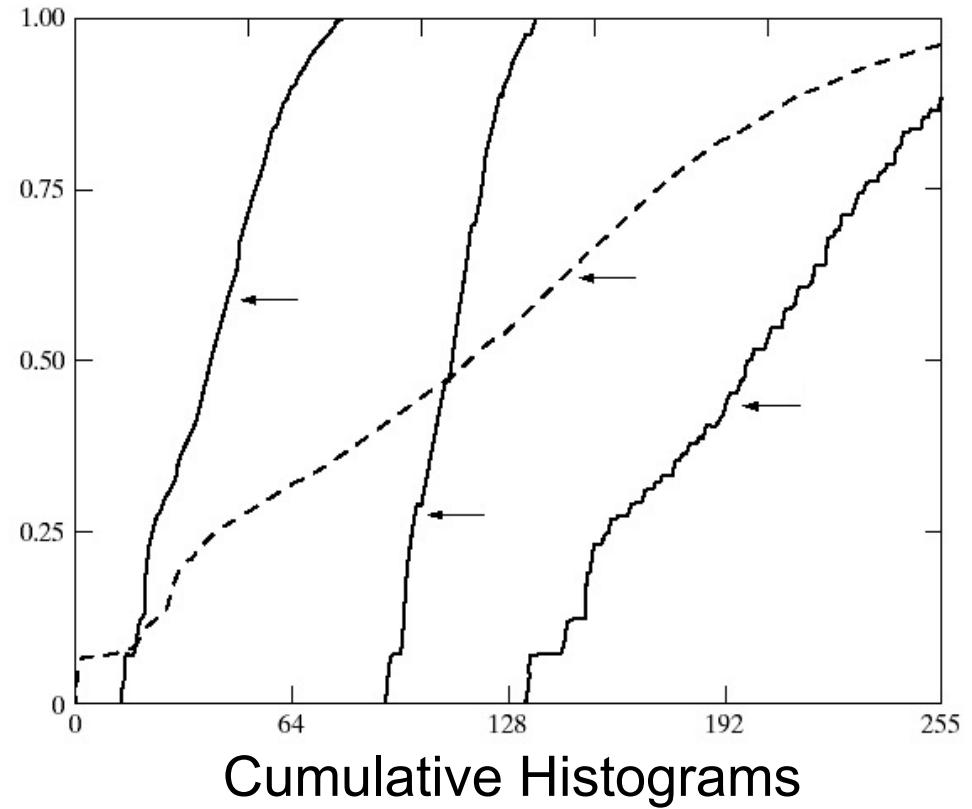
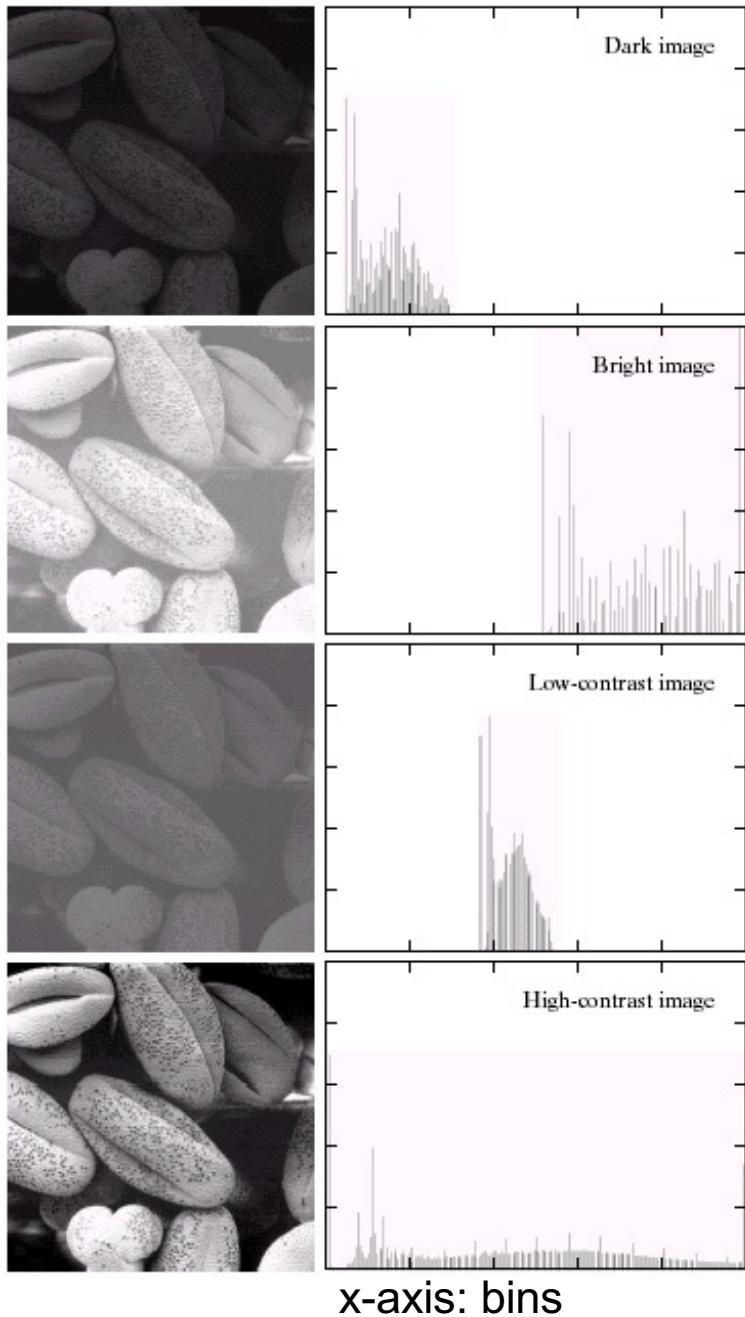
Contrast Stretching



a b
c d

FIGURE 3.10
Contrast stretching.
(a) Form of transformation function. (b) A low-contrast image. (c) Result of contrast stretching. (d) Result of thresholding.
(Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Image Histograms



$$s = T(r)$$

a b

FIGURE 3.15 Four basic image types: dark, light, low contrast, high contrast, and their corresponding histograms. (Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Histogram Equalization

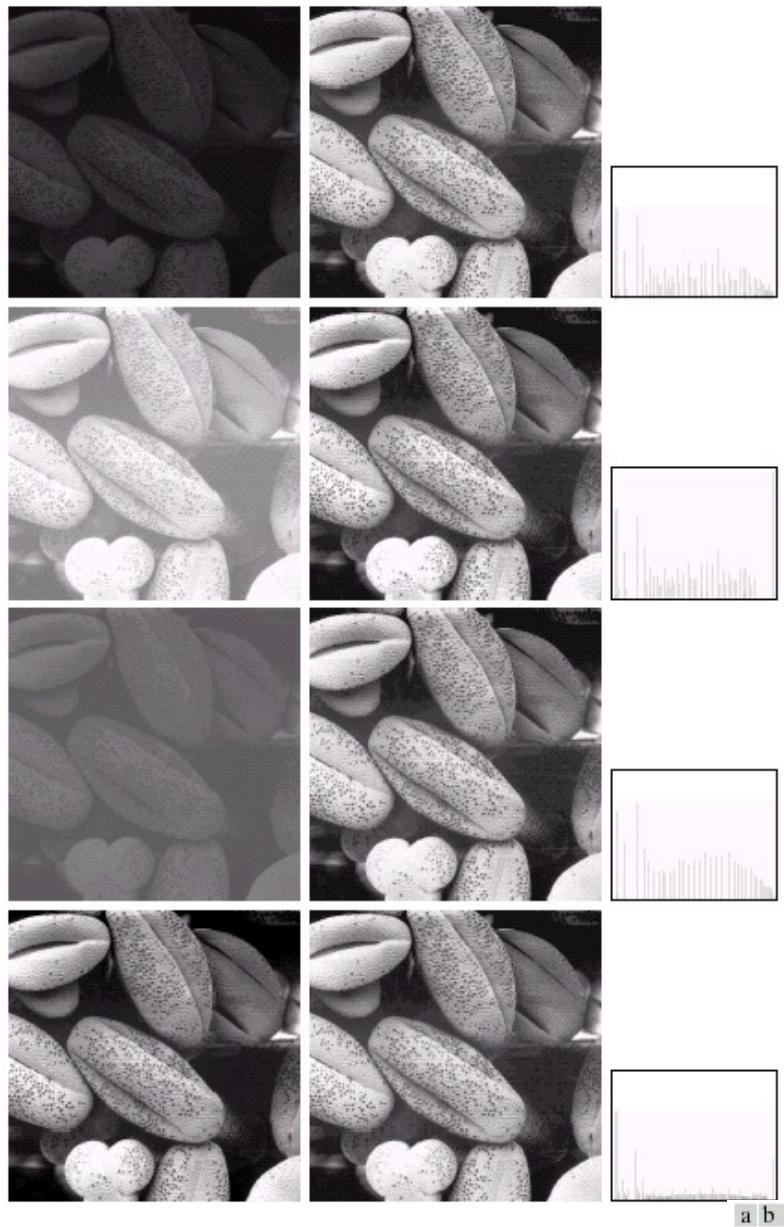
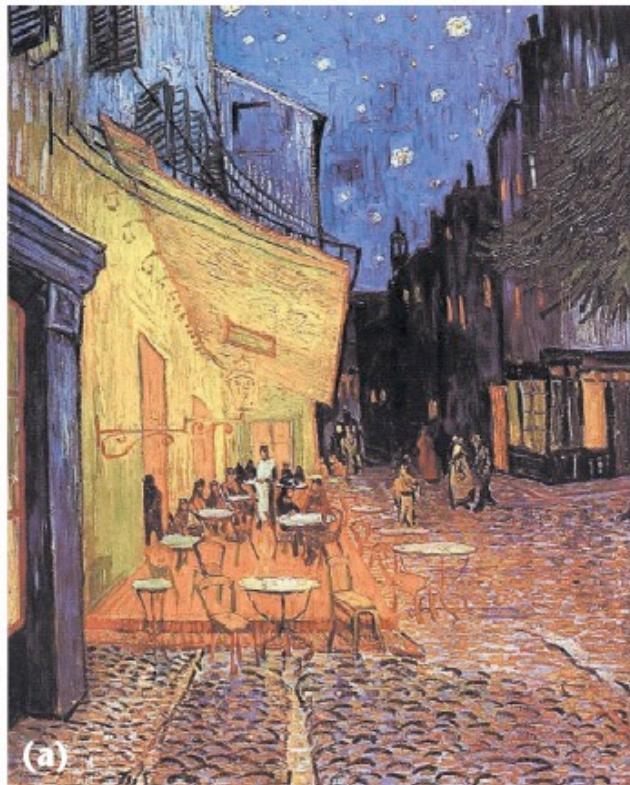


FIGURE 3.17 (a) Images from Fig. 3.15. (b) Results of histogram equalization. (c) Corresponding histograms.

Color Transfer [Reinhard, et al, 2001]



(a)

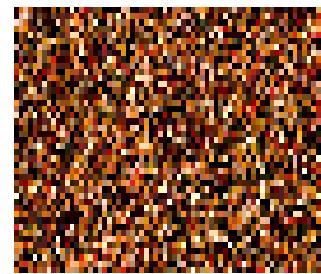


(b)

Erik Reinhard, Michael Ashikhmin, Bruce Gooch, Peter Shirley, [Color Transfer between Images](#). *IEEE Computer Graphics and Applications*, 21(5), pp. 34–41. September 2001.

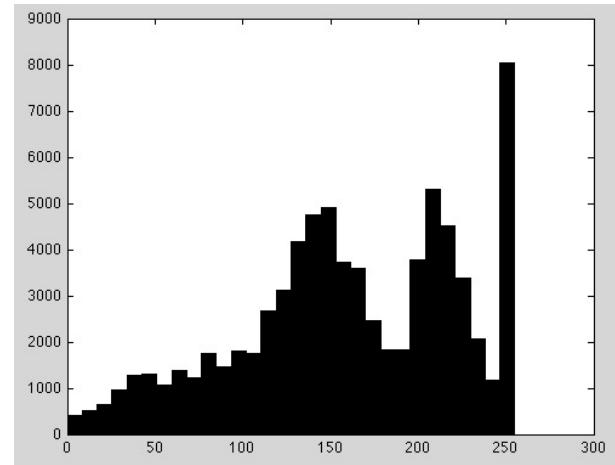
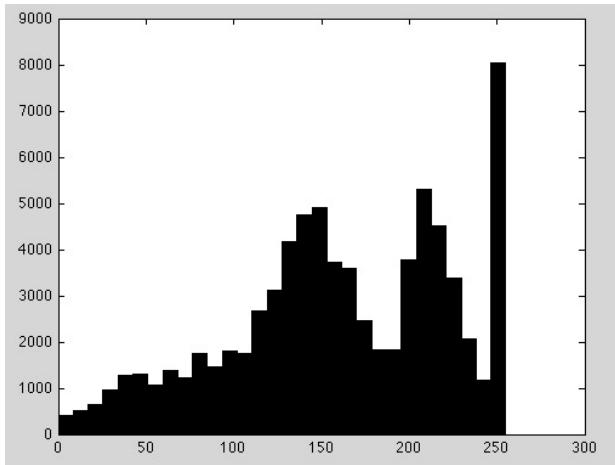
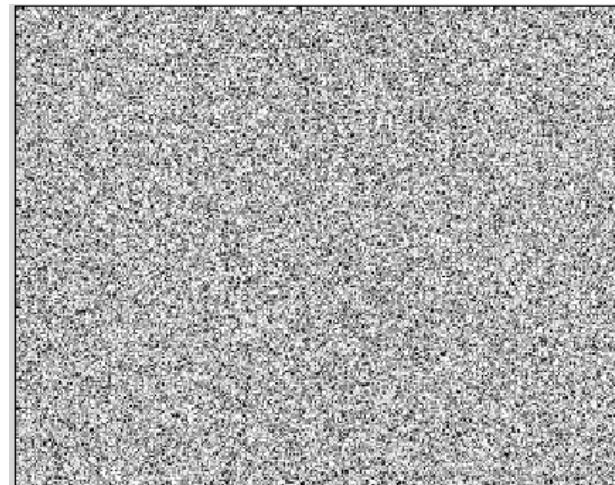
Limitations of Point Processing

Q: What happens if I reshuffle all pixels within the image?

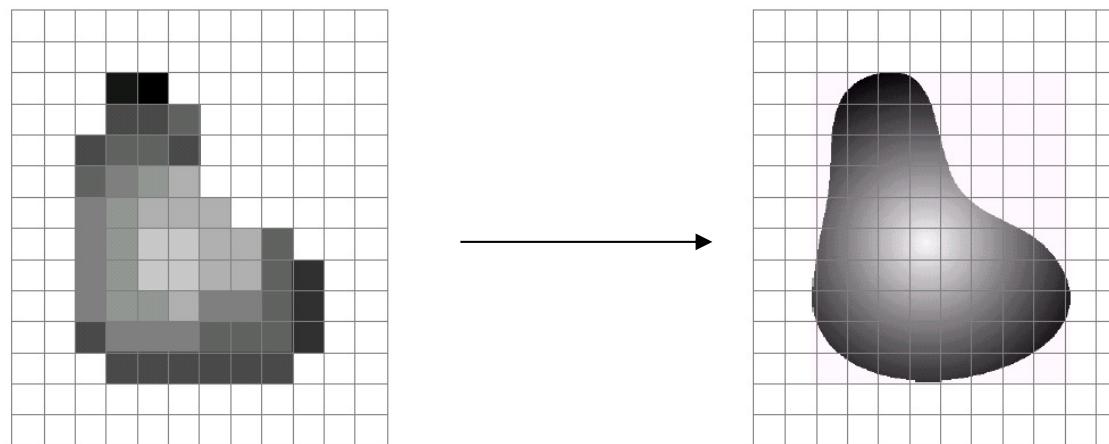
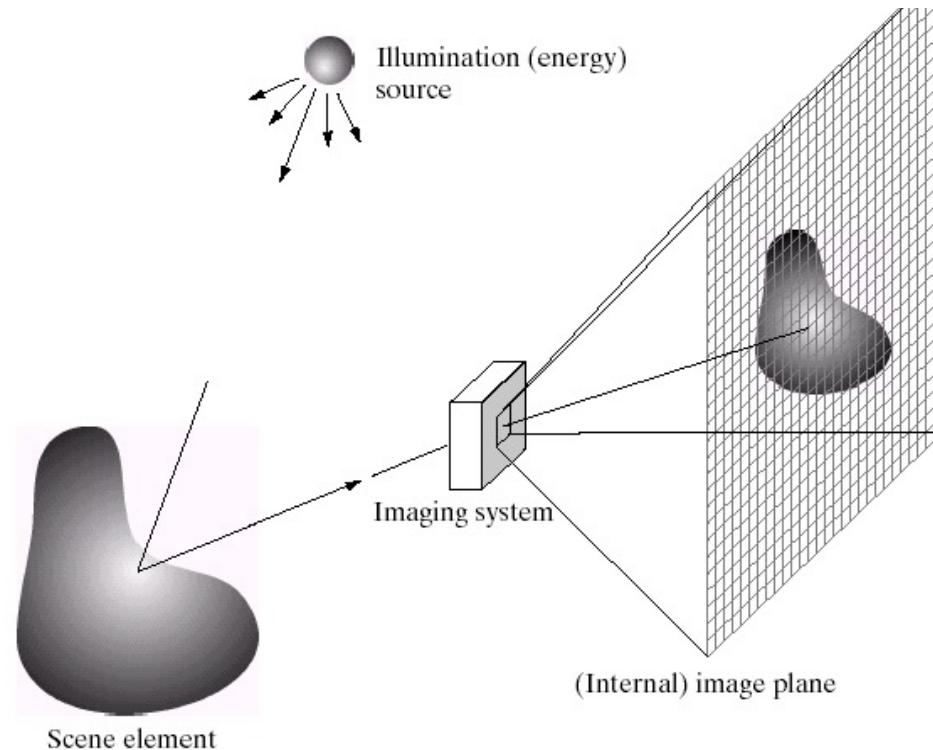


A: Its histogram won't change. No point processing will be affected...

Limitations of Point Processing...



Sampling and Reconstruction

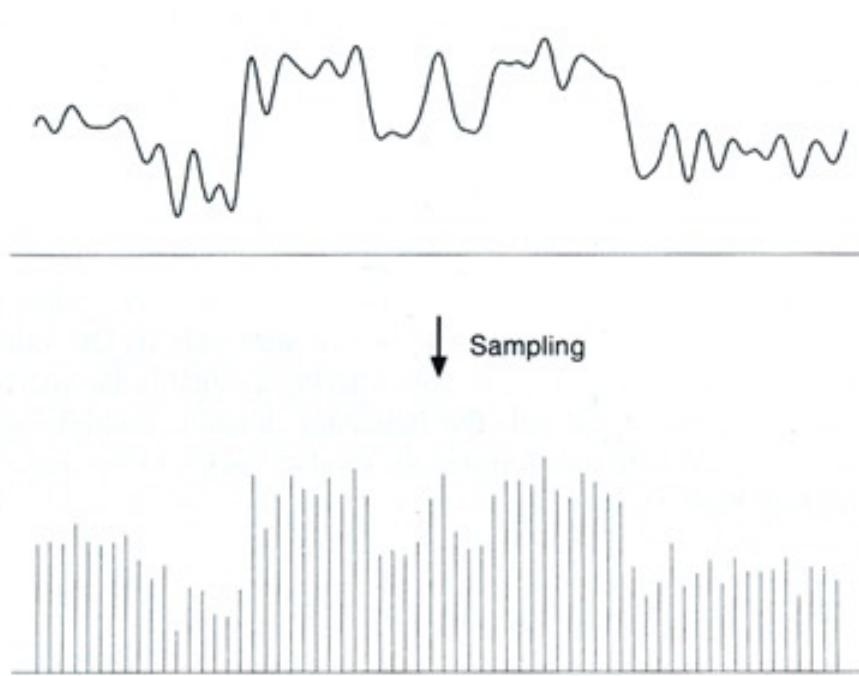


Sampled representations

How to store and compute with continuous functions?

Common scheme for representation: samples

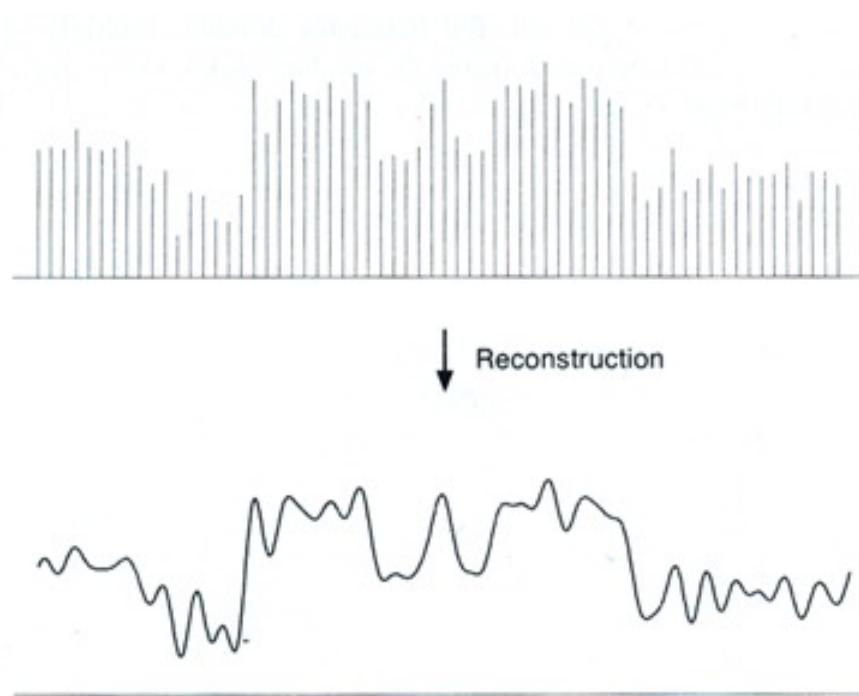
- write down the function's values at many points



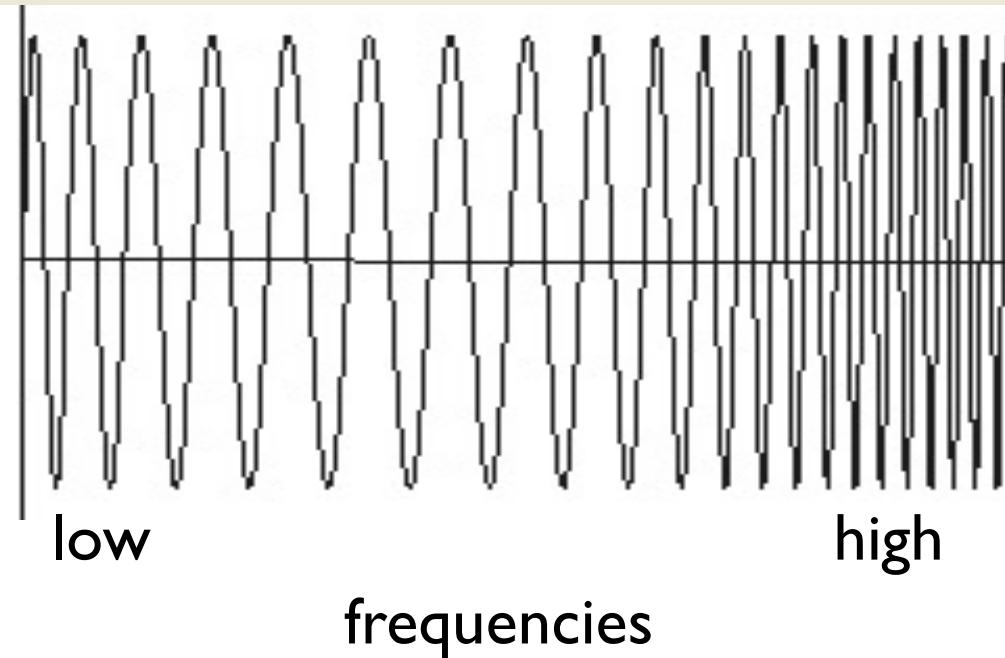
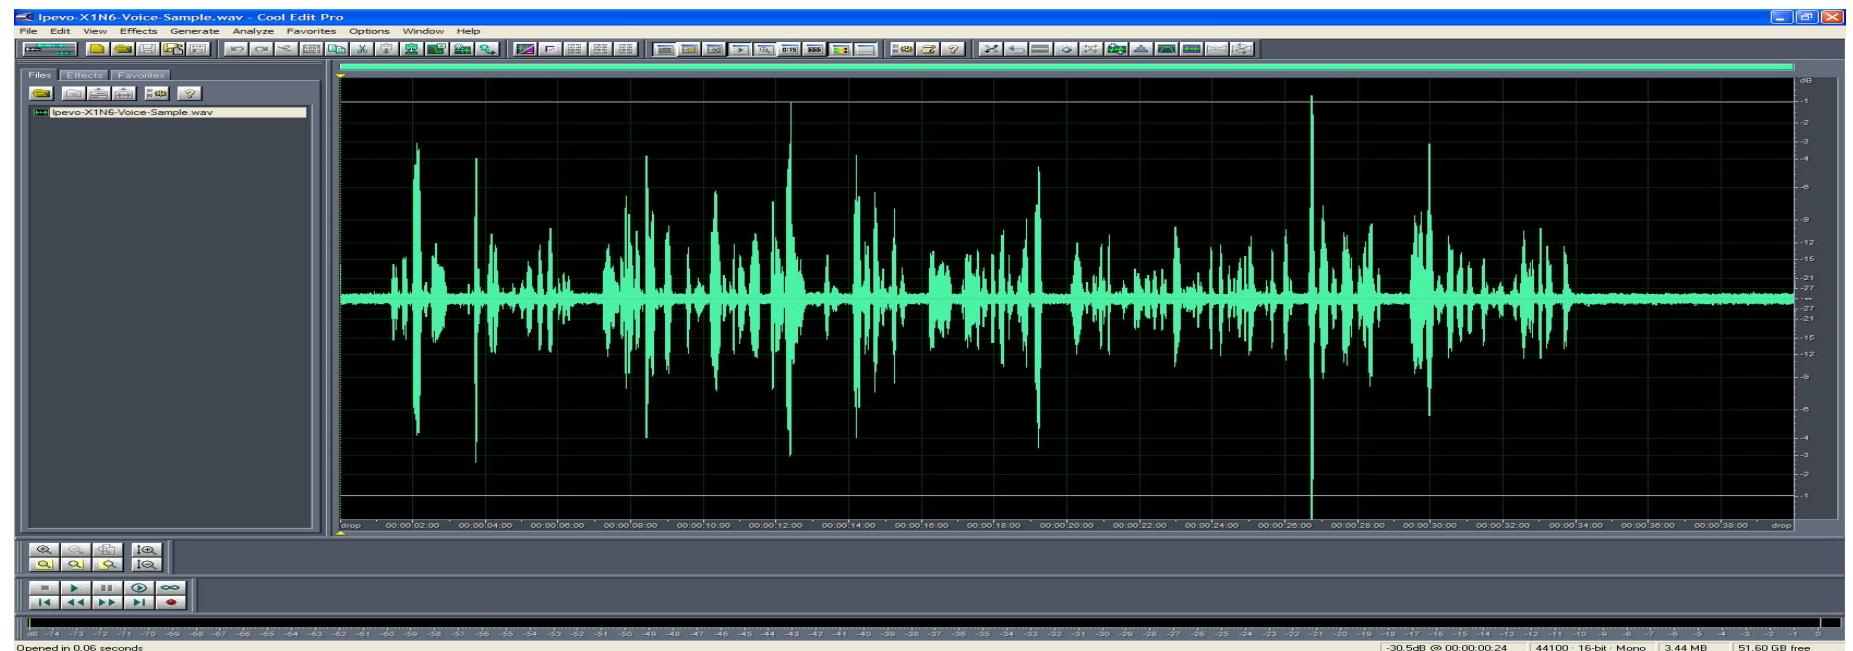
Reconstruction

Making samples back into a continuous function

- for output (need realizable method)
- for analysis or processing (need mathematical method)
- amounts to “guessing” what the function did in between



1D Example: Audio

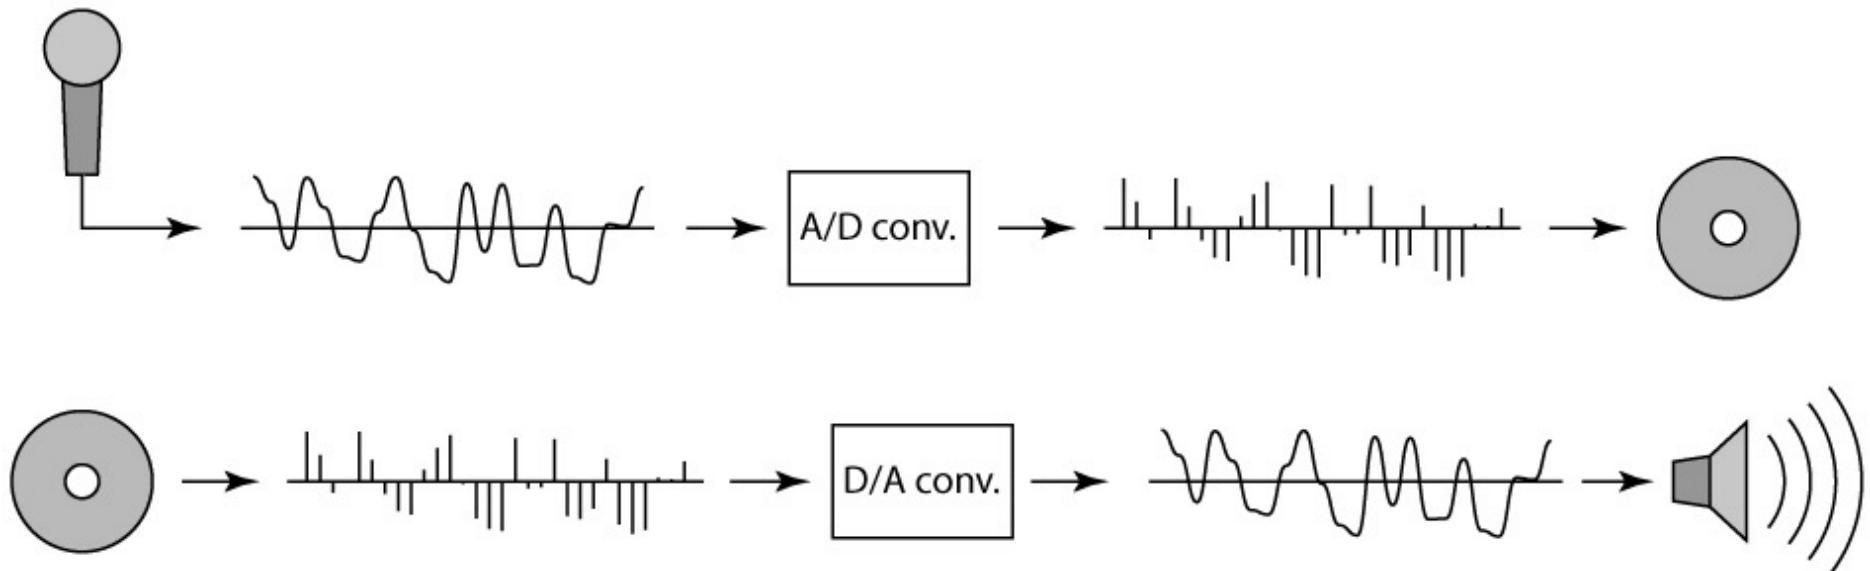


Sampling in digital audio

Recording: sound to analog to samples to disc

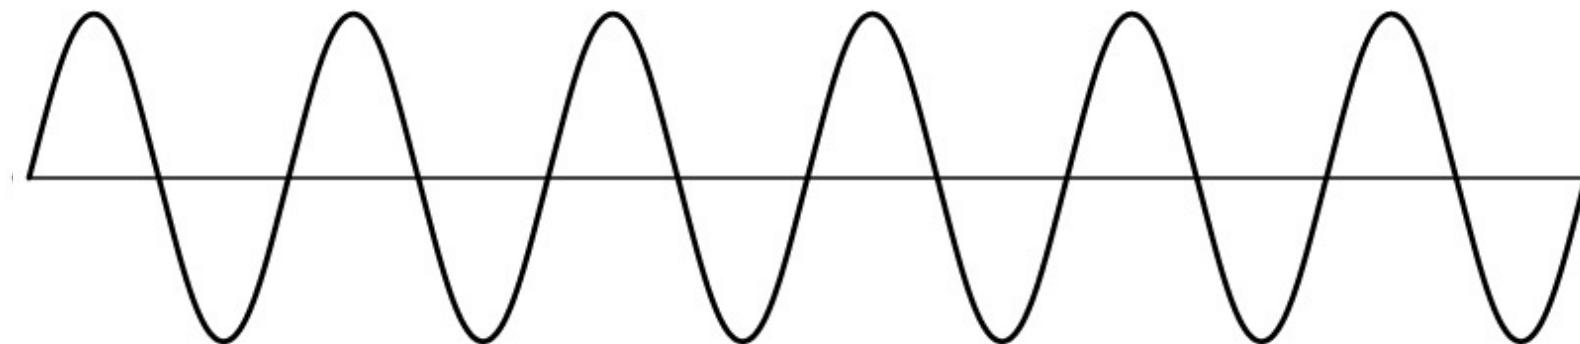
Playback: disc to samples to analog to sound again

- how can we be sure we are filling in the gaps correctly?



Sampling and Reconstruction

Simple example: a sign wave

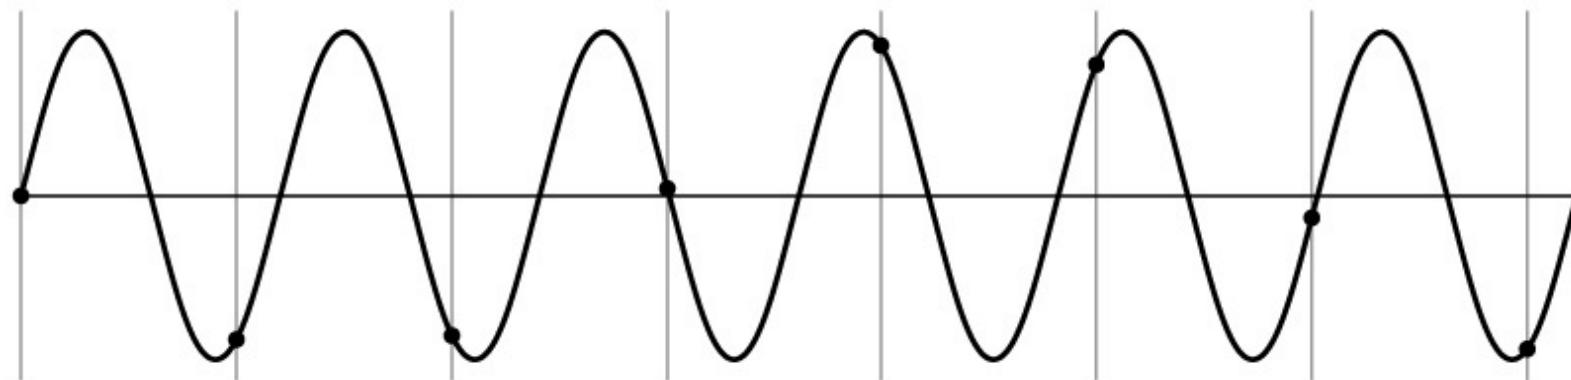


Undersampling

What if we “missed” things between the samples?

Simple example: undersampling a sine wave

- unsurprising result: information is lost

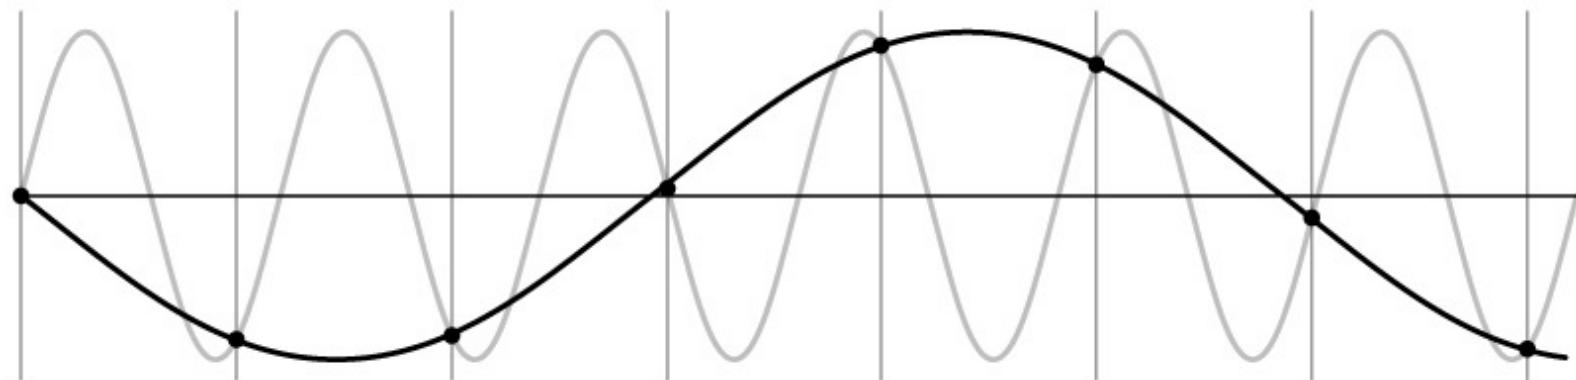


Undersampling

What if we “missed” things between the samples?

Simple example: undersampling a sine wave

- unsurprising result: information is lost
- surprising result: indistinguishable from lower frequency

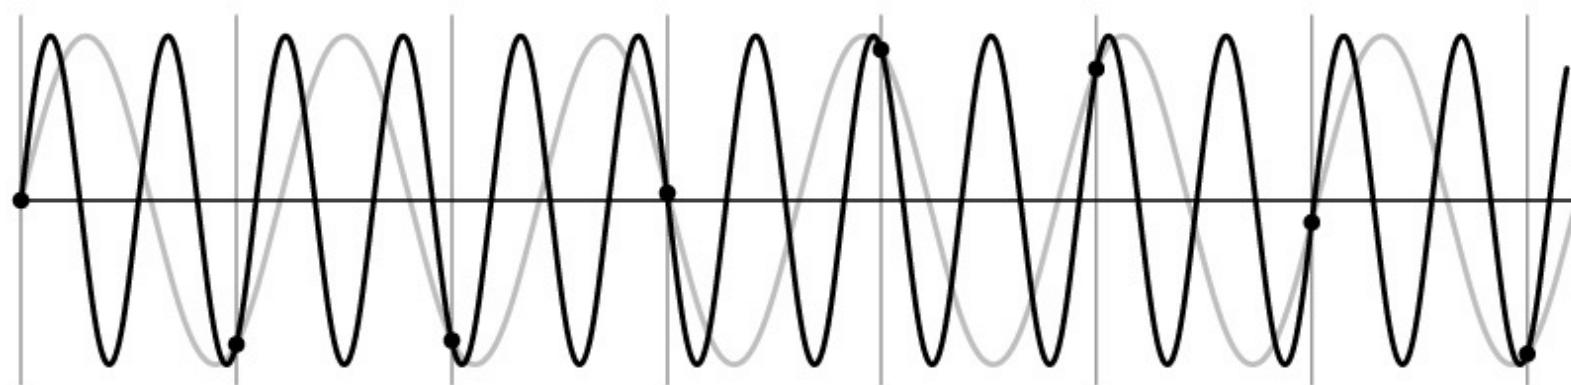


Undersampling

What if we “missed” things between the samples?

Simple example: undersampling a sine wave

- unsurprising result: information is lost
- surprising result: indistinguishable from lower frequency
- also, was always indistinguishable from higher frequencies
- aliasing: signals “traveling in disguise” as other frequencies



Aliasing!

Alias = Pseudonym, False Identity

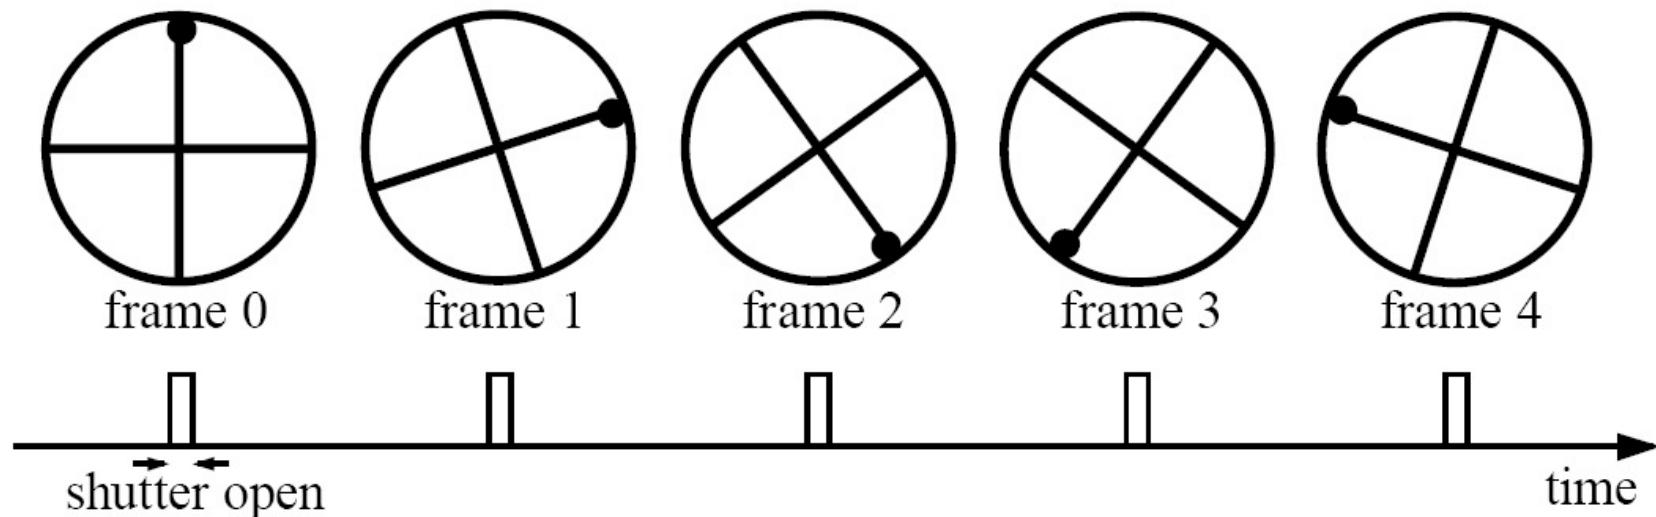


Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = $1/30$ sec. for video, $1/24$ sec. for film):



Without dot, wheel appears to be rotating slowly backwards!
(counterclockwise)



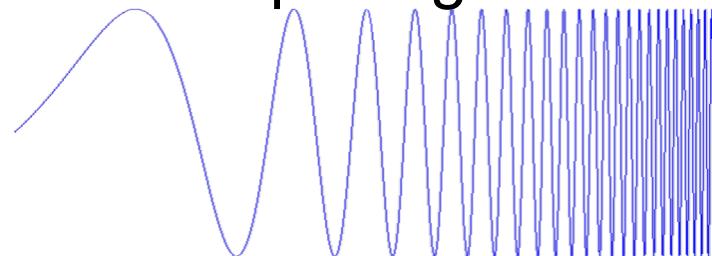
Aliasing in images



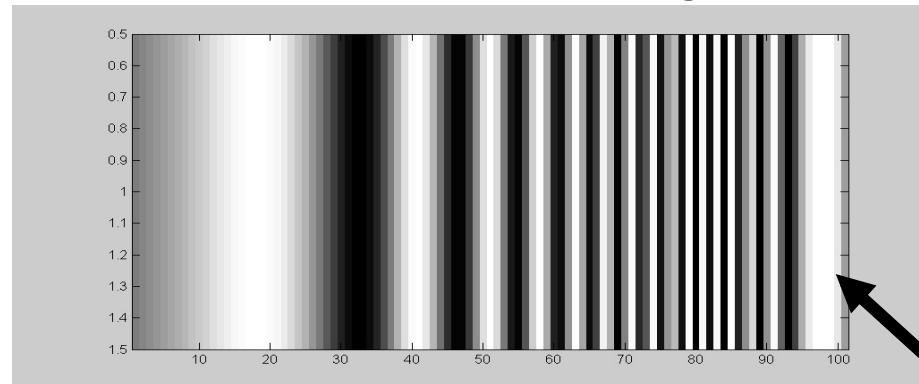
Disintegrating textures

What's happening?

Input signal:



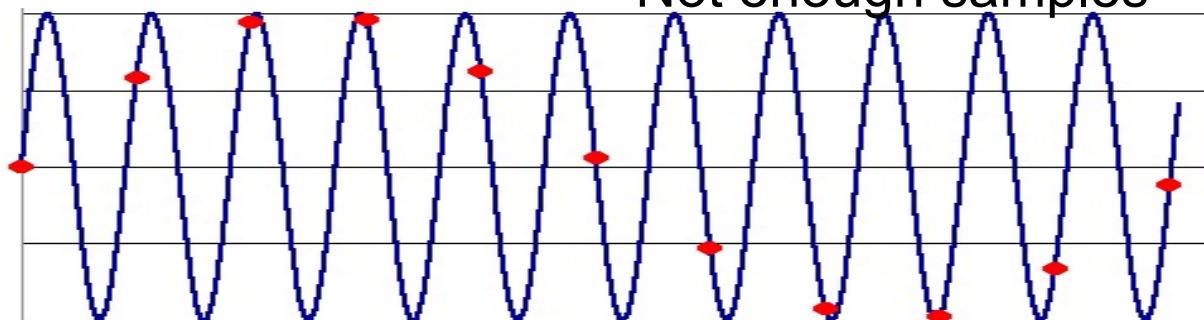
Plot as image:



```
x = 0:.05:5; imagesc(sin((2.^x).*x))
```

Alias!

Not enough samples



Antialiasing

What can we do about aliasing?

Sample more often

- Join the Mega-Pixel craze of the photo industry
- But this can't go on forever

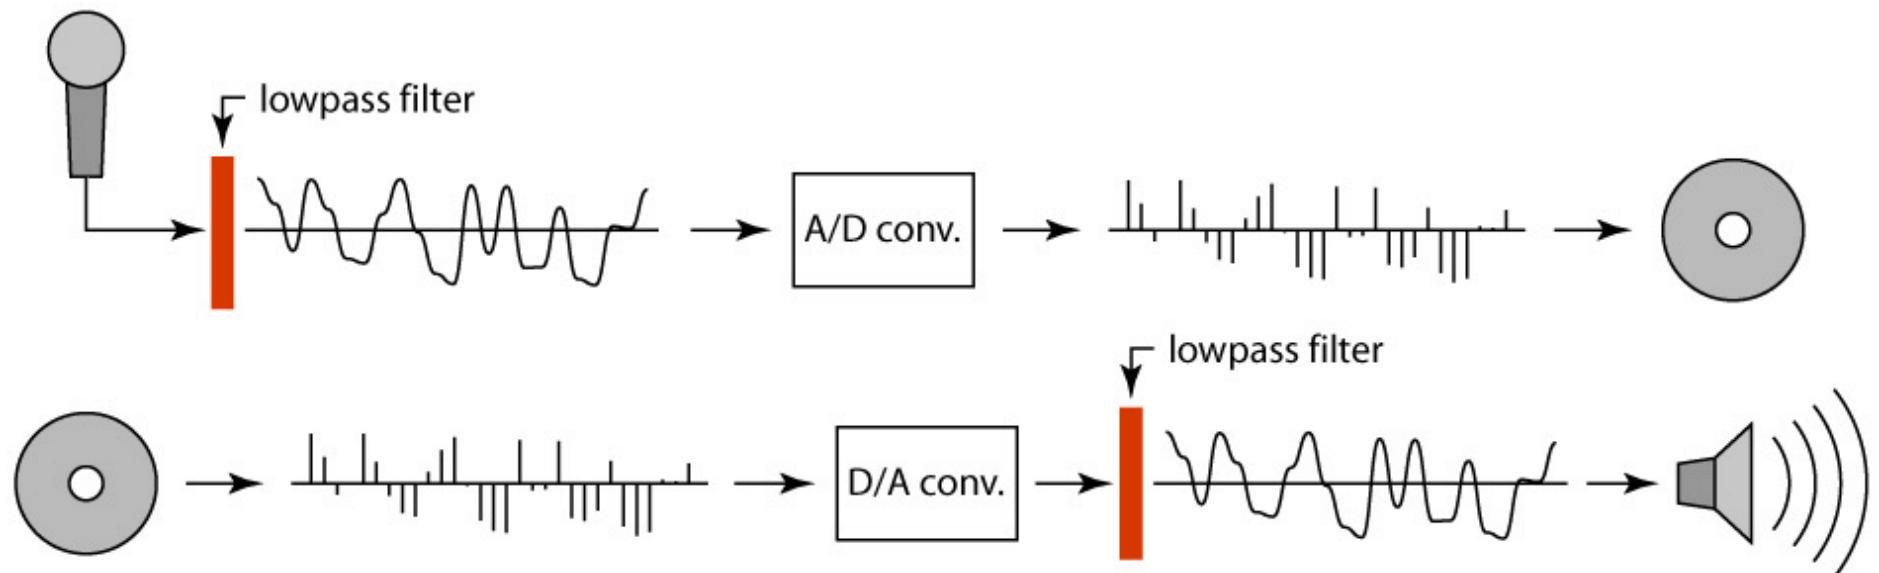
Make the signal less “wiggly”

- Get rid of some high frequencies
- Will loose information
- But it's better than aliasing

Preventing aliasing

Introduce lowpass filters:

- remove high frequencies leaving only safe, low frequencies
- choose lowest frequency in reconstruction (disambiguate)



Linear filtering: a key idea

Transformations on signals; e.g.:

- bass/treble controls on stereo
- blurring/sharpening operations in image editing
- smoothing/noise reduction in tracking

Key properties

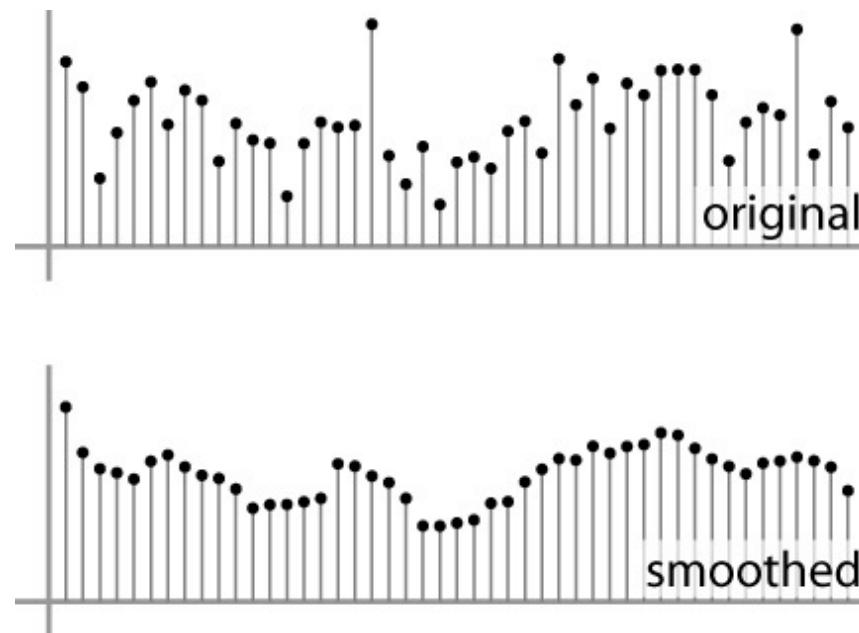
- linearity: $\text{filter}(f + g) = \text{filter}(f) + \text{filter}(g)$
- shift invariance: behavior invariant to shifting the input
 - delaying an audio signal
 - sliding an image around

Can be modeled mathematically by *convolution*

Moving Average

basic idea: define a new function by averaging over a sliding window

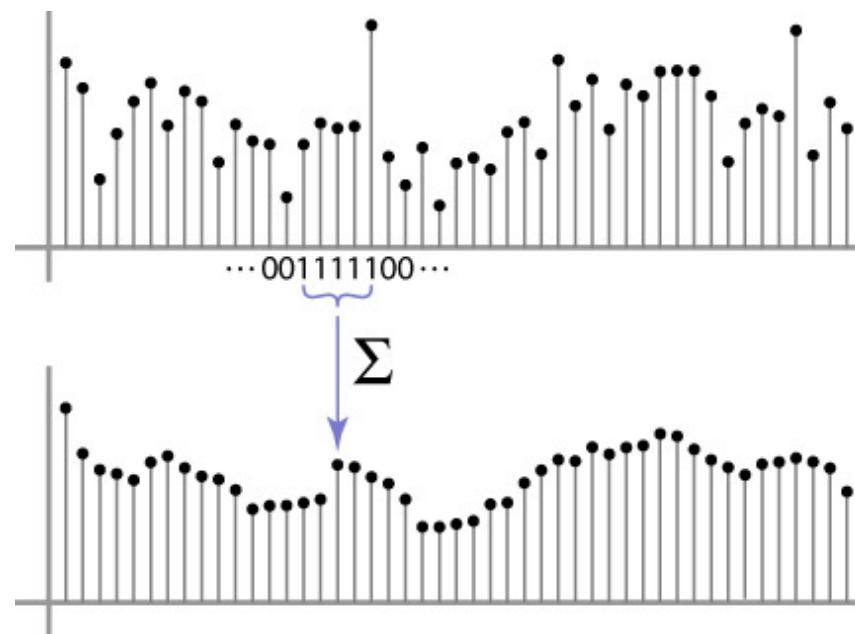
a simple example to start off: smoothing



Moving Average

Can add weights to our moving average

Weights $[..., 0, 1, 1, 1, 1, 1, 0, ...] / 5$



Cross-correlation

Let F be the image, H be the kernel (of size $2k+1 \times 2k+1$), and G be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

- Can think of as a “dot product” between local neighborhood and kernel for each pixel

In 2D: box filter

$$h[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Image filtering

$$h[\cdot, \cdot]$$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

$$g[m, n] = \sum_{k, l} h[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

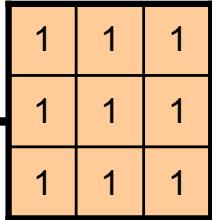
$$g[\cdot, \cdot]$$

0	10									

$$g[m, n] = \sum_{k,l} h[k, l] f[m+k, n+l]$$

Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$


$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20						

Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30					

Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

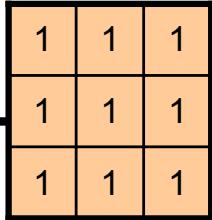
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30				

Credit: S. Seitz

Image filtering

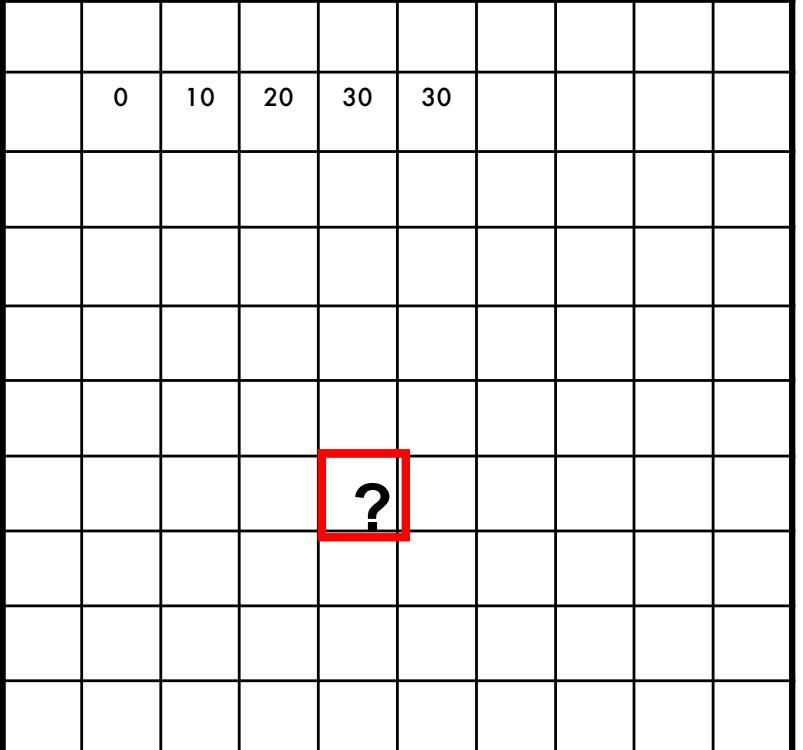
$$h[\cdot, \cdot] \frac{1}{9}$$


$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30				



Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30				

Credit: S. Seitz

Image filtering

$$h[\cdot, \cdot] \quad \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$g[m, n] = \sum_{k, l} h[k, l] f[m + k, n + l]$$

Credit: S. Seitz

Box Filter

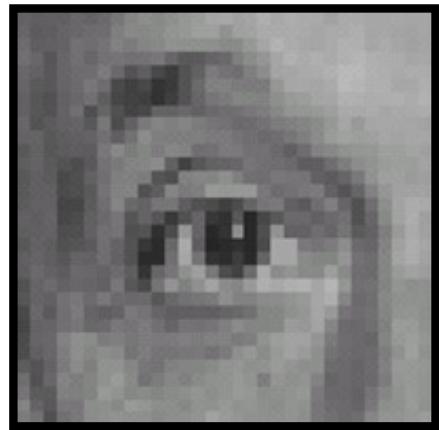
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$h[\cdot, \cdot]$$

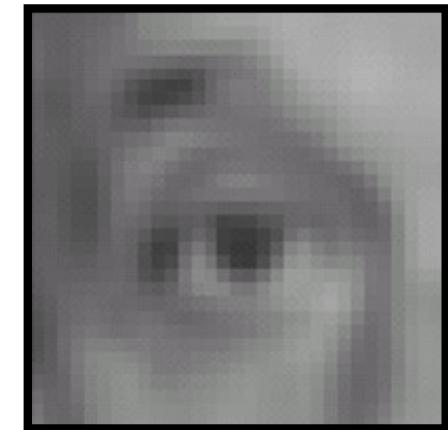
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Linear filters: examples



Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Blur (with a mean filter)

Practice with linear filters



0	0	0
0	1	0
0	0	0

?

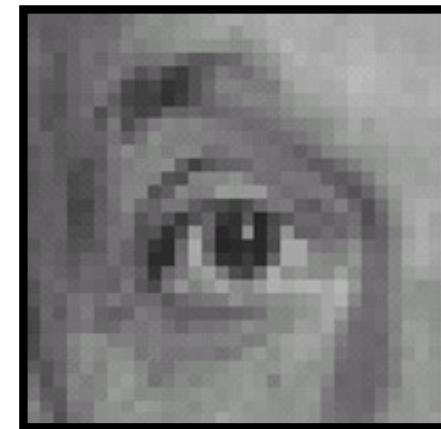
Original

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



0	0	0
0	0	1
0	0	0

?

Original

Practice with linear filters



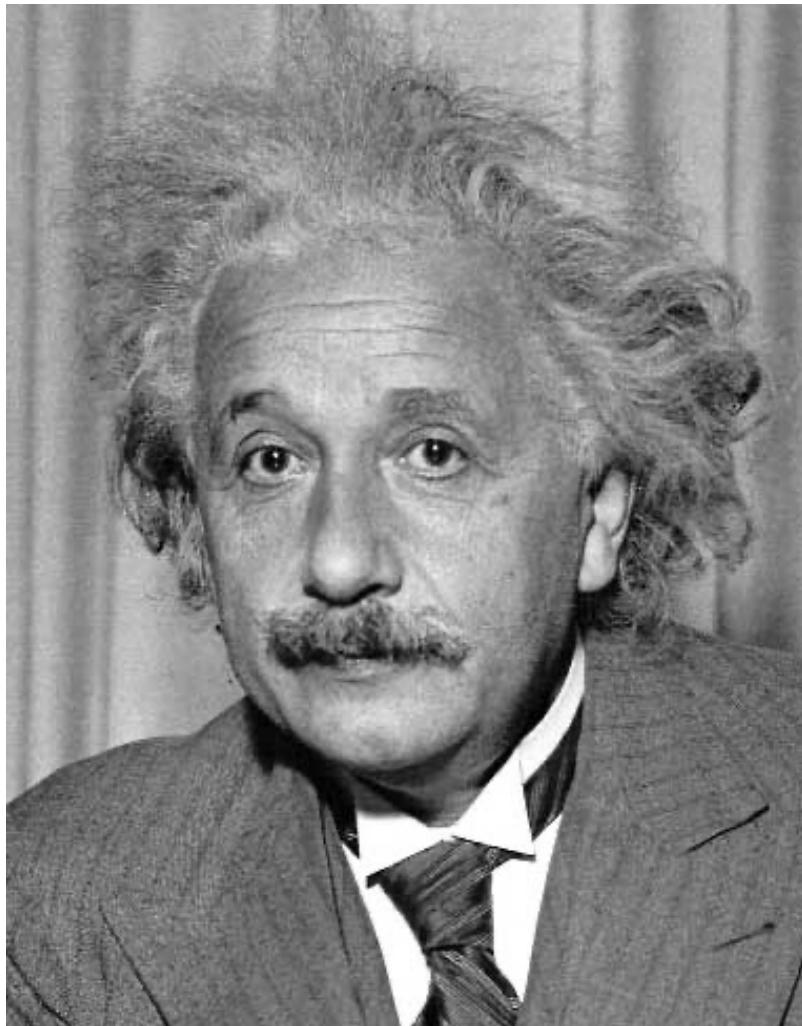
Original

0	0	0
0	0	1
0	0	0



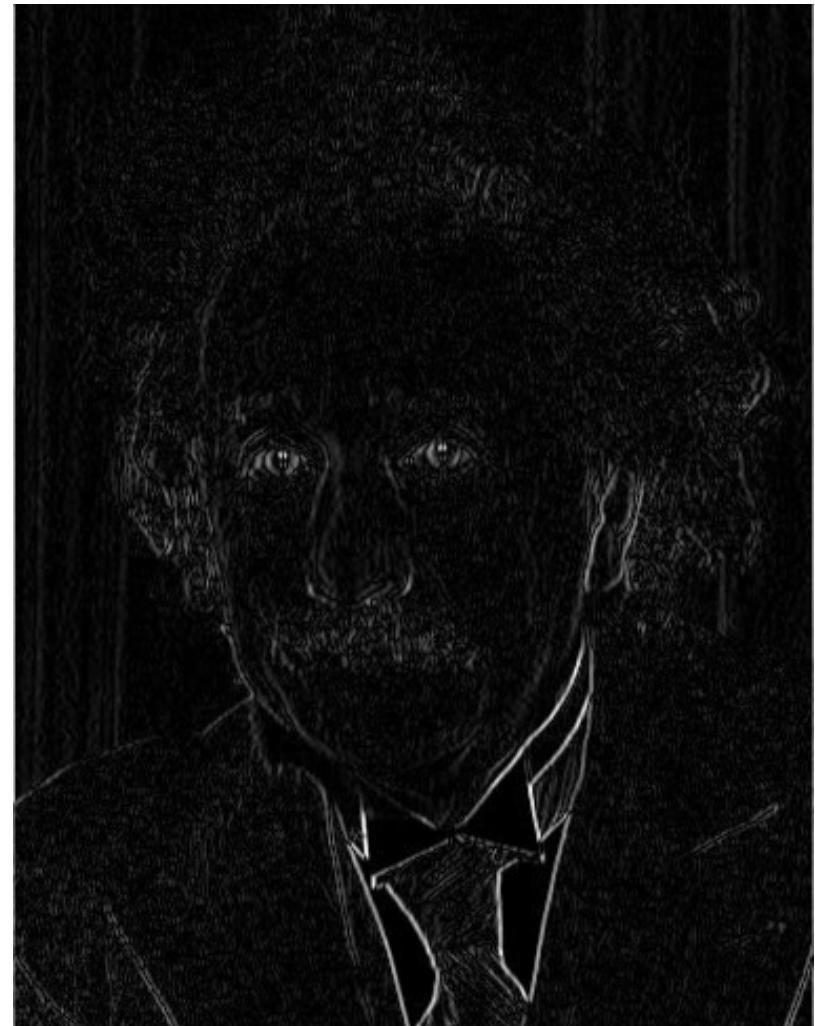
Shifted left
By 1 pixel

Other filters



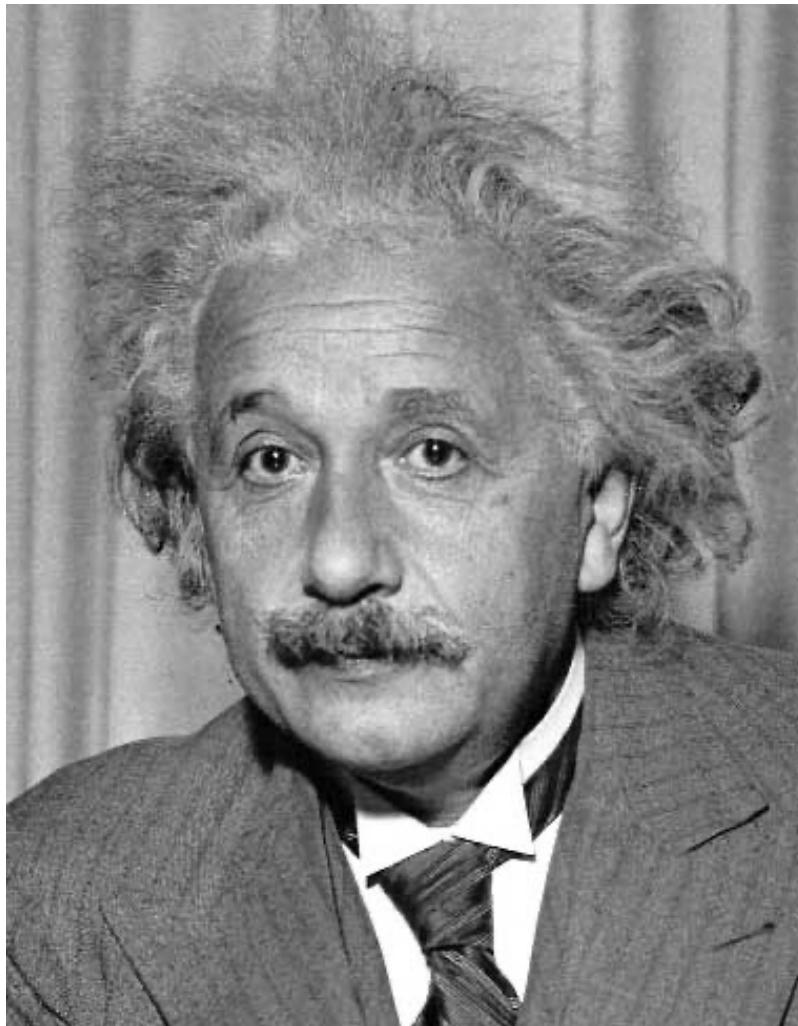
1	0	-1
2	0	-2
1	0	-1

Sobel



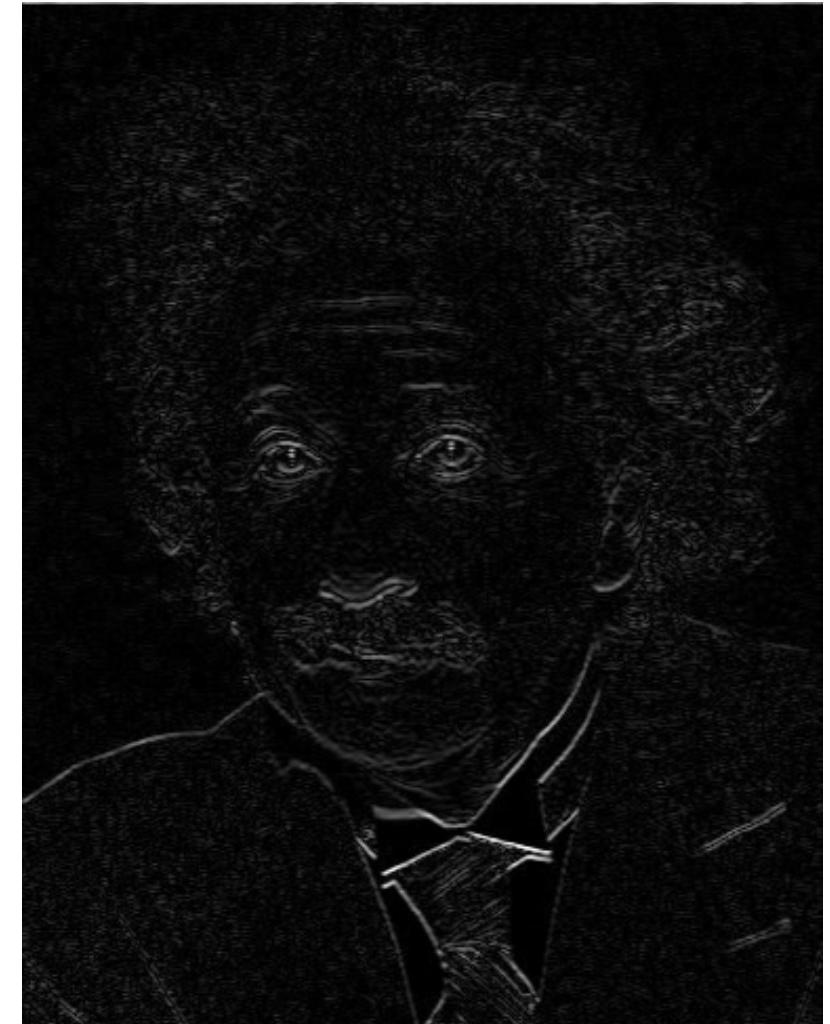
Vertical Edge
(absolute value)

Other filters



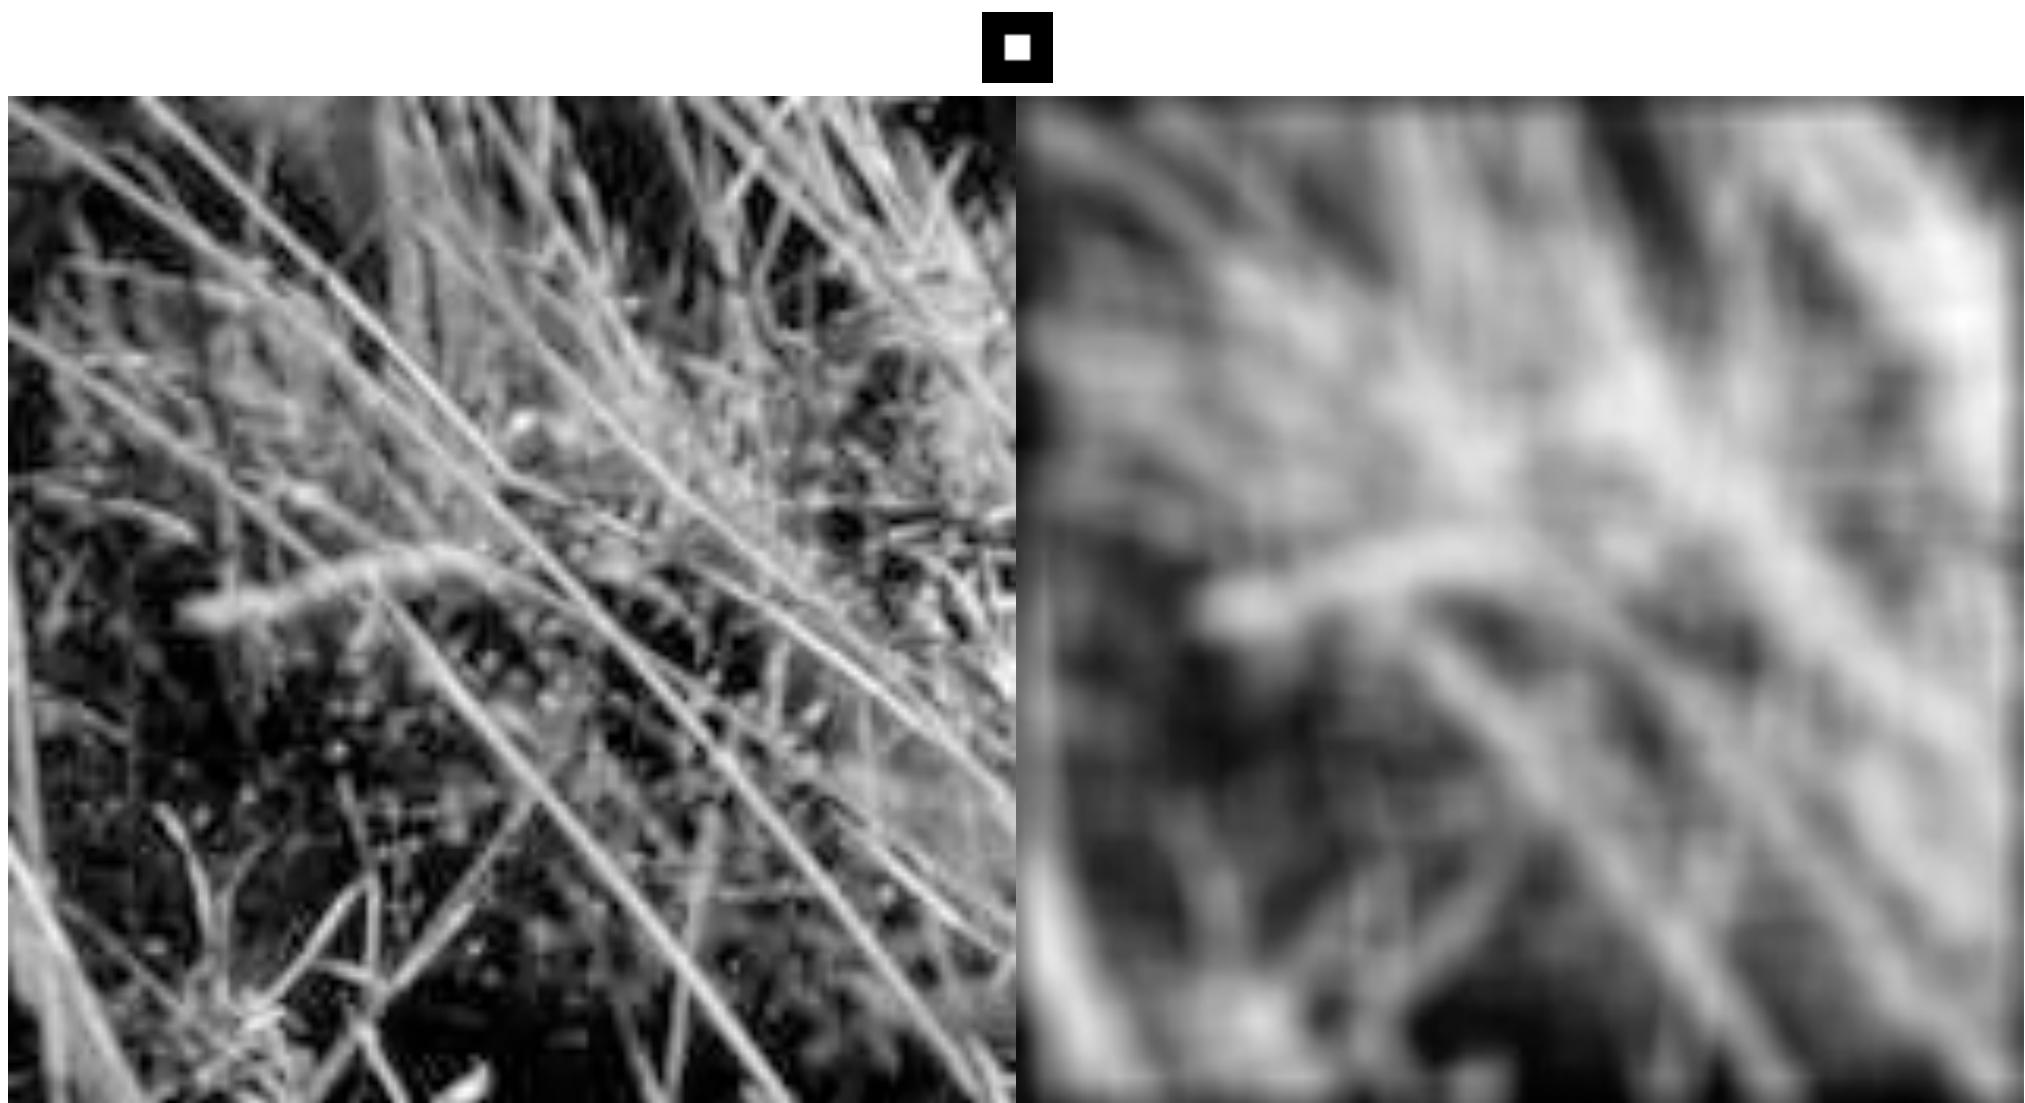
1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

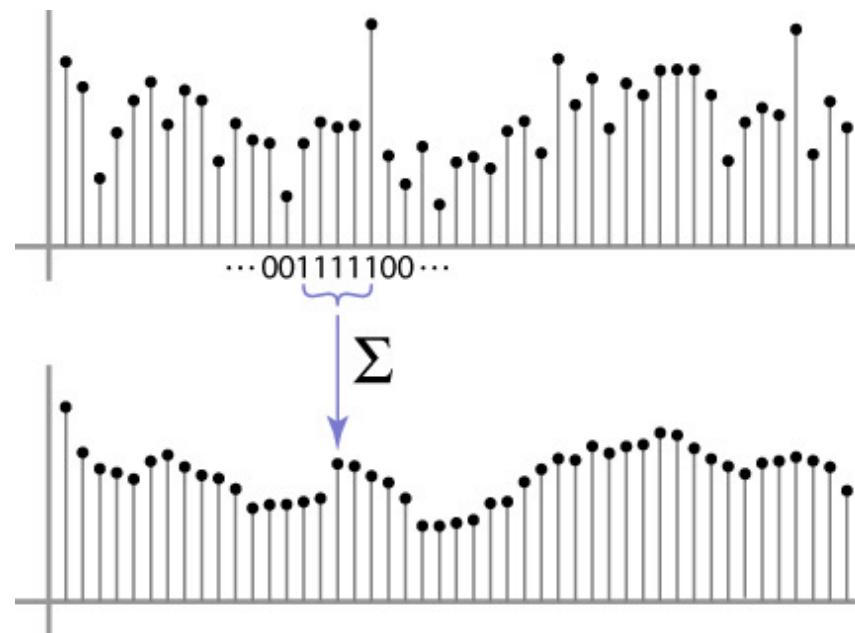
Back to the box filter



Moving Average

Can add weights to our moving average

Weights $[..., 0, 1, 1, 1, 1, 1, 0, ...] / 5$



Weighted Moving Average

bell curve (gaussian-like) weights $[..., 1, 4, 6, 4, 1, ...]$

