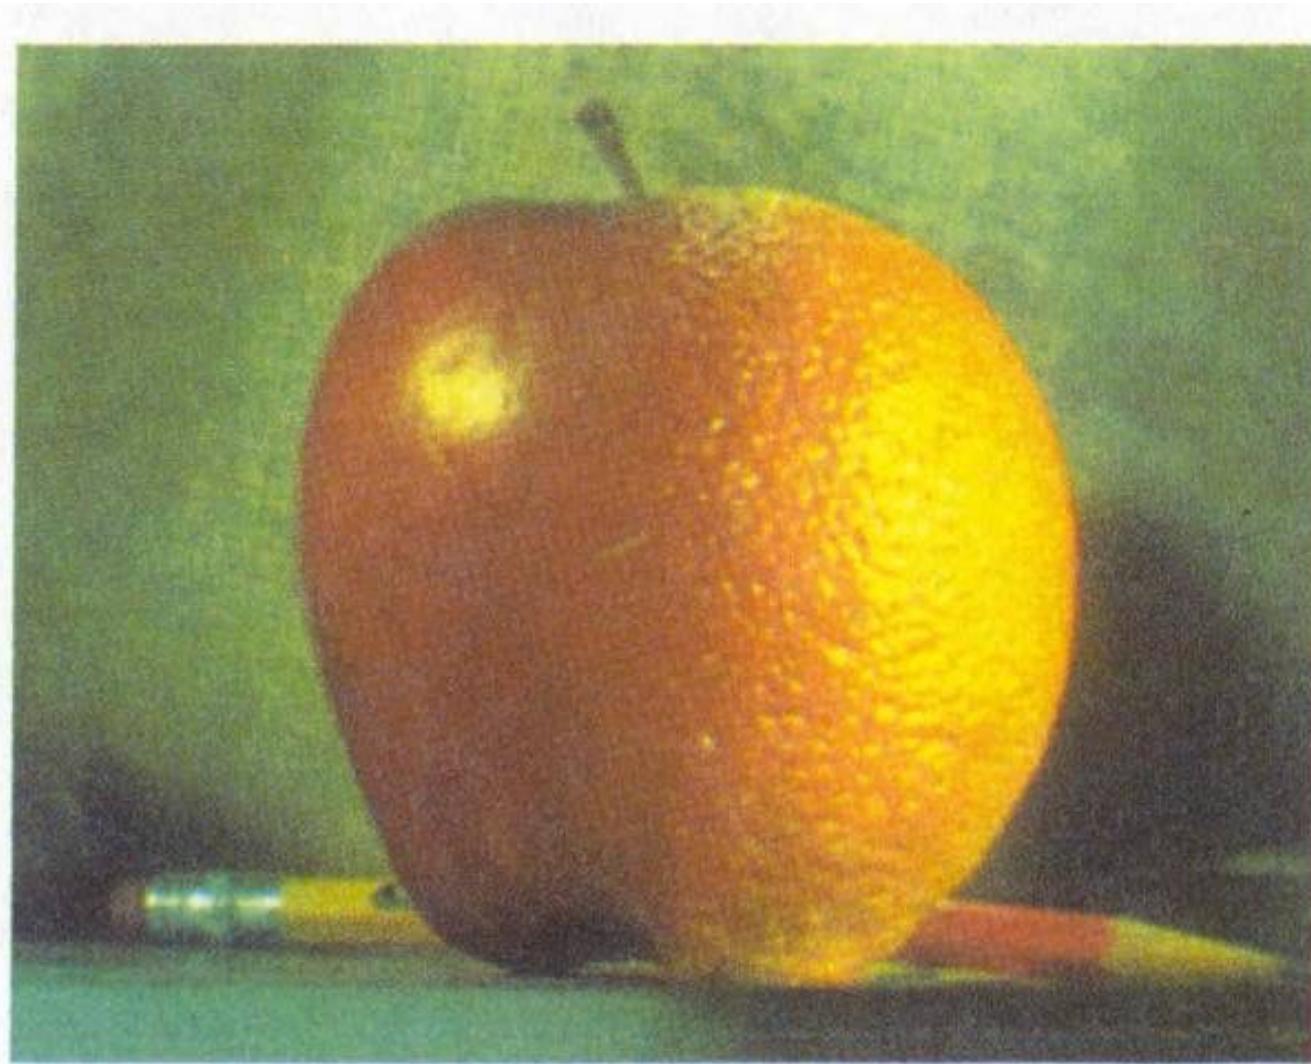


# Pyramid Blending, Templates, NL Filters

---



CS194: Intro to Comp. Vision and Comp. Photo  
Angjoo Kanazawa & Alexei Efros, UC Berkeley, Fall 2022

## 5 min recap to watch

---

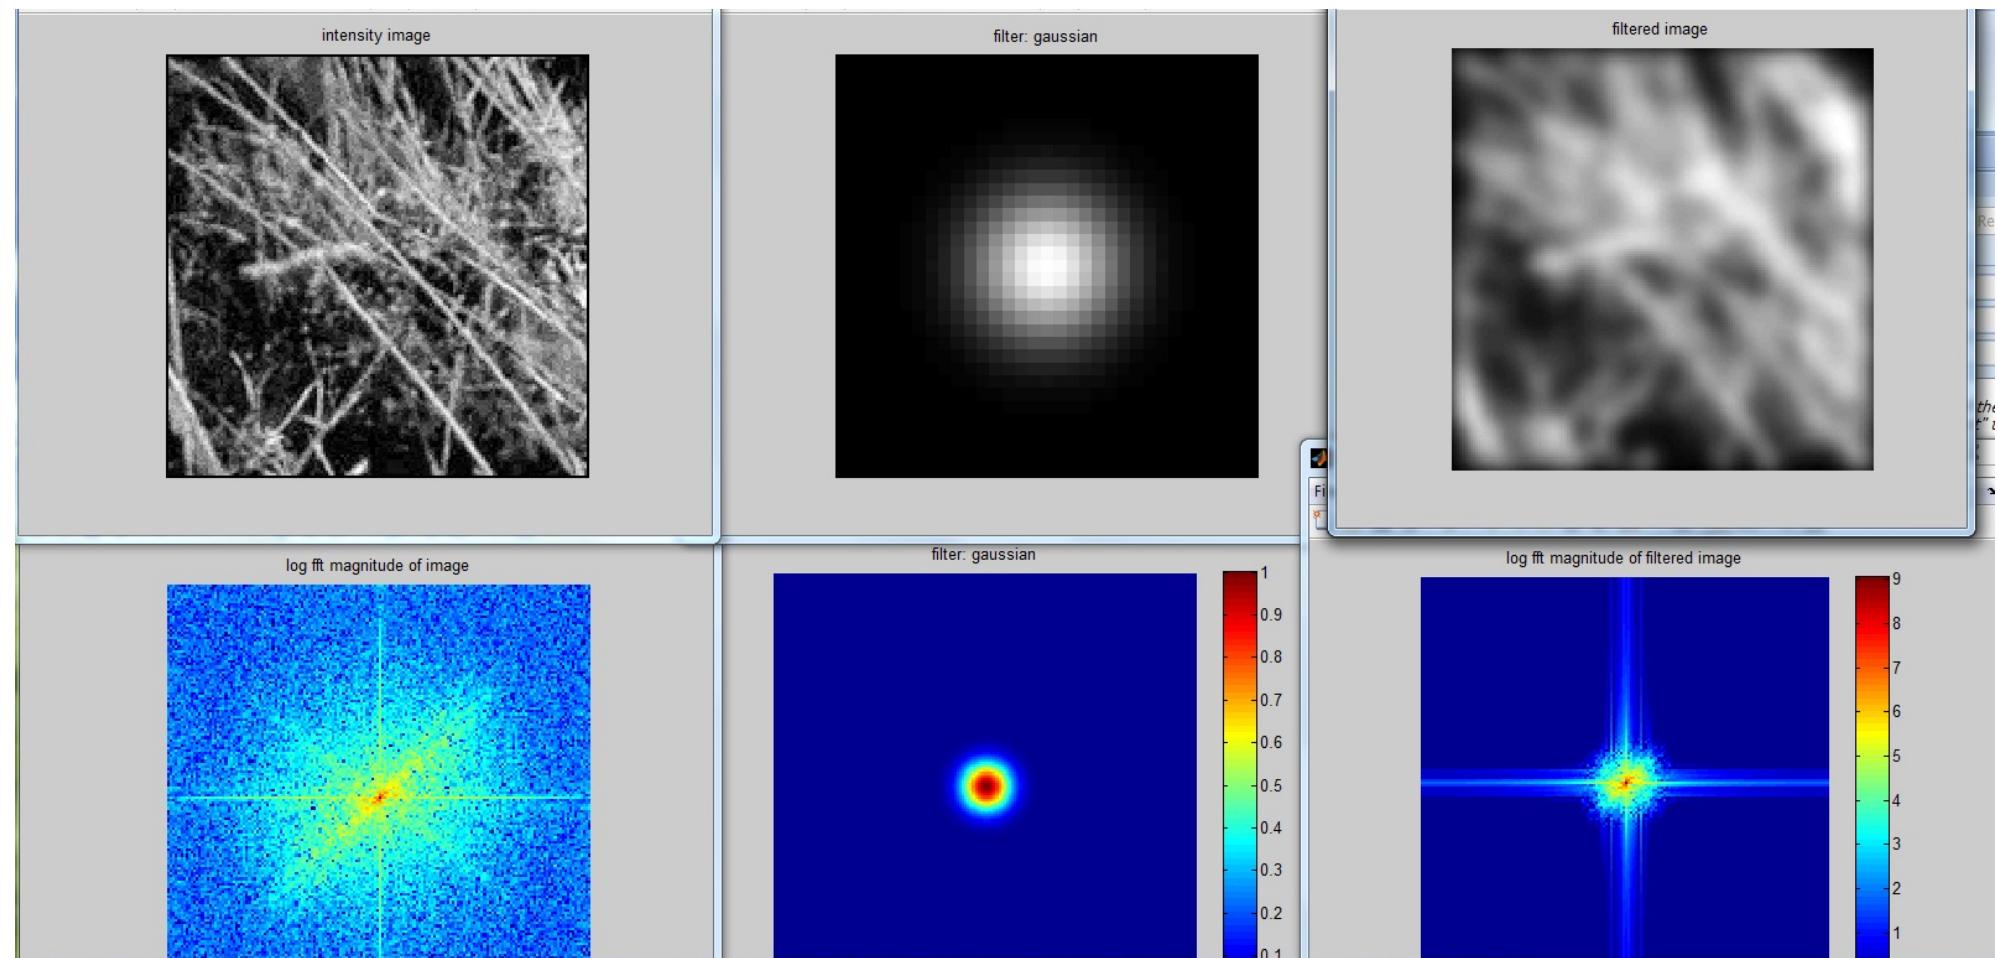
Fourier Transform in 5 minutes: The Case of  
the Splotched Van Gogh, Part 3

<https://www.youtube.com/watch?v=JciZYrh36LY>

(on the class website)

# Gaussian is not perfect

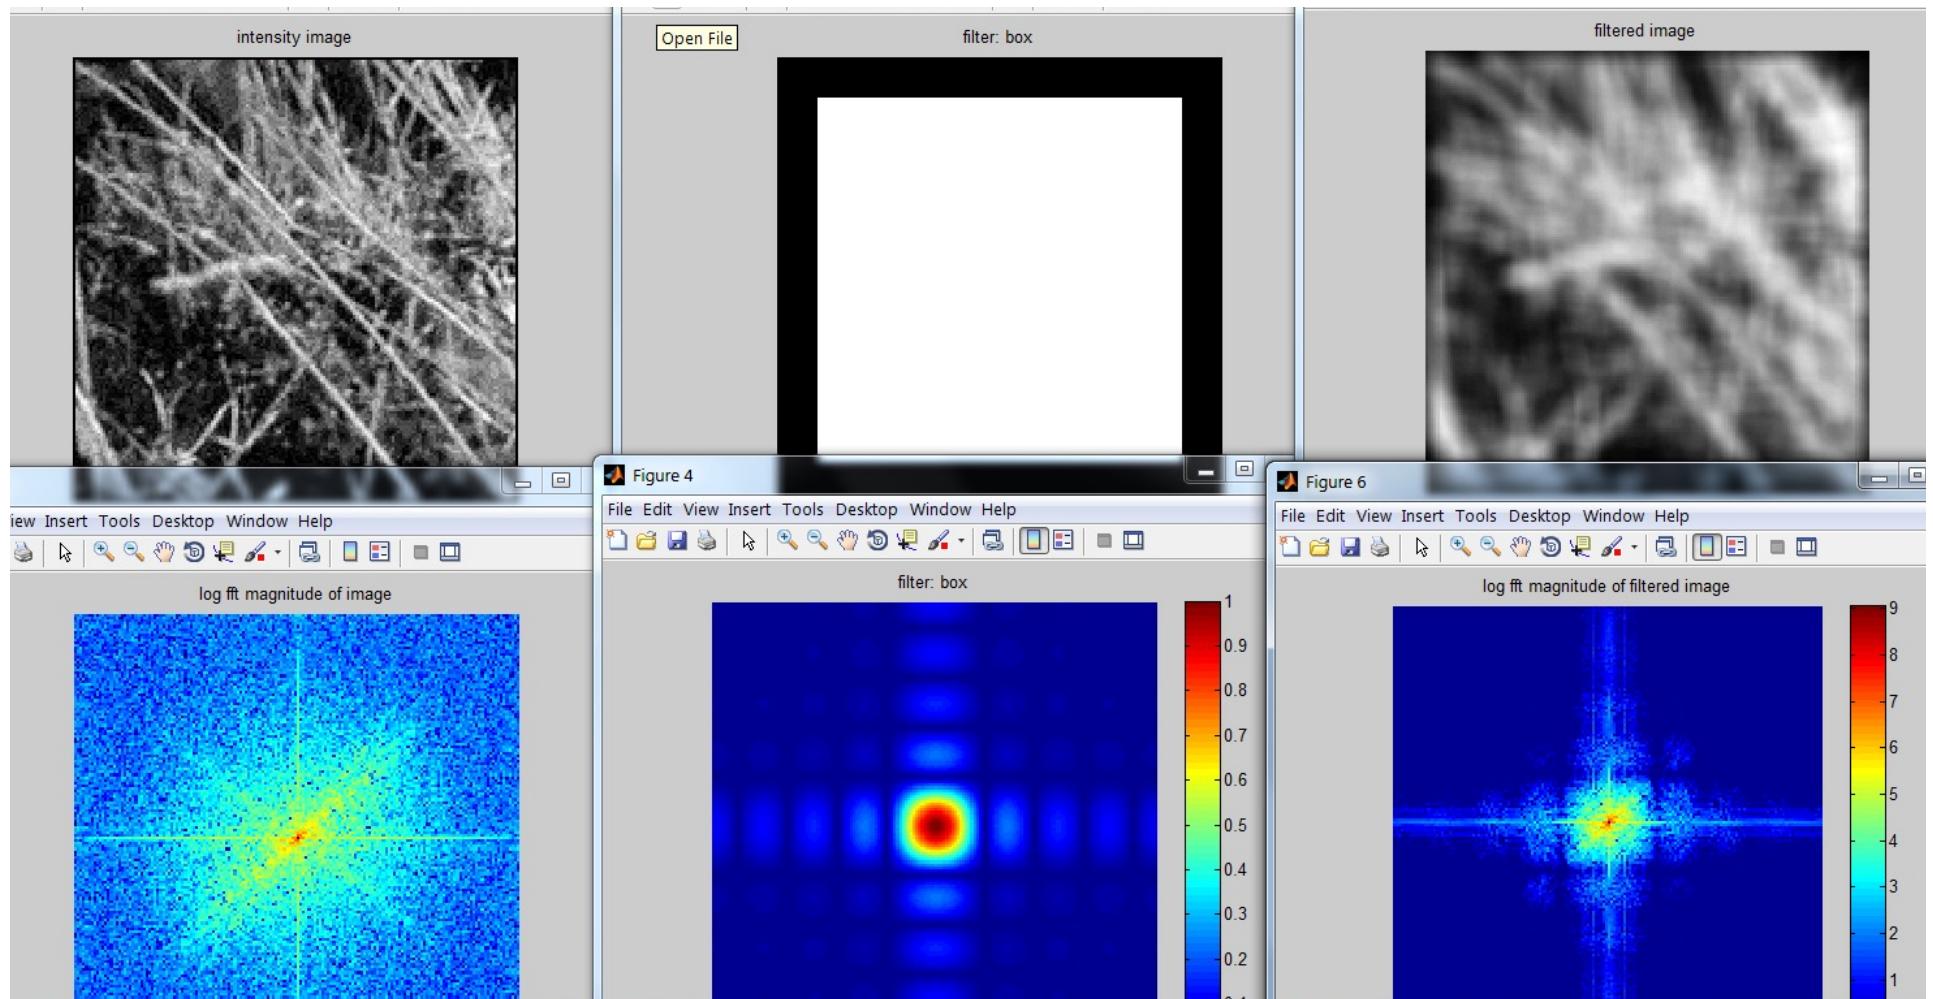
## Gaussian



# But better than box filter!

---

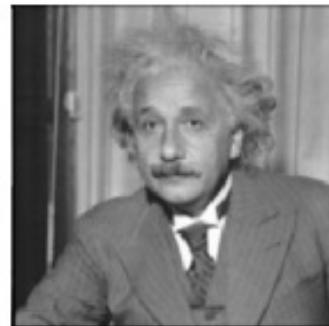
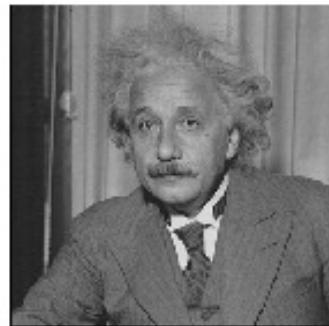
## Box Filter



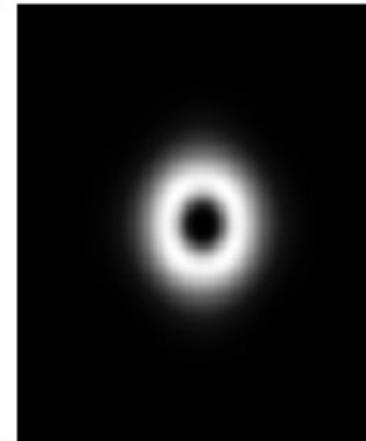
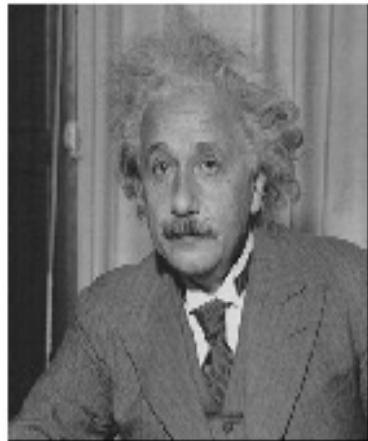
# Low-pass, Band-pass, High-pass filters

---

low-pass:

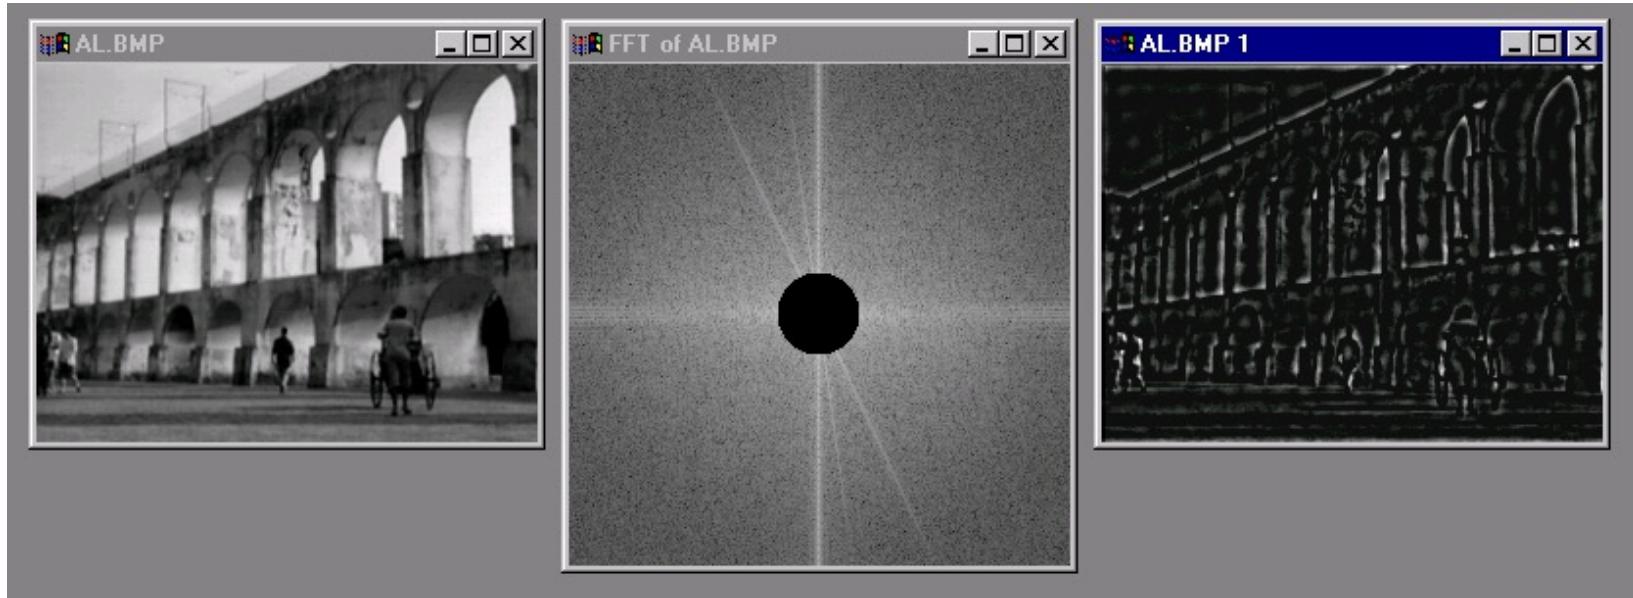


High-pass / band-pass:



# Edges in images

---



# Low Pass vs. High Pass filtering

---

Image

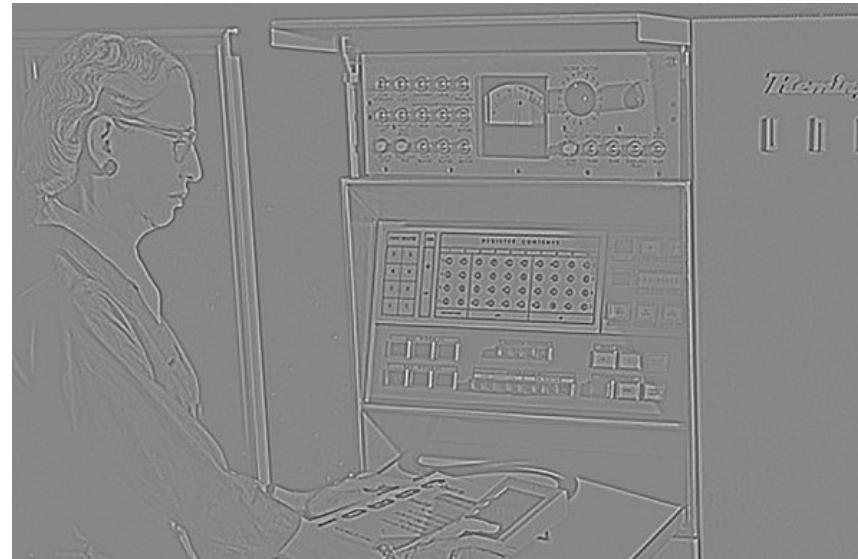


Smoothed



Details

=



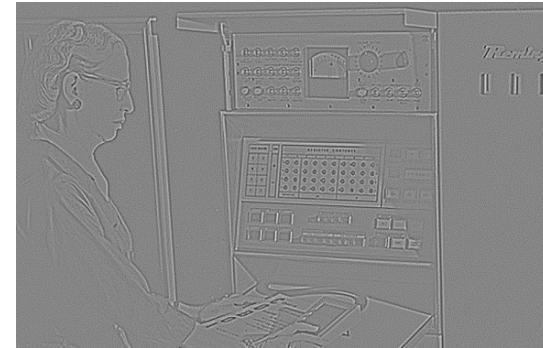
# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=1$

=



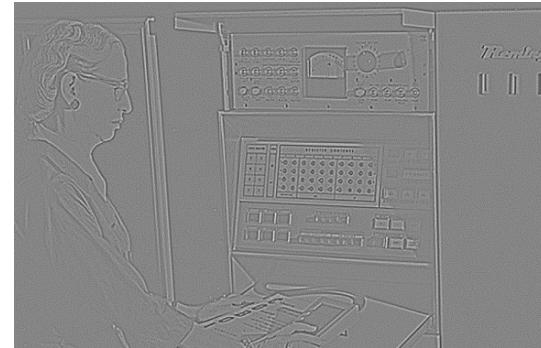
# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=0$

=



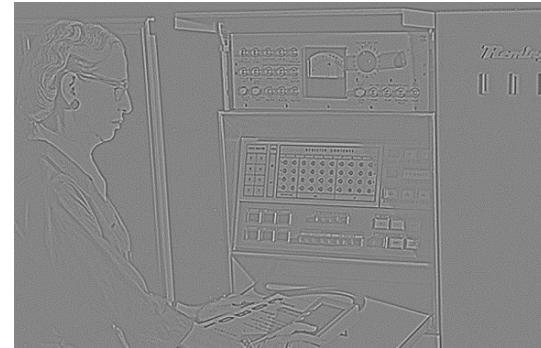
# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=2$

=



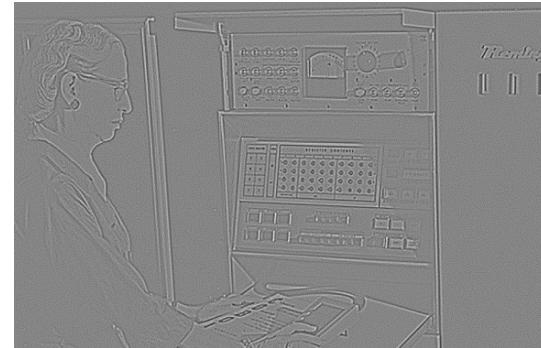
# Filtering – Sharpening

---

Image



Details



$+ \alpha$

“Sharpened”  $\alpha=0$

=



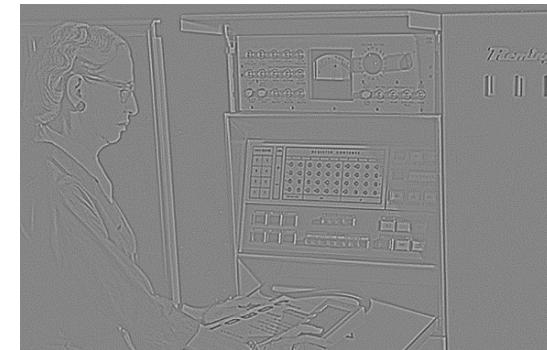
# Filtering – Extreme Sharpening

---

Image



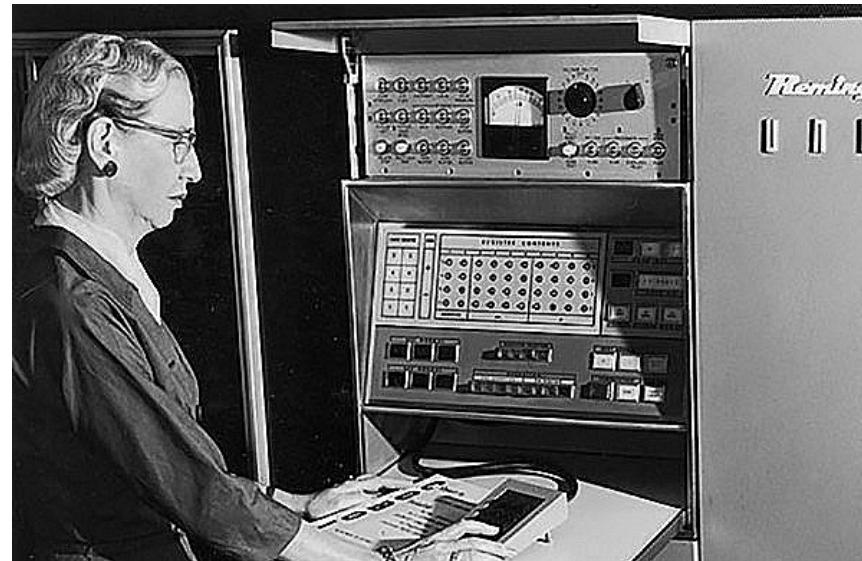
Details



$+ \alpha$

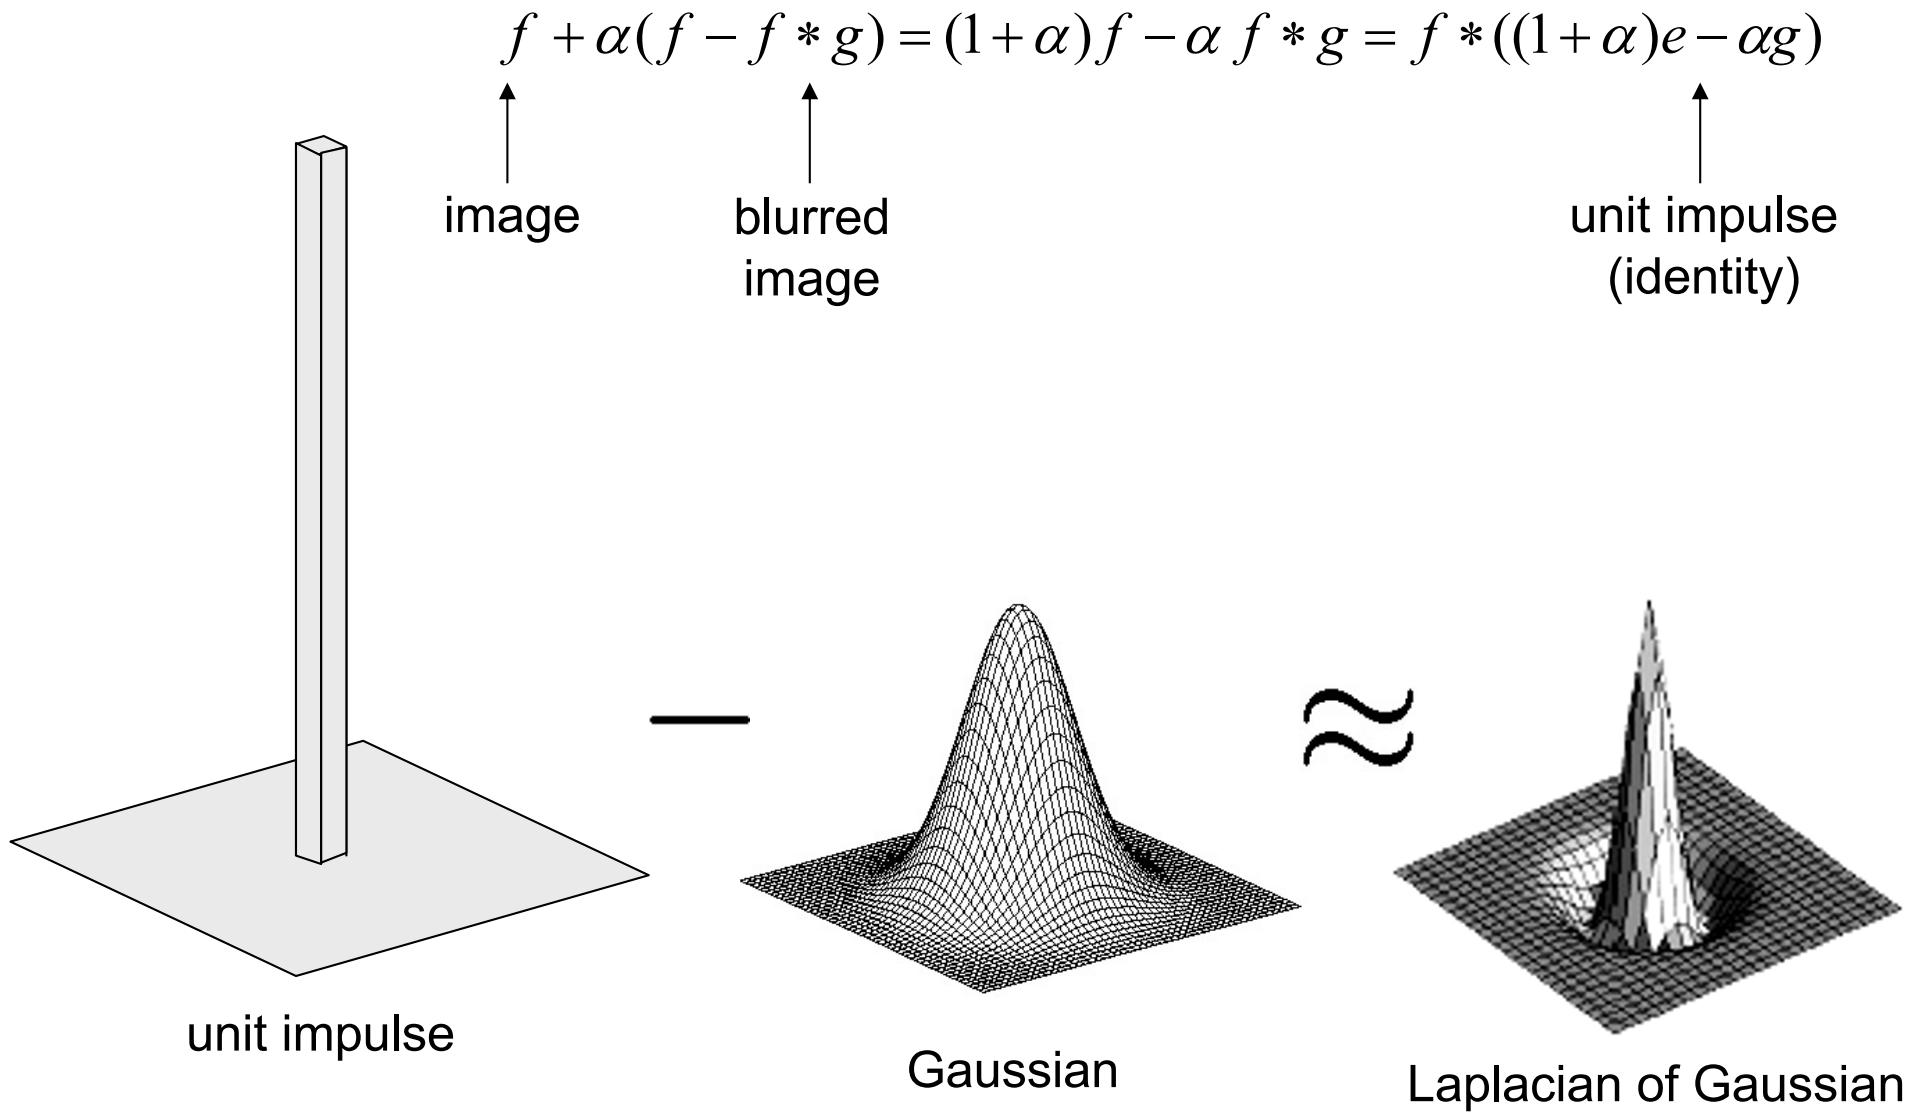
“Sharpened”  $\alpha=10$

=

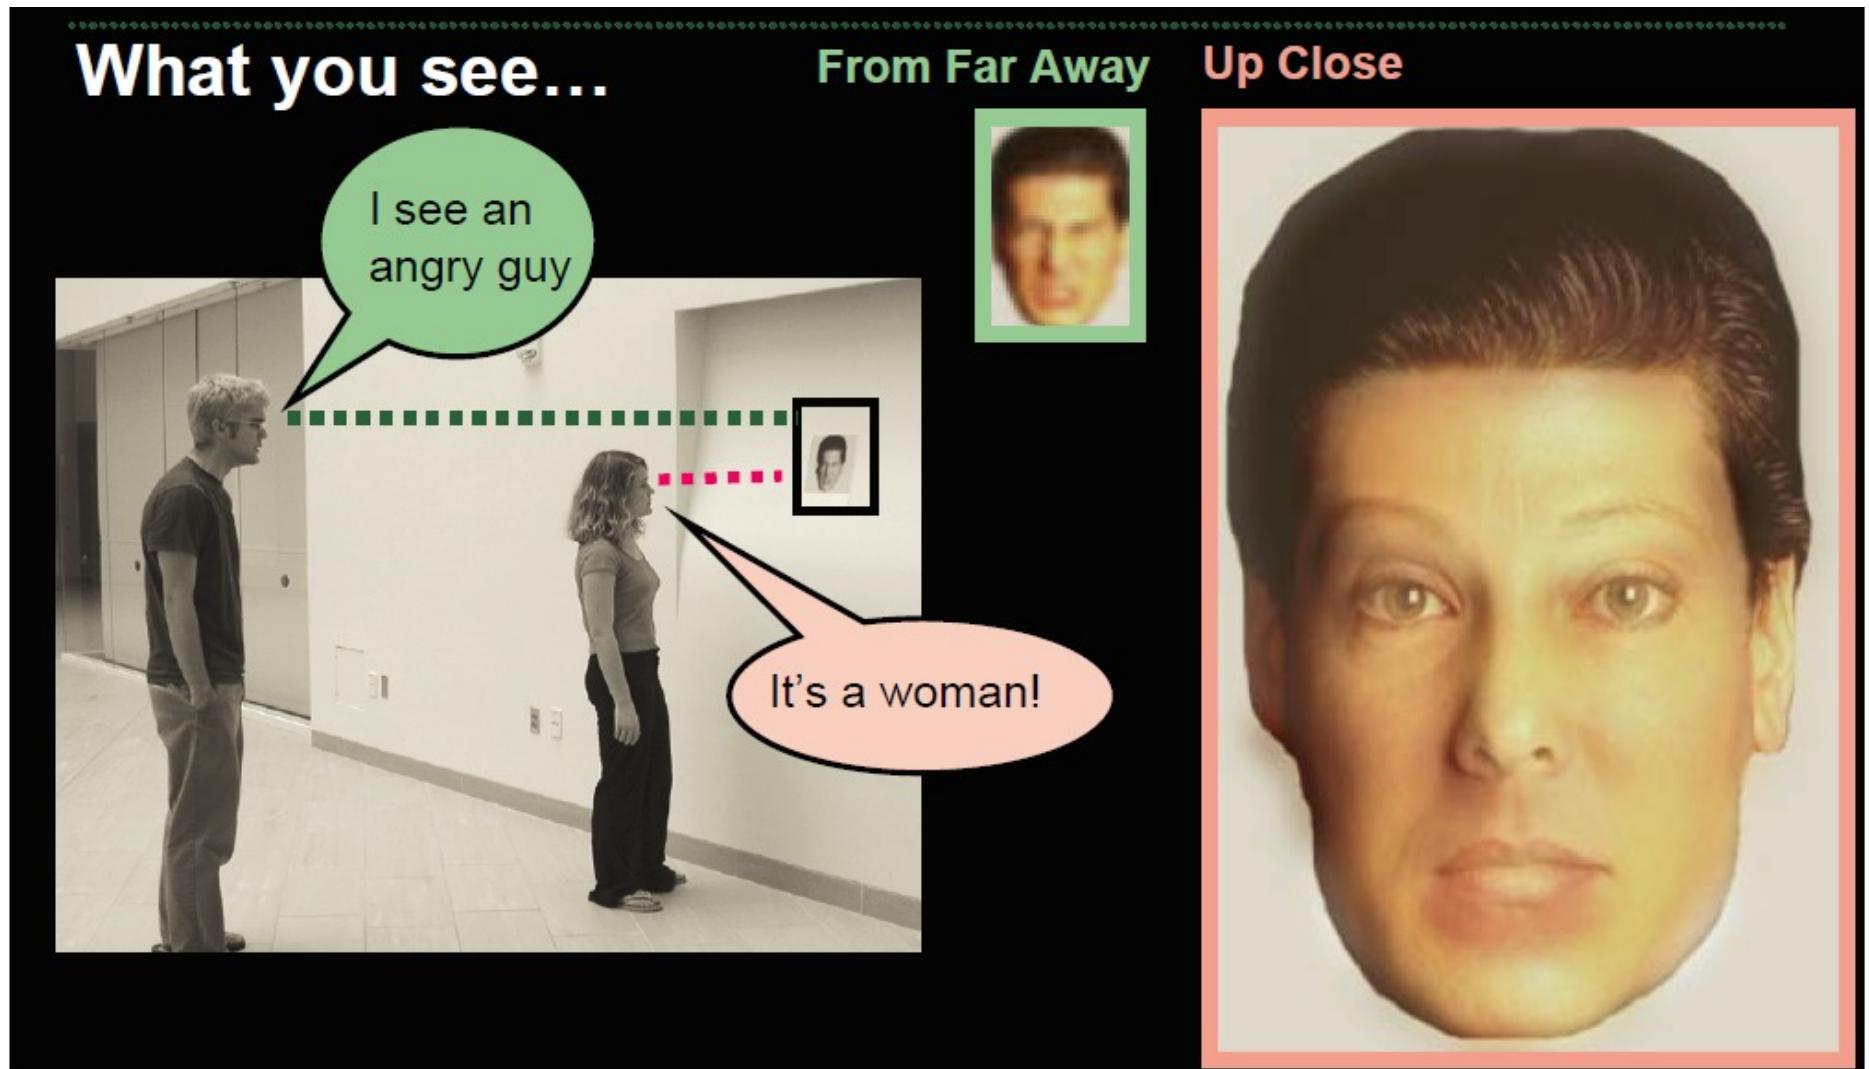


# Unsharp mask filter

---



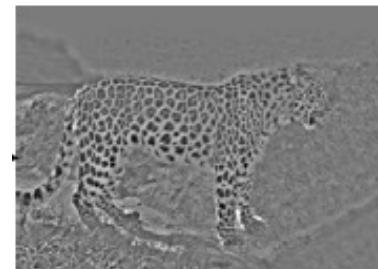
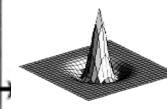
# application: Hybrid Images



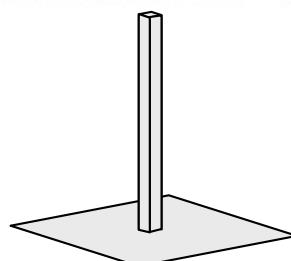
# Application: Hybrid Images

Gaussian Filter

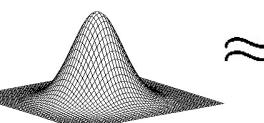
A. Oliva, A. Torralba, P.G. Schyns,  
[“Hybrid Images,” SIGGRAPH 2006](#)



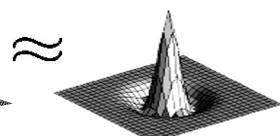
Laplacian Filter



unit impulse



Gaussian

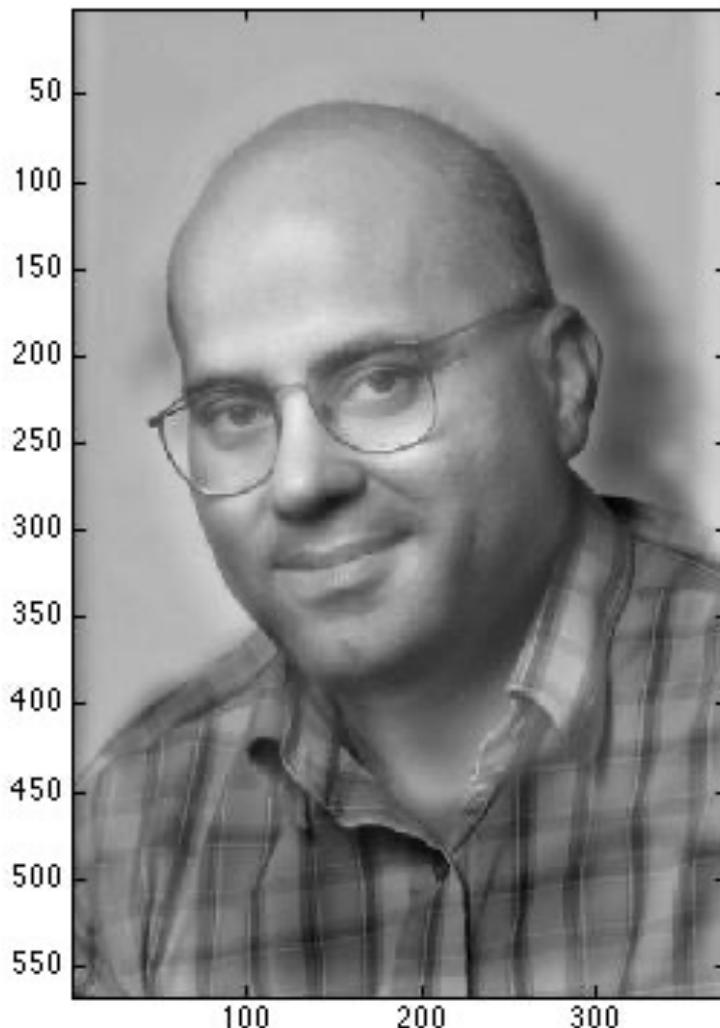


$\approx$  Laplacian of Gaussian

# Yestaryear's homework

---

(CS194-26: Riyaz Faizullabhoy)

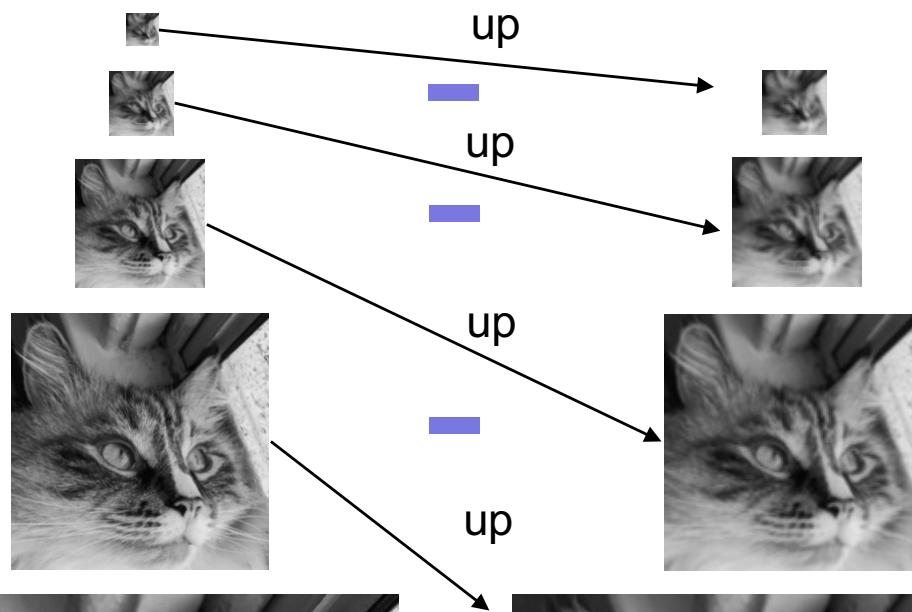


Prof. Jitendros Papadimalik

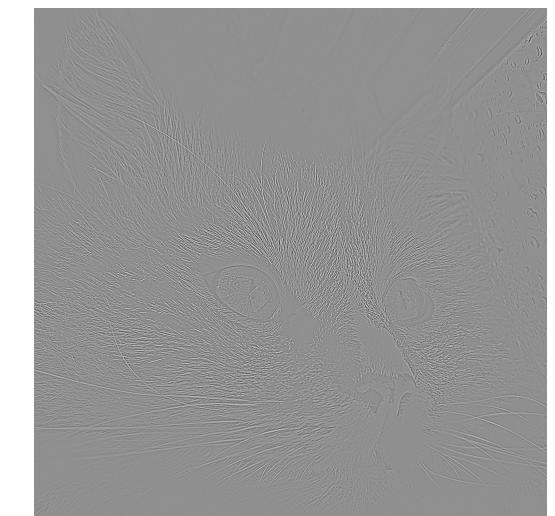
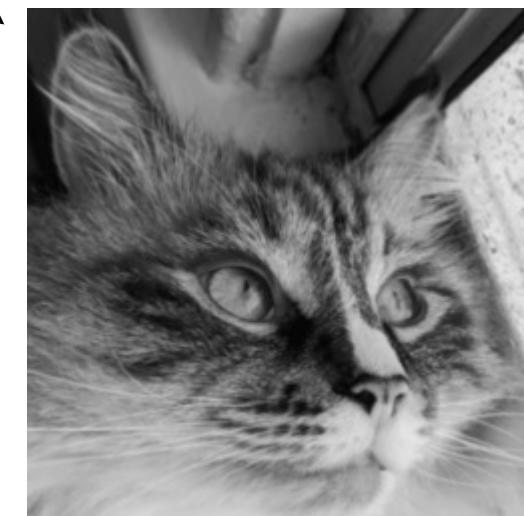
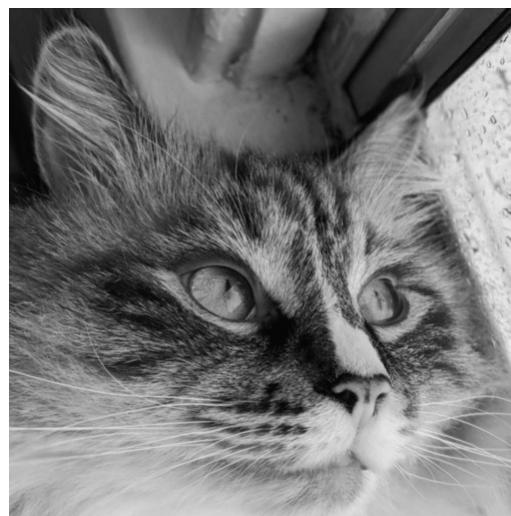
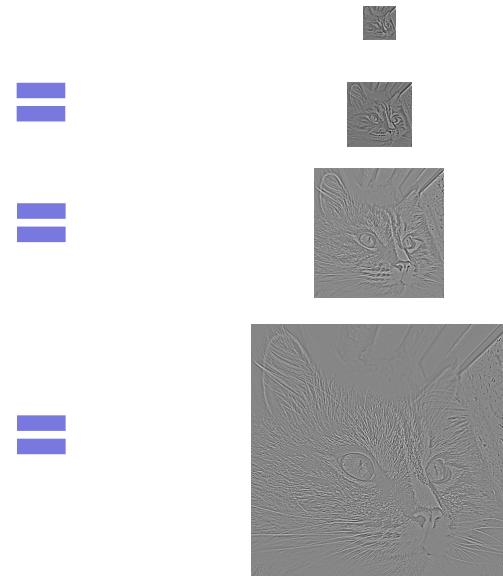
# Band-pass filtering in spatial domain

Gaussian Pyramid  
(low-pass images)

:



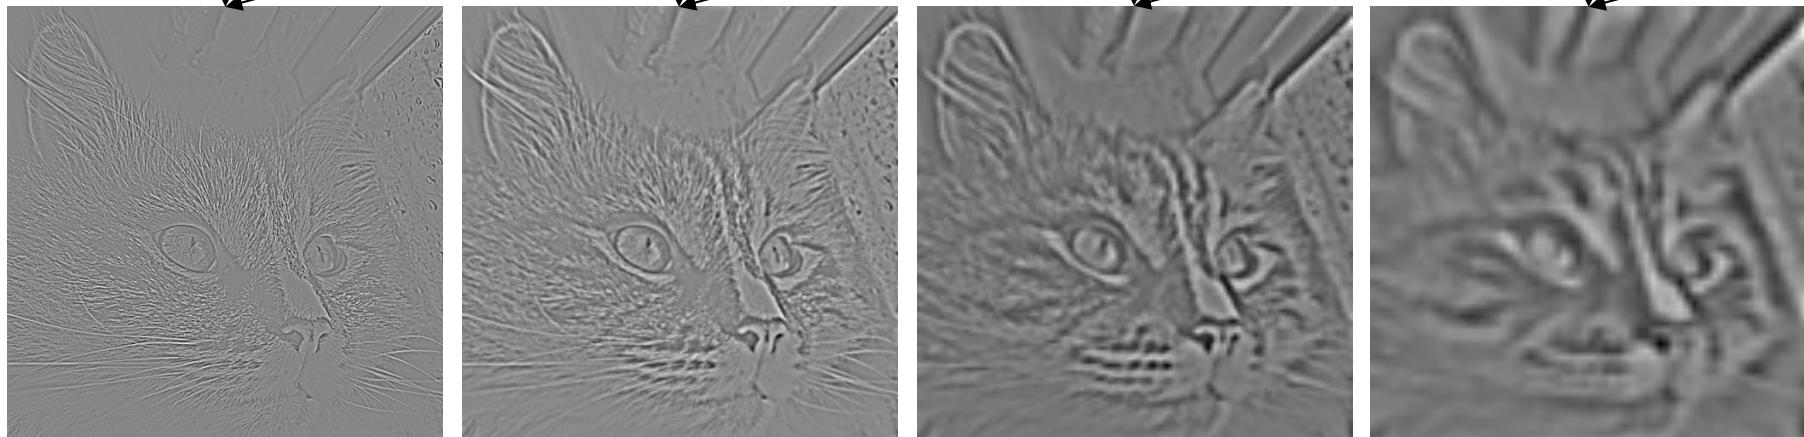
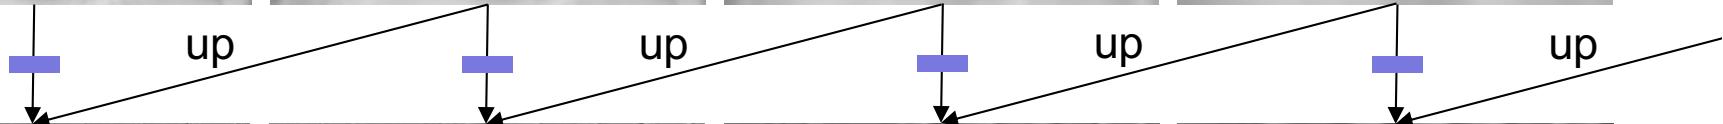
Laplacian Pyramid  
(sub-band images)



# As a stack

---

## Gaussian Pyramid (low-pass images)

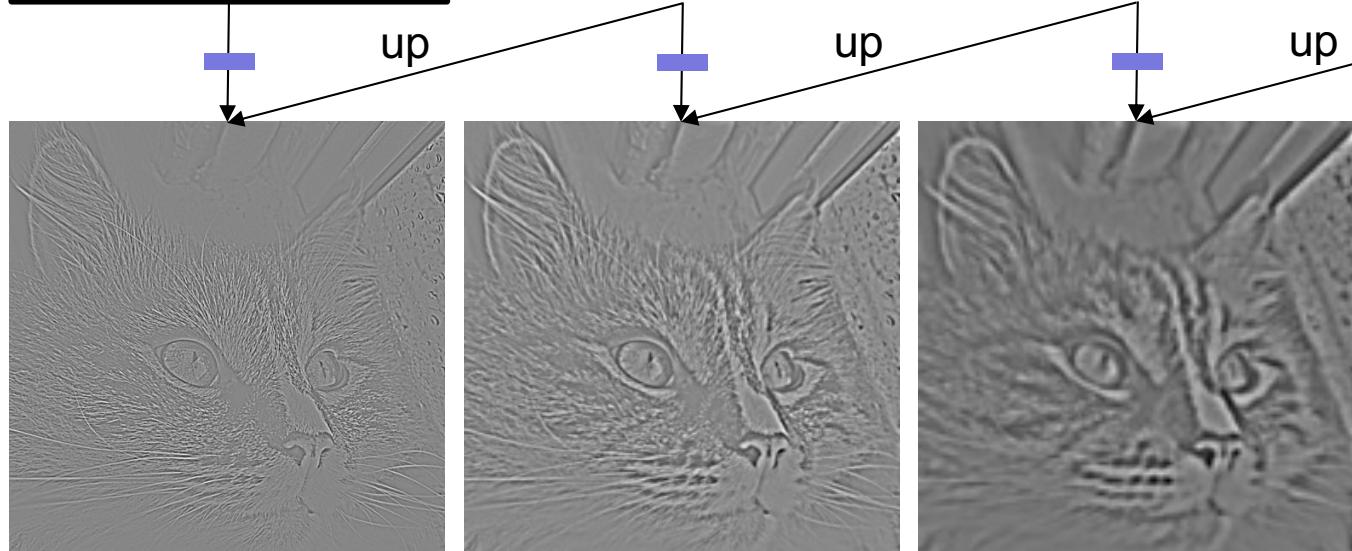
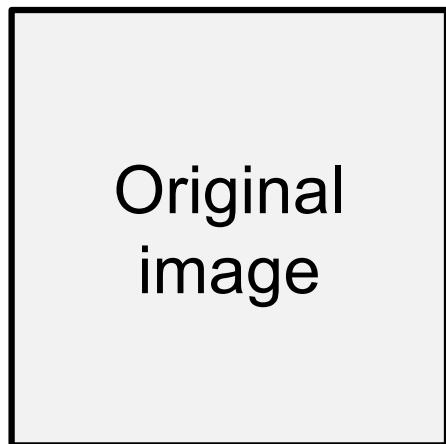


## Laplacian Pyramid (sub-band images)

Created from Gaussian pyramid by subtraction

# Laplacian Pyramid

---



Need this!  
(Lowest Freq)



How can we reconstruct (collapse) this pyramid into the original image?

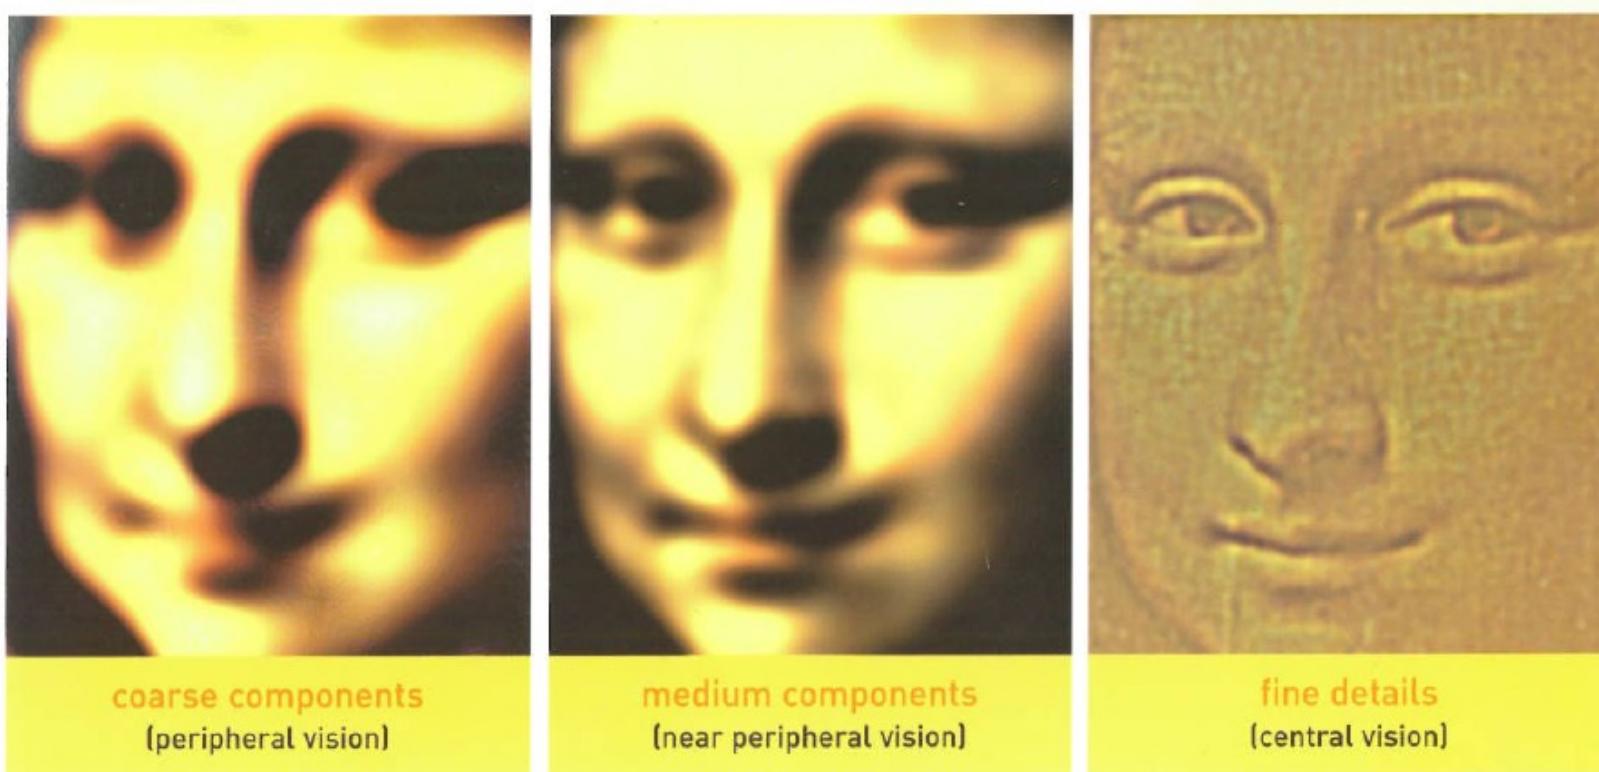
# Da Vinci and The Laplacian Pyramid

---



# Da Vinci and The Laplacian Pyramid

---

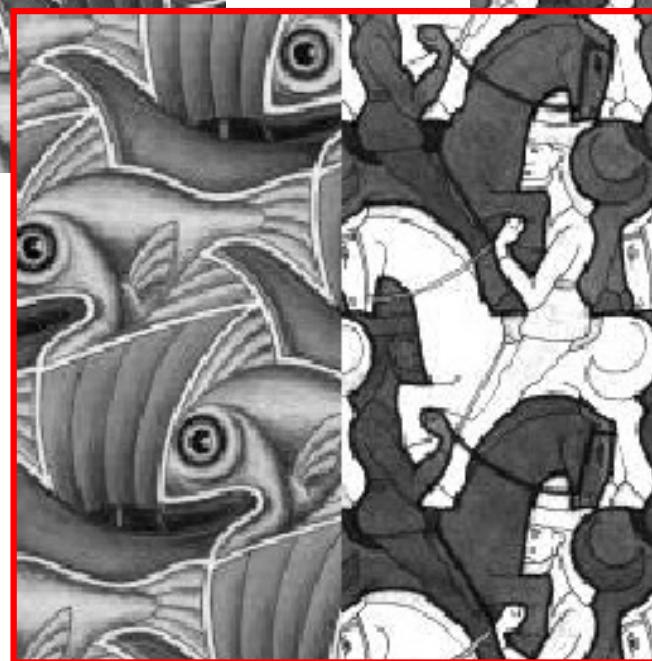
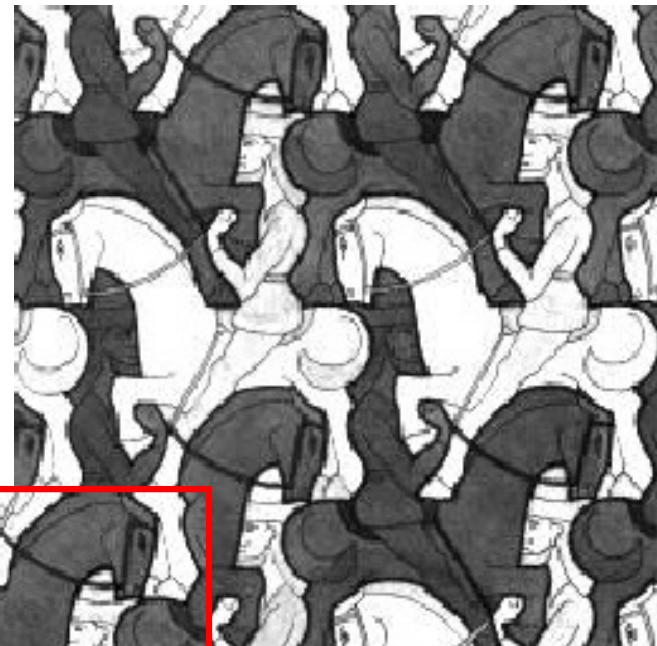
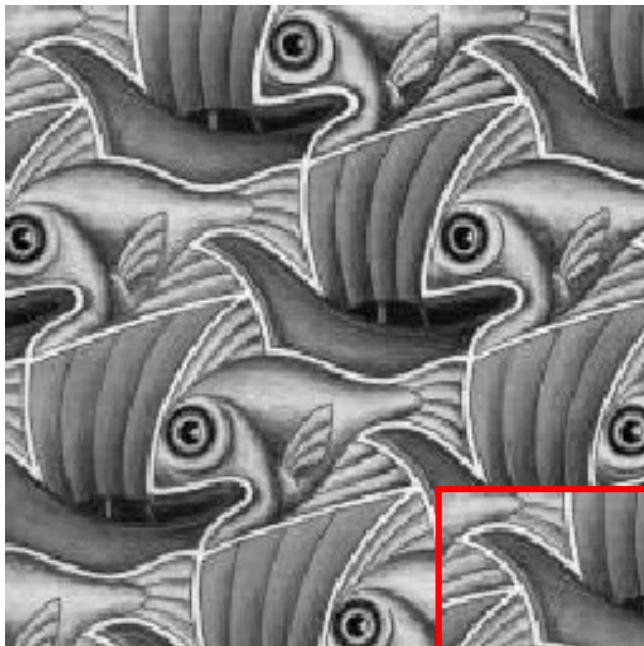


Leonardo playing with peripheral vision

[Livingstone, Vision and Art: The Biology of Seeing](#)

# Blending

---

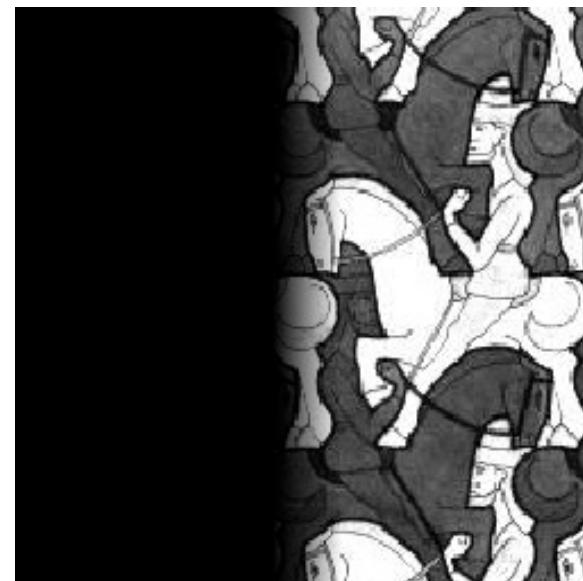


# Alpha Blending / Feathering

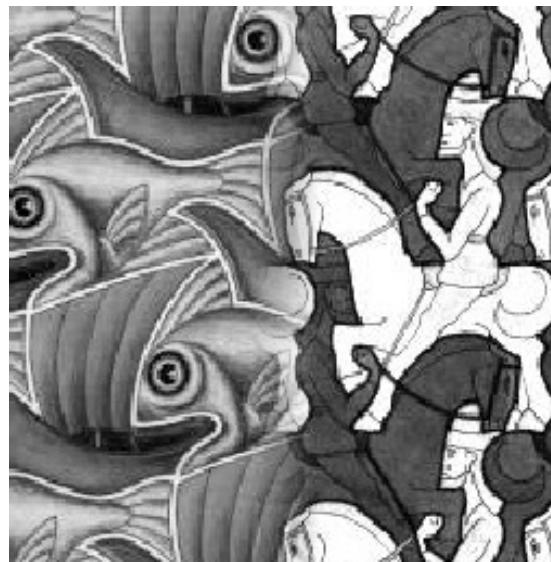
---



+



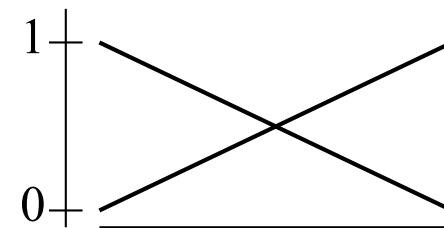
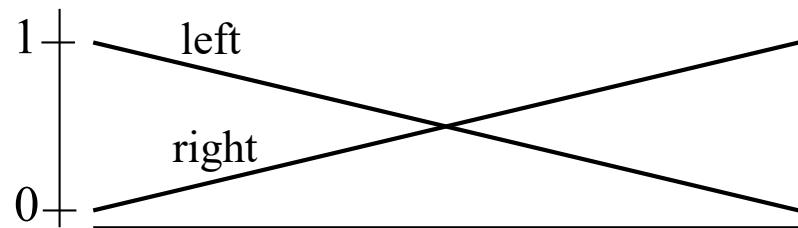
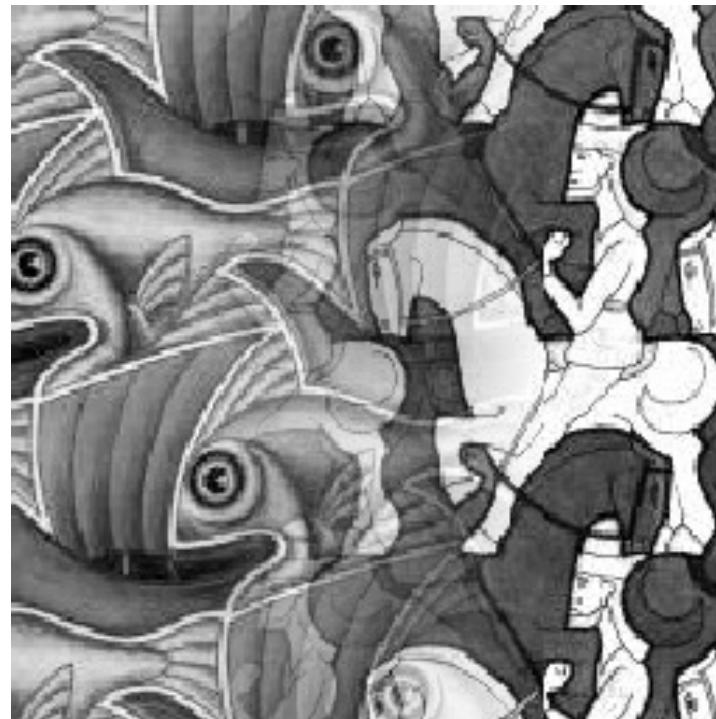
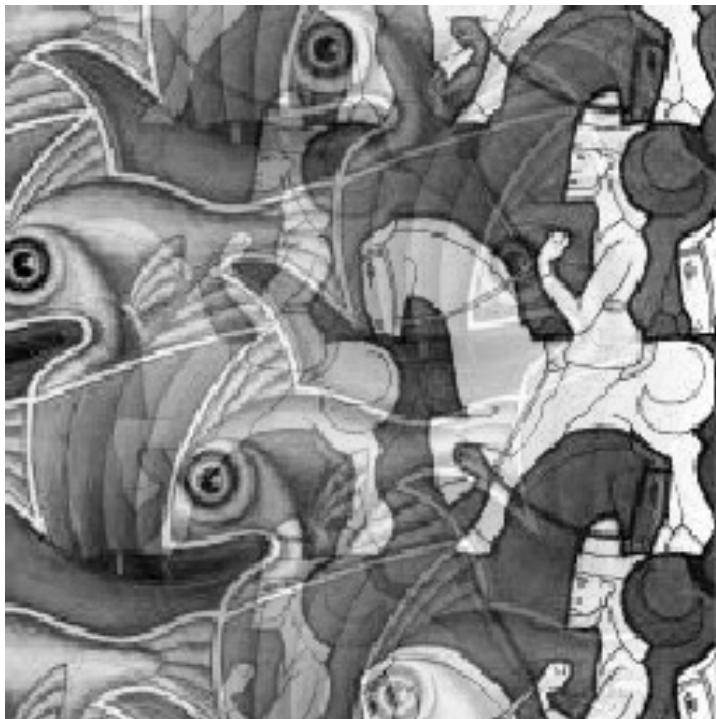
=



$$I_{\text{blend}} = \alpha I_{\text{left}} + (1-\alpha) I_{\text{right}}$$

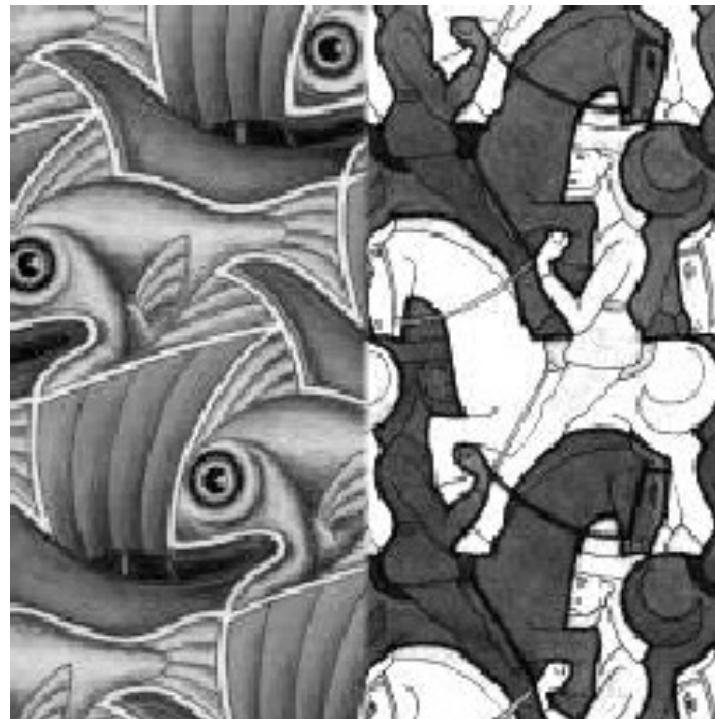
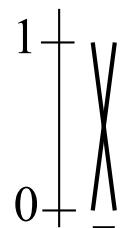
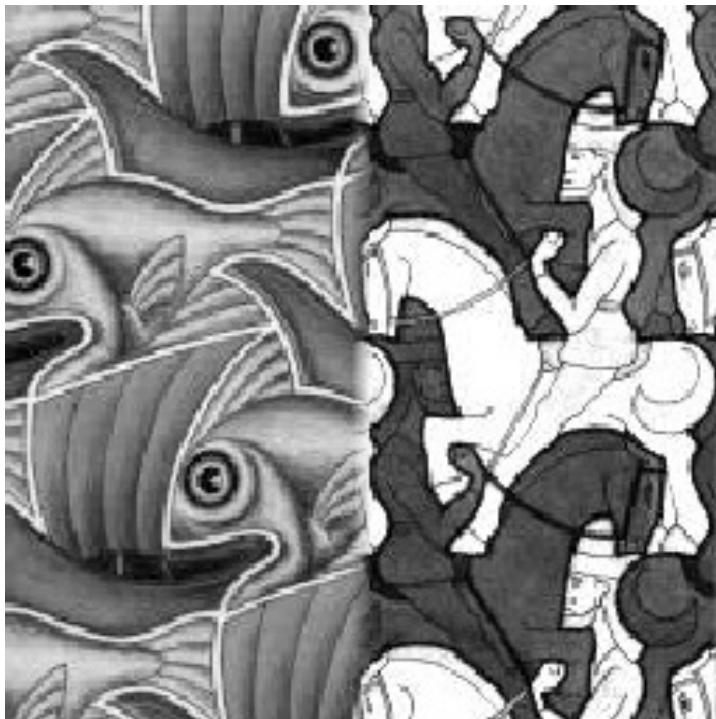
# Affect of Window Size

---



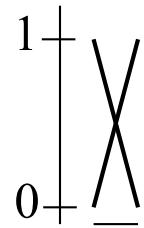
# Affect of Window Size

---



# Good Window Size

---



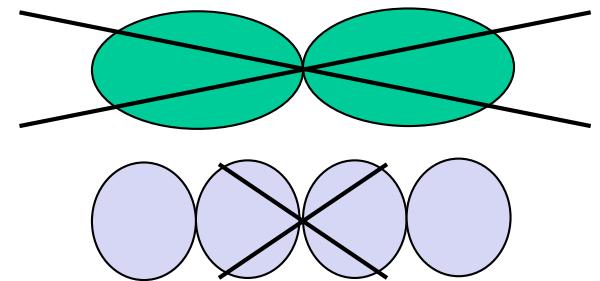
“Optimal” Window: smooth but not ghosted

# What is the Optimal Window?

---

To avoid seams

- window = size of largest prominent feature

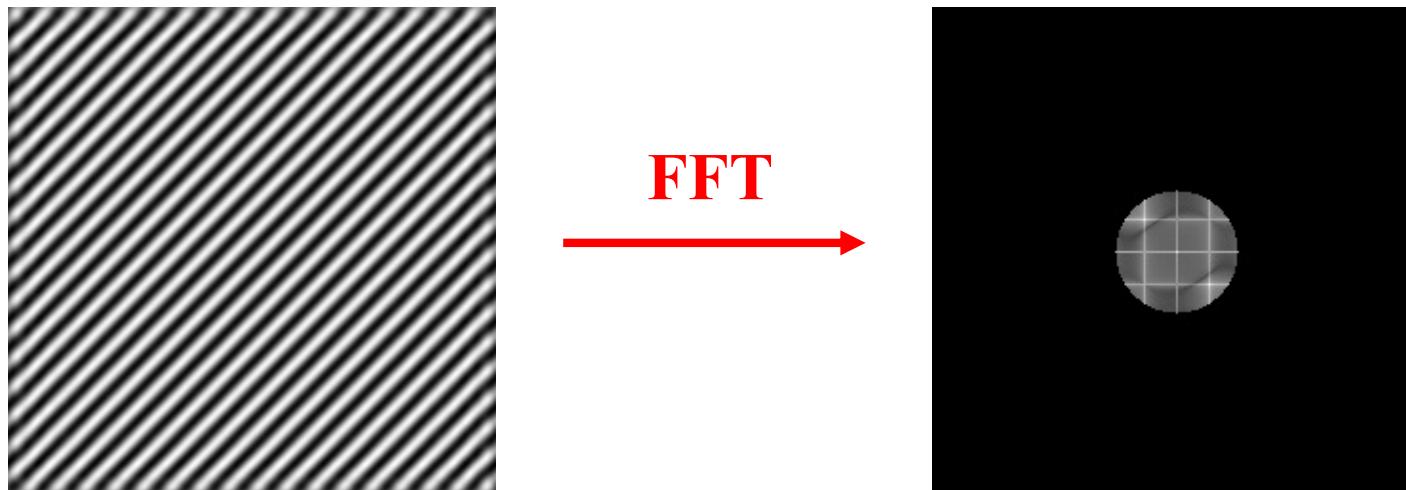


To avoid ghosting

- window  $\leq 2 \times$  size of smallest prominent feature

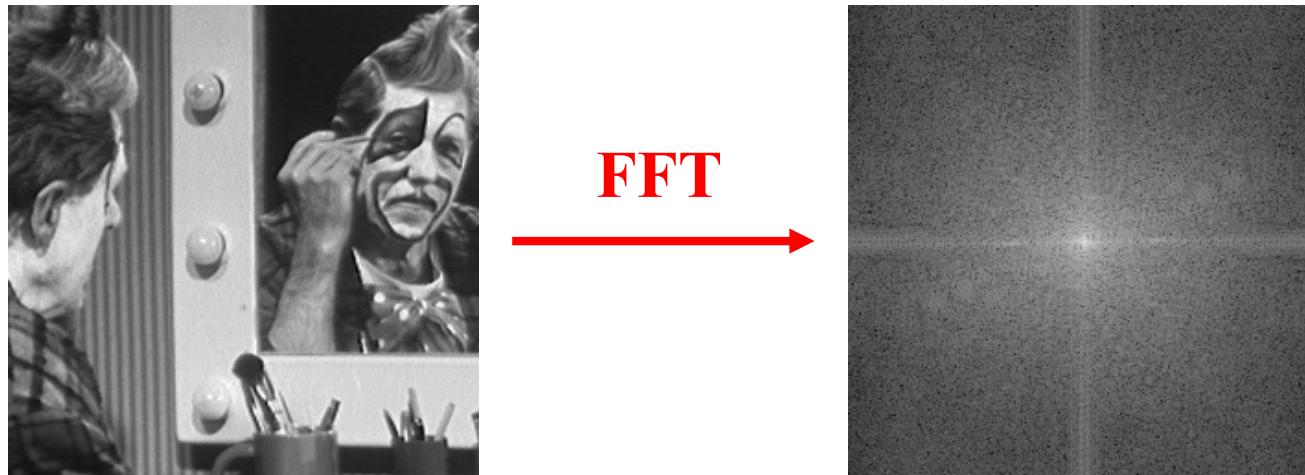
Natural to cast this in the *Fourier domain*

- largest frequency  $\leq 2 \times$  size of smallest frequency
- image frequency content should occupy one “octave” (power of two)



# What if the Frequency Spread is Wide

---



## Idea (Burt and Adelson)

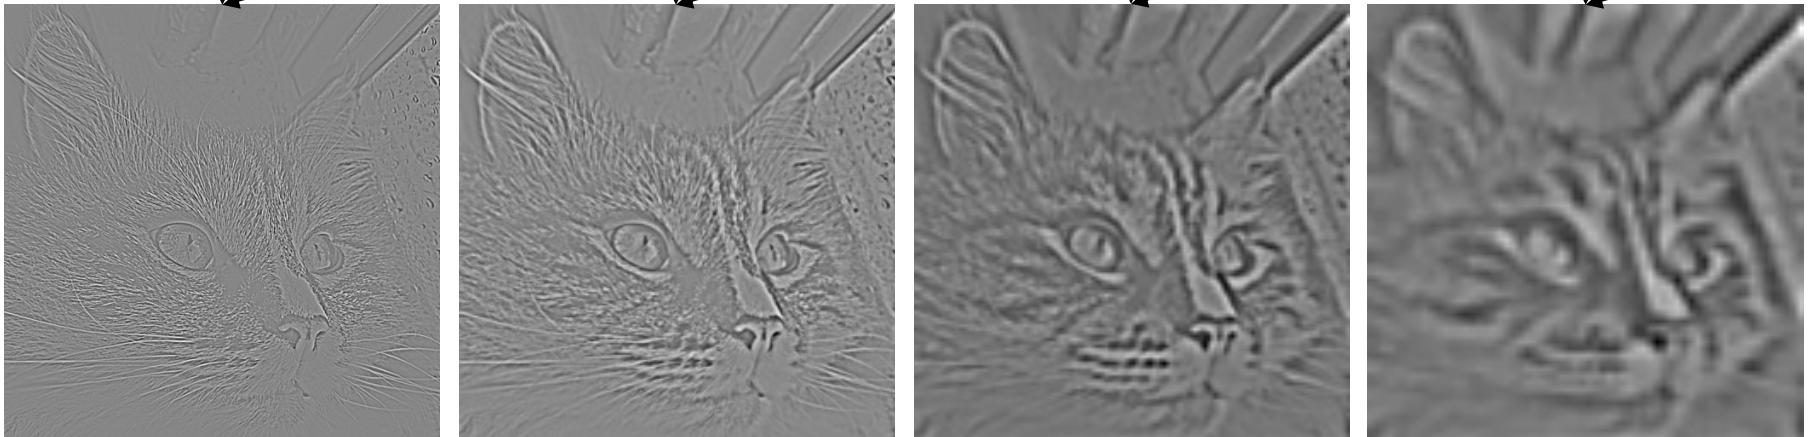
- Compute  $F_{\text{left}} = \text{FFT}(I_{\text{left}})$ ,  $F_{\text{right}} = \text{FFT}(I_{\text{right}})$
- Decompose Fourier image into octaves (bands)
  - $F_{\text{left}} = F_{\text{left}}^1 + F_{\text{left}}^2 + \dots$
- Feather corresponding octaves  $F_{\text{left}}^i$  with  $F_{\text{right}}^i$ 
  - Can compute inverse FFT and feather in spatial domain
- Sum feathered octave images in frequency domain

Better implemented in *spatial domain*

# As a stack

---

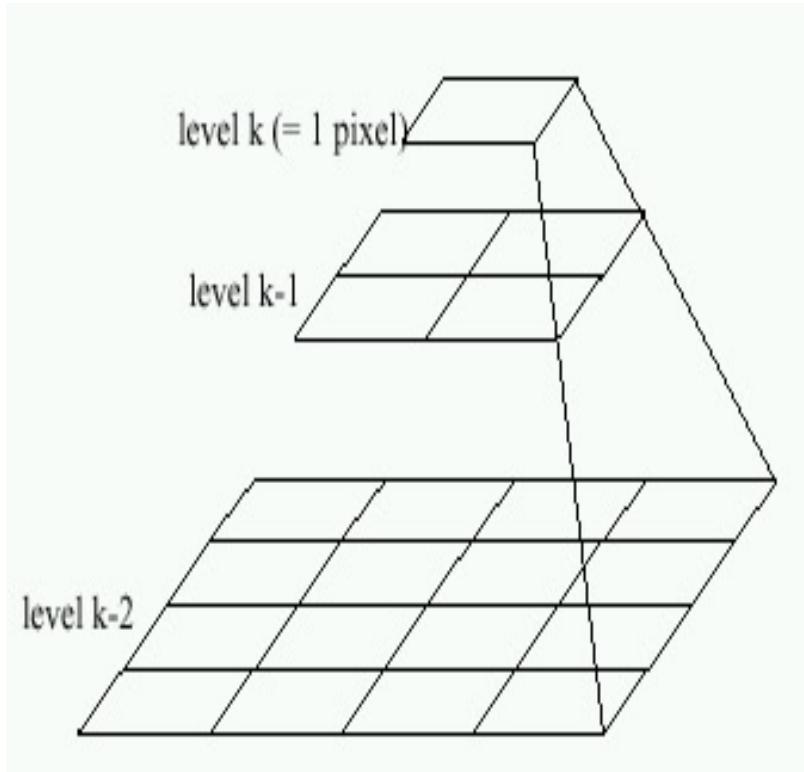
## Gaussian Pyramid (low-pass images)



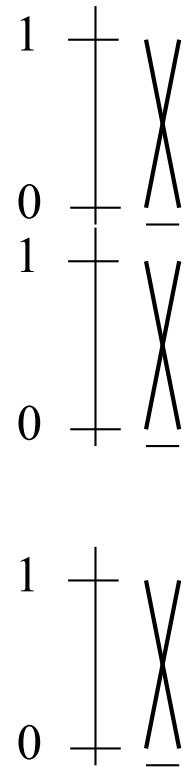
Bandpass Images

# Pyramid Blending

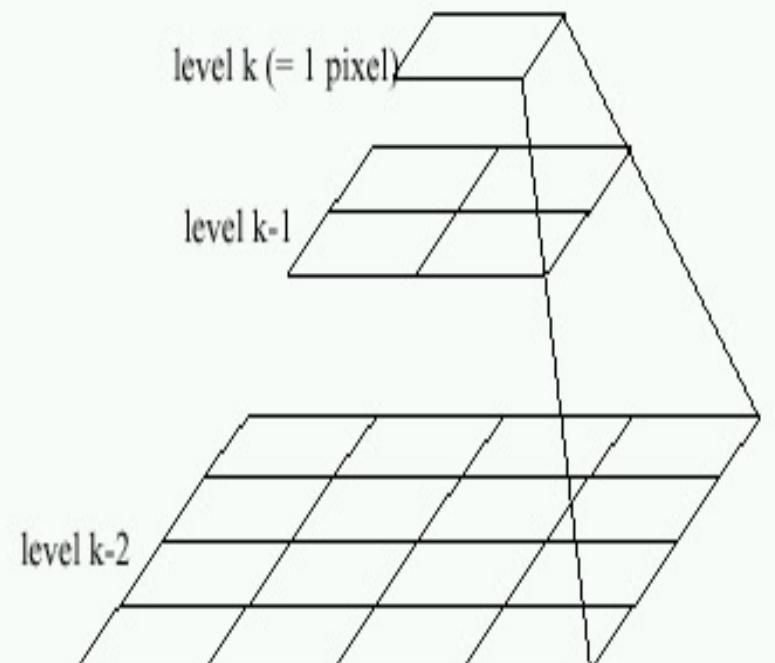
---



Left pyramid



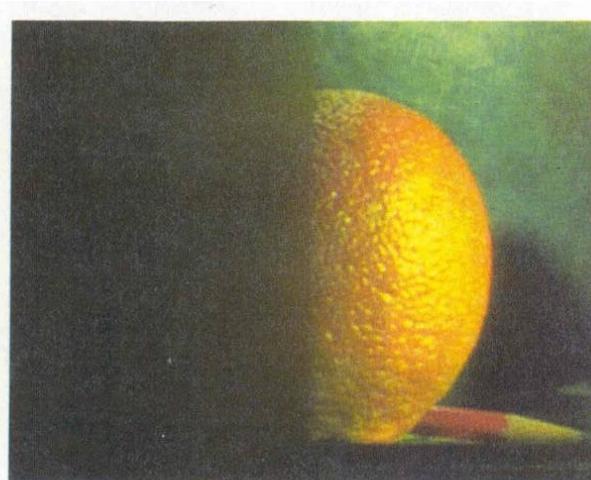
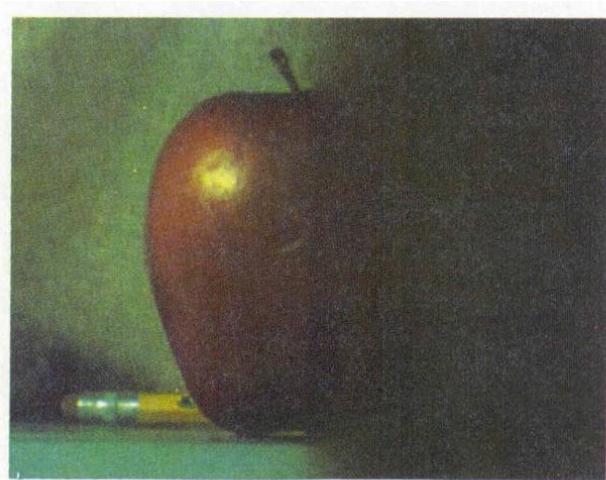
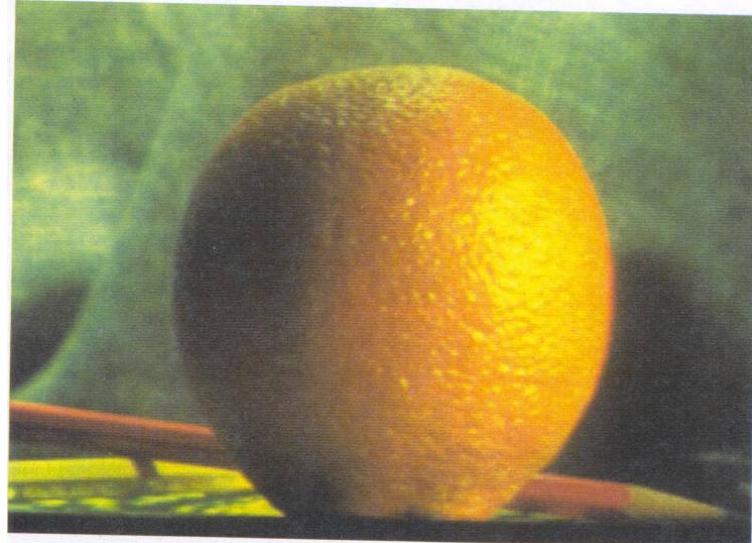
blend



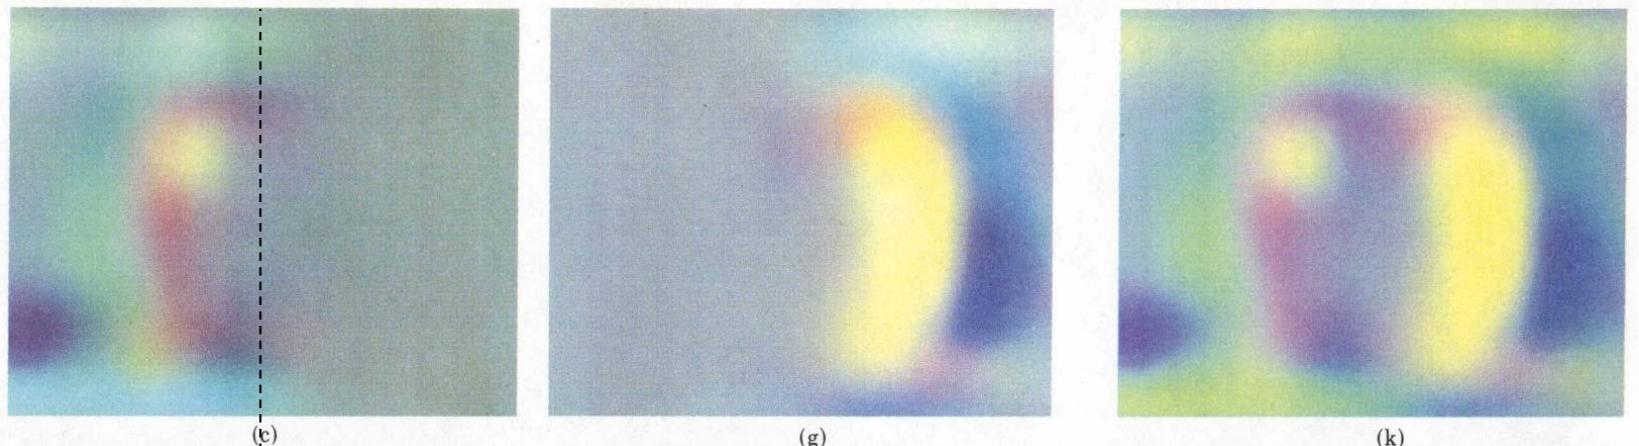
Right pyramid

# Pyramid Blending

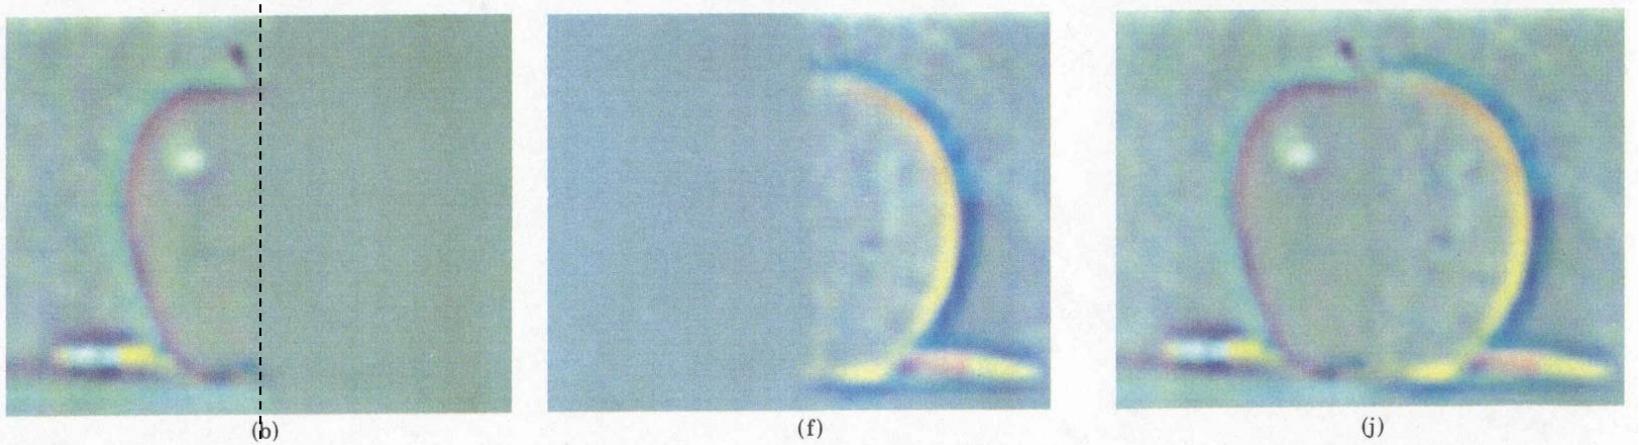
---



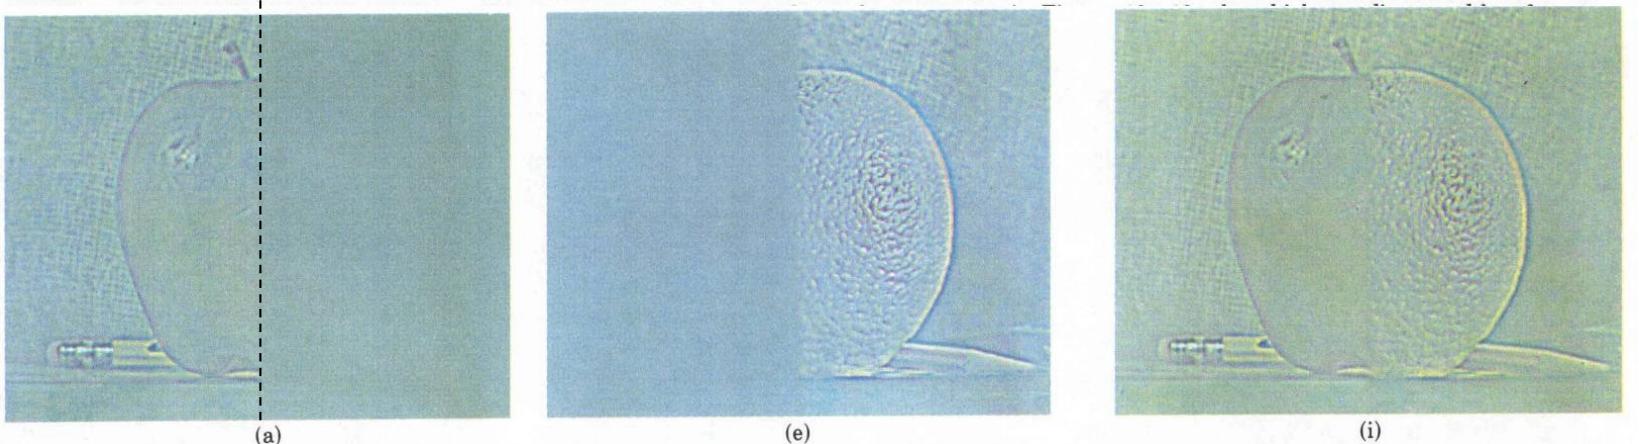
laplacian  
level  
4



laplacian  
level  
2



laplacian  
level  
0



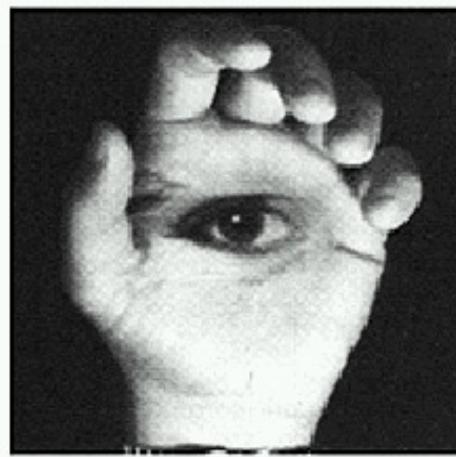
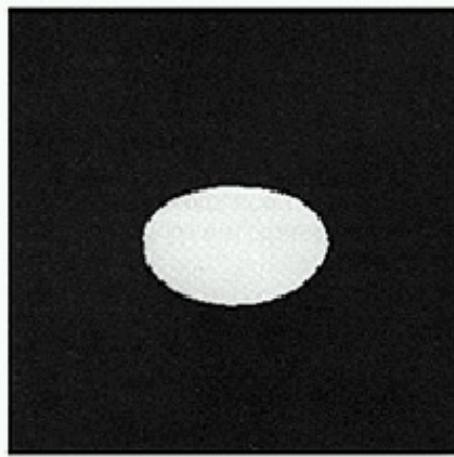
left pyramid

right pyramid

blended pyramid

# Blending Regions

---



# Laplacian Pyramid: Blending

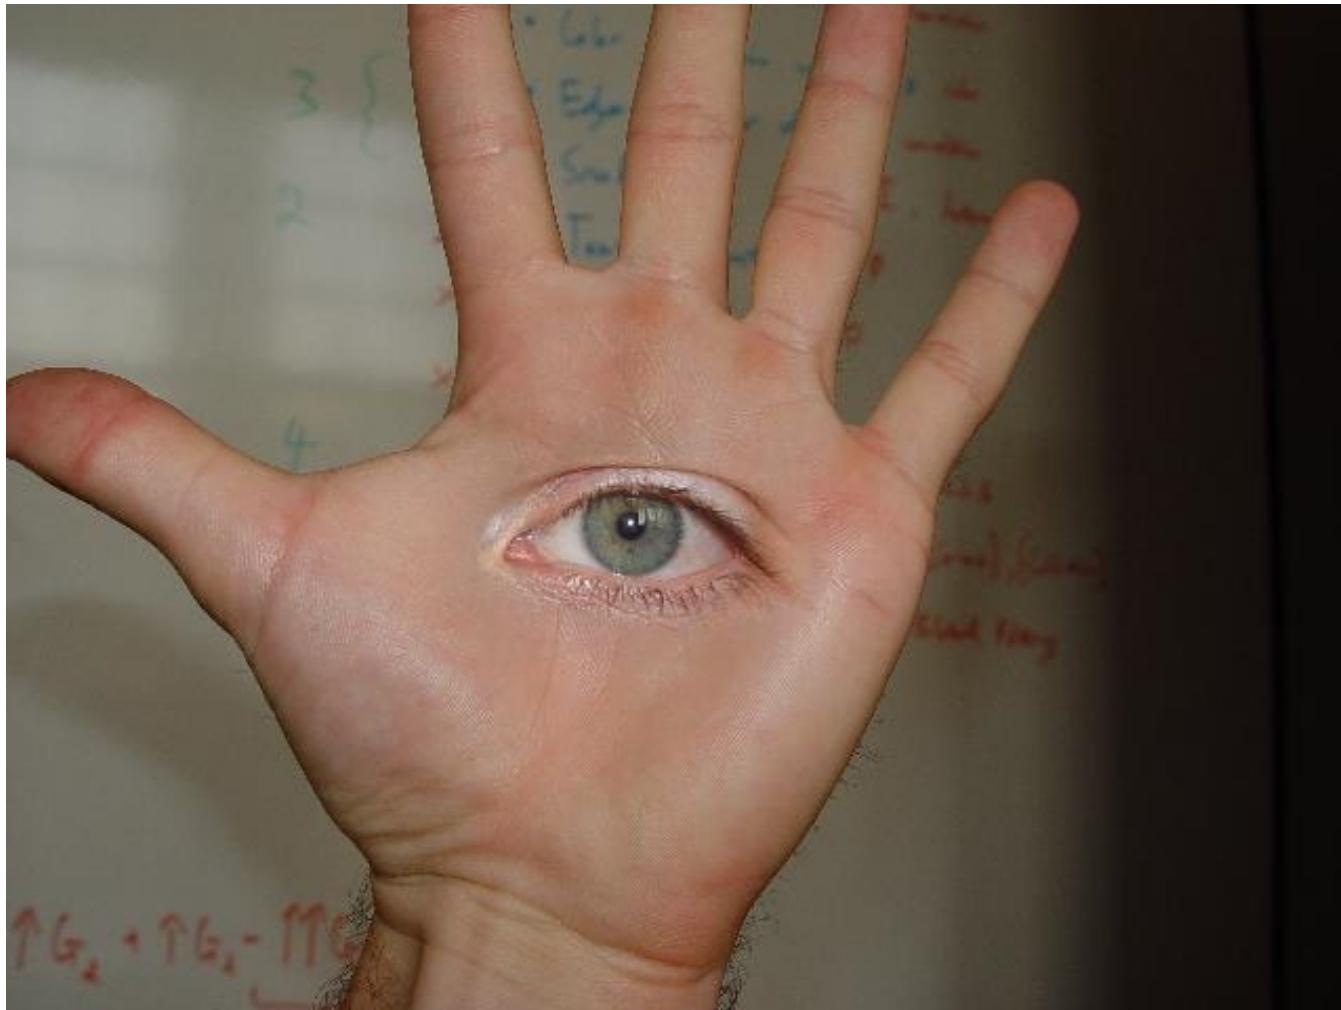
---

## General Approach:

1. Build Laplacian pyramids  $LA$  and  $LB$  from images  $A$  and  $B$
2. Build a Gaussian pyramid  $GR$  from selected region  $R$
3. Form a combined pyramid  $LS$  from  $LA$  and  $LB$  using nodes of  $GR$  as weights:
  - $LS(i,j) = GR(i,j) * LA(i,j) + (1 - GR(i,j)) * LB(i,j)$
4. Collapse the  $LS$  pyramid to get the final blended image

# Horror Photo

---



© david dmartin (Boston College)

# Results from this class (fall 2005)

---



© Chris Cameron

# Simplification: Two-band Blending

---

Brown & Lowe, 2003

- Only use two bands – high freq. and low freq. – without downsampling
- Blends low freq. smoothly
- Blend high freq. with no smoothing: use binary alpha



# 2-band “Laplacian Stack” Blending

---



Low frequency ( $\lambda > 2$  pixels)



High frequency ( $\lambda < 2$  pixels)

# Linear Blending



# 2-band Blending



# Side note: Image Compression

---

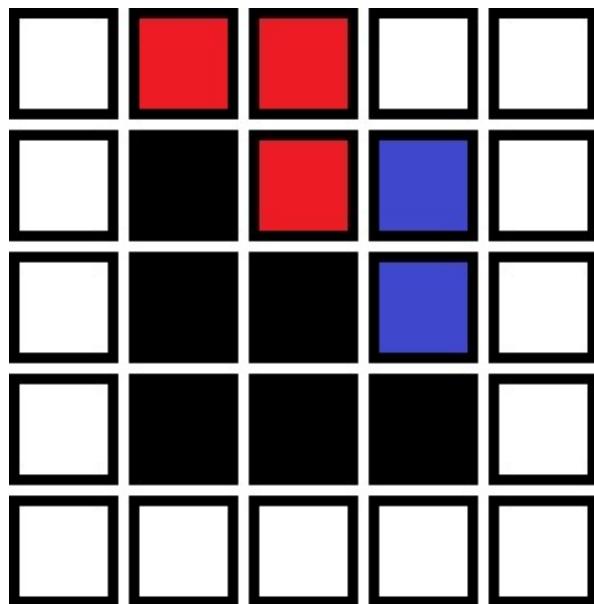


89k

# Lossless Compression (e.g. Huffman coding)

---

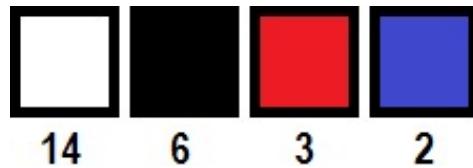
Input image:



Pixel code:

color	freq.	bit code
White	14	0
Black	6	10
Red	3	110
Blue	2	111

Pixel histogram:

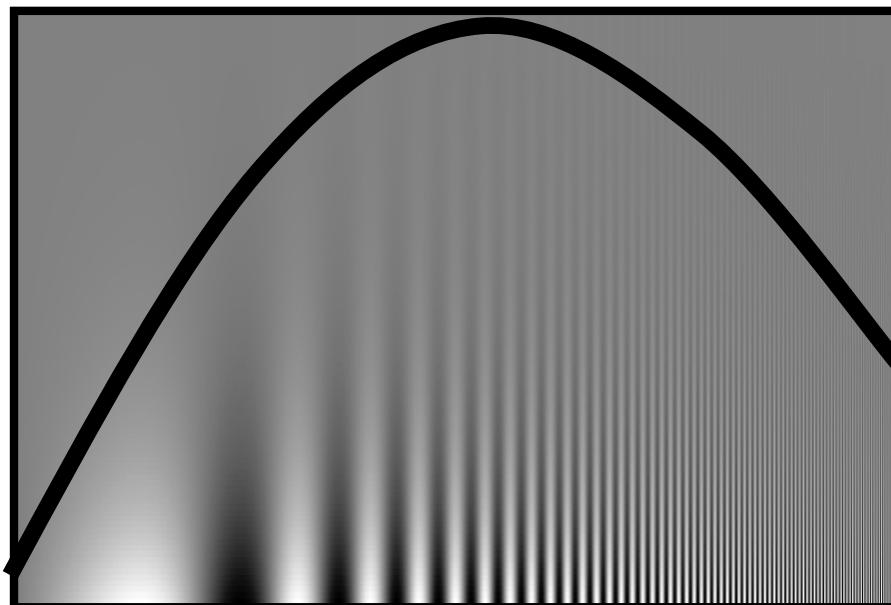


Compressed image:

0 110 110 0 0  
0 10 110 111 0  
...

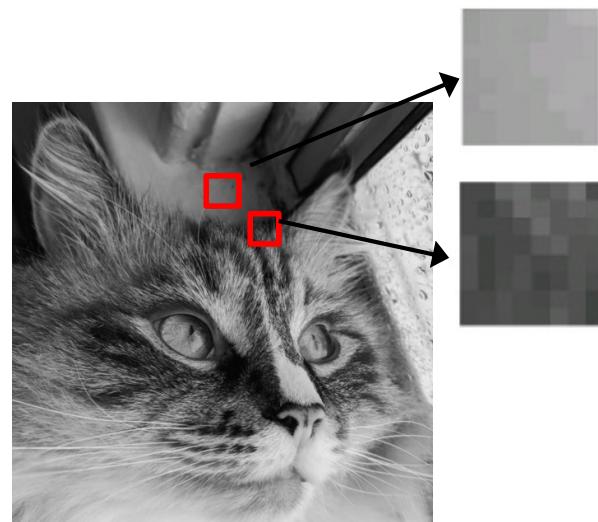
# Lossless Compression not enough

---

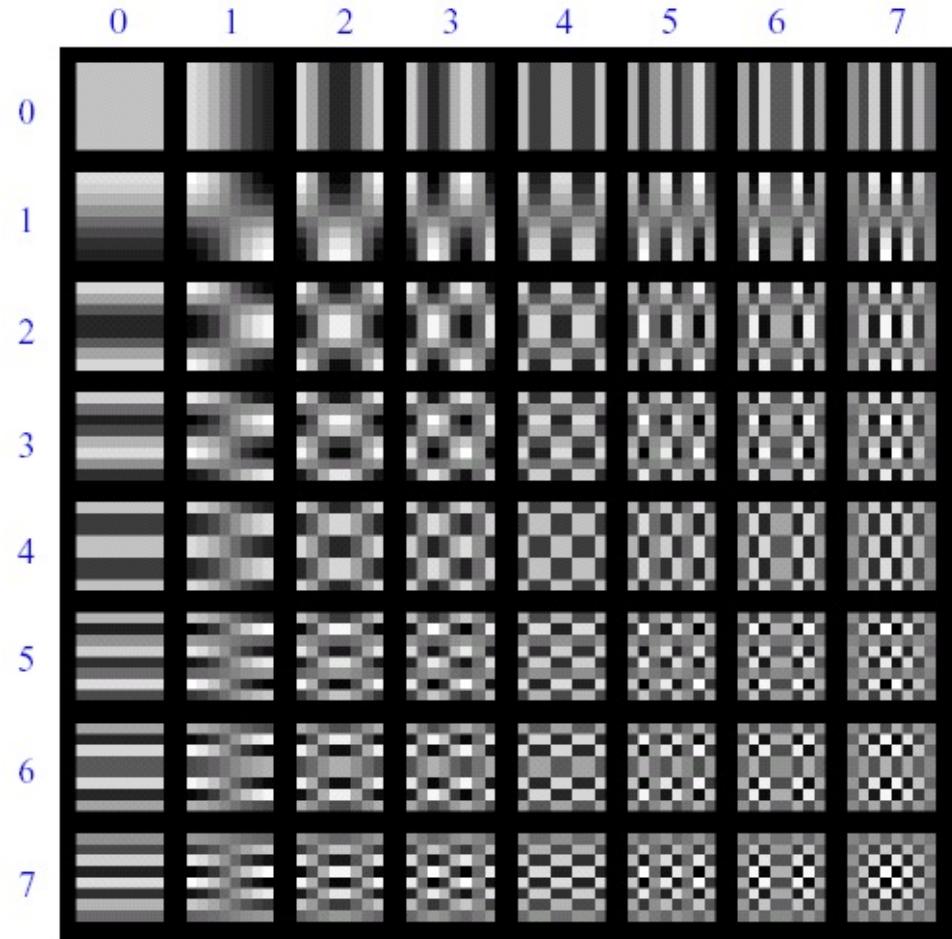


# Lossy Image Compression (JPEG)

---



cut up into 8x8 blocks



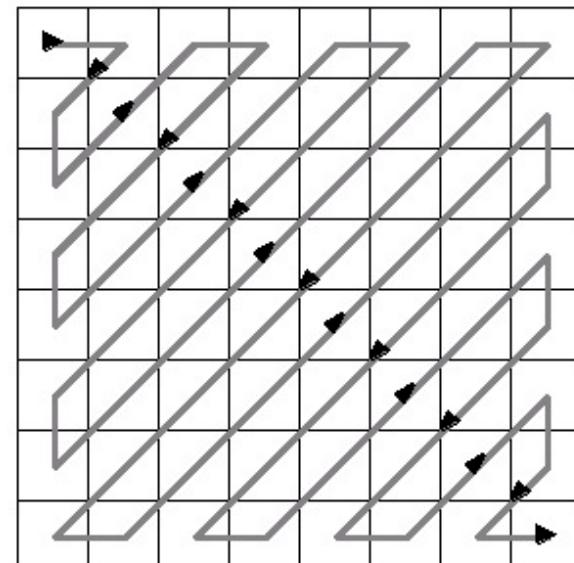
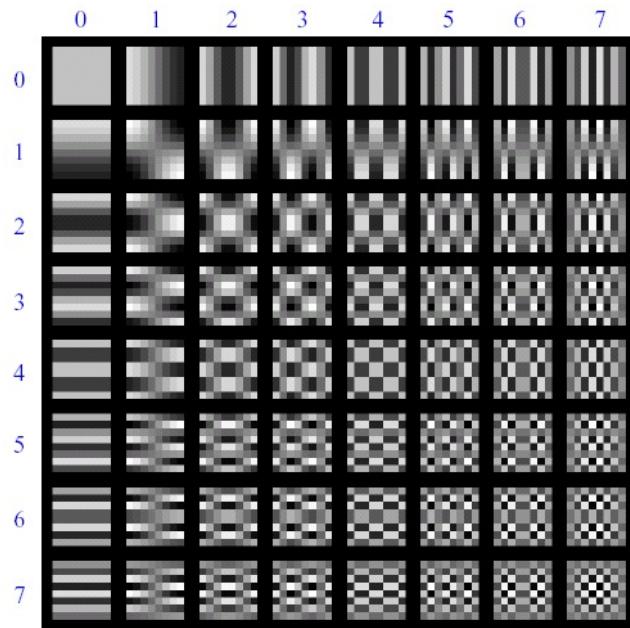
Block-based Discrete Cosine Transform (DCT)

# Using DCT in JPEG

---

The first coefficient  $B(0,0)$  is the DC component,  
the average intensity

The top-left coeffs represent low frequencies,  
the bottom right – high frequencies



# Image compression using DCT

---

## Quantize

- More coarsely for high frequencies (tend to have smaller values anyway)
- Many quantized high frequency values will be zero

## Encode

- Can decode with inverse dct

### Filter responses

$$G = \begin{bmatrix} -415.38 & -30.19 & -61.20 & 27.24 & 56.13 & -20.10 & -2.39 & 0.46 \\ 4.47 & -21.86 & -60.76 & 10.25 & 13.15 & -7.09 & -8.54 & 4.88 \\ -46.83 & 7.37 & 77.13 & -24.56 & -28.91 & 9.93 & 5.42 & -5.65 \\ -48.53 & 12.07 & 34.10 & -14.76 & -10.24 & 6.30 & 1.83 & 1.95 \\ 12.12 & -6.55 & -13.20 & -3.95 & -1.88 & 1.75 & -2.79 & 3.14 \\ -7.73 & 2.91 & 2.38 & -5.94 & -2.38 & 0.94 & 4.30 & 1.85 \\ -1.03 & 0.18 & 0.42 & -2.42 & -0.88 & -3.02 & 4.12 & -0.66 \\ -0.17 & 0.14 & -1.07 & -4.19 & -1.17 & -0.10 & 0.50 & 1.68 \end{bmatrix}$$

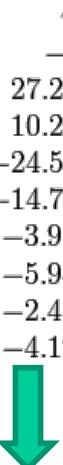
$\xrightarrow{u}$

### Quantized values

$$B = \begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -3 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

### Quantization table

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$



# JPEG Compression Summary

---

Subsample color by factor of 2

- People have bad resolution for color

Split into blocks (8x8, typically), subtract 128

For each block

- a. Compute DCT coefficients
- b. Coarsely quantize
  - Many high frequency components will become zero
- c. Encode (e.g., with Huffman coding)

Spatial dimension of color channels are reduced by 2  
(lecture 2)!

<http://en.wikipedia.org/wiki/YCbCr>  
<http://en.wikipedia.org/wiki/JPEG>

# Block size in JPEG

---

## Block size

- small block
  - faster
  - correlation exists between neighboring pixels
- large block
  - better compression in smooth regions
- It's 8x8 in standard JPEG

# JPEG compression comparison

---



89k



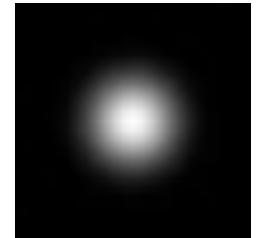
12k

# Review: Smoothing vs. derivative filters

---

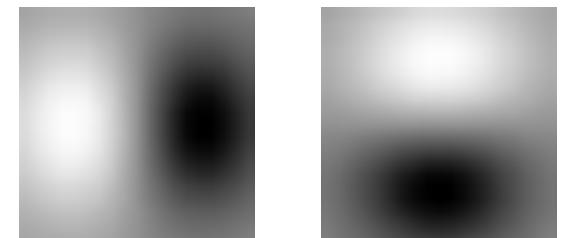
## Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
  - **One**: constant regions are not affected by the filter



## Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
  - **Zero**: no response in constant regions
- High absolute value at points of high contrast



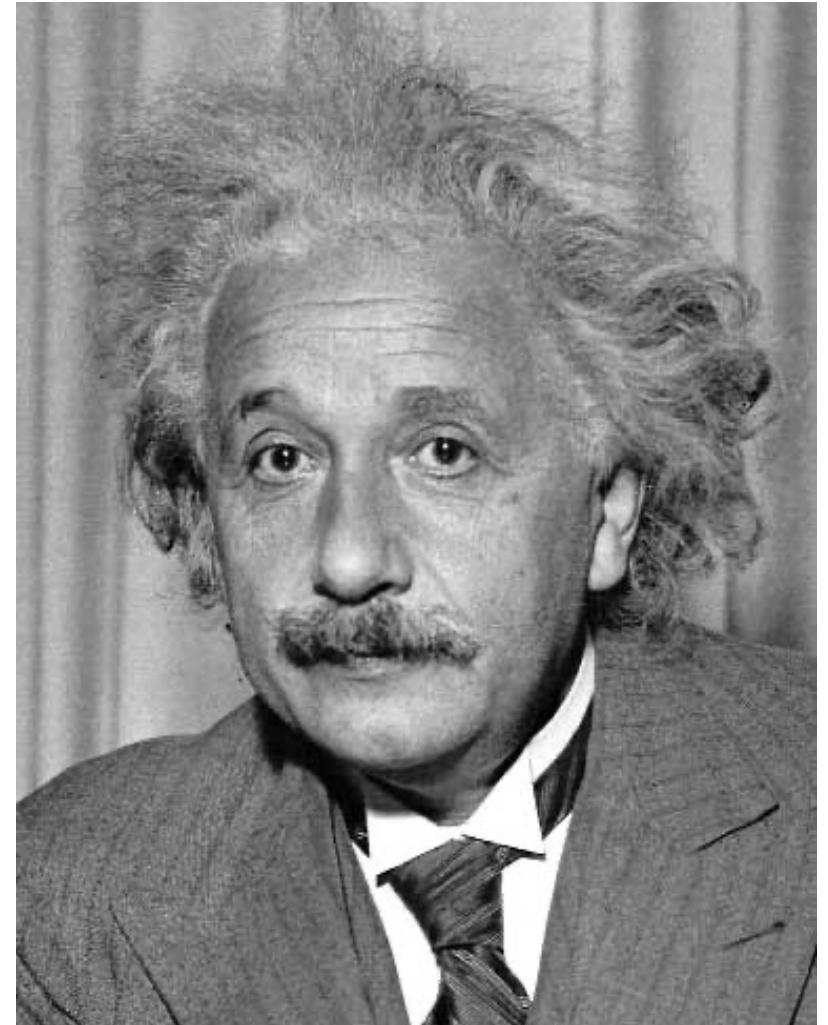
# Template matching

---

Goal: find  in image

Main challenge: What is a good similarity or distance measure between two patches?

- Correlation
- Zero-mean correlation
- Sum Square Difference
- Normalized Cross Correlation



# Matching with filters

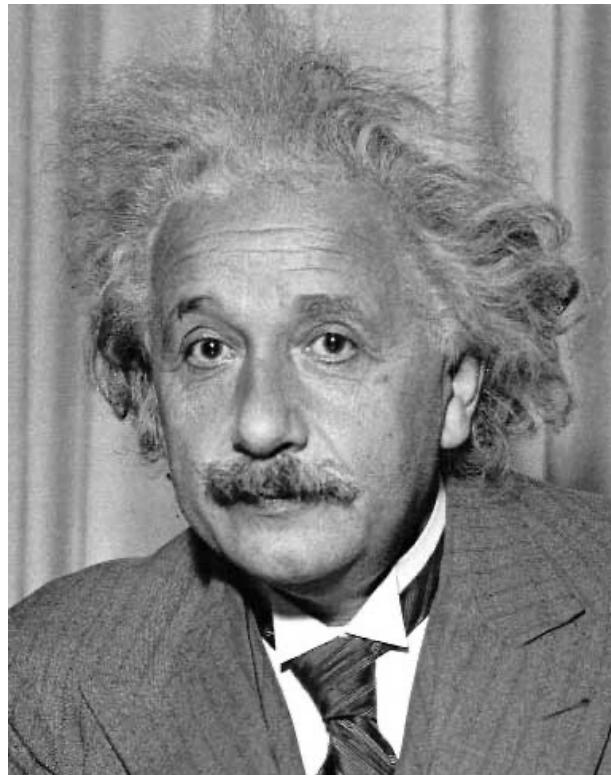
---

Goal: find  in image

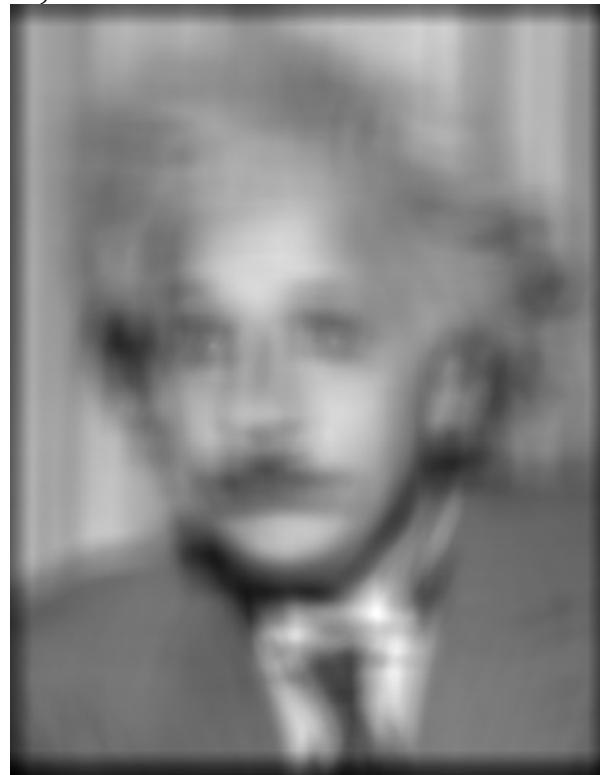
Method 0: filter the image with eye patch

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$

f = image  
g = filter



Input



Filtered Image

What went wrong?

Side by Derek Hoiem

# Matching with filters

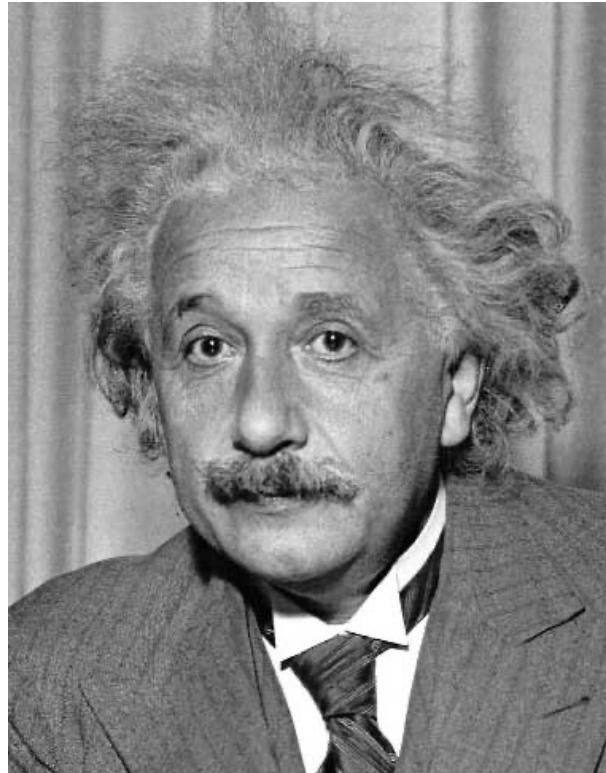
Goal: find  in image

$f$  = image  
 $g$  = filter

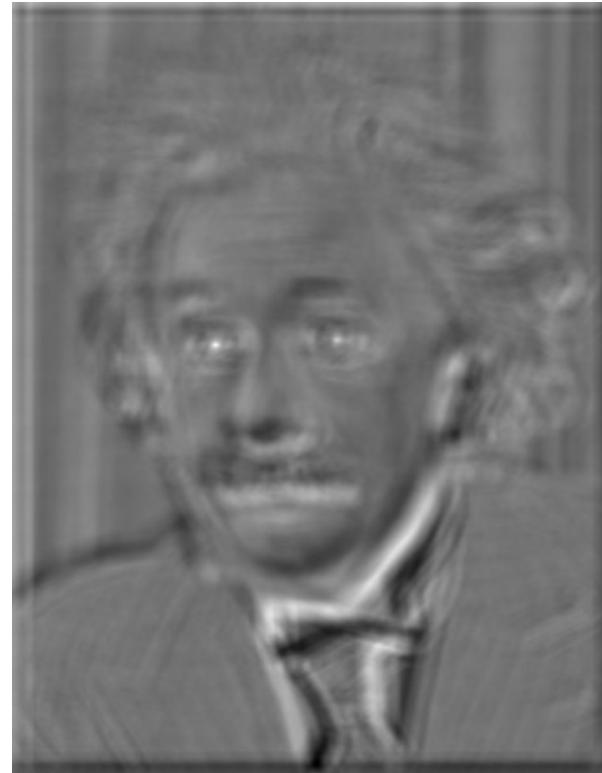
Method 1: filter the image with zero-mean eye

$$h[m, n] = \sum_{k, l} (g[k, l] - \bar{g})(f[m + k, n + l])$$

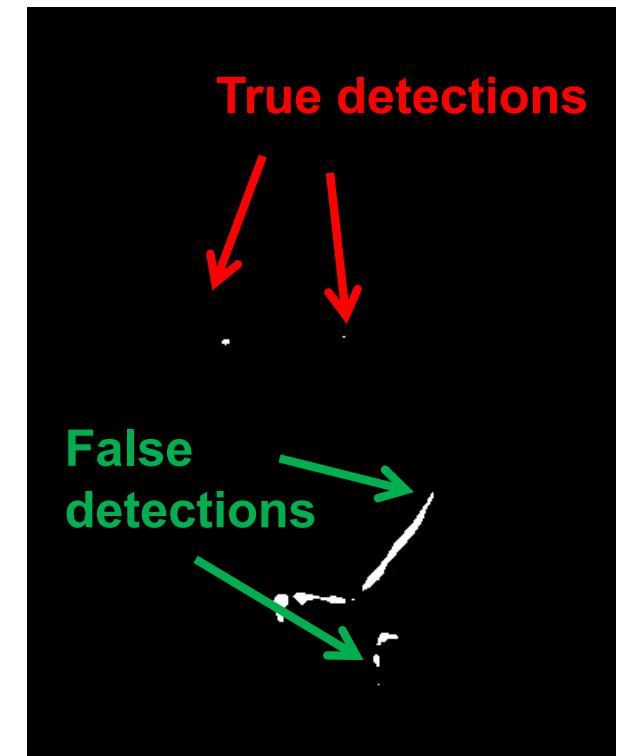
mean of g



Input



Filtered Image (scaled)



Thresholded Image

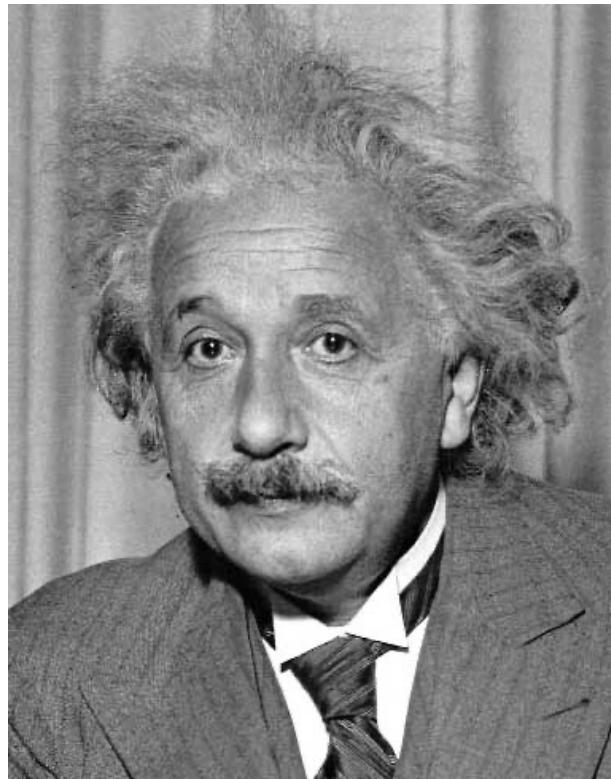
# Matching with filters

---

Goal: find  in image

Method 2: SSD (L2)

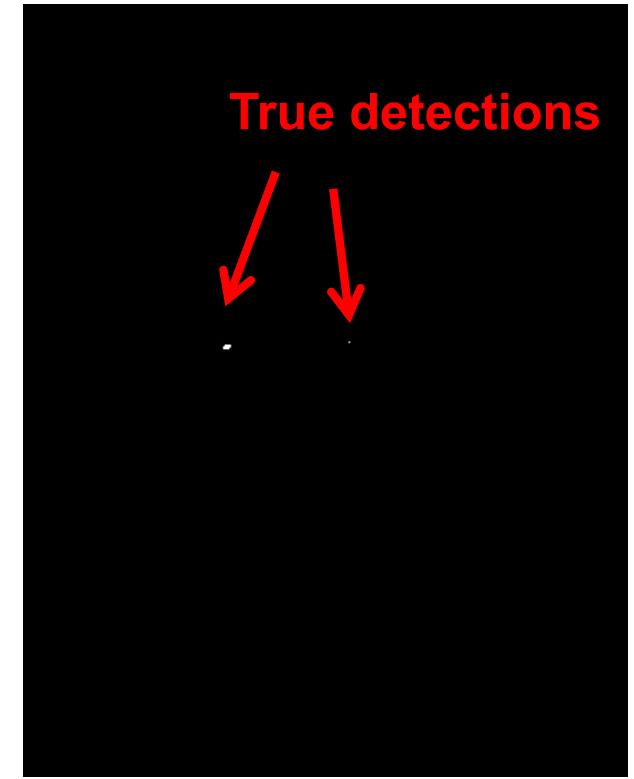
$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



Input



1 -  $\sqrt{\text{SSD}}$



True detections

# Matching with filters

---

Can SSD be implemented with linear filters?

$$h[m, n] = \sum_{k,l} (g[k, l] - f[m + k, n + l])^2$$

# Matching with filters

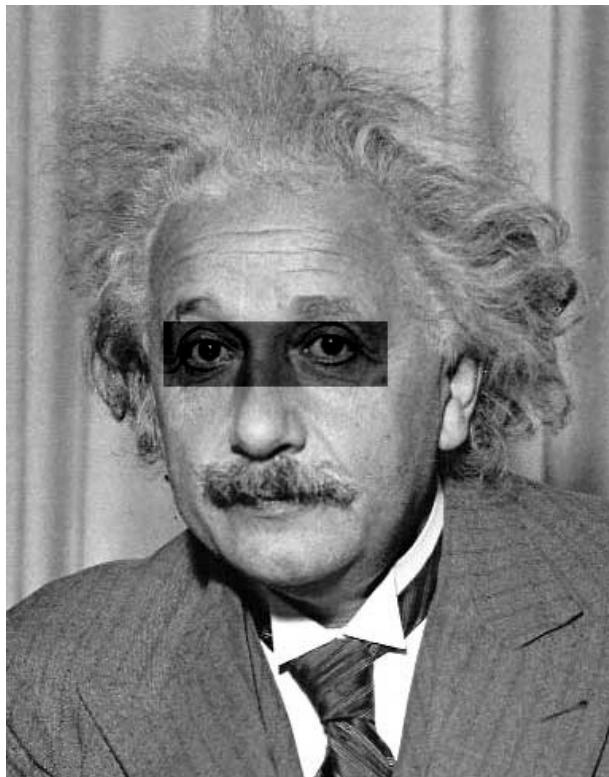
---

Goal: find  in image

What's the potential downside of SSD?

Method 2: SSD

$$h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$$



Input



1- sqrt(SSD)

Side by Derek Hoiem

# Matching with filters

---

Goal: find  in image

Method 3: Normalized cross-correlation

$$h[m, n] = \frac{\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})}{\left( \sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m + k, n + l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

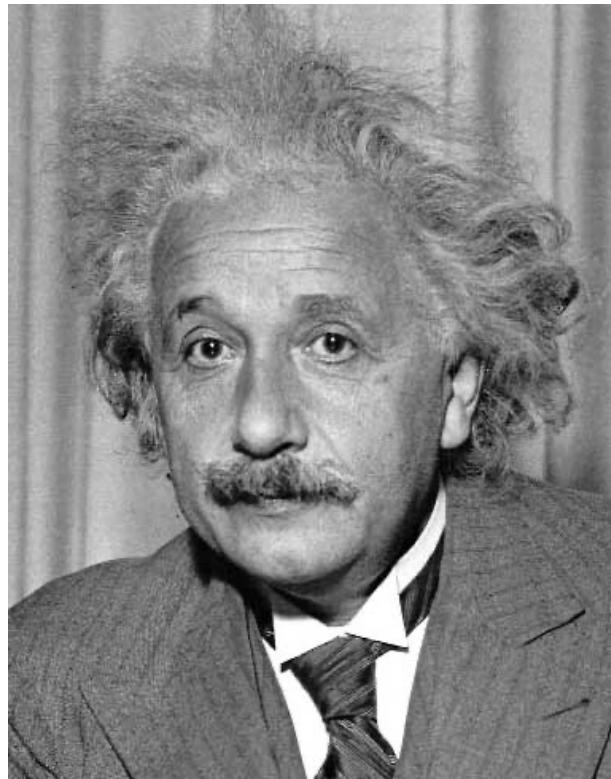
mean template  
↓  
 $\sum_{k,l}$   
mean image patch  
↓

# Matching with filters

---

Goal: find  in image

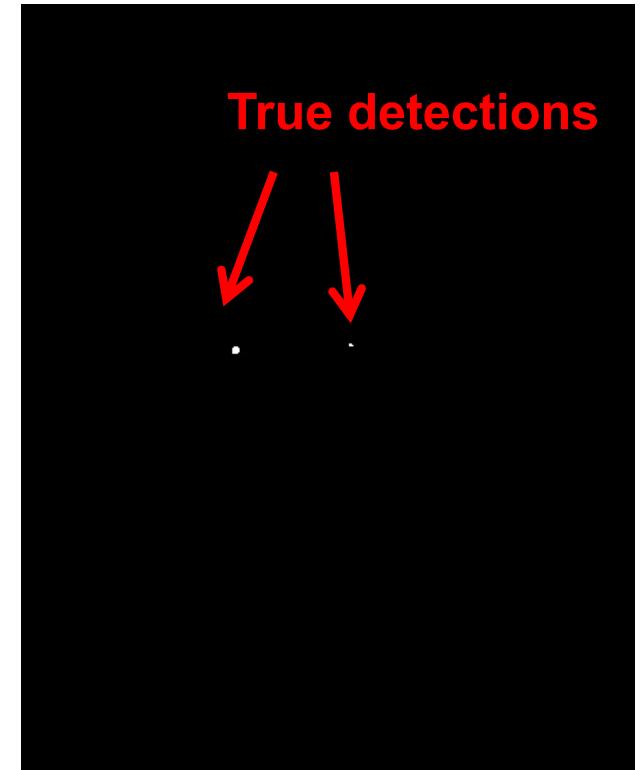
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



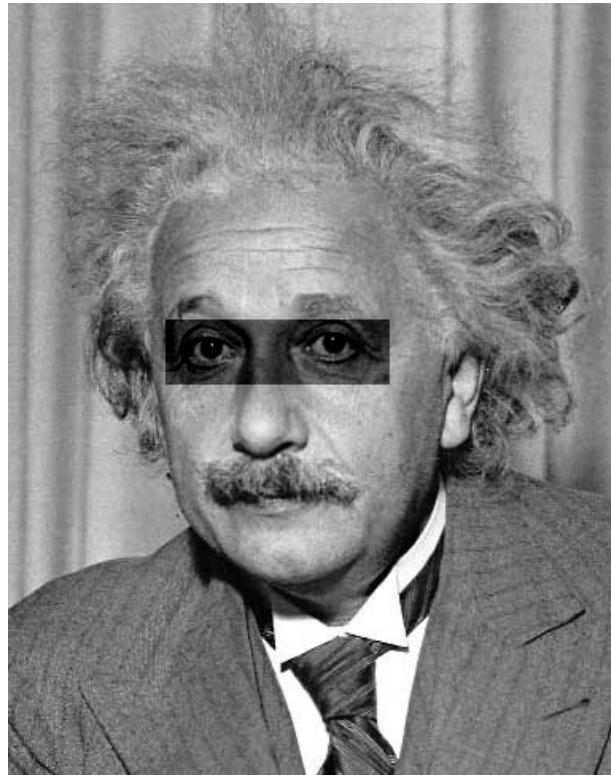
True detections  
↓  
↓  
Thresholded Image

# Matching with filters

---

Goal: find  in image

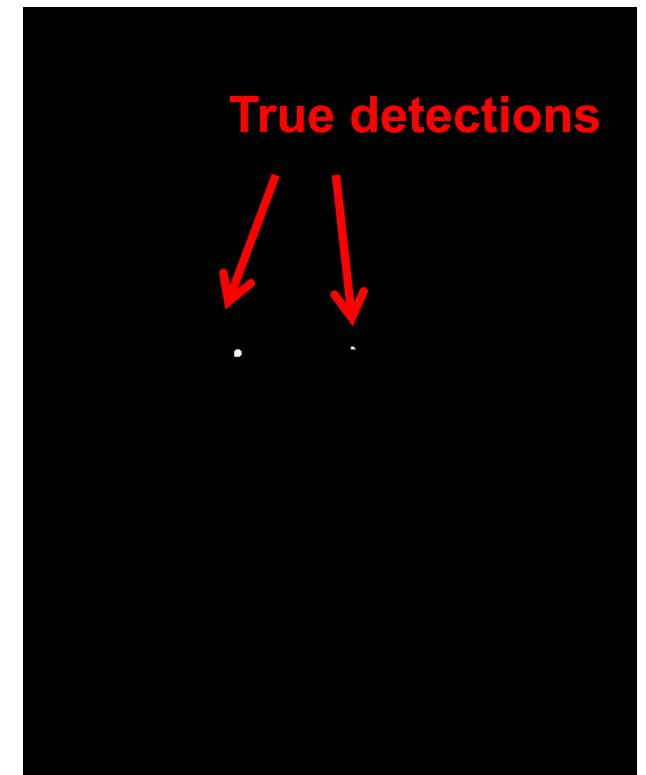
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



True detections  
↓  
↓  
Thresholded Image

# Q: What is the best method to use?

A: Depends

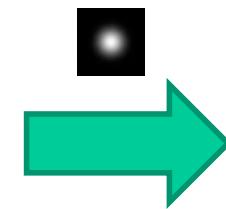
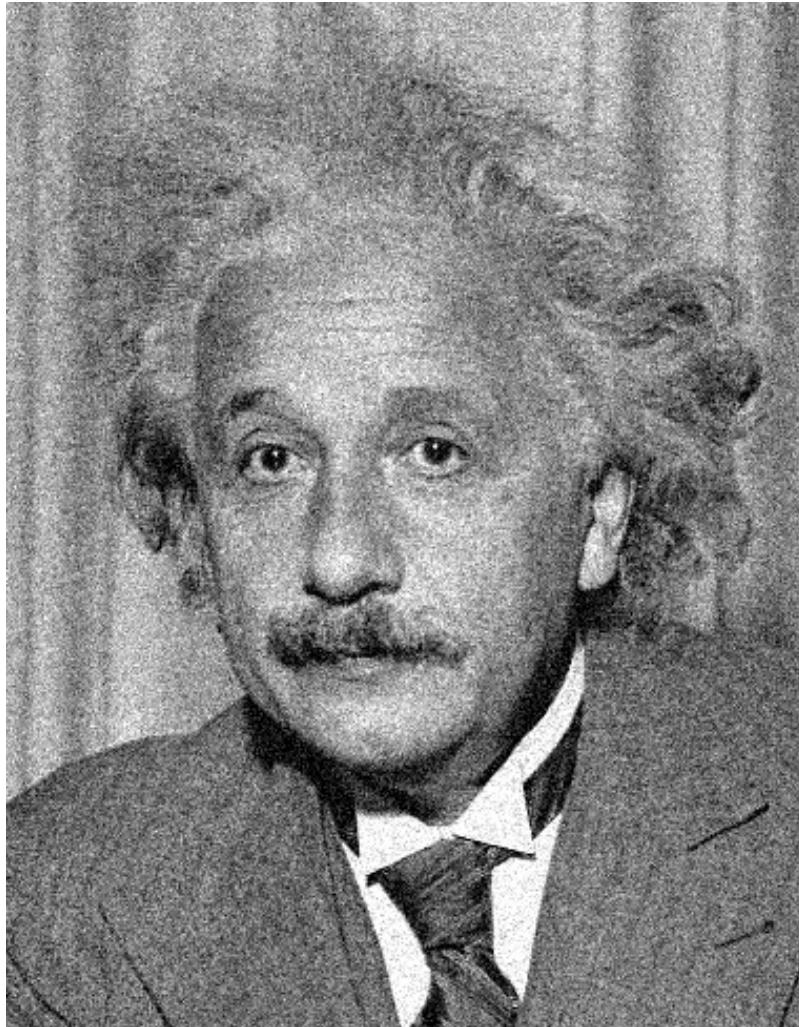
Zero-mean filter: fastest but not a great matcher

SSD: next fastest, sensitive to overall intensity

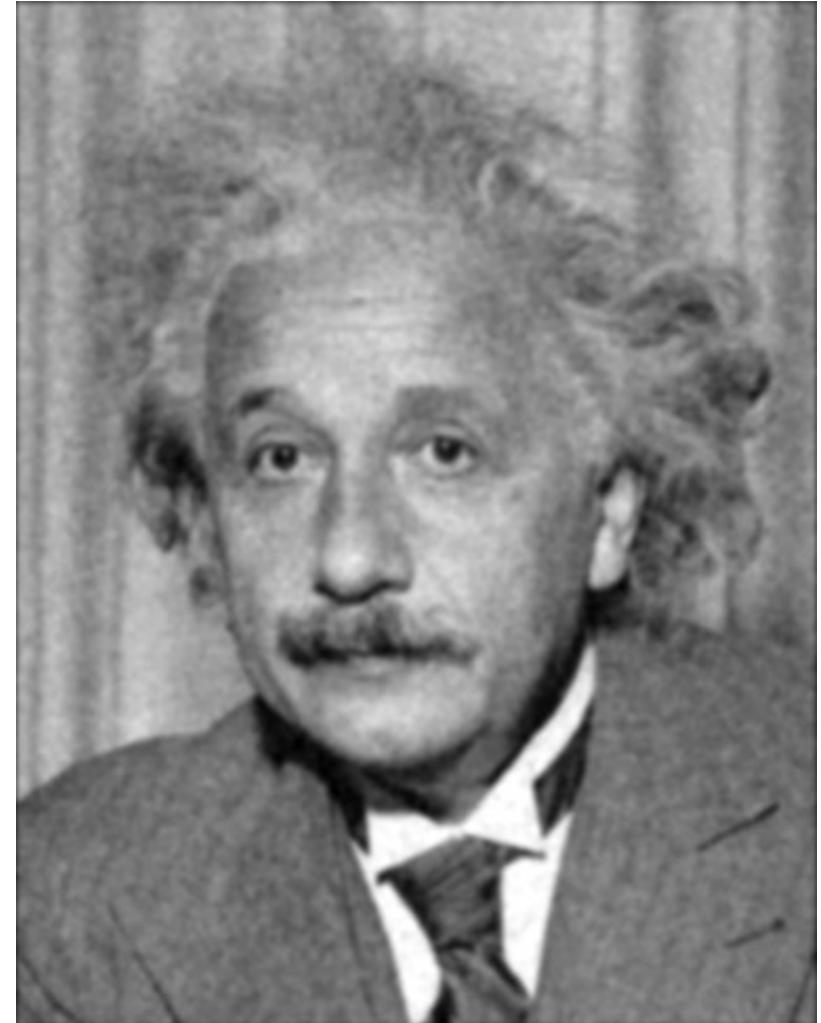
Normalized cross-correlation: slowest, invariant to local average intensity and contrast

# Denoising

---

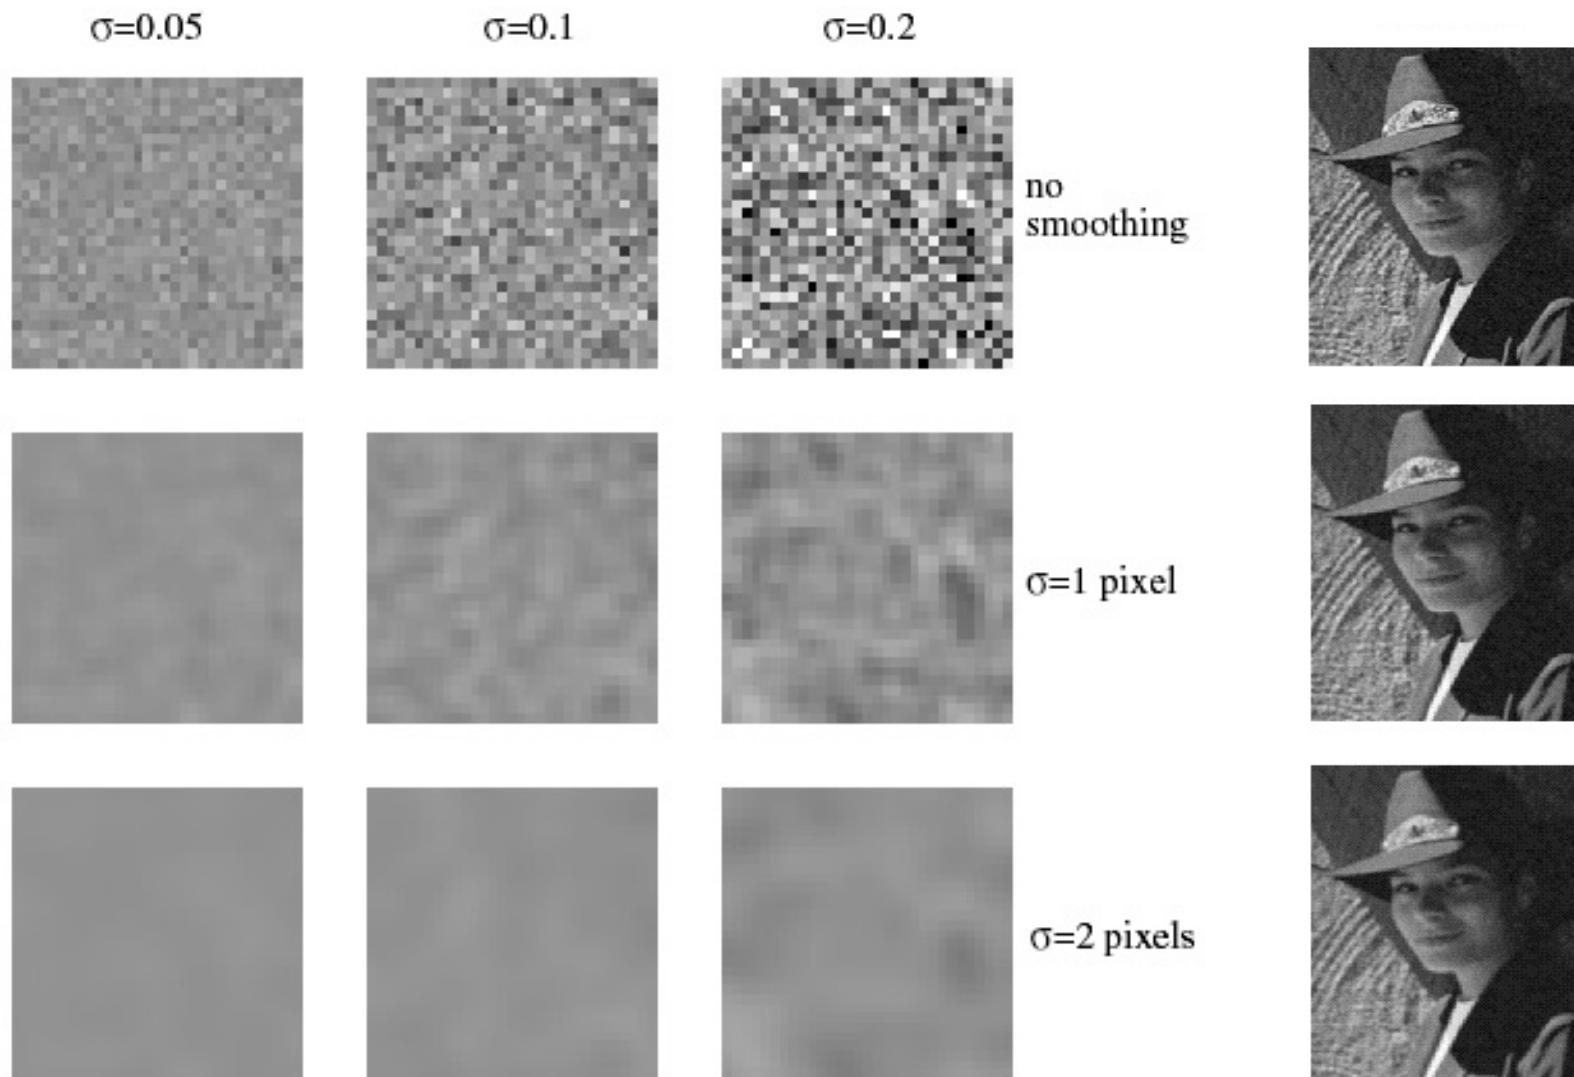


Gaussian  
Filter



Additive Gaussian Noise

# Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Source: S. Lazebnik

# Reducing salt-and-pepper noise by Gaussian smoothing

---

3x3



5x5

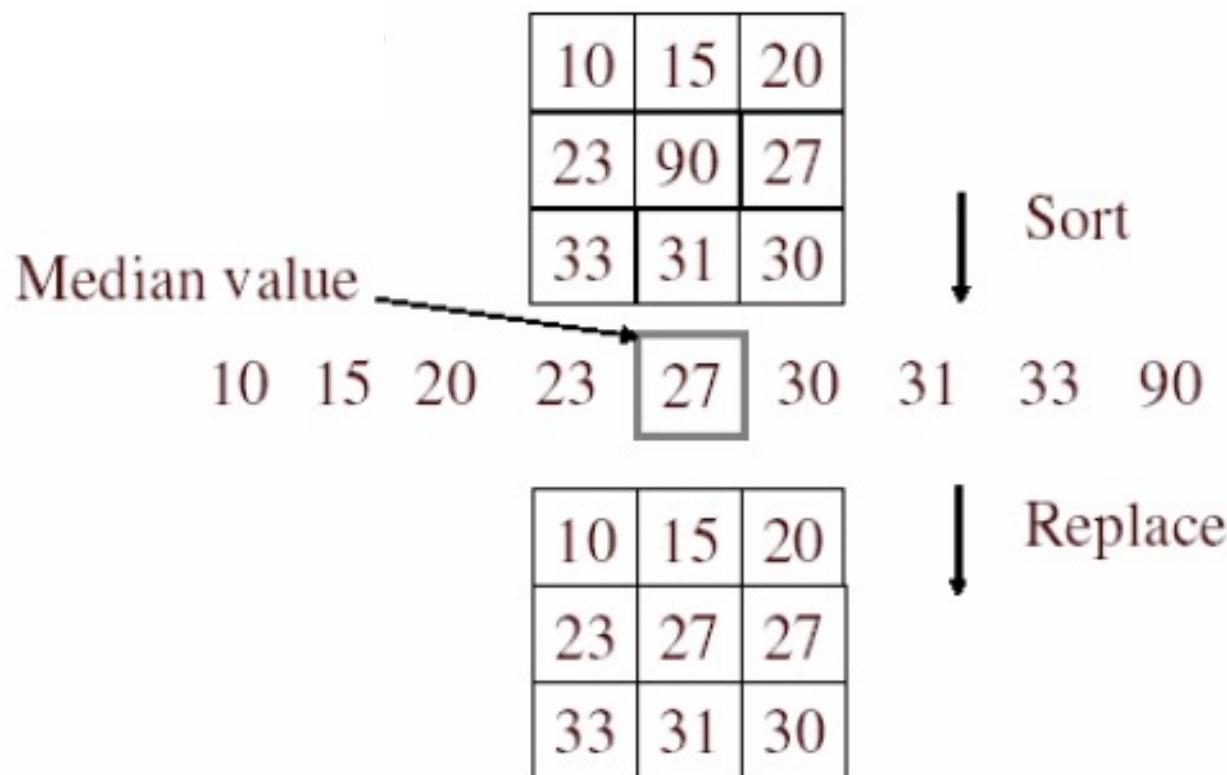


7x7



# Alternative idea: Median filtering

A **median filter** operates over a window by selecting the median intensity in the window



- Is median filtering linear?

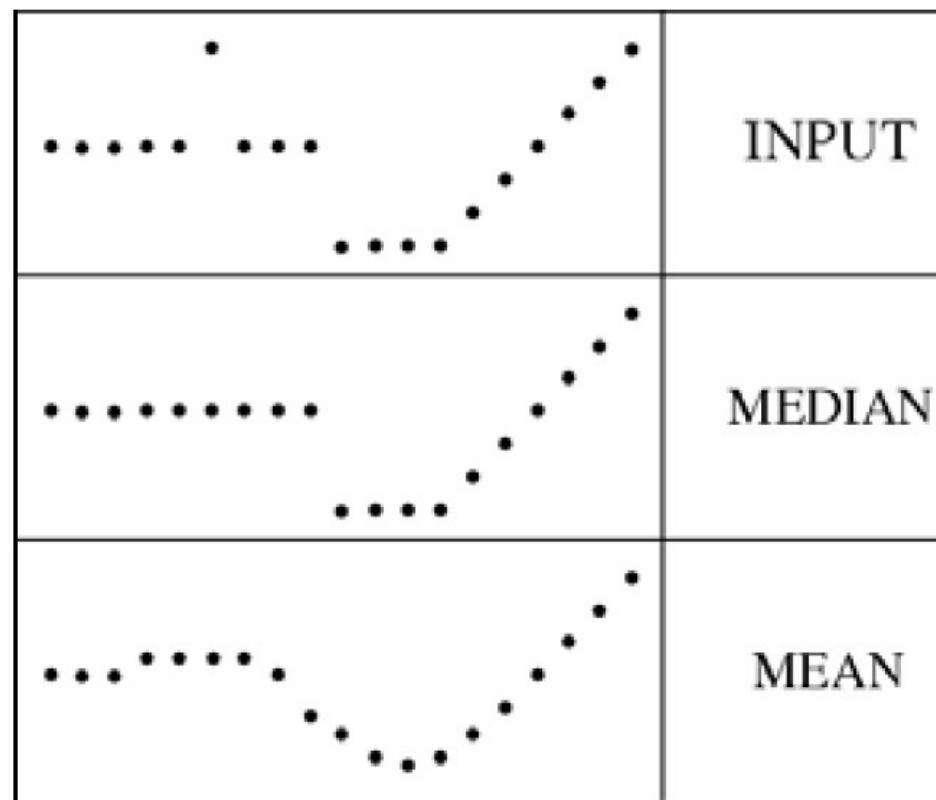
# Median filter

---

What advantage does median filtering have over Gaussian filtering?

- Robustness to outliers

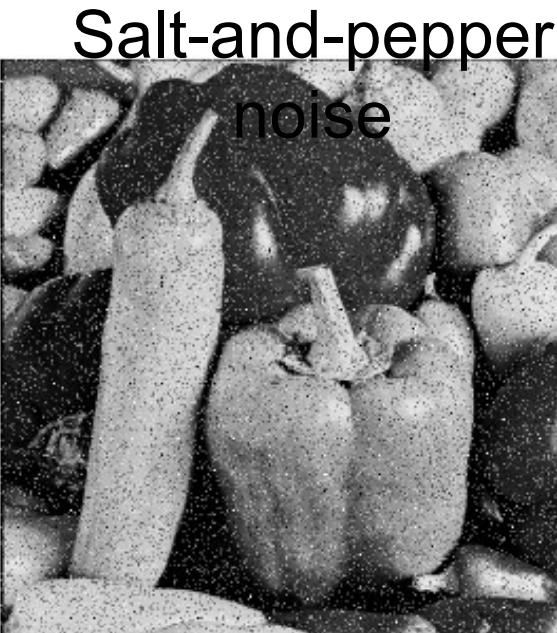
filters have width 5 :



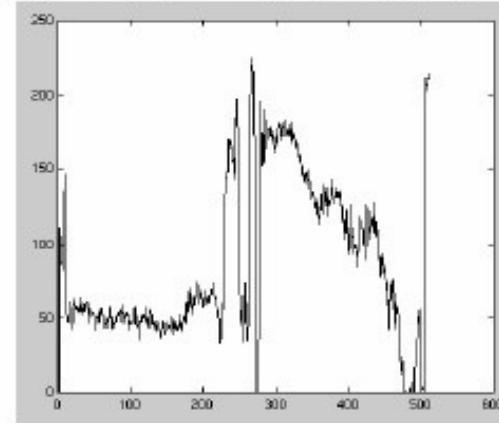
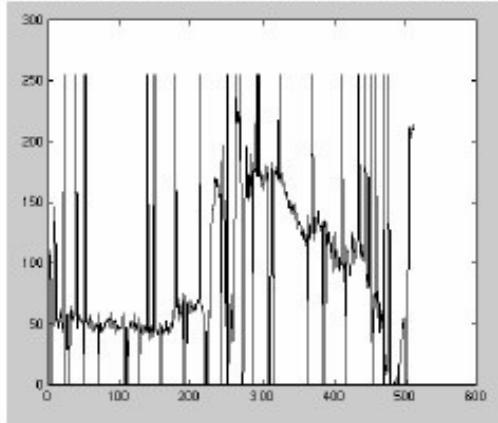
Source: K. Grauman

# Median filter

---



Median filtered



MATLAB: `medfilt2(image, [h w])`

Source: M. Hebert

# Median vs. Gaussian filtering

---

3x3



5x5



7x7



Gaussian

Median



# A Gentle Introduction to Bilateral Filtering and its Applications

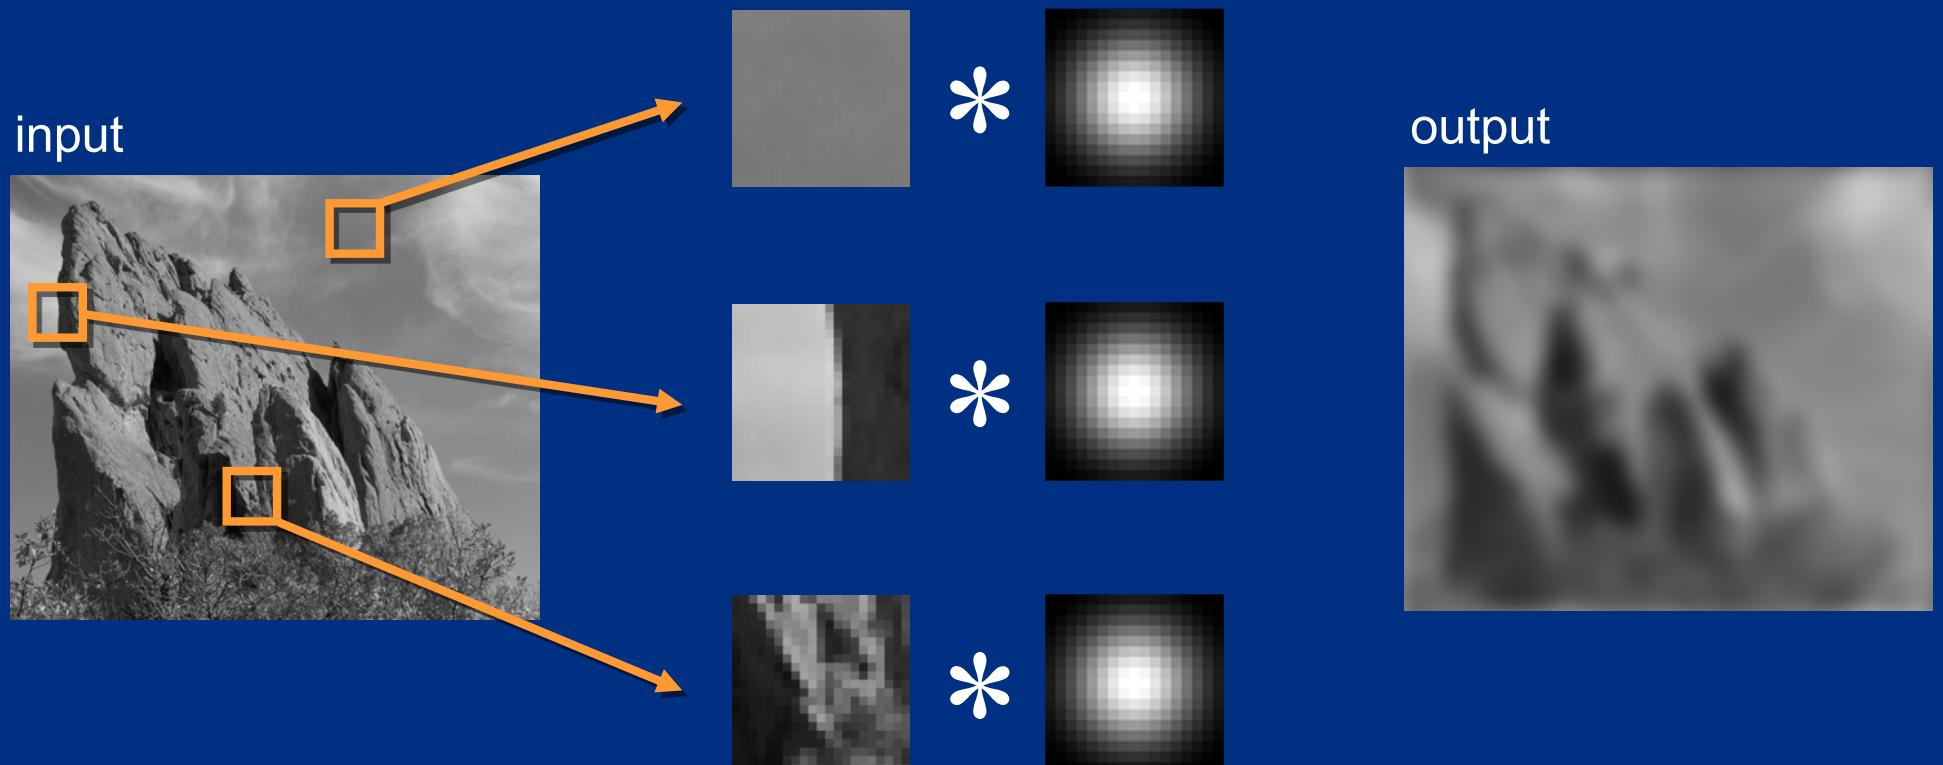


SIGGRAPH2007

## “Fixing the Gaussian Blur”: the Bilateral Filter

*Sylvain Paris – MIT CSAIL*

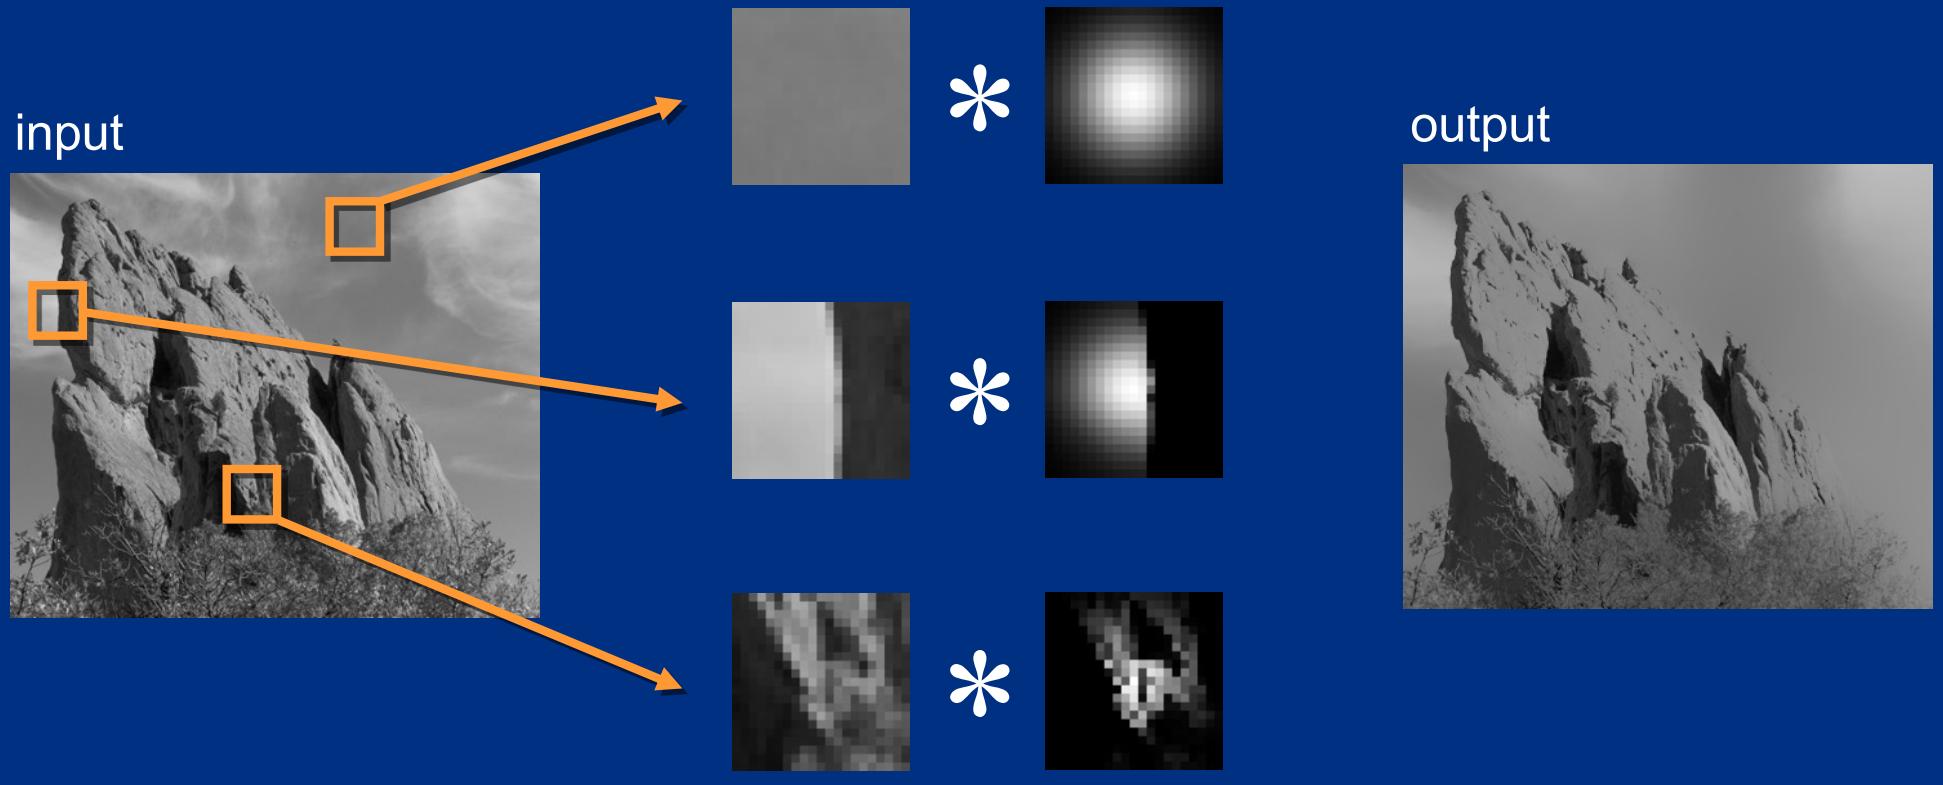
# Blur Comes from Averaging across Edges



Same Gaussian kernel everywhere.

# Bilateral Filter [Aurich 95, Smith 97, Tomasi 98]

## No Averaging across Edges



The kernel shape depends on the image content.

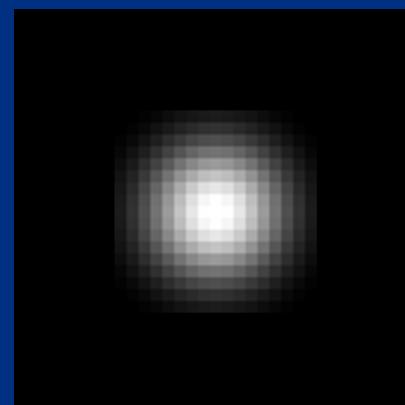
# Bilateral Filter Definition: an Additional Edge Term

Same idea: **weighted average of pixels.**

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

new  
not new  
new

normalization factor  
**space weight**  
**range weight**

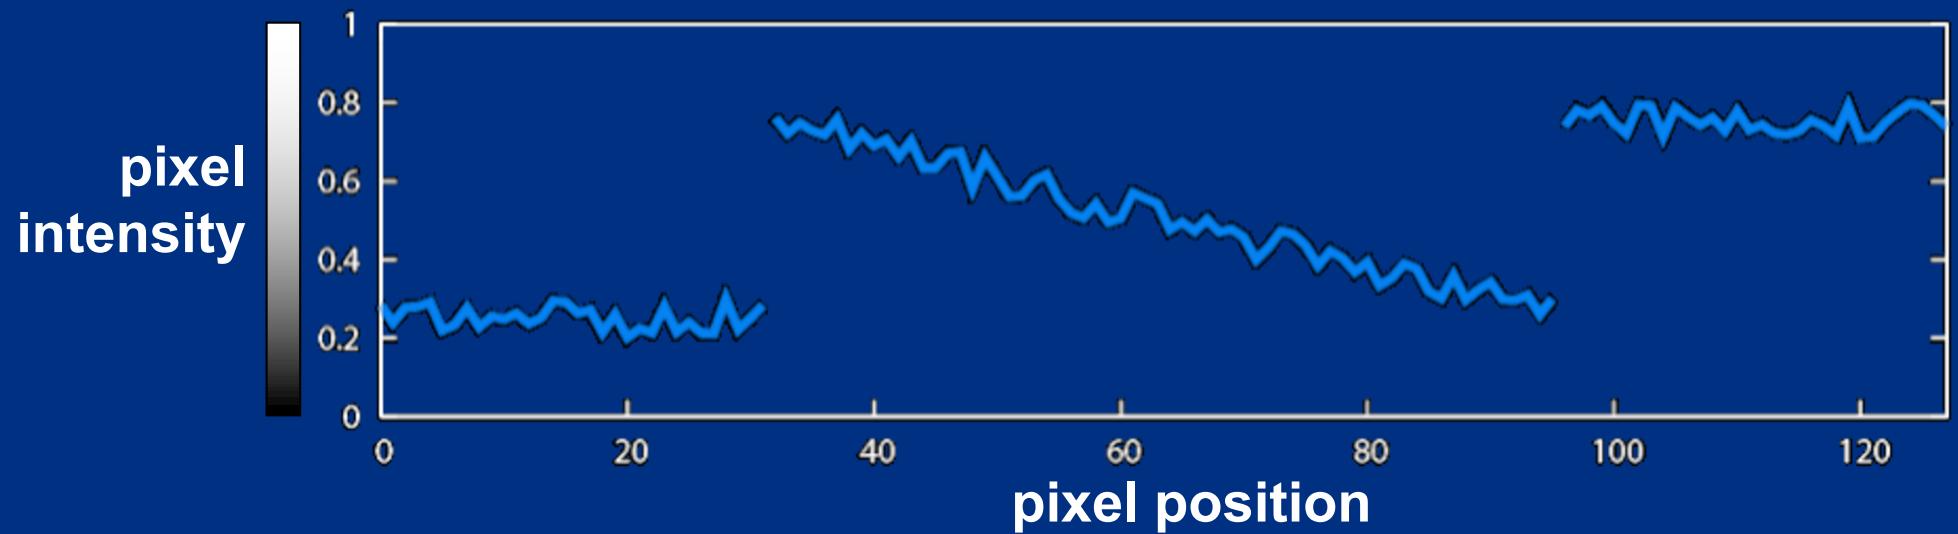


# Illustration a 1D Image

- 1D image = line of pixels

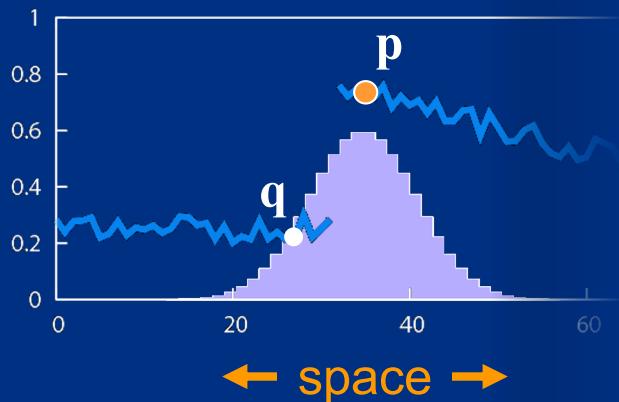


- Better visualized as a plot



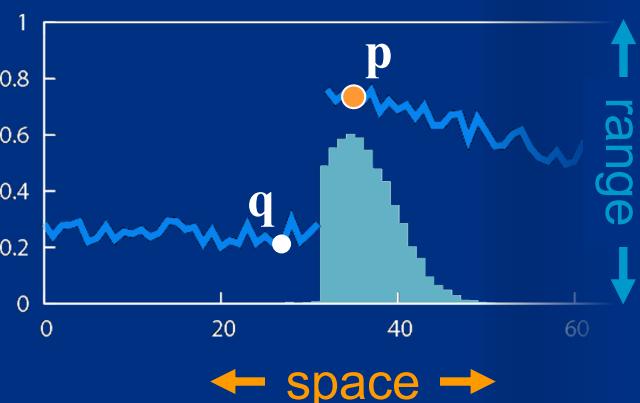
# Gaussian Blur and Bilateral Filter

# Gaussian blur



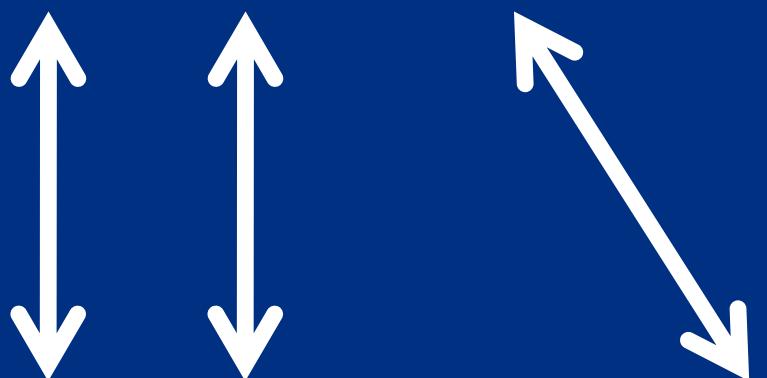
## Bilateral filter

[Aurich 95, Smith 97, Tomasi 98]



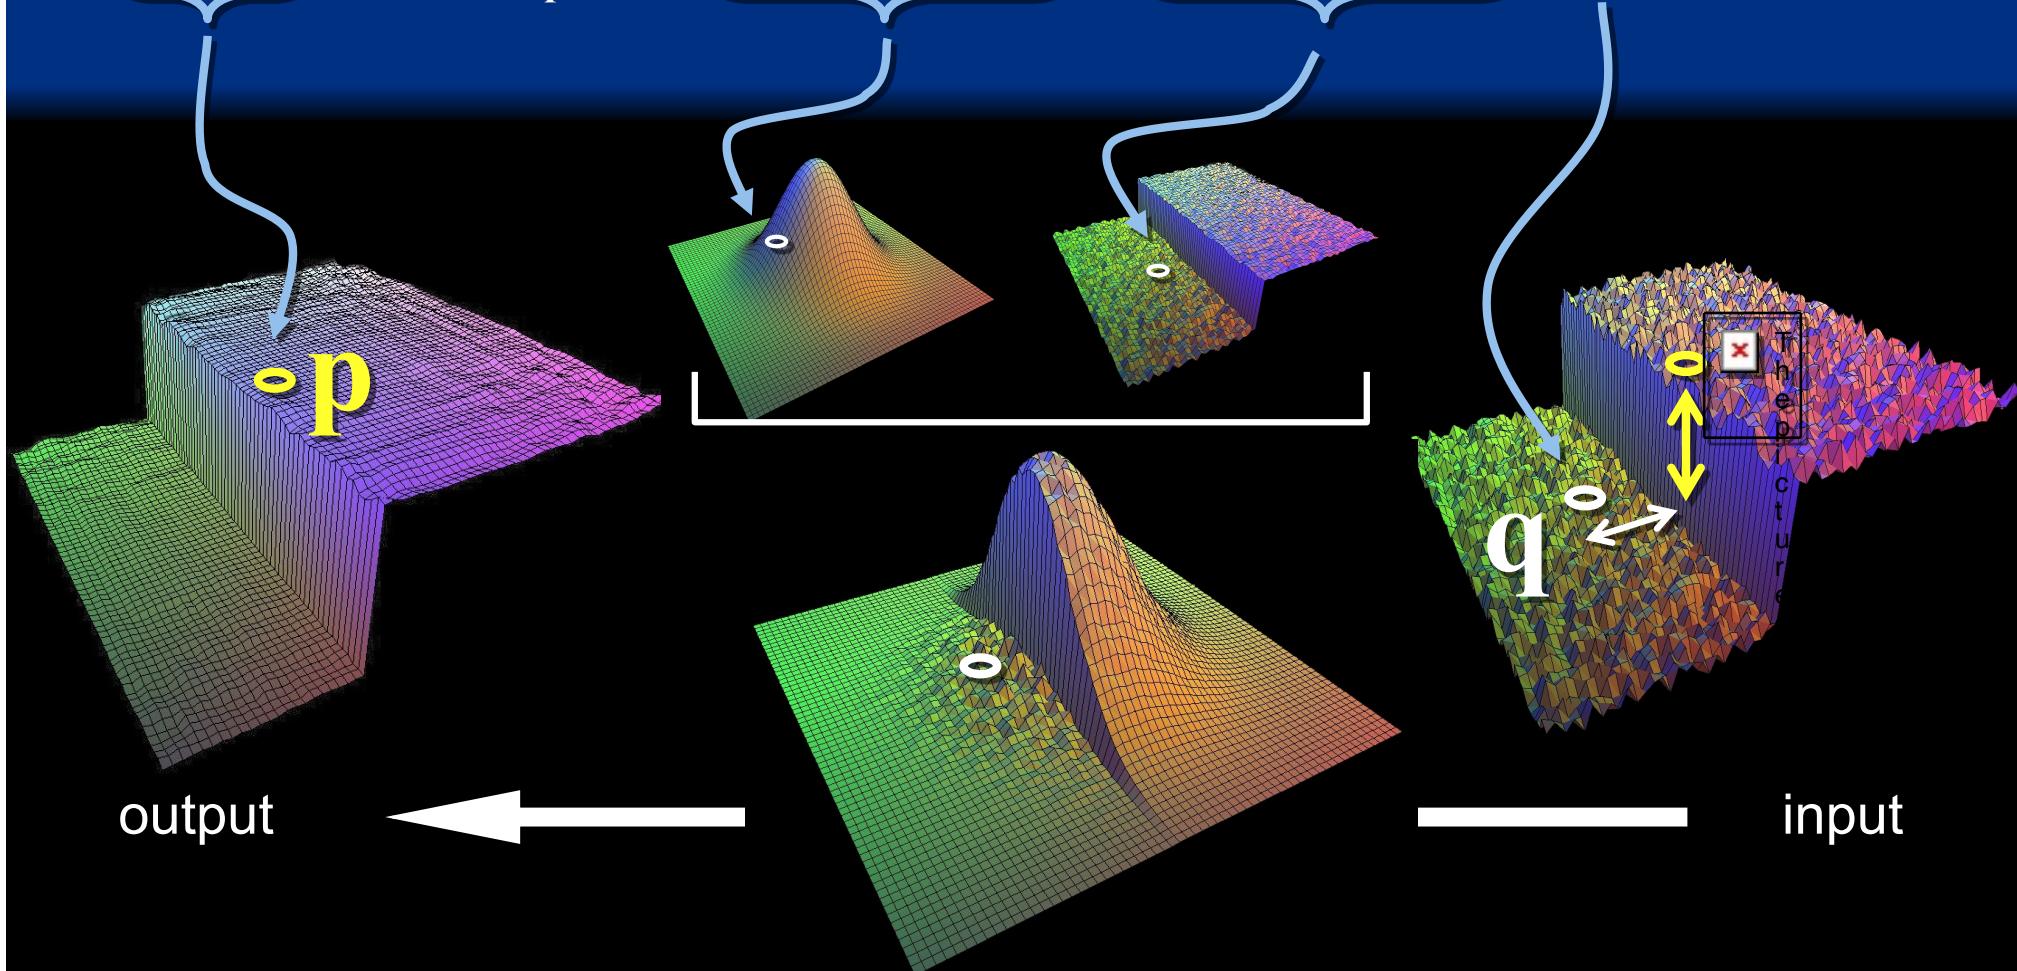
$$GB[I]_p = \sum_{q \in S} G_\sigma(\| p - q \|) I_q$$

space



# Bilateral Filter on a Height Field

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) \underbrace{G_{\sigma_r}(|I_p - I_q|)}_{I_q}$$



reproduced  
from [Durand 02]

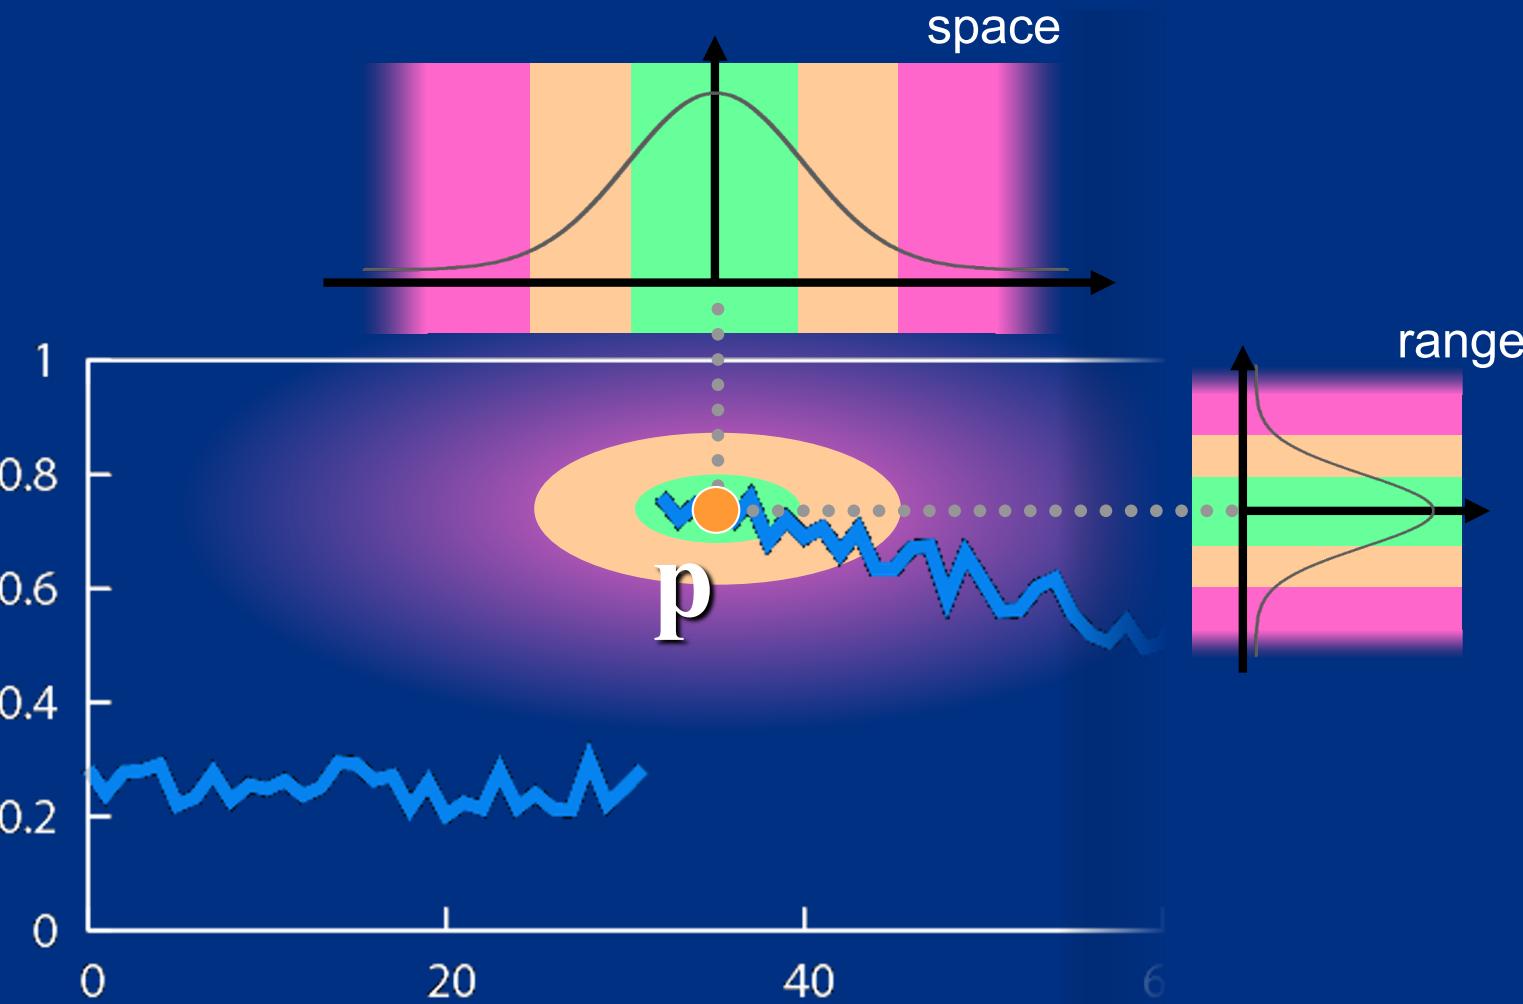
# Space and Range Parameters

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) G_{\sigma_r}(|I_p - I_q|) I_q$$


- space  $\sigma_s$  : spatial extent of the kernel, size of the considered neighborhood.
- range  $\sigma_r$  : “minimum” amplitude of an edge

# Influence of Pixels

Only pixels close in space and in range are considered.





input

## Exploring the Parameter Space

$\sigma_r = 0.1$



$\sigma_r = 0.25$



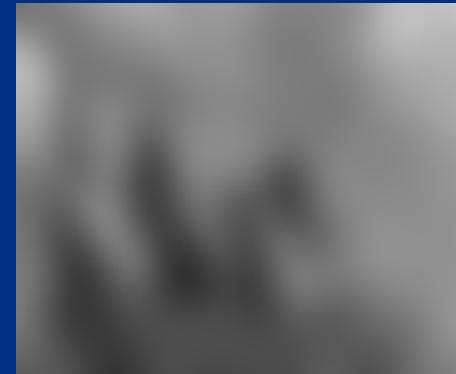
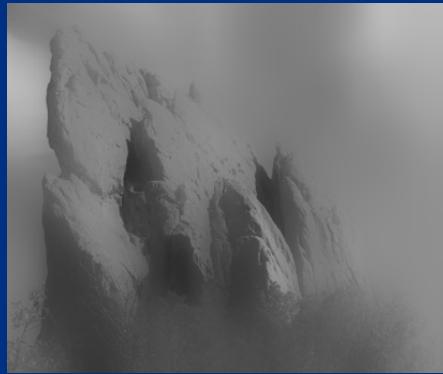
$\sigma_r = \infty$   
(Gaussian blur)



$\sigma_s = 2$



$\sigma_s = 6$



$\sigma_s = 18$

# Varying the Range Parameter

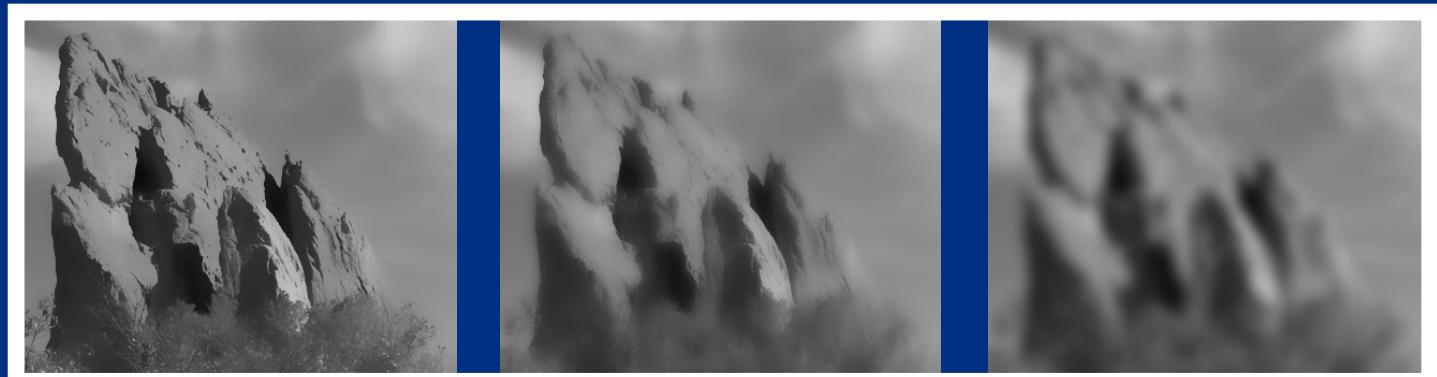


input

$\sigma_s = 2$



$\sigma_s = 6$



$\sigma_s = 18$



$\sigma_r = \infty$   
(Gaussian blur)

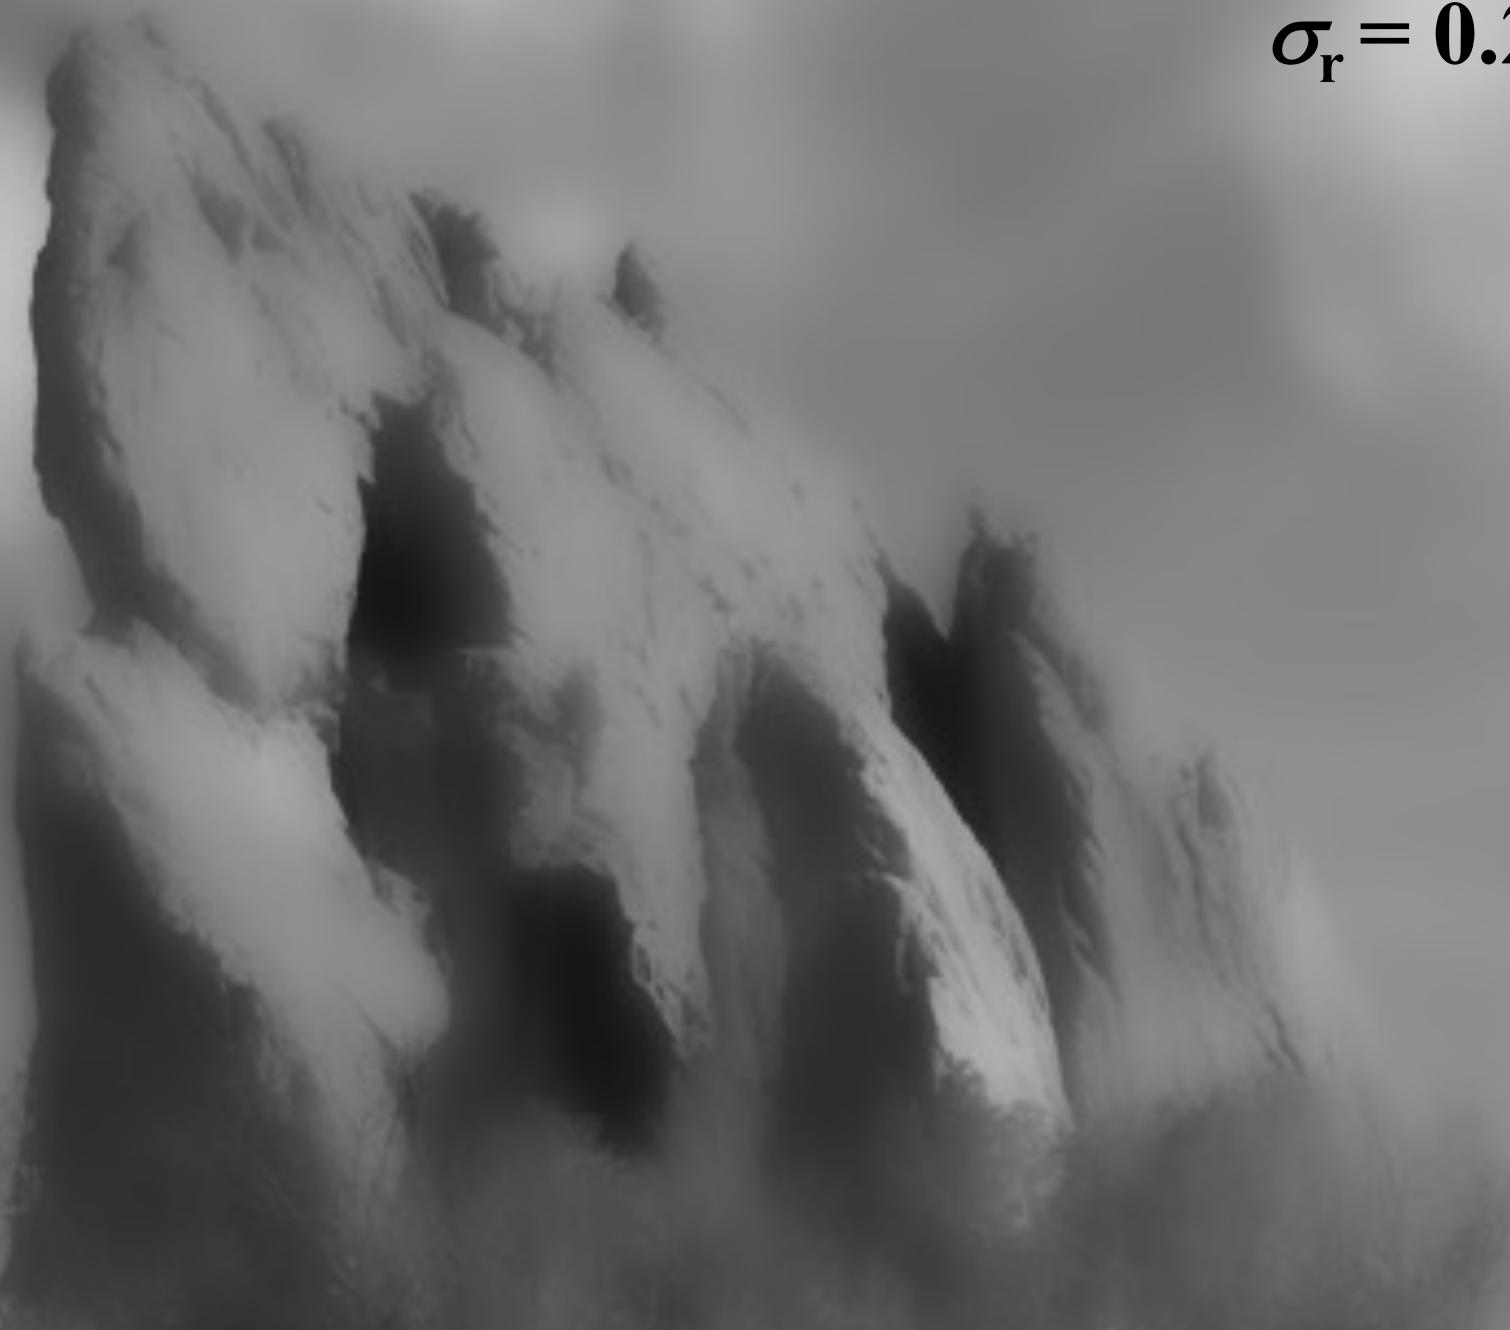
input



$\sigma_r = 0.1$



$\sigma_r = 0.25$





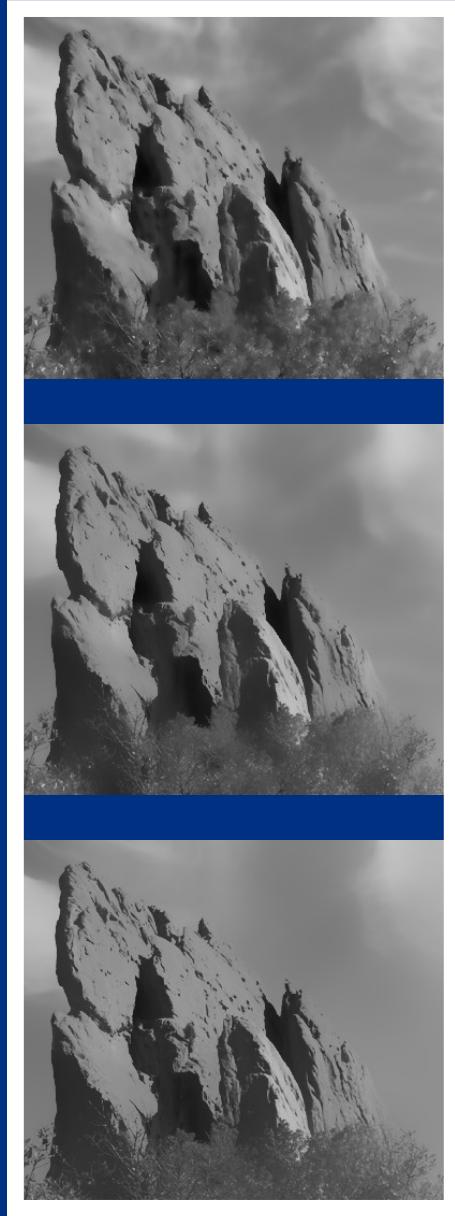
$\sigma_r = \infty$   
**(Gaussian blur)**

## Varying the Space Parameter



input

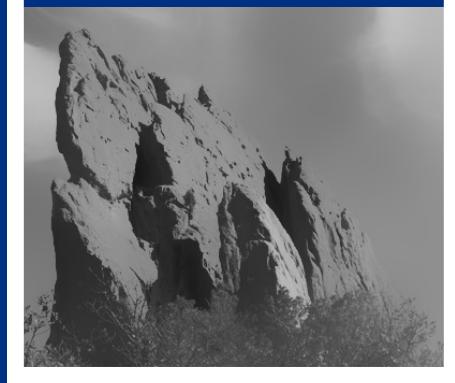
$\sigma_s = 2$



$\sigma_s = 6$



$\sigma_s = 18$



$\sigma_r = 0.1$

$\sigma_r = 0.25$

$\sigma_r = \infty$   
(Gaussian blur)



input



$\sigma_s = 2$ 

$\sigma_s = 6$ 

$\sigma_s = 18$ 